

Simulated data is all you need: Bayesian parameter inference for scientific simulators with SBI

Jan Boelts (Teusen) | EuroSciPy2024 | August 30th, 2024 | Szczecin

A joint
initiative

UNTER
NEHMER
TUM

* IPAI

About me

- AI Researcher at appliedAI

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb

About



About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb



About

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb



About

- Non-profit (for-impact) organization
- **transferlab.ai**

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb



About

2

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb



About

- Non-profit (for-impact) organization
- **transferlab.ai**
 - research
 - education

About me

- AI Researcher at appliedAI
- Background
 - Cognitive Science
 - Computational Neuroscience
 - “Machine Learning for Science”
- Maintainer of **sbi**
- **X, Mastodon:** @janfiete
- **GitHub:** @janfb



About

- Non-profit (for-impact) organization
- **transferlab.ai**
 - research
 - education
 - open source tools

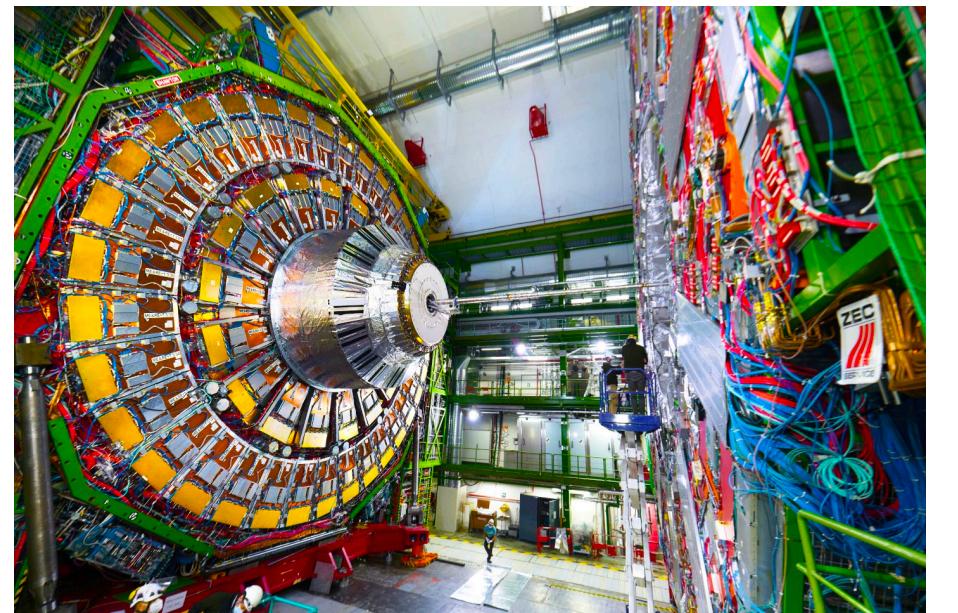
Outline

1. **Why** simulators?
2. **Why** Bayesian parameter inference?
3. **How** to do Bayesian inference for simulators?
4. **Demo:** How can you apply SBI to your simulator?

Outline

1. **Why** simulators?
2. **Why** Bayesian parameter inference?
3. **How** to do Bayesian inference for simulators?
4. **Demo:** How can you apply SBI to your simulator?

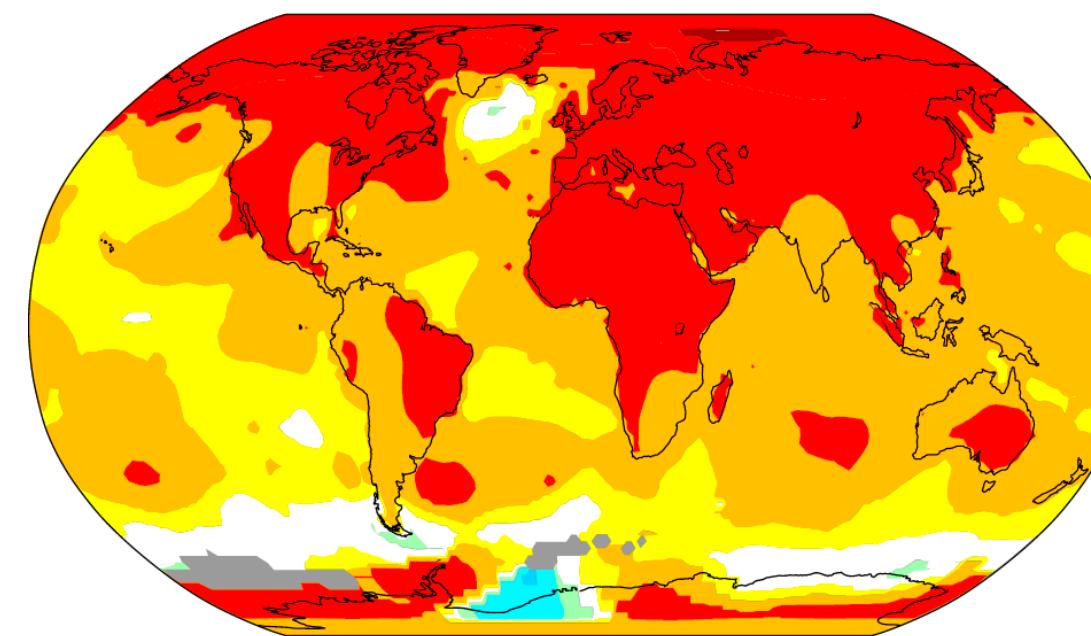
Simulators are used all across science



particle physics



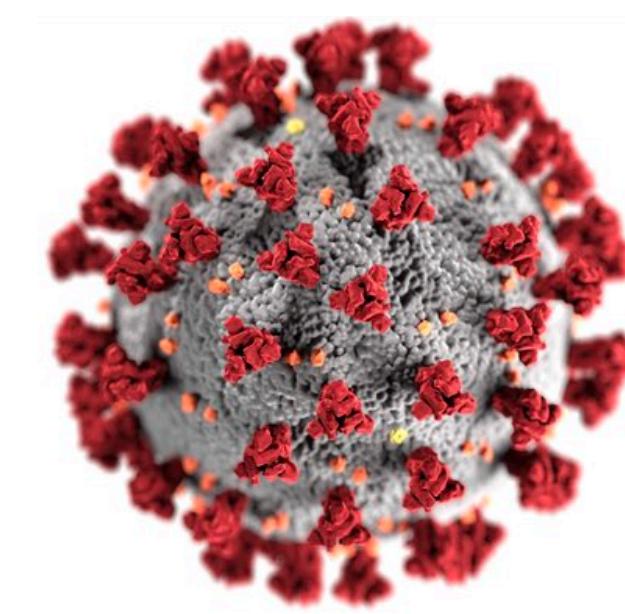
neuroscience



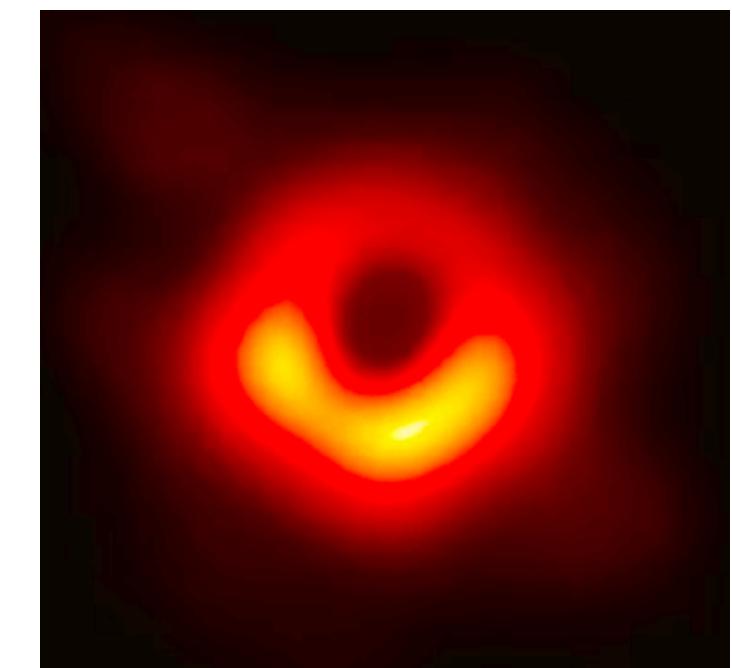
climate science



genomics



epidemiology



astrophysics

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

Examples:

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

Examples:

- Explicit **mechanistic** models, e.g., ODEs, PDEs: “simulator as numerical solver”

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

Examples:

- Explicit **mechanistic** models, e.g., ODEs, PDEs: “simulator as numerical solver”
- Implicit models: “simulator as black-box computer simulation”

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

Examples:

- Explicit **mechanistic** models, e.g., ODEs, PDEs: “simulator as numerical solver”
- Implicit models: “simulator as black-box computer simulation”
- Anything in-between: input parameter θ , output simulated data x

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

Examples:

- Explicit **mechanistic** models, e.g., ODEs, PDEs: “simulator as numerical solver”
- Implicit models: “simulator as black-box computer simulation”
- Anything in-between: input parameter θ , output simulated data x

$$x = \text{simulator}(\theta)$$

What is a simulator?

Simulator: A computer program that takes input parameters θ and returns simulated data x

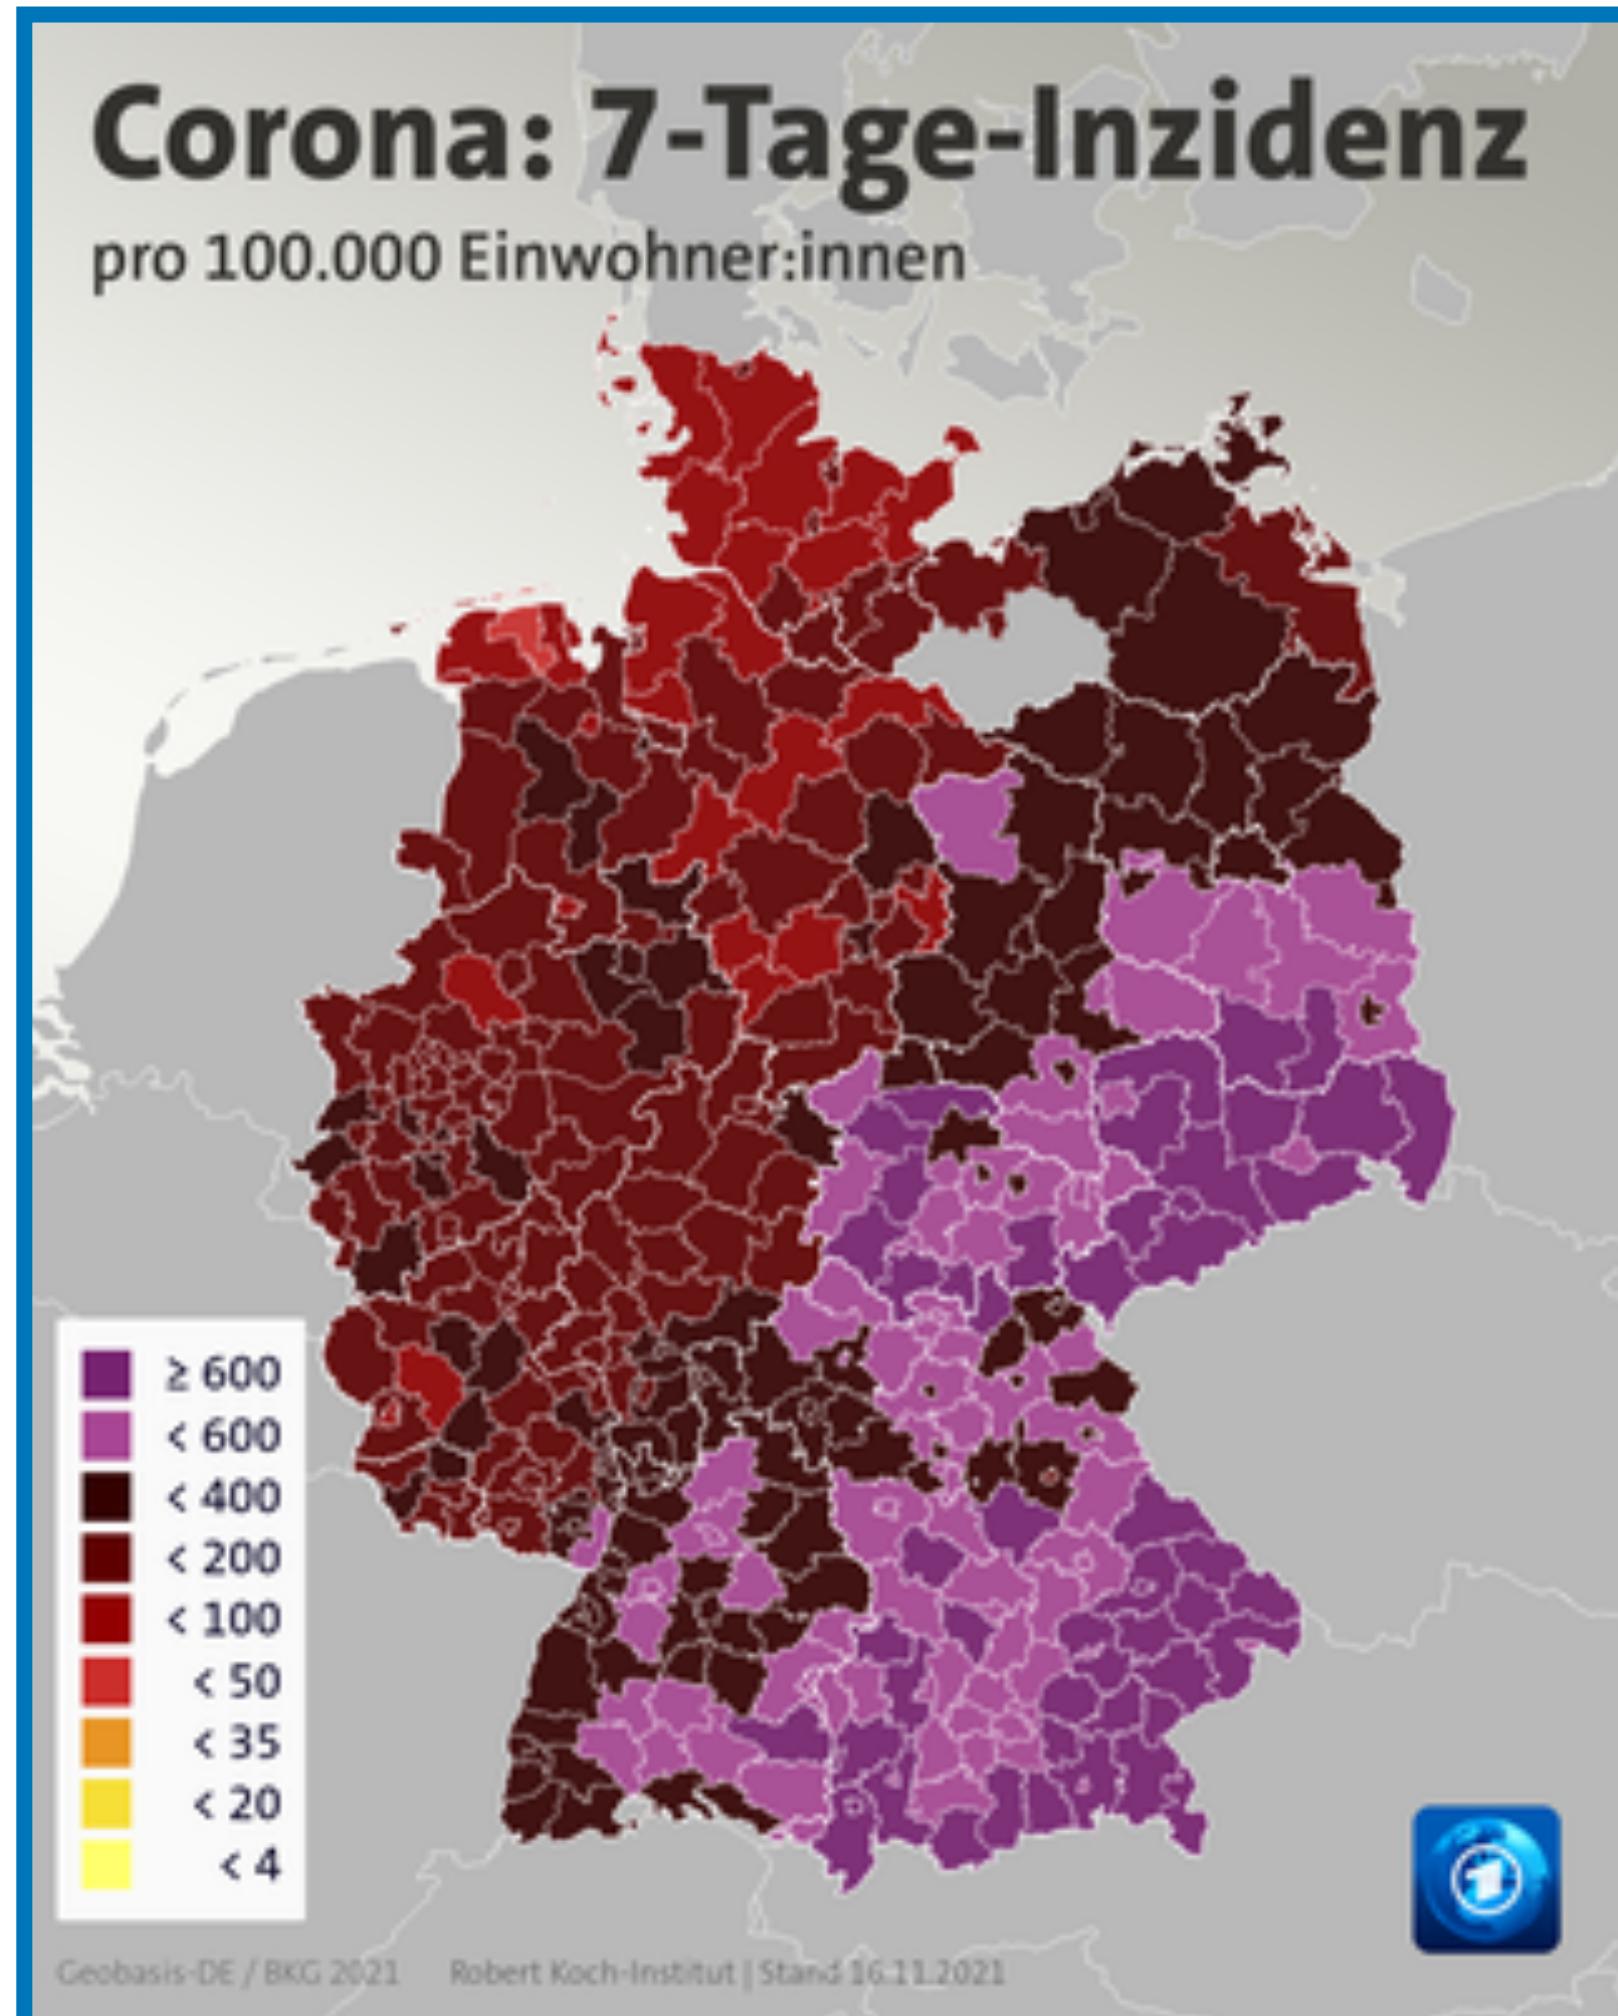
Examples:

- Explicit **mechanistic** models, e.g., ODEs, PDEs: “simulator as numerical solver”
- Implicit models: “simulator as black-box computer simulation”
- Anything in-between: input parameter θ , output simulated data x

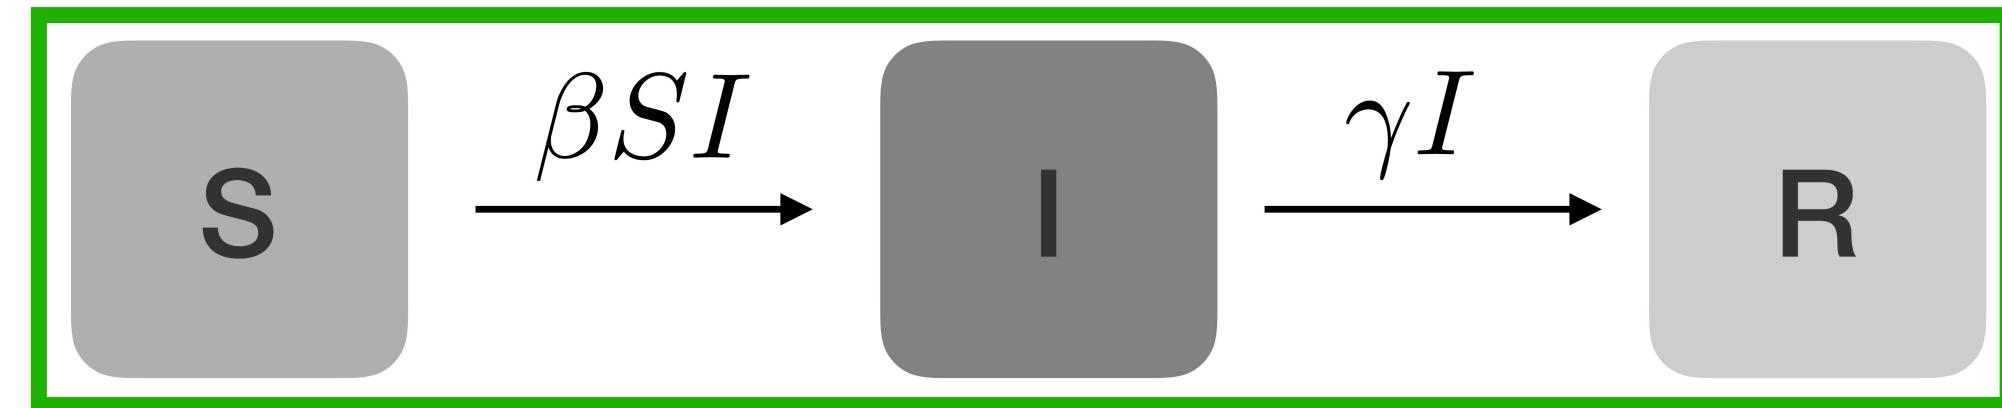
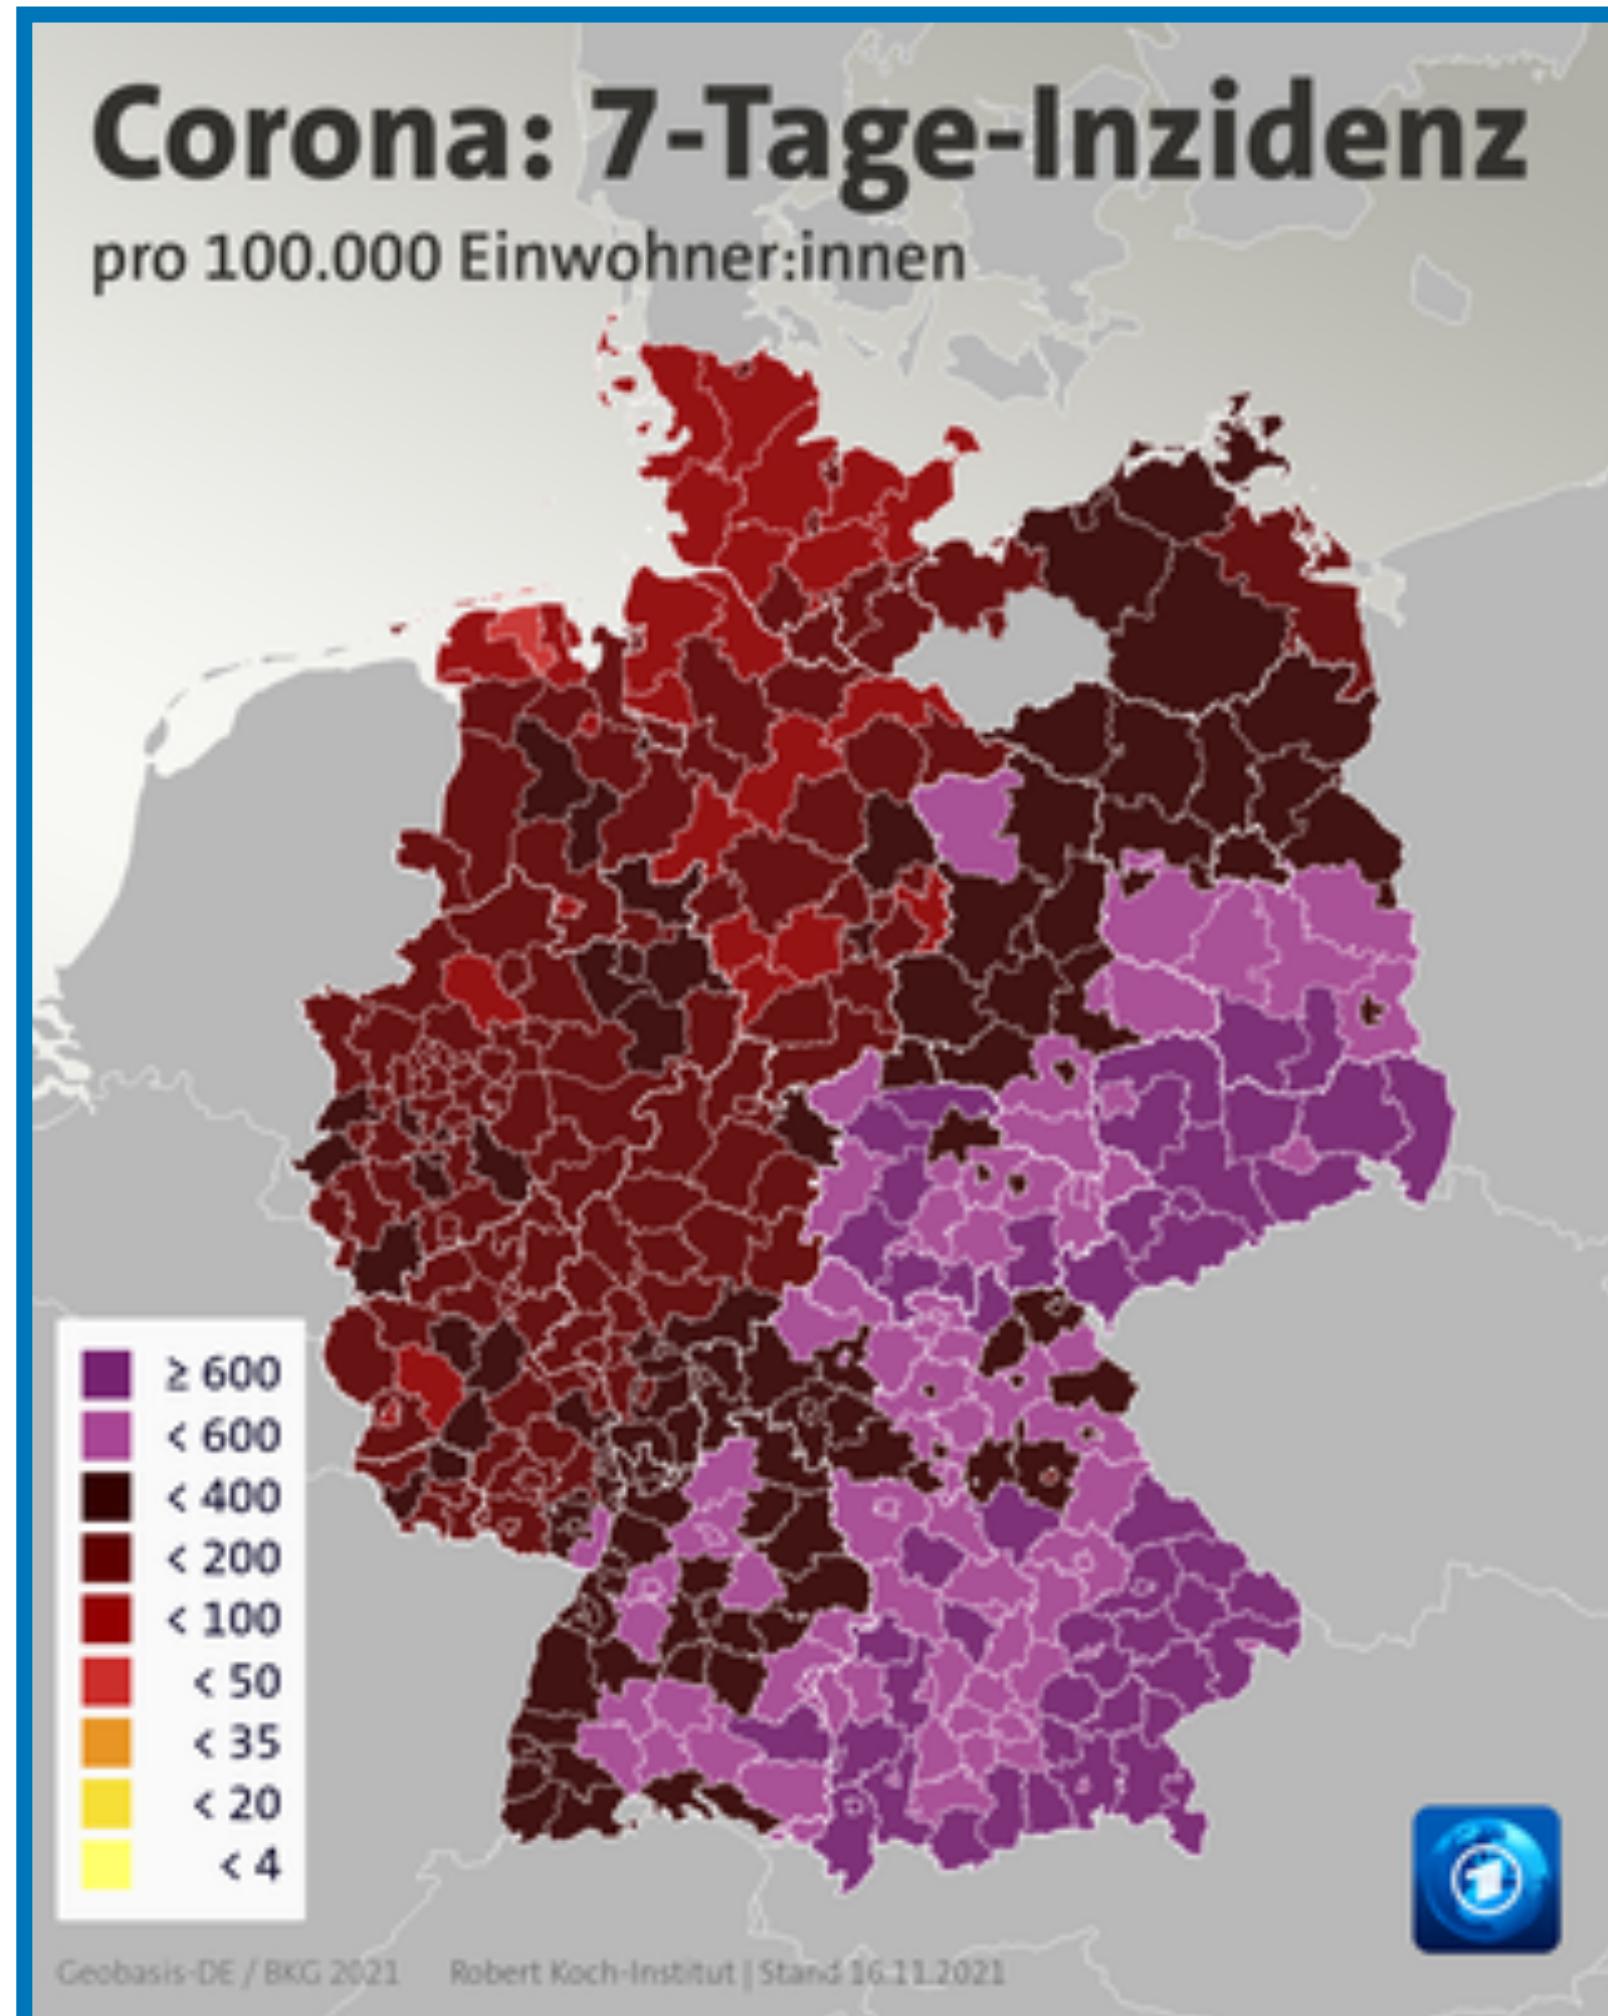
$$x = \text{simulator}(\theta)$$

- **no access** to the inner workings (no likelihoods, no gradients)

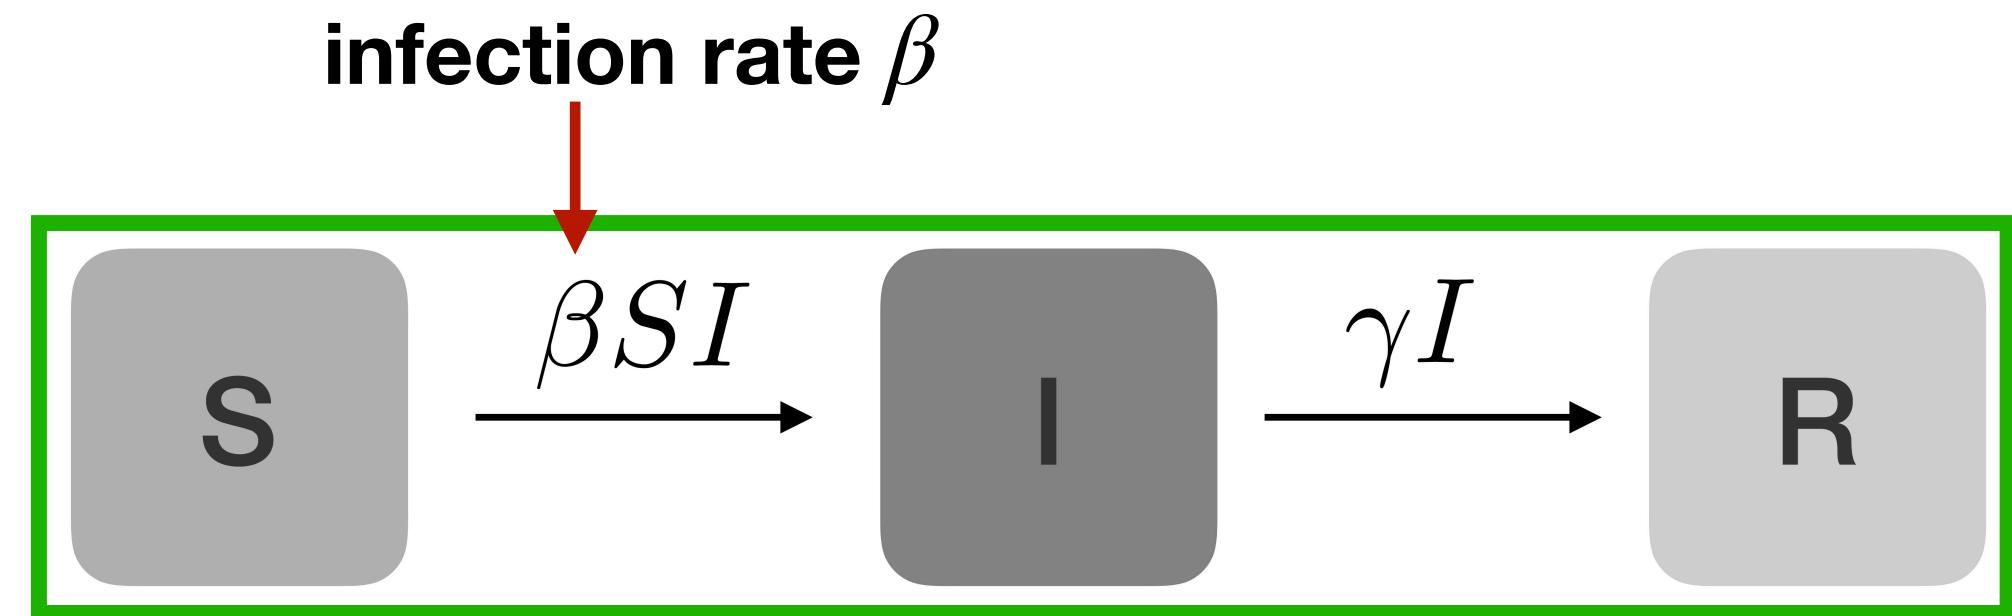
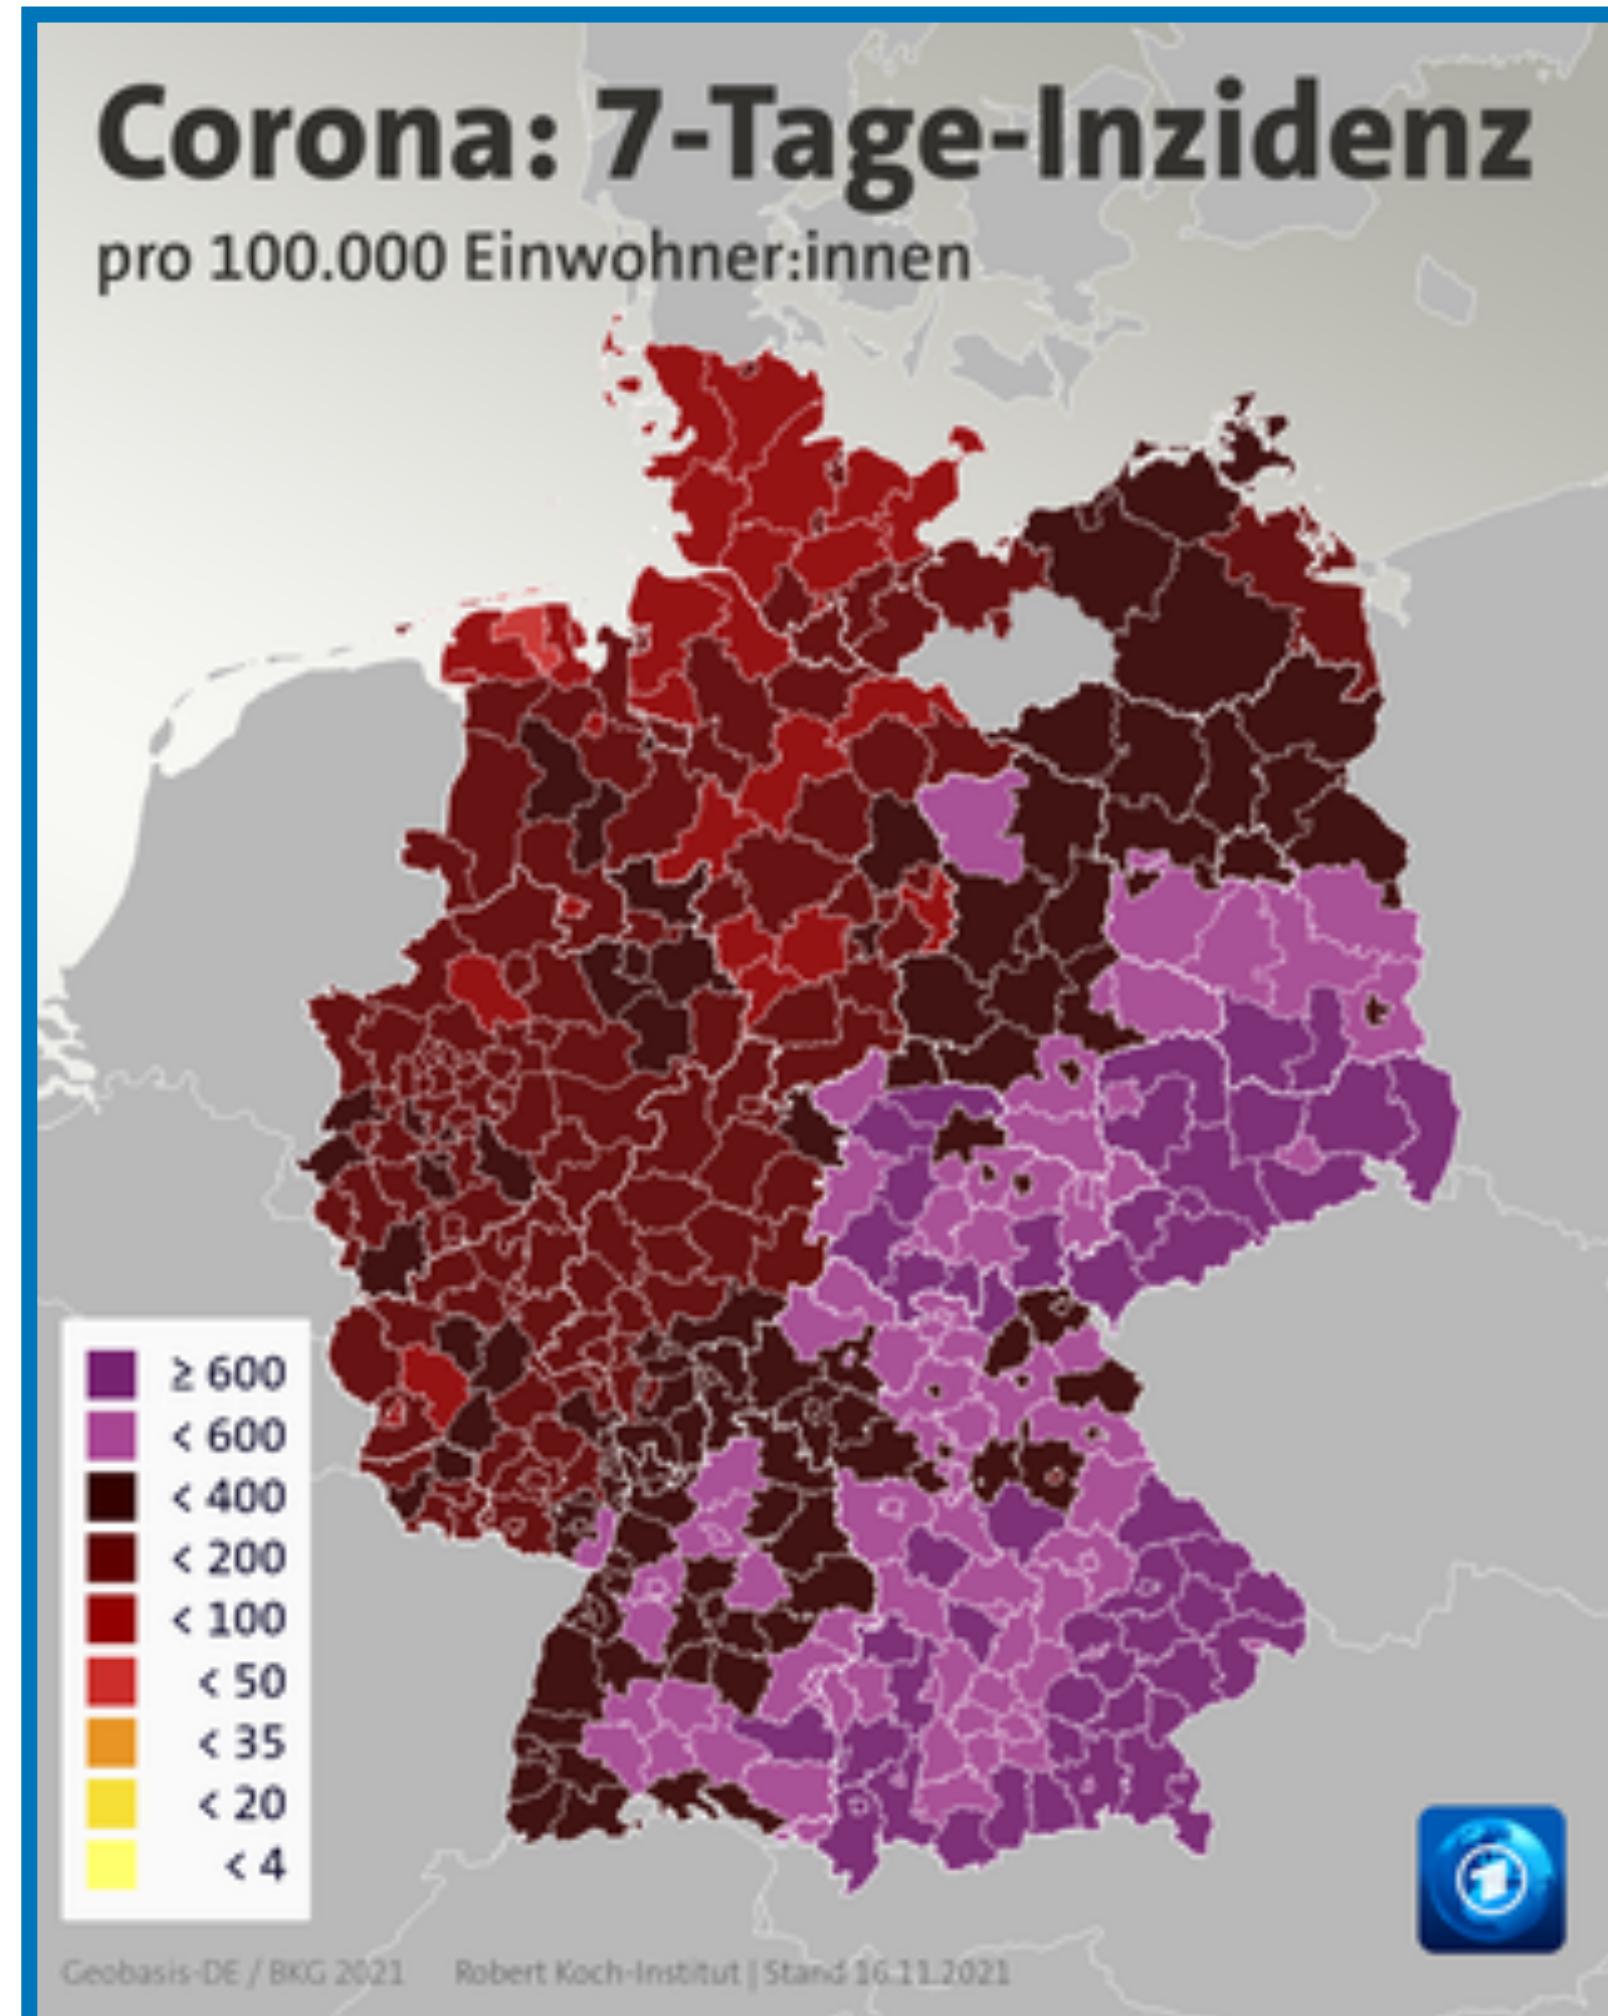
Example: The SIR Simulator for COVID-19



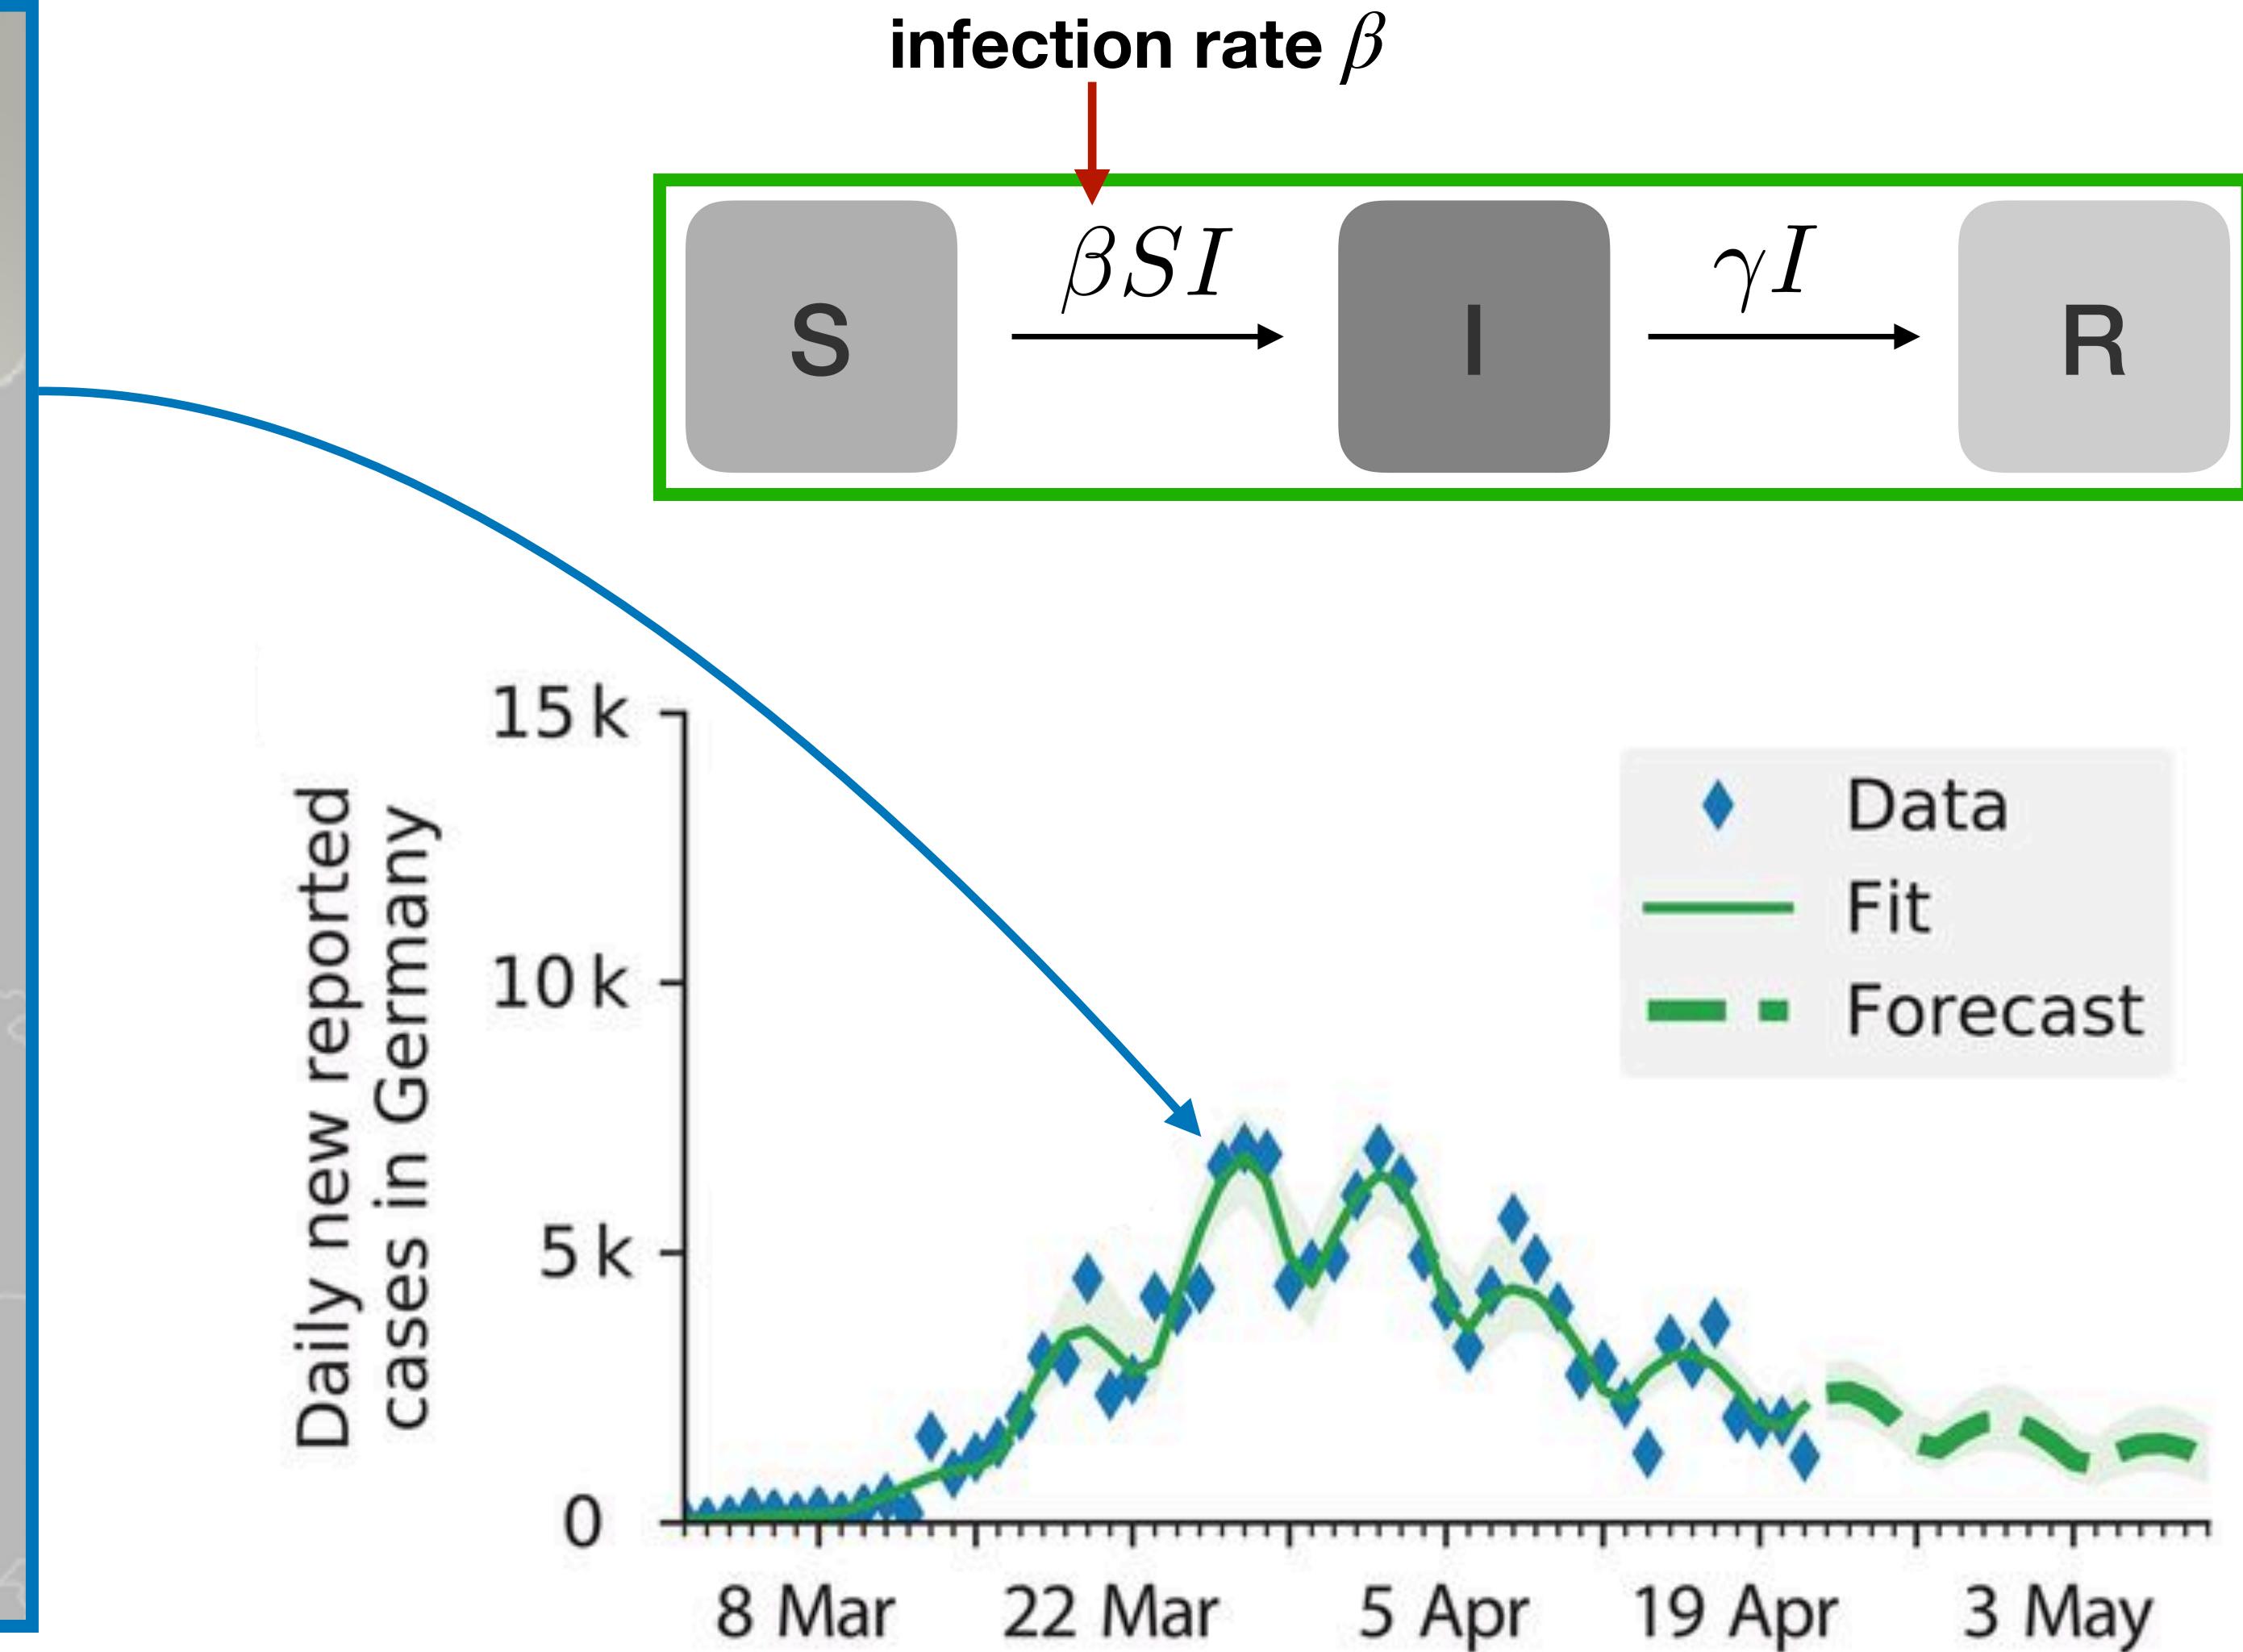
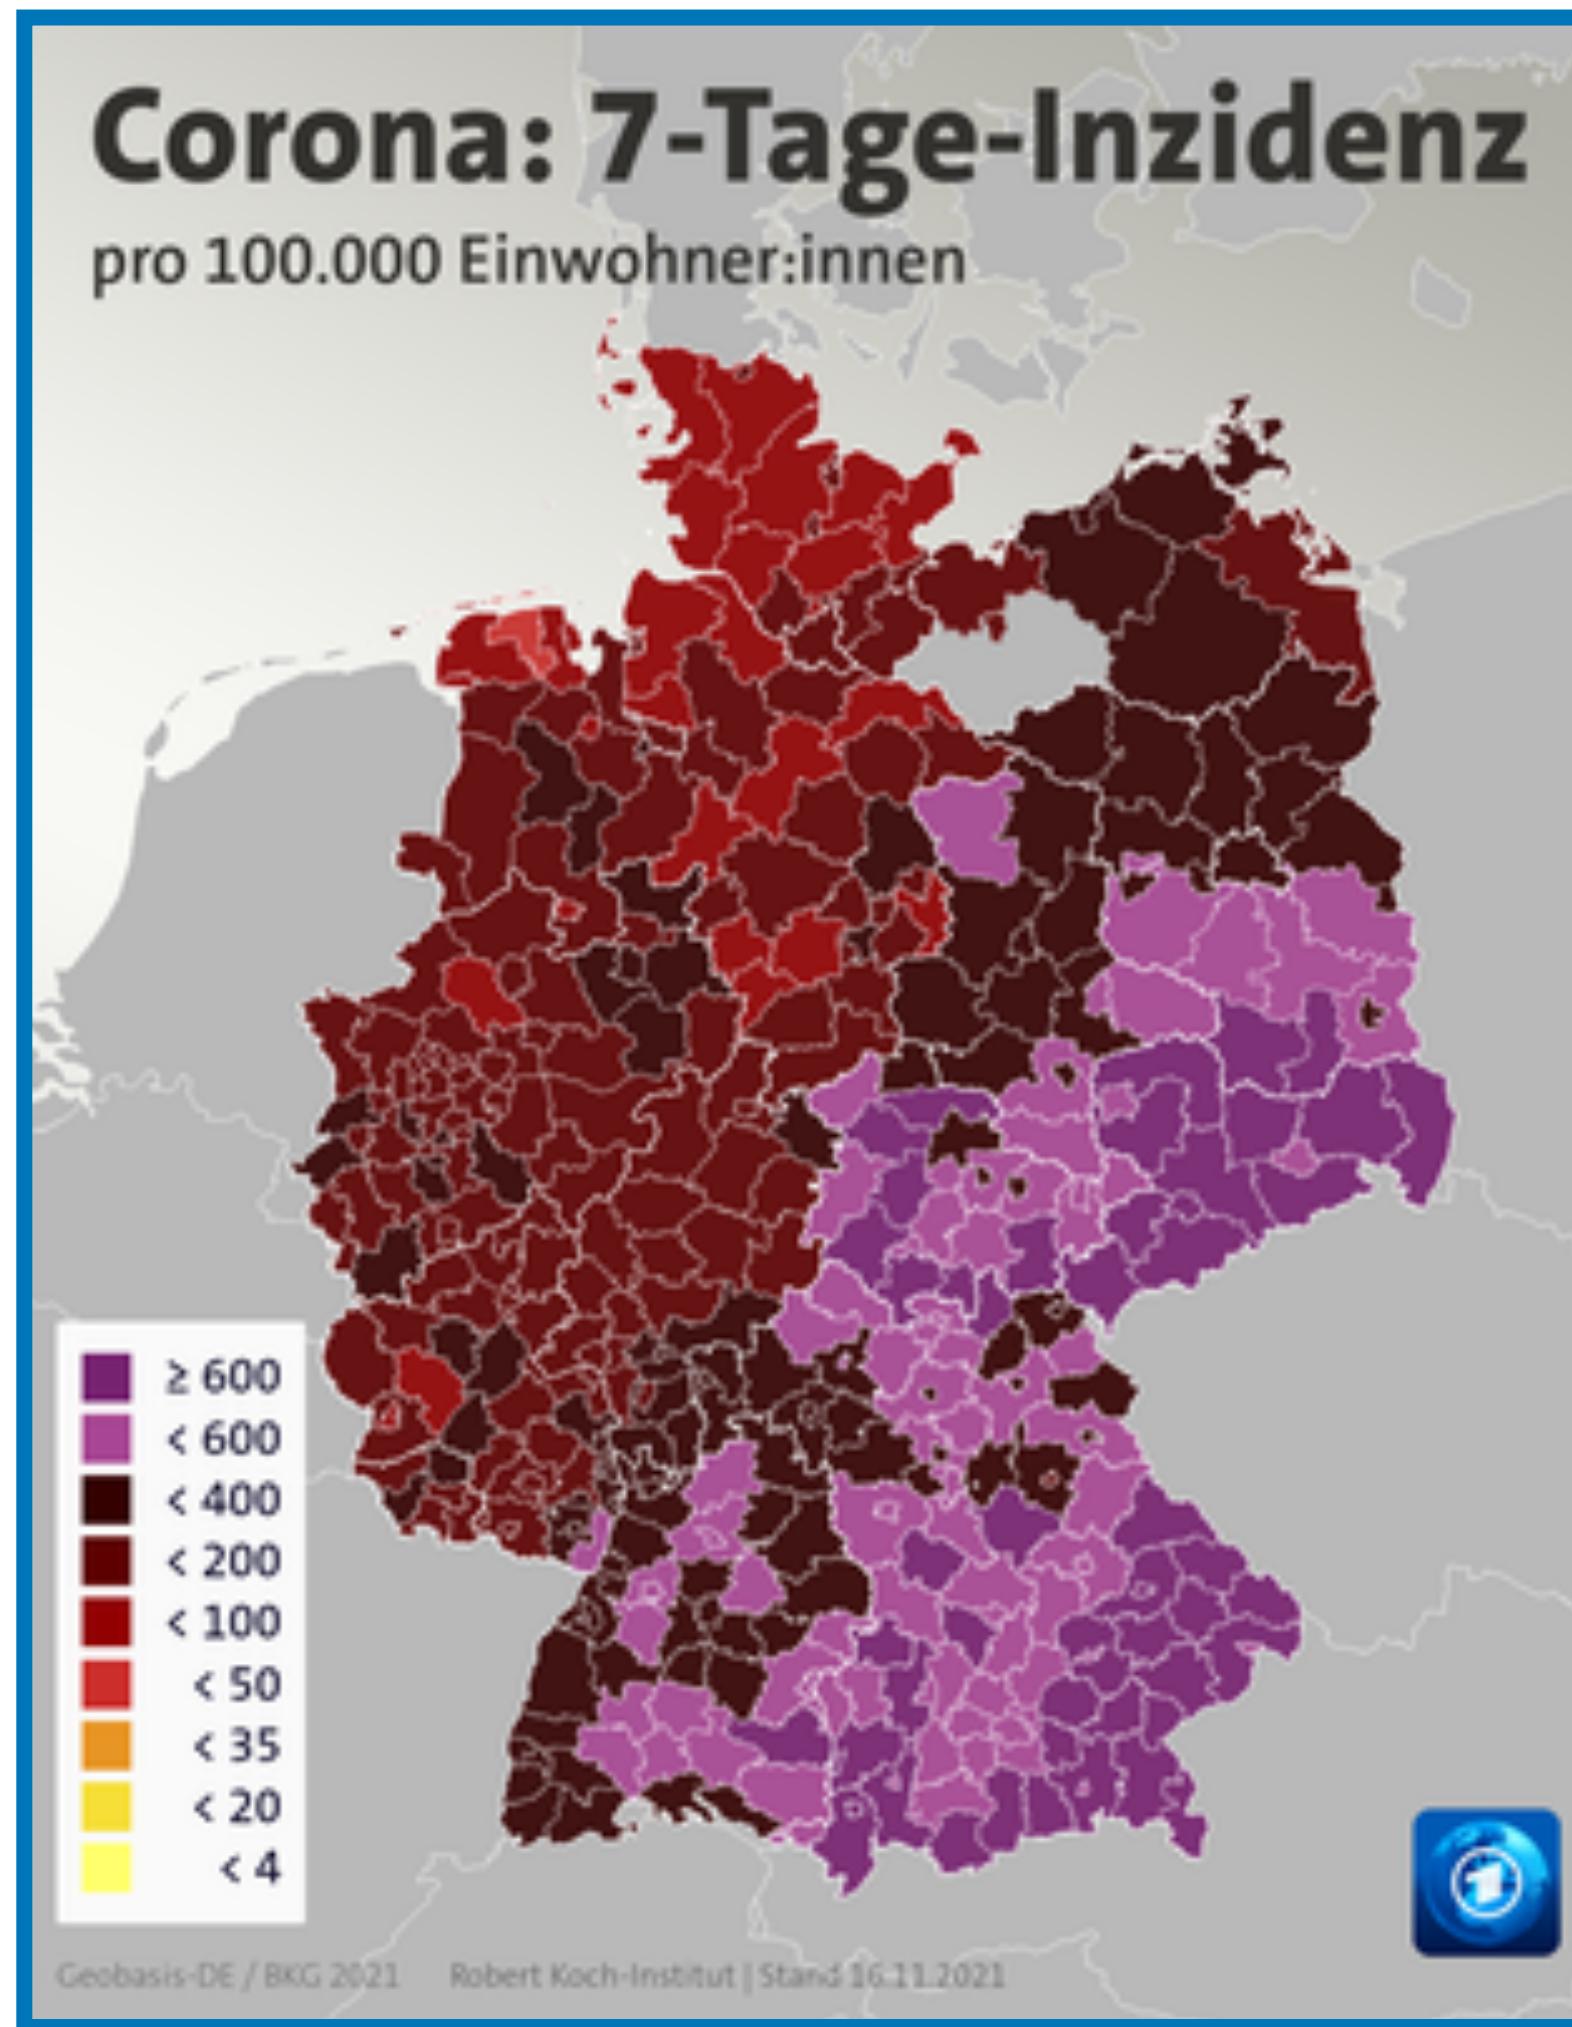
Example: The SIR Simulator for COVID-19



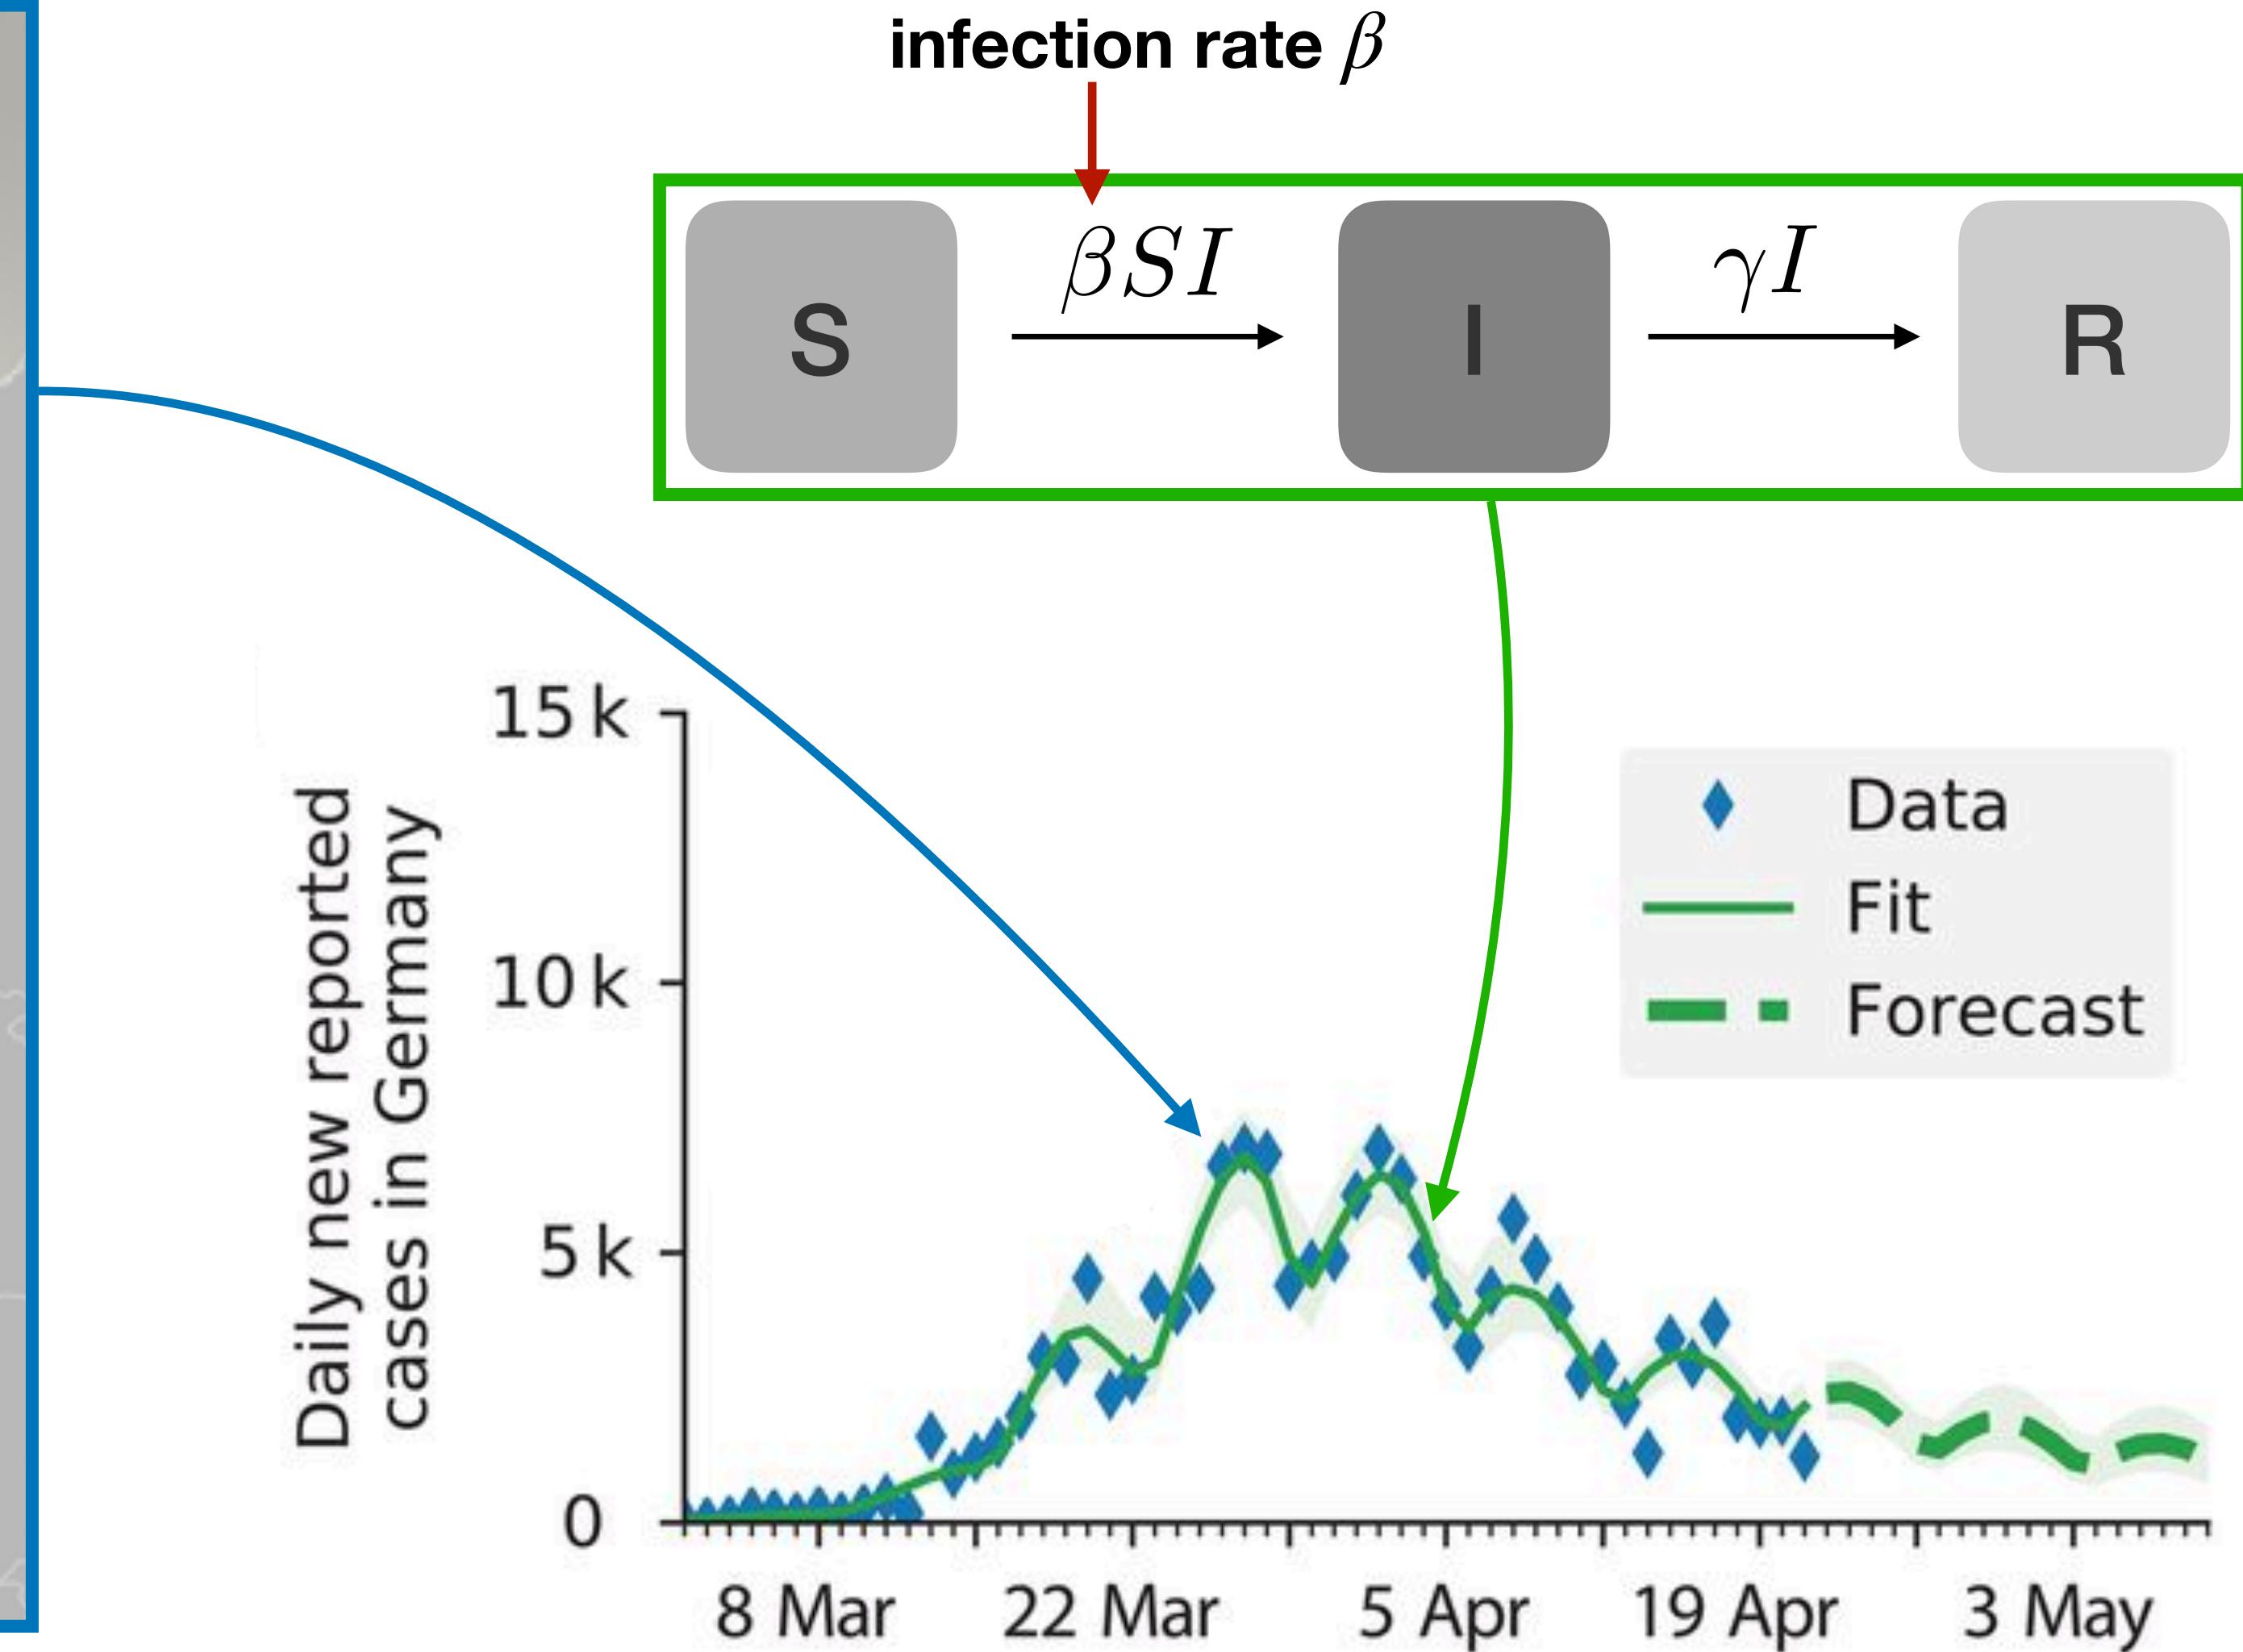
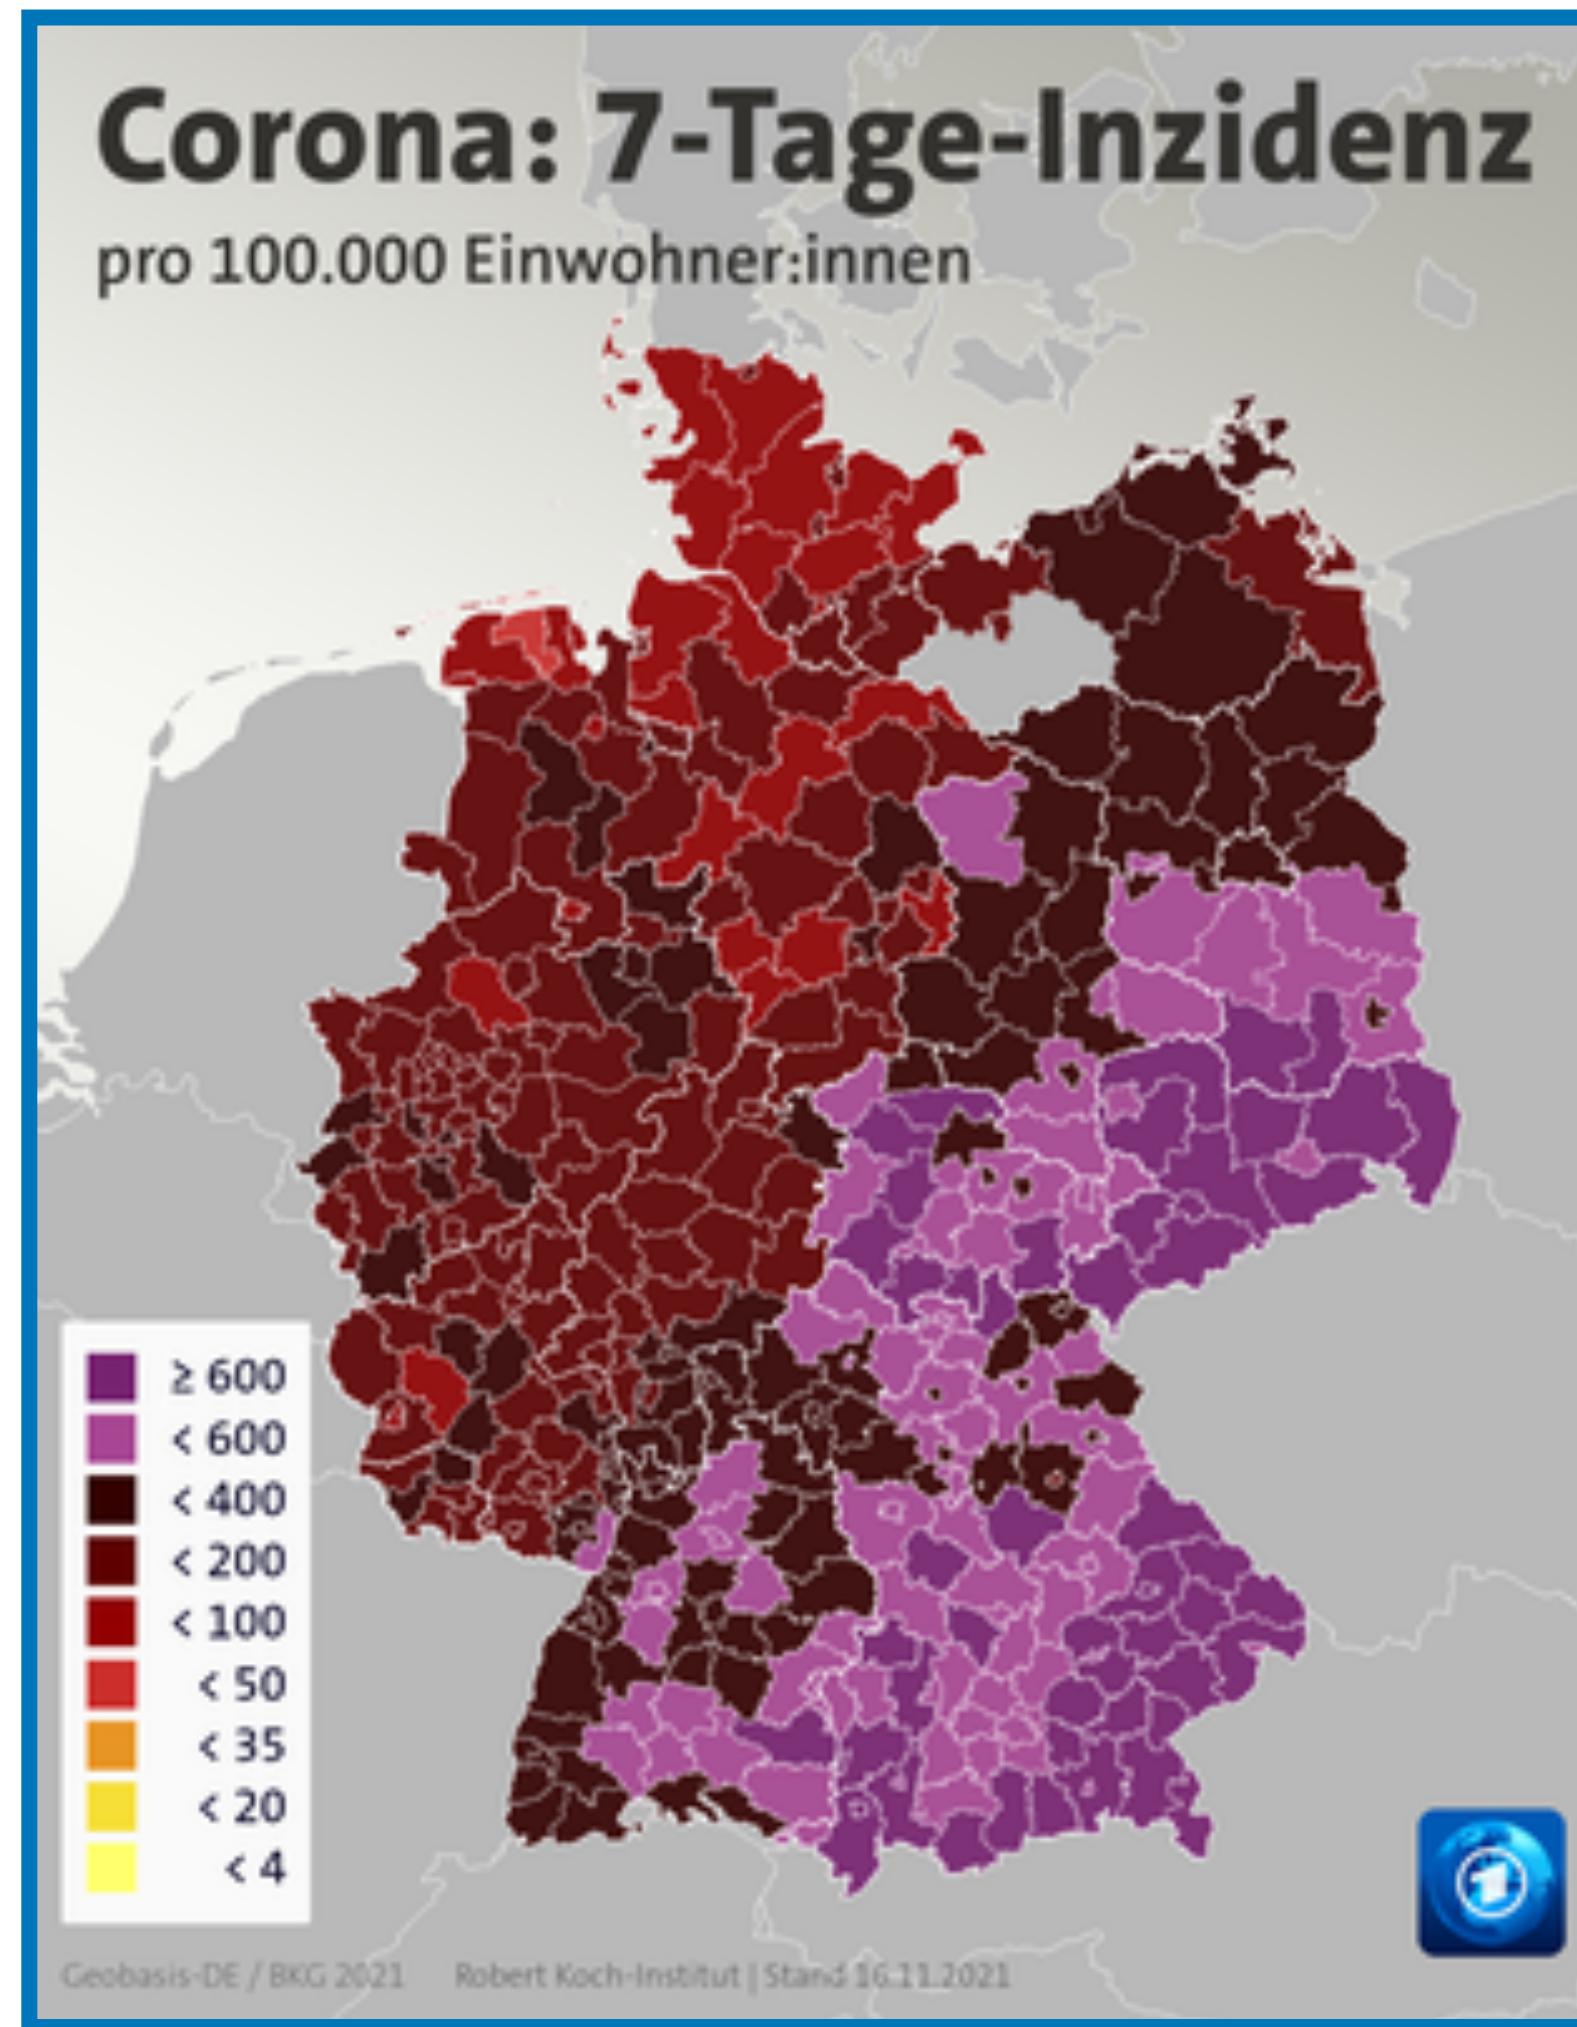
Example: The SIR Simulator for COVID-19



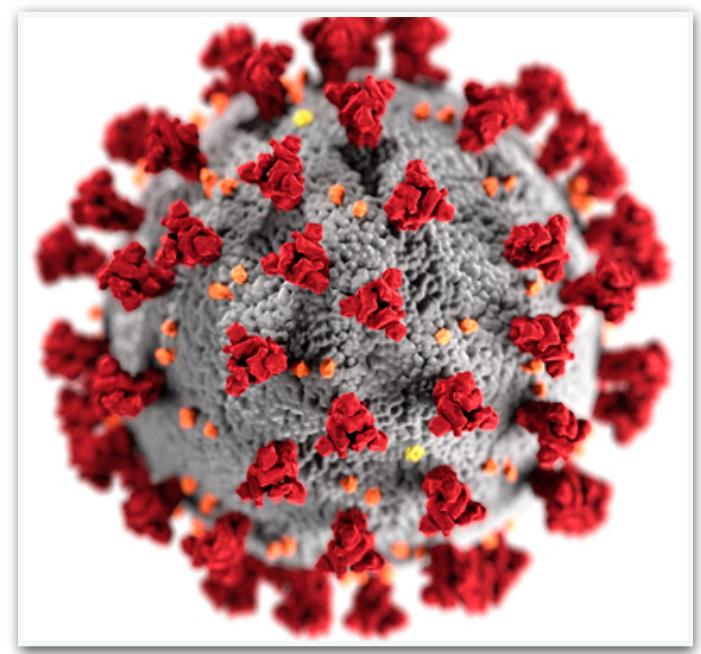
Example: The SIR Simulator for COVID-19



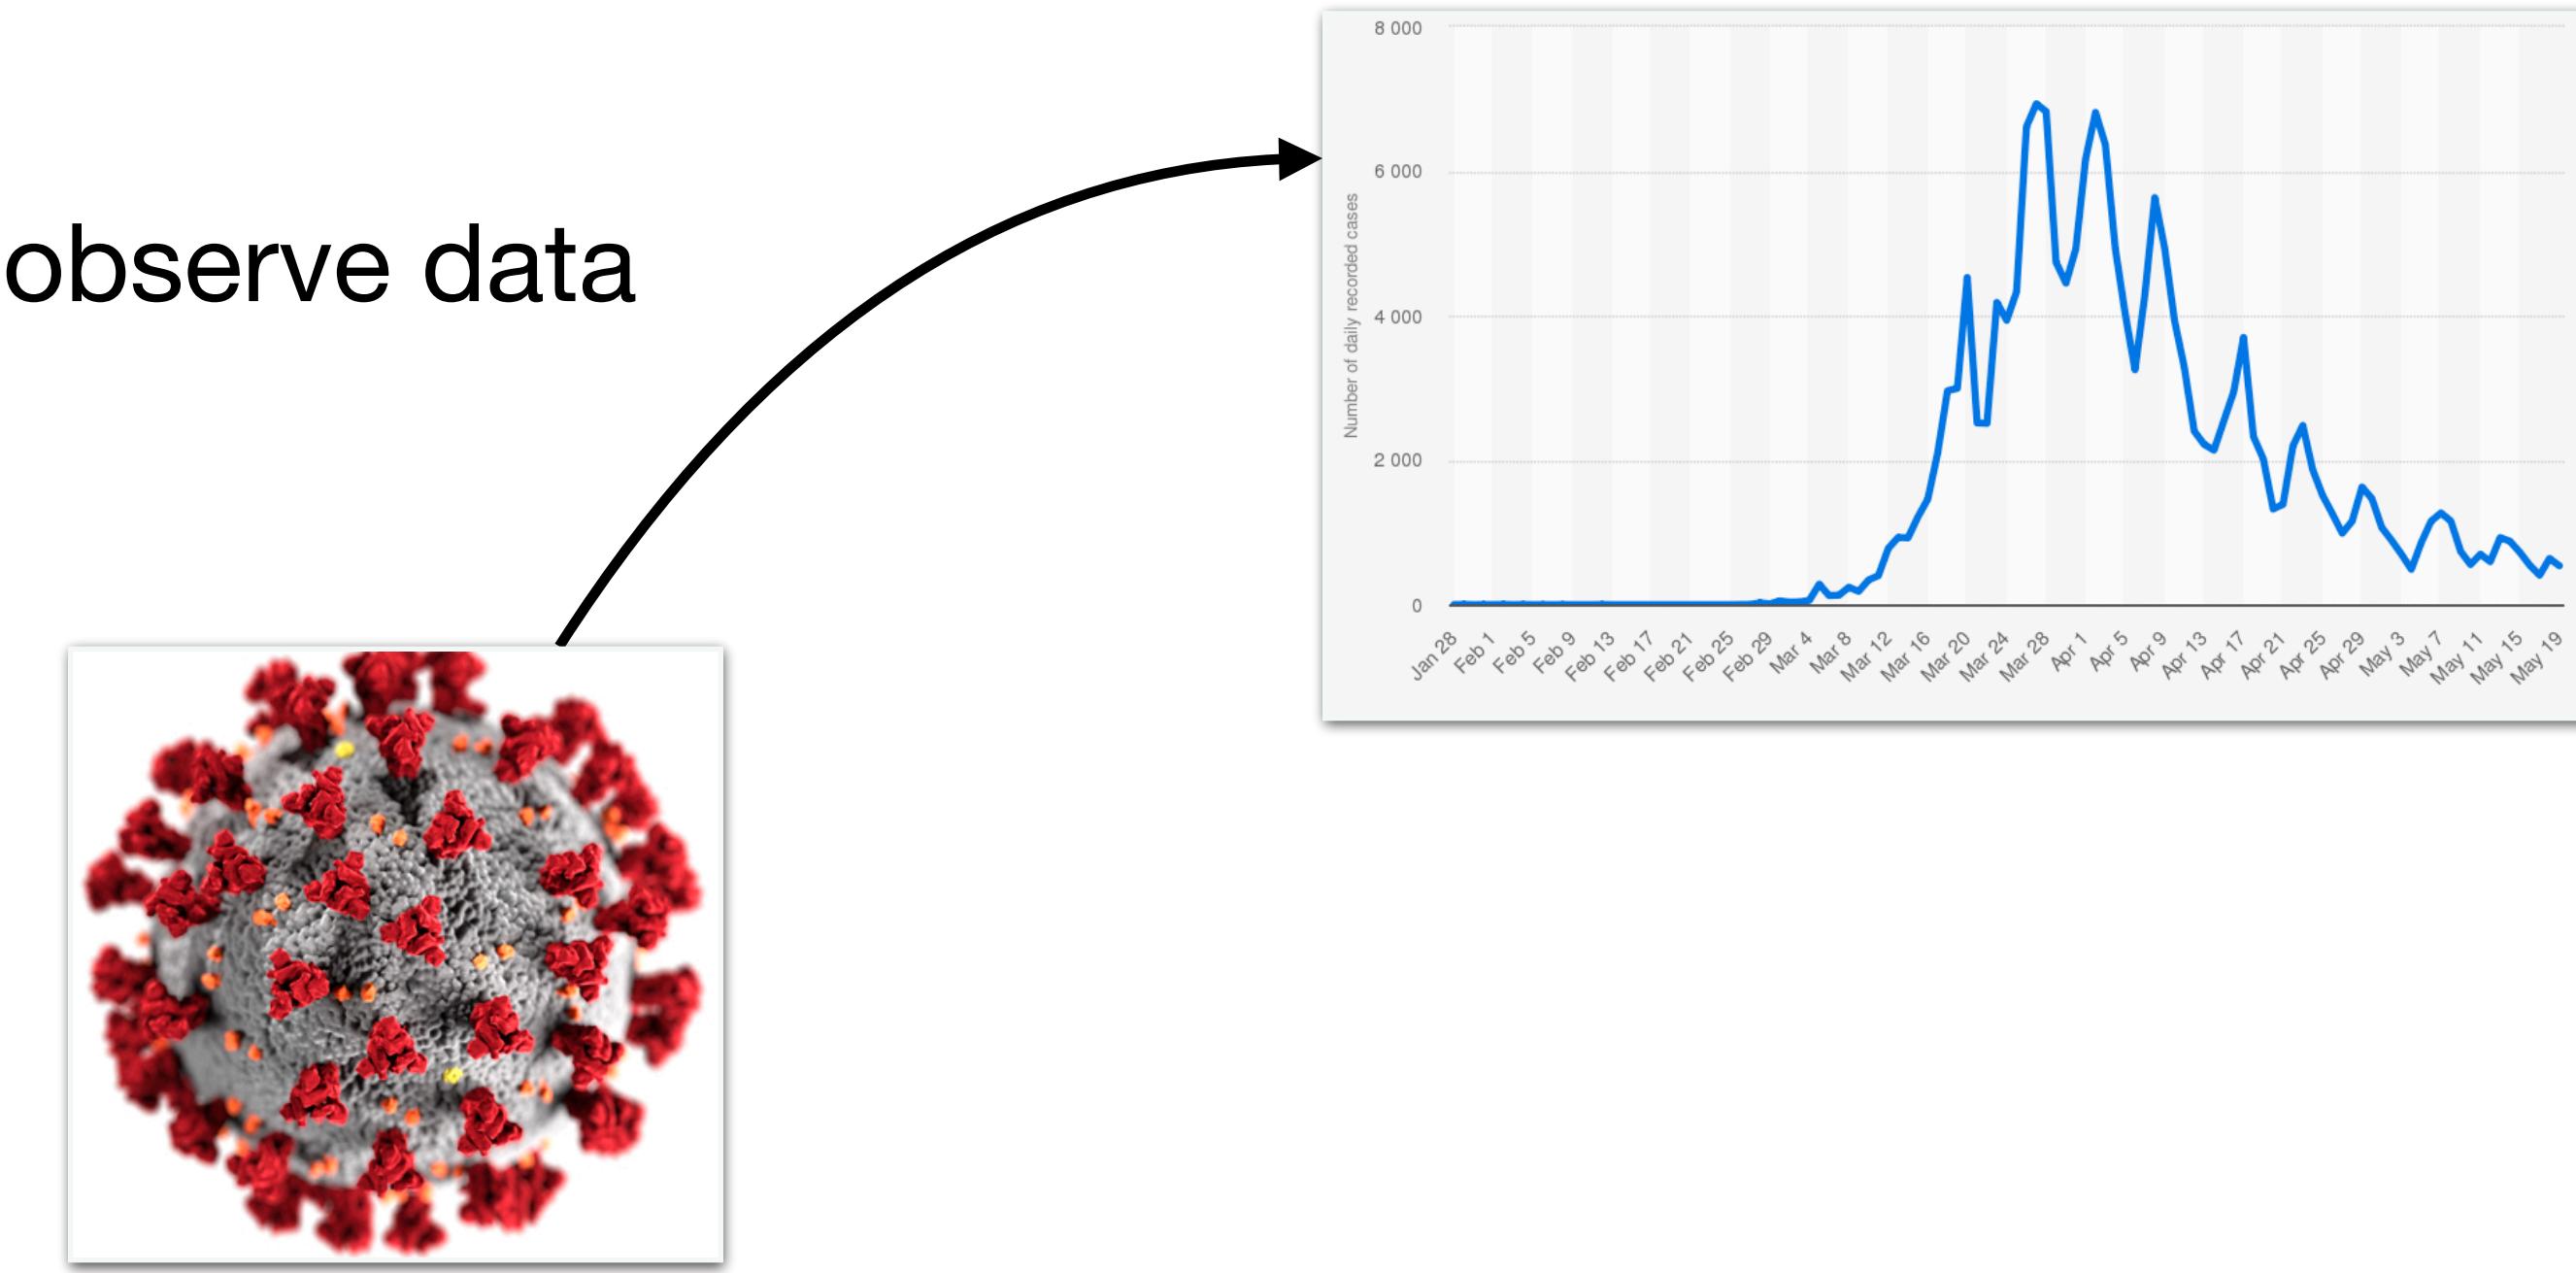
Example: The SIR Simulator for COVID-19



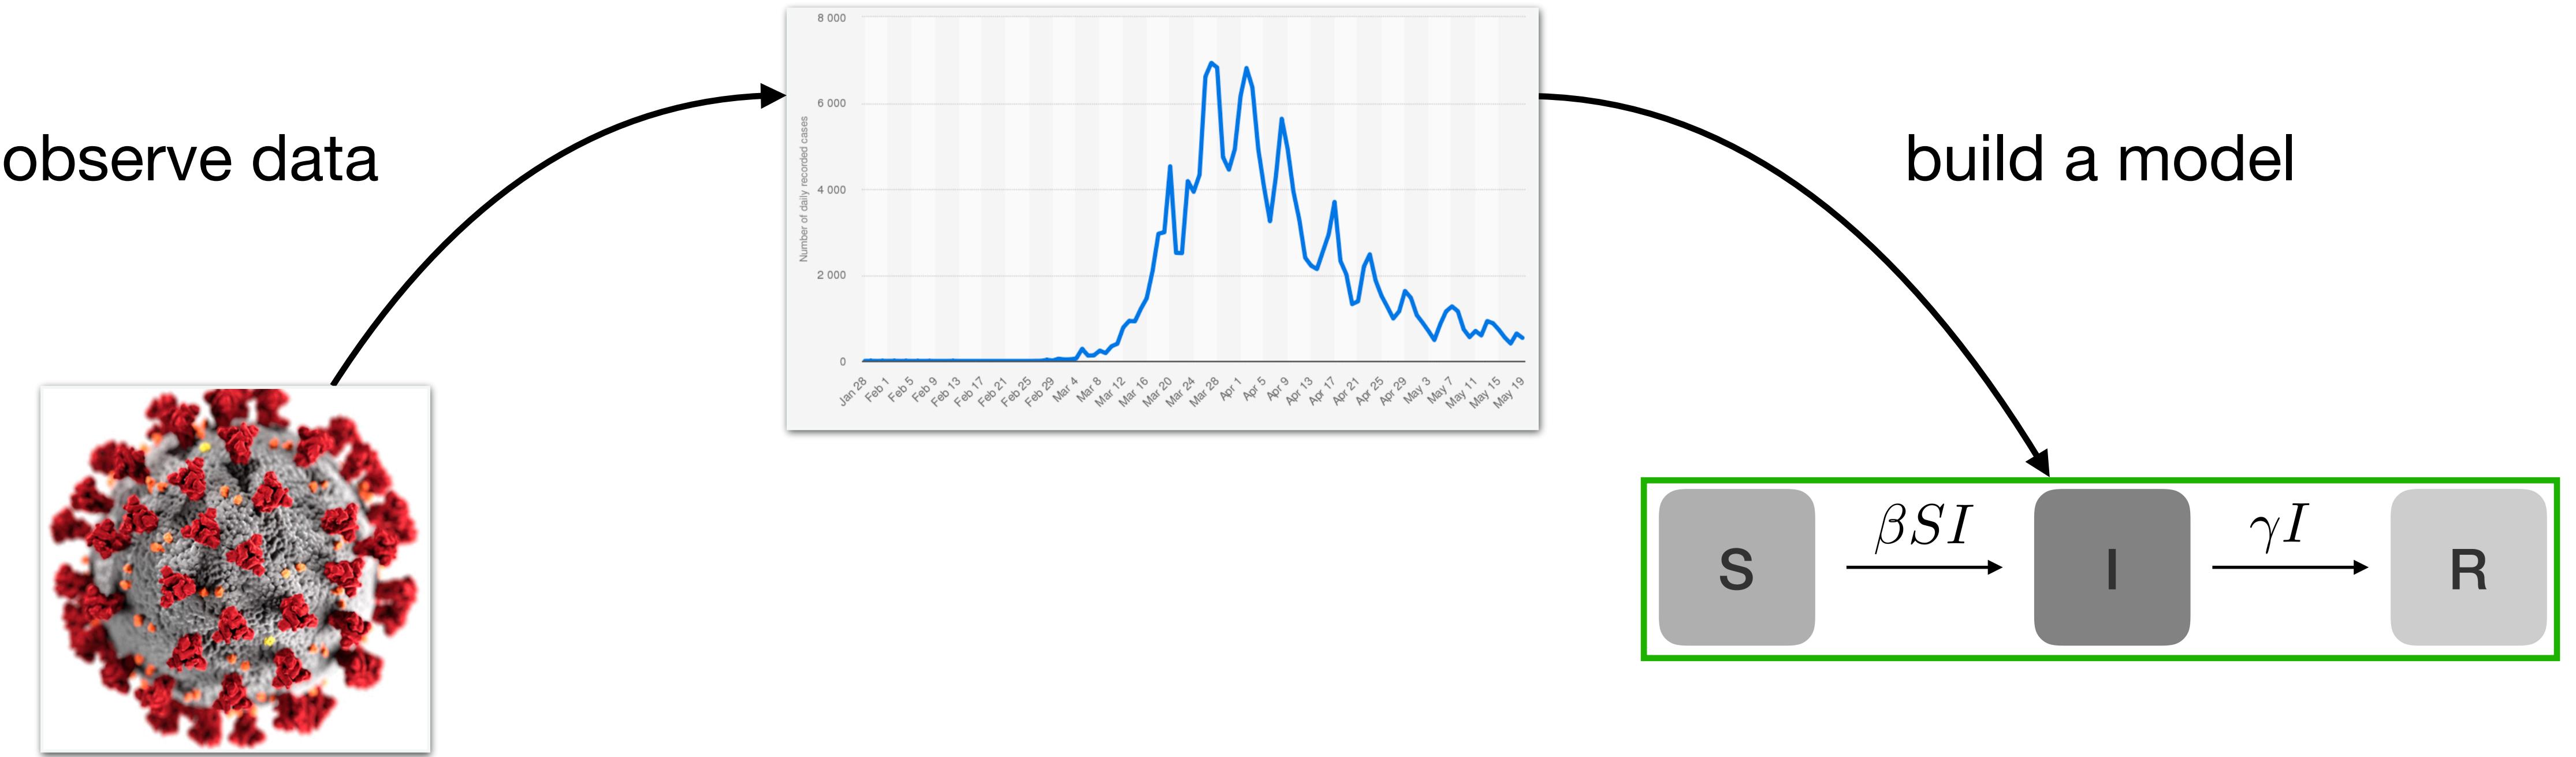
The simulation-based modeling approach



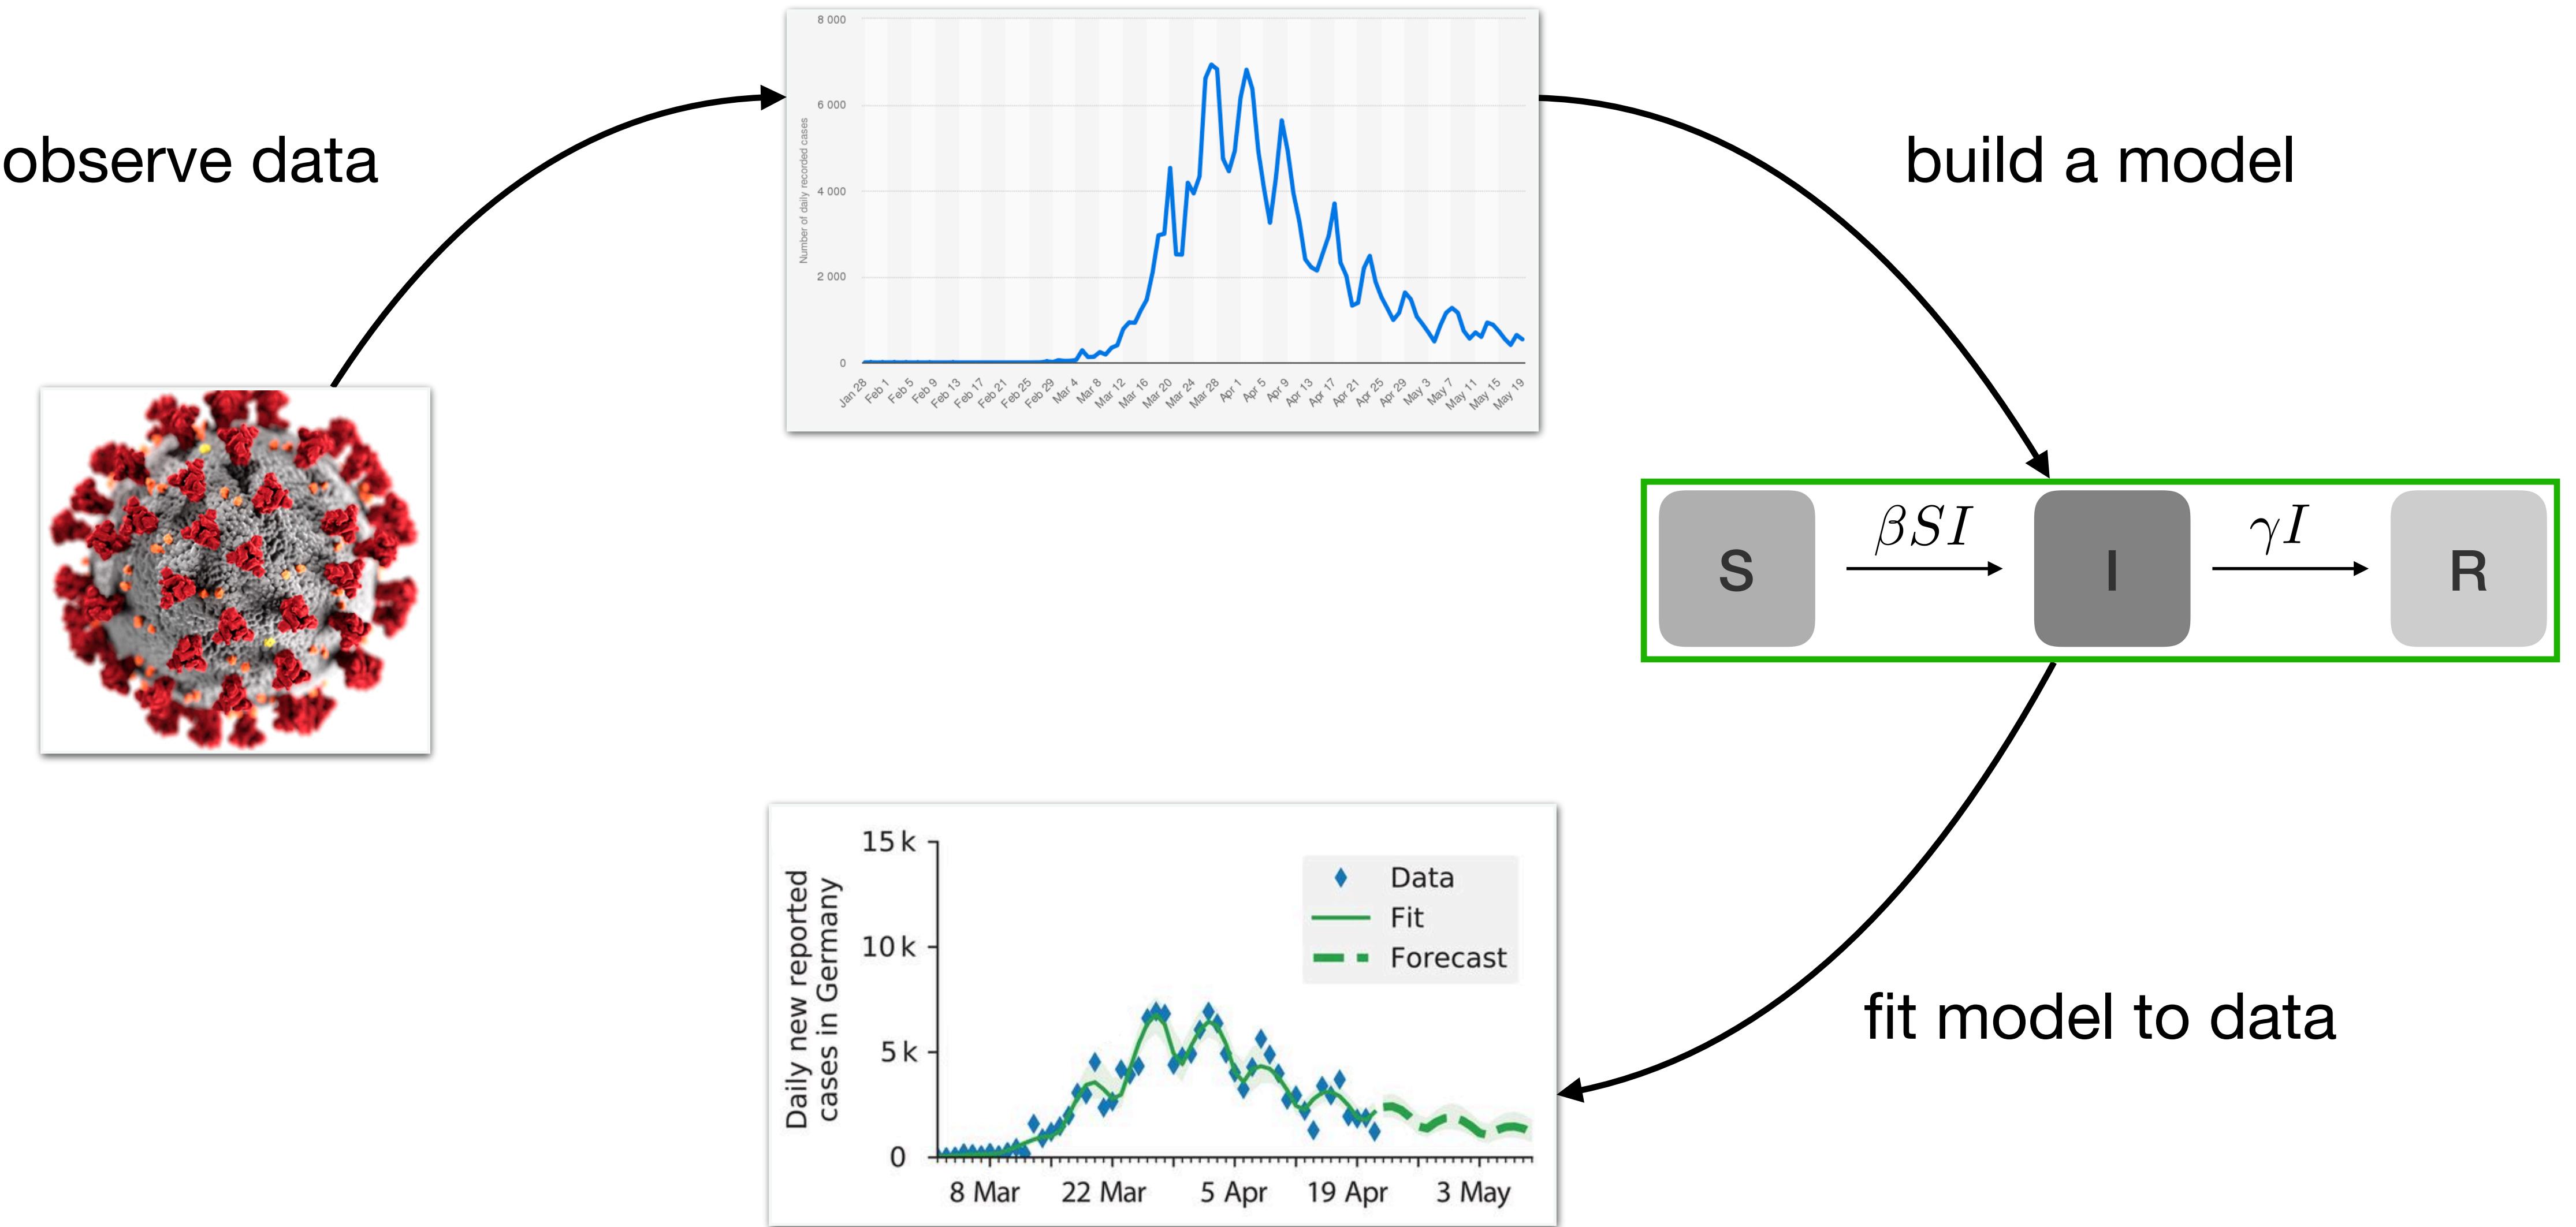
The simulation-based modeling approach



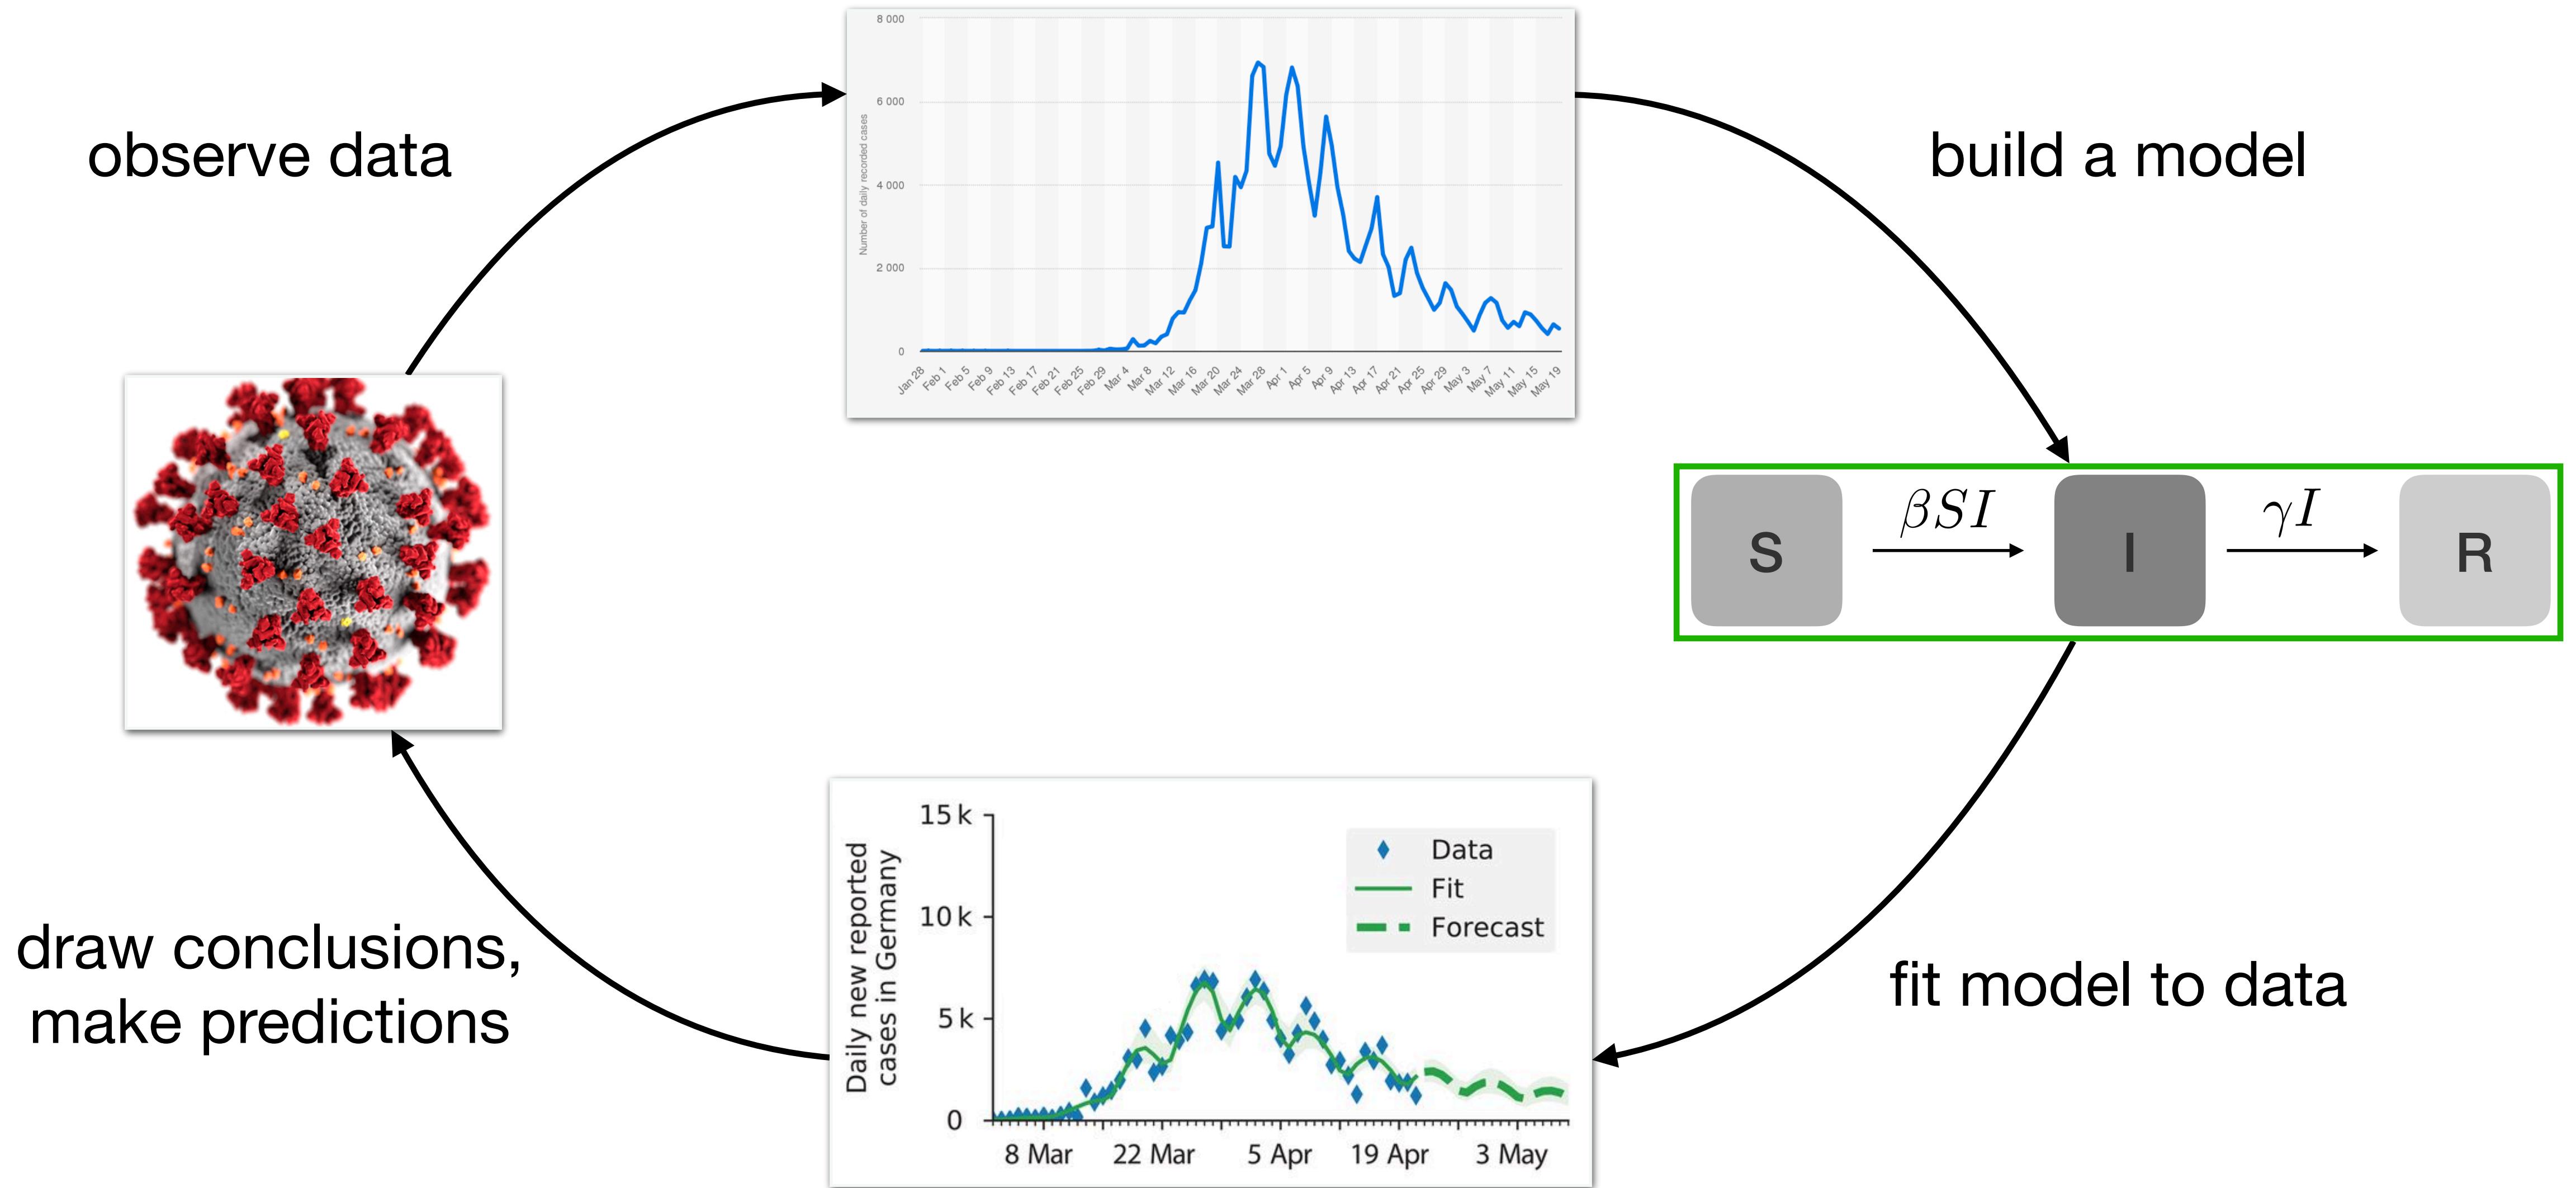
The simulation-based modeling approach



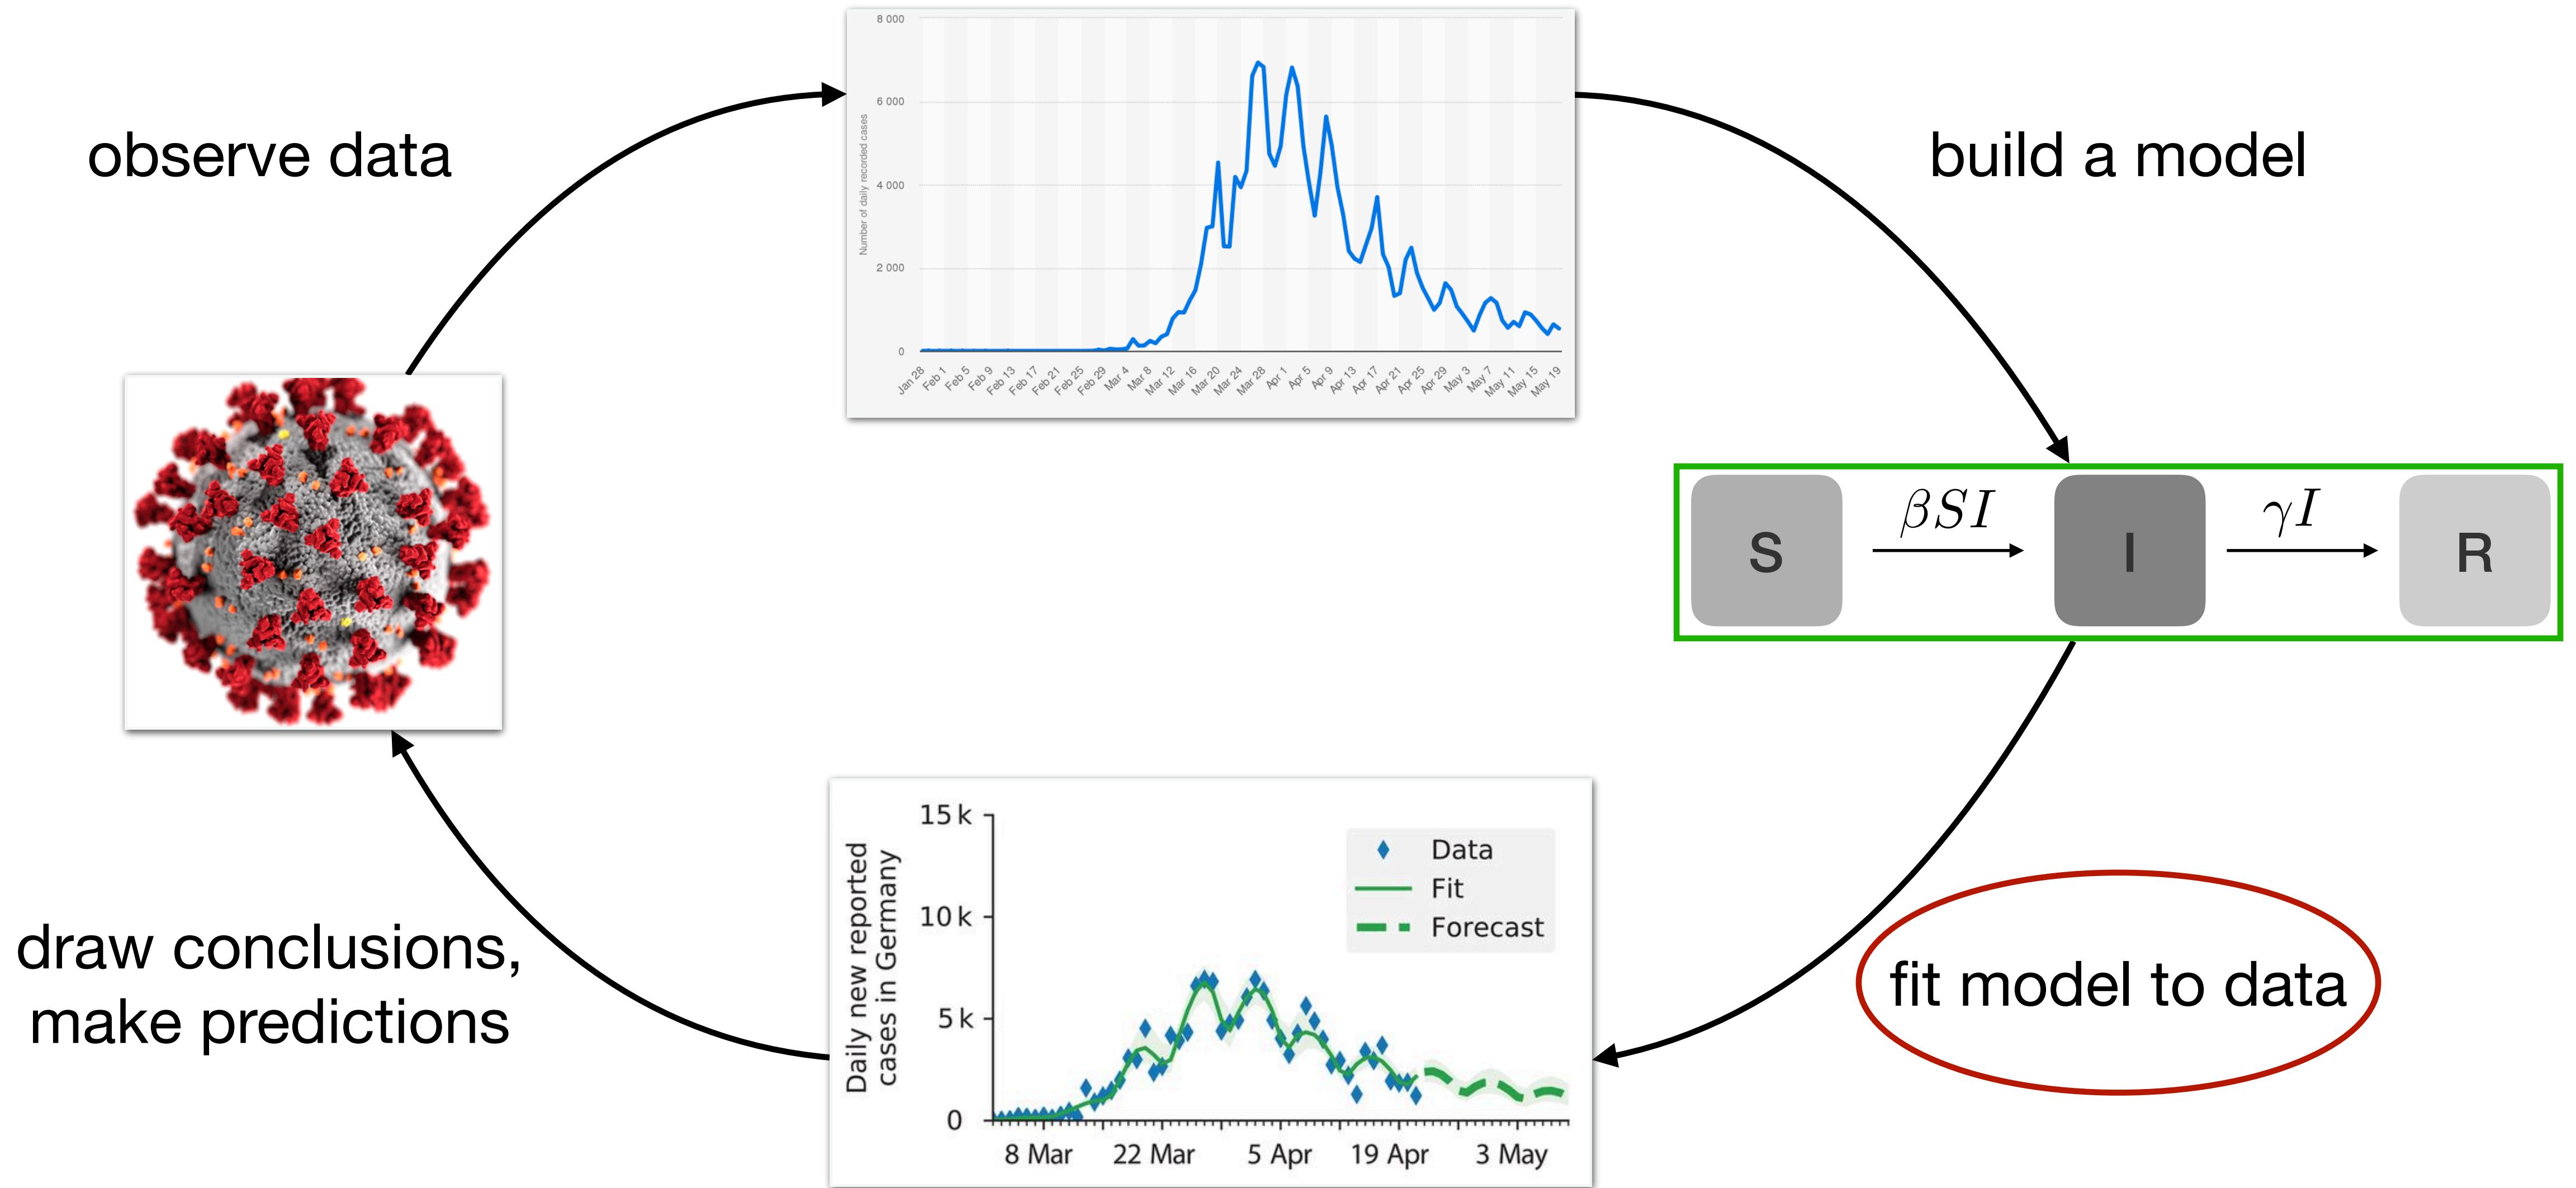
The simulation-based modeling approach



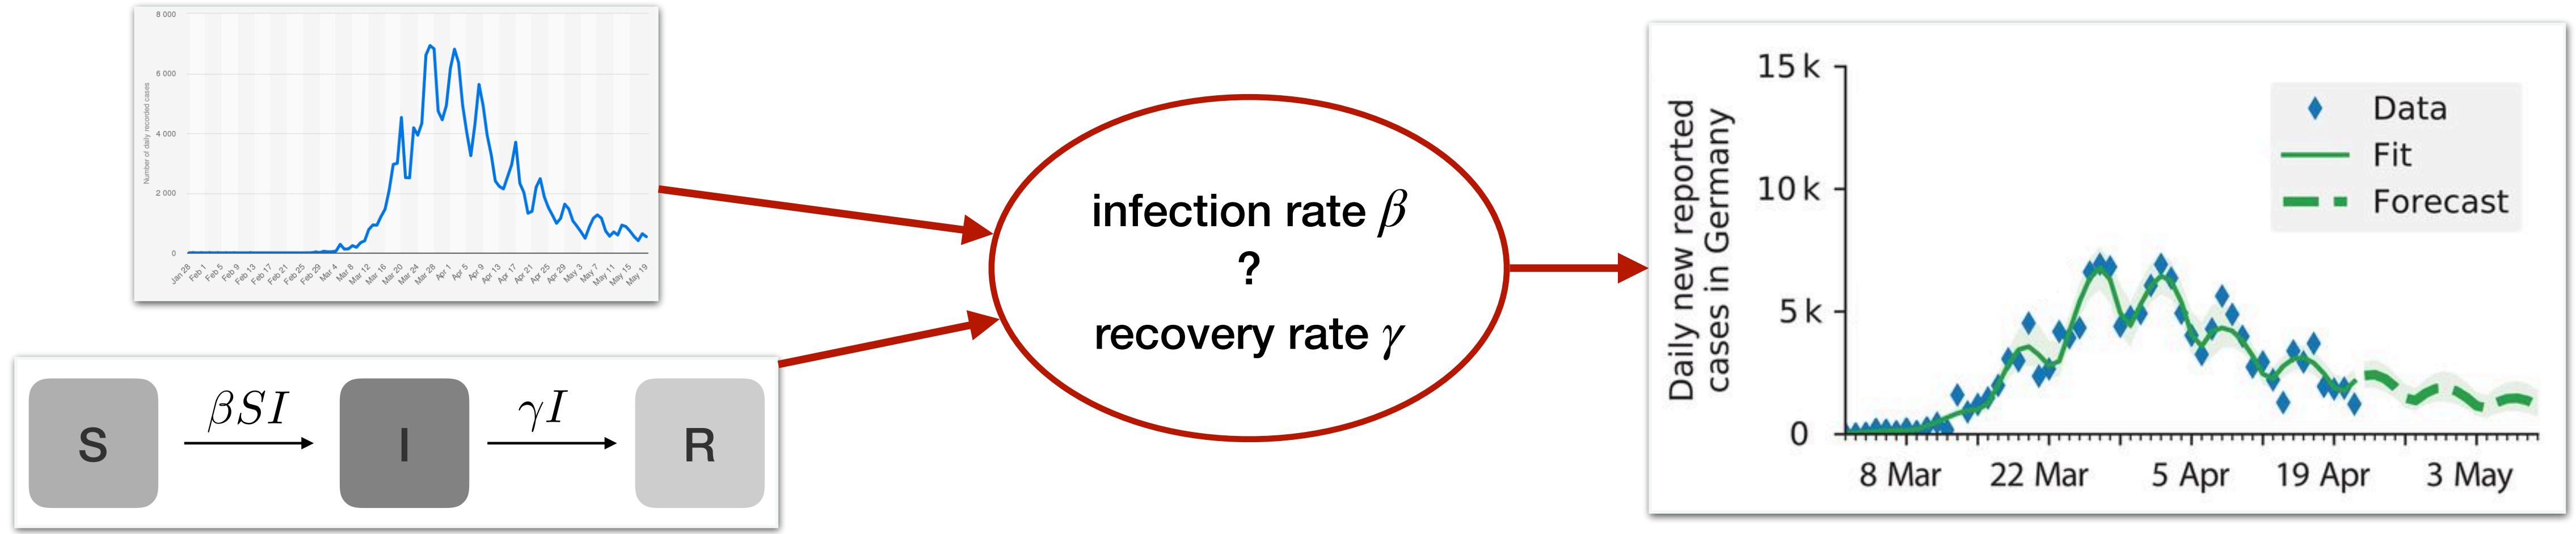
The simulation-based modeling approach



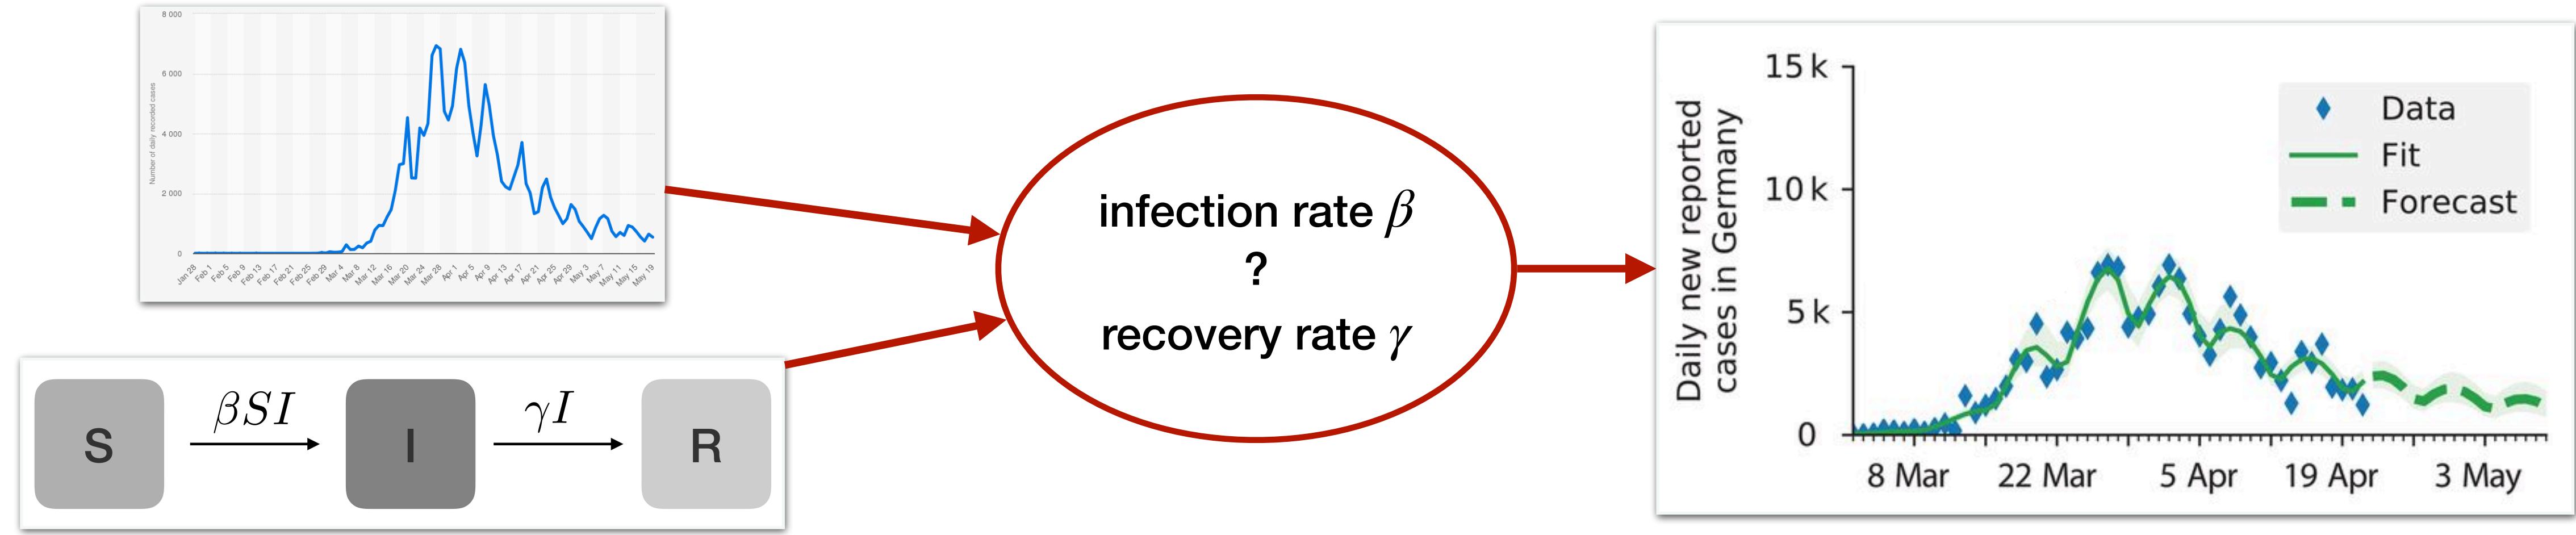
The simulation-based modeling approach



Central challenge: finding model parameters

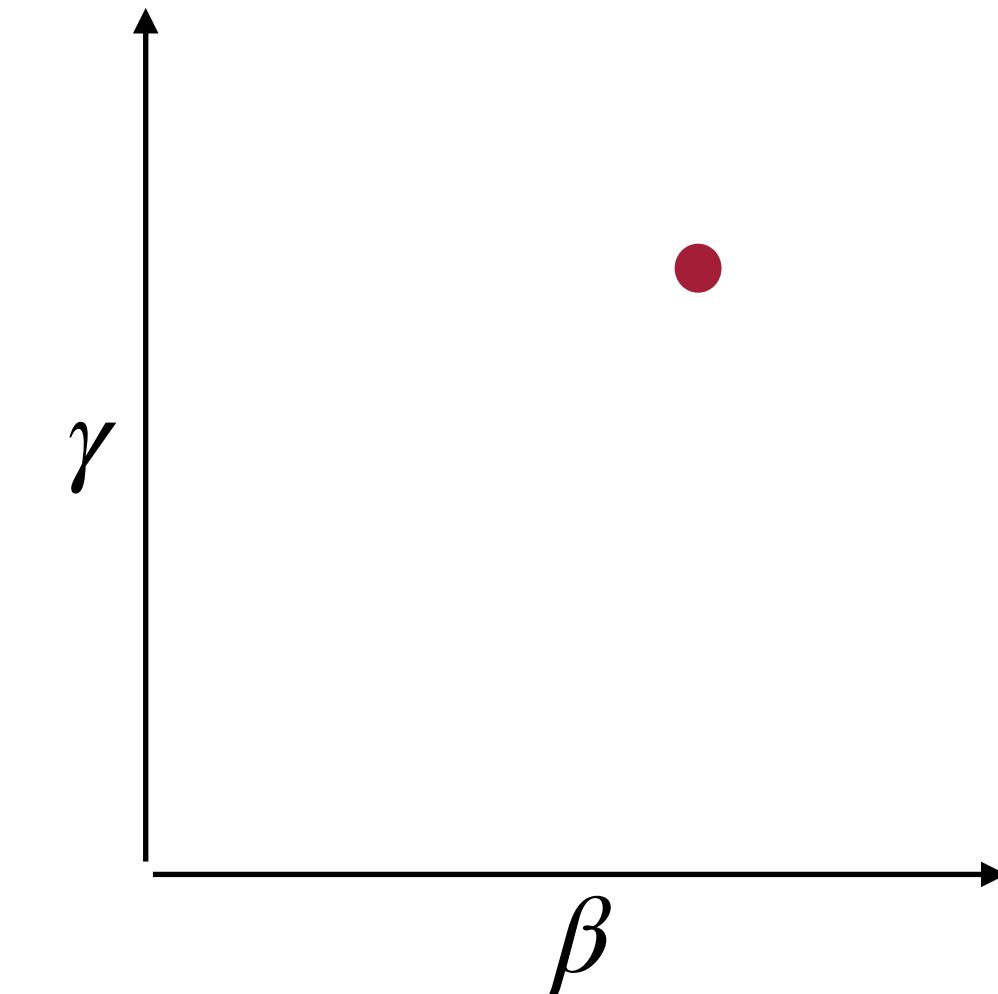


Central challenge: finding model parameters

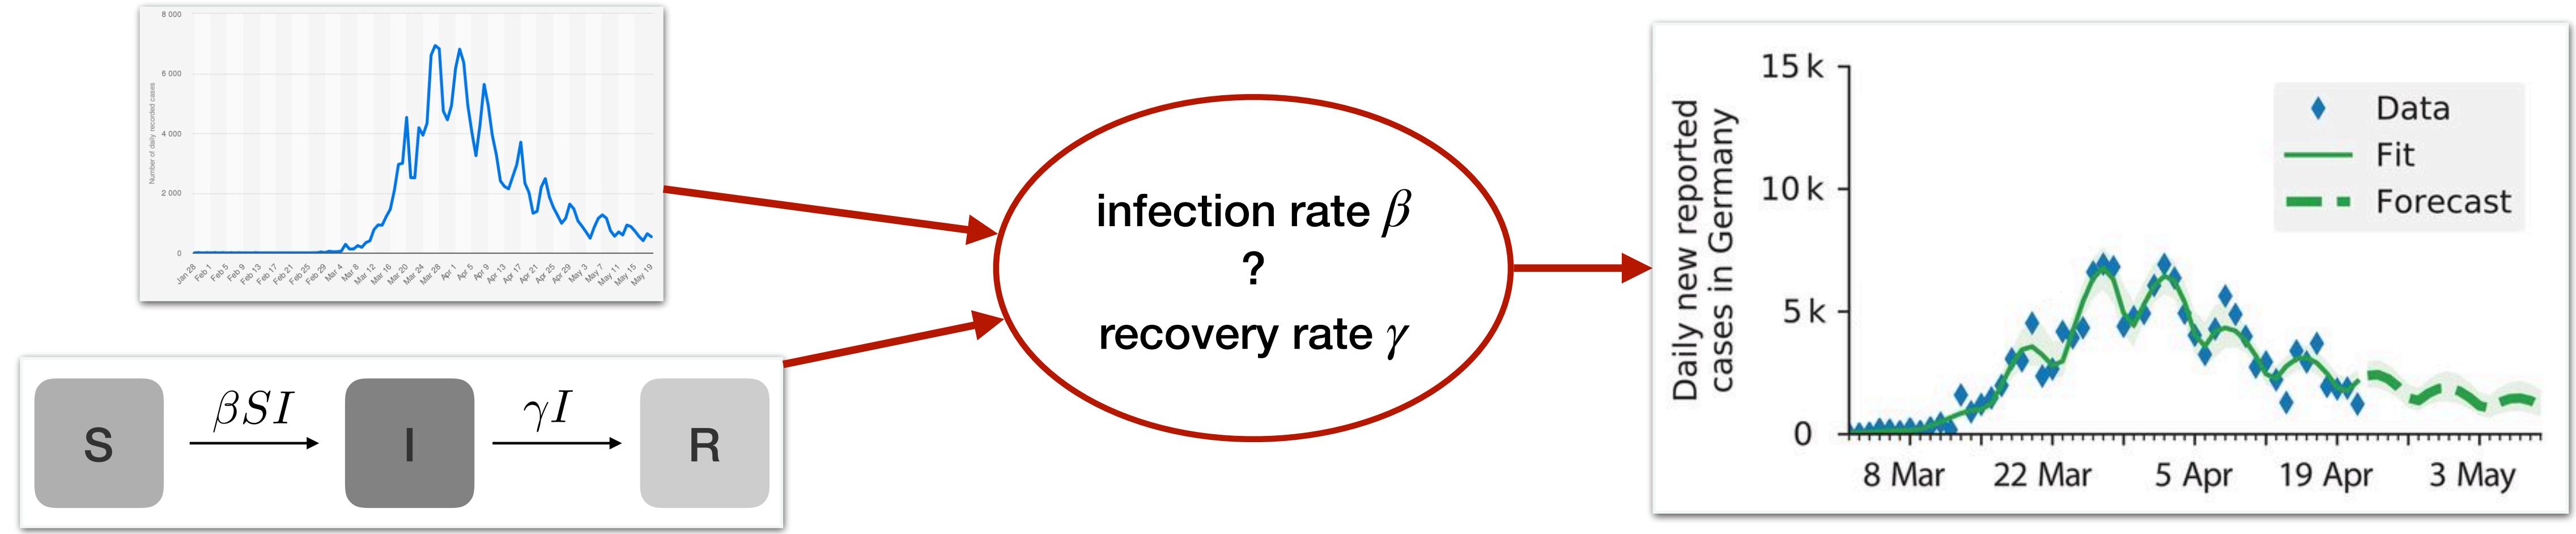


Naive approach

- find single best-fitting parameters
 - hand tuning, grid search
 - optimization, genetic algorithms

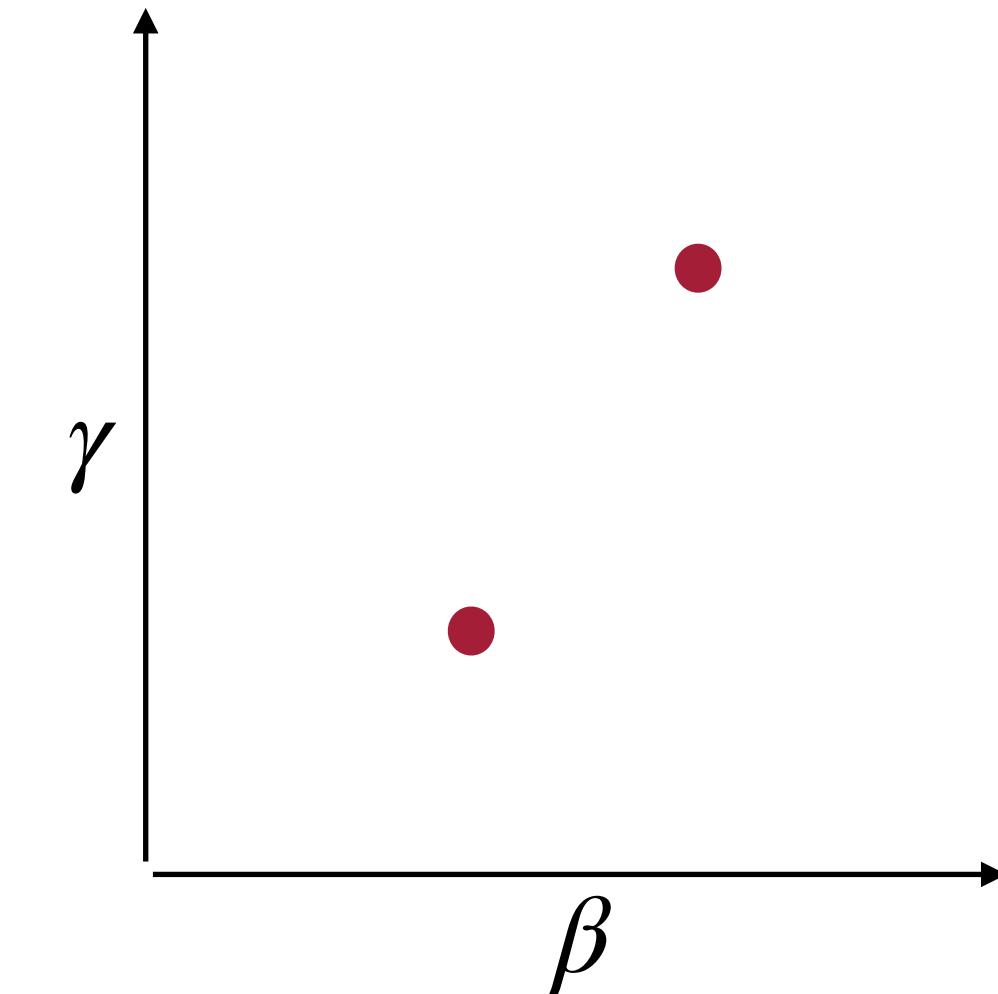


Central challenge: finding model parameters

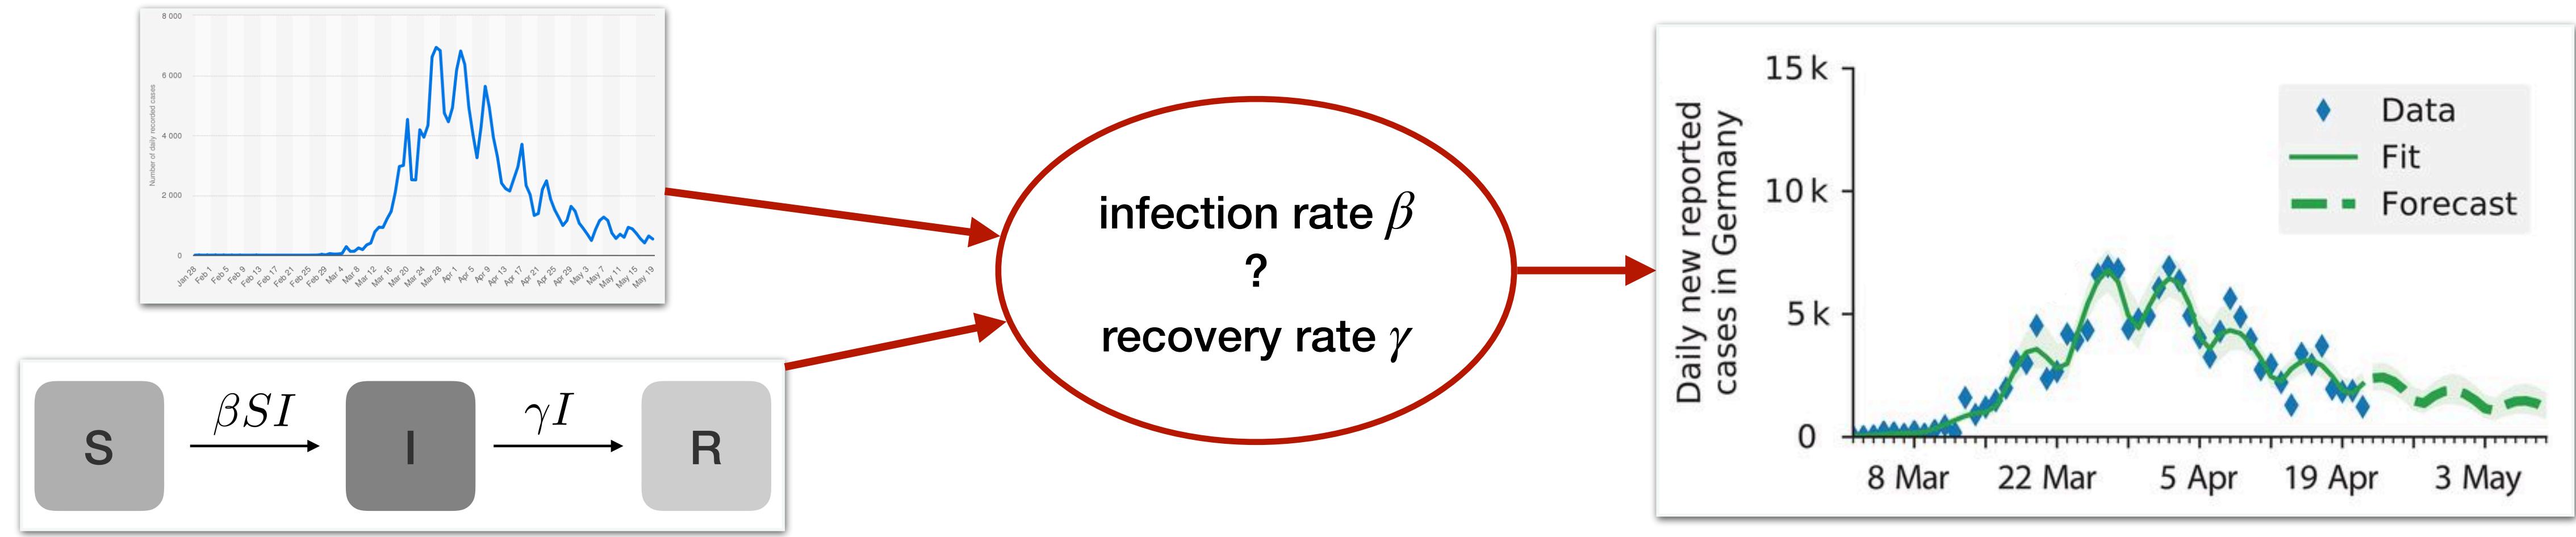


Naive approach

- find single best-fitting parameters
 - hand tuning, grid search
 - optimization, genetic algorithms

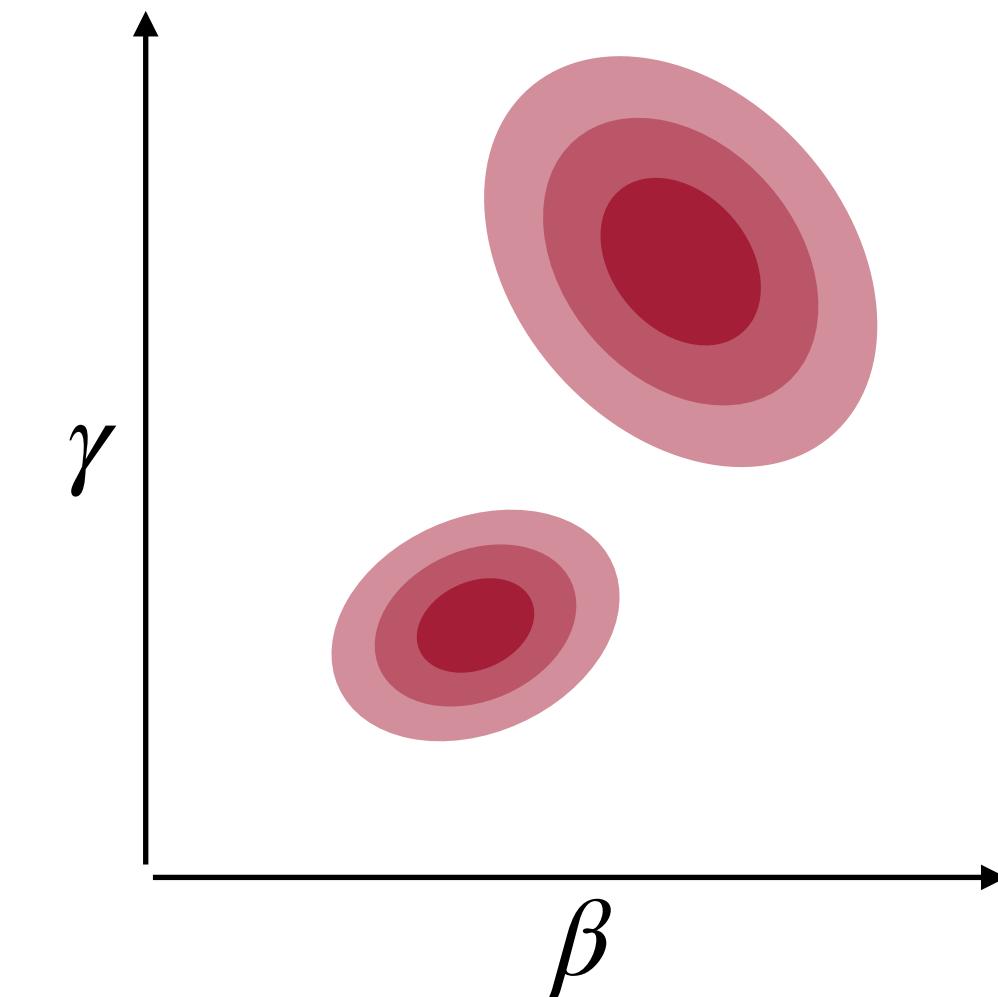


Central challenge: finding model parameters

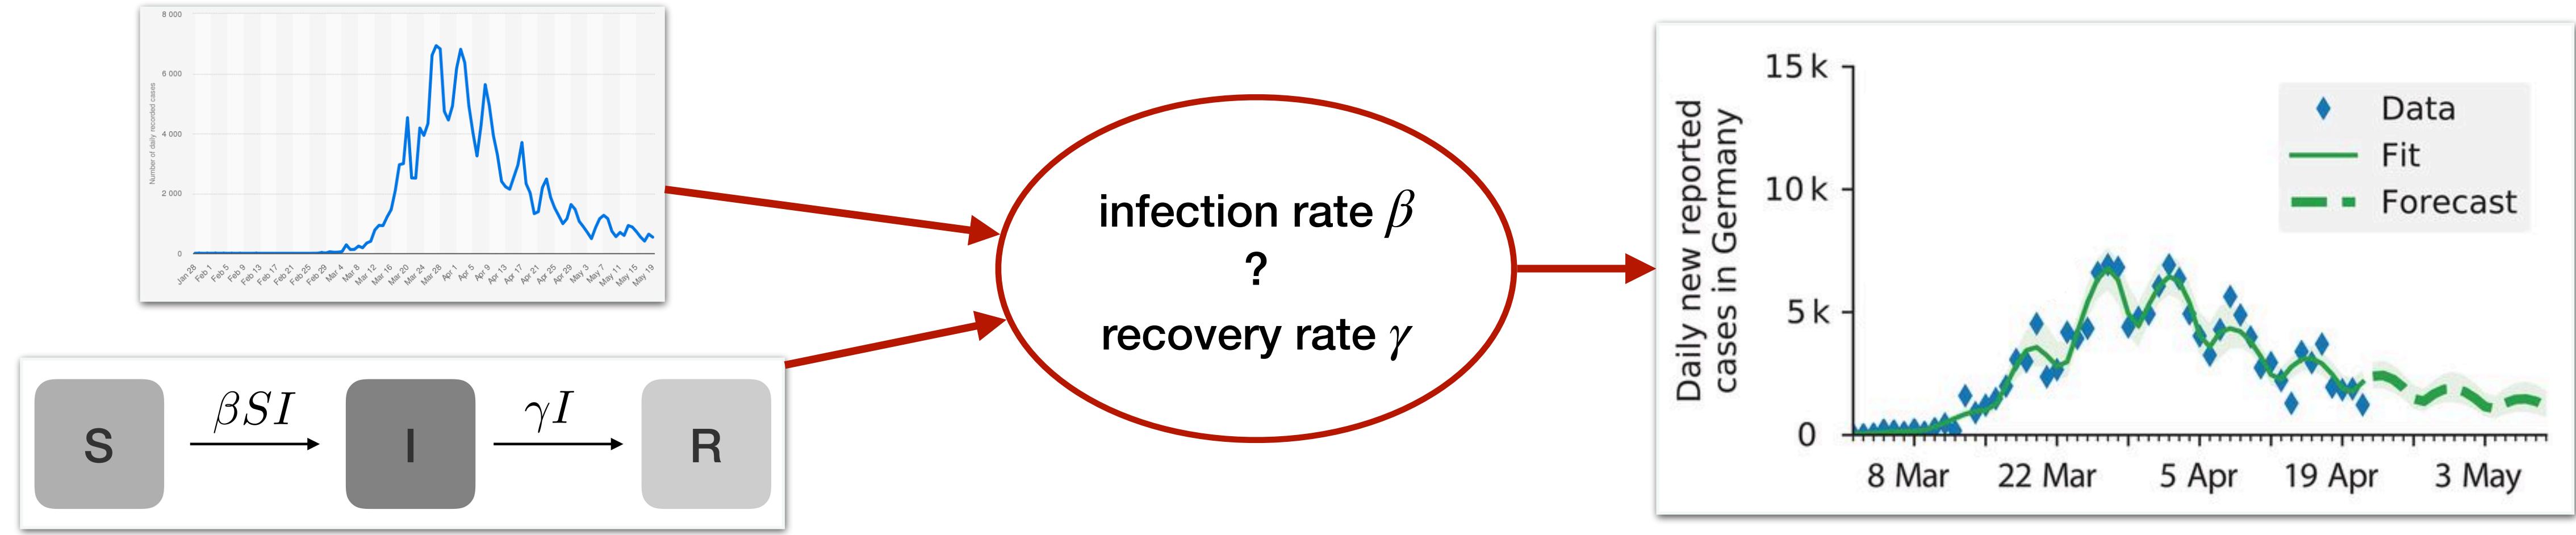


Naive approach

- find single best-fitting parameters
 - hand tuning, grid search
 - optimization, genetic algorithms

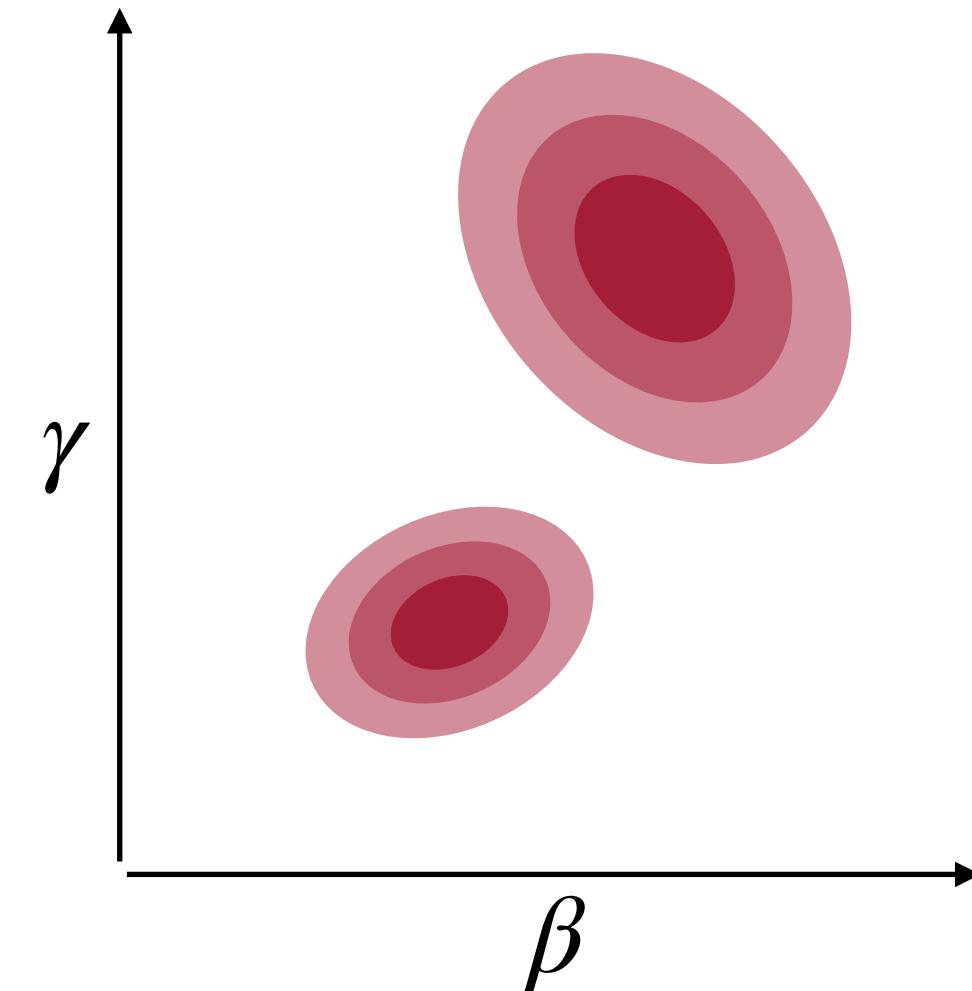


Central challenge: finding model parameters



Naive approach

- find single best-fitting parameters
 - hand tuning, grid search
 - optimization, genetic algorithms



Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters

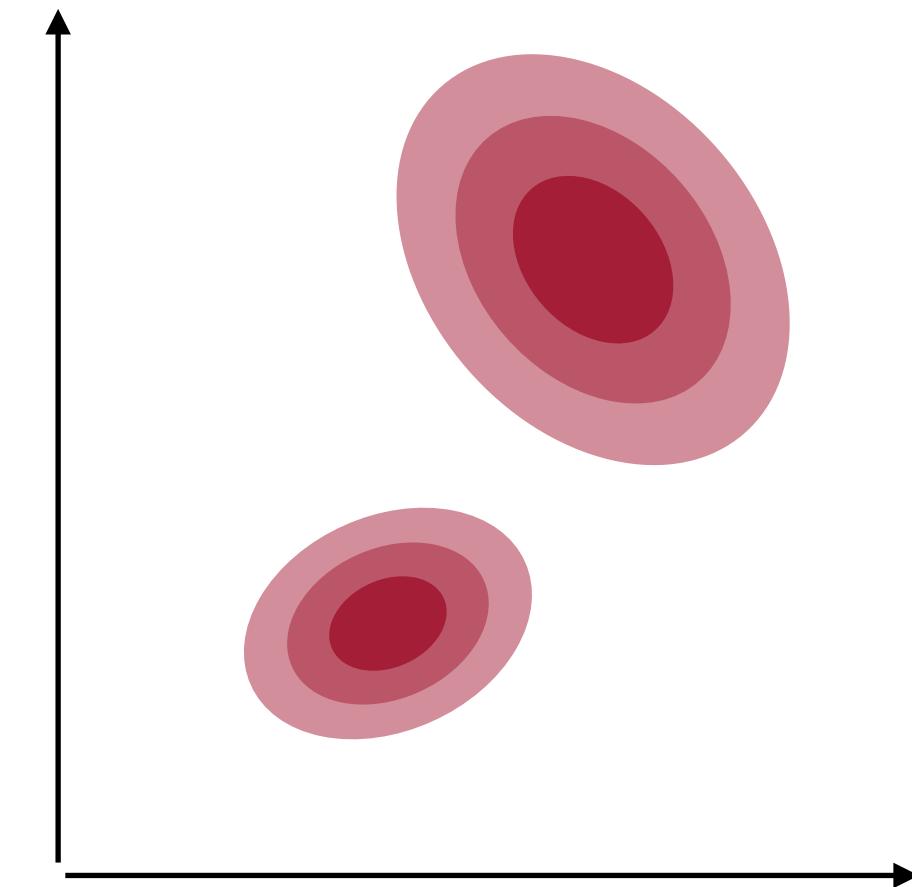
Outline

1. **Why** simulators?
2. **Why** Bayesian parameter inference?
3. **How** to do Bayesian inference for simulators?
4. **How** can you apply SBI to your simulator?

Bayesian parameter inference

Ideal approach

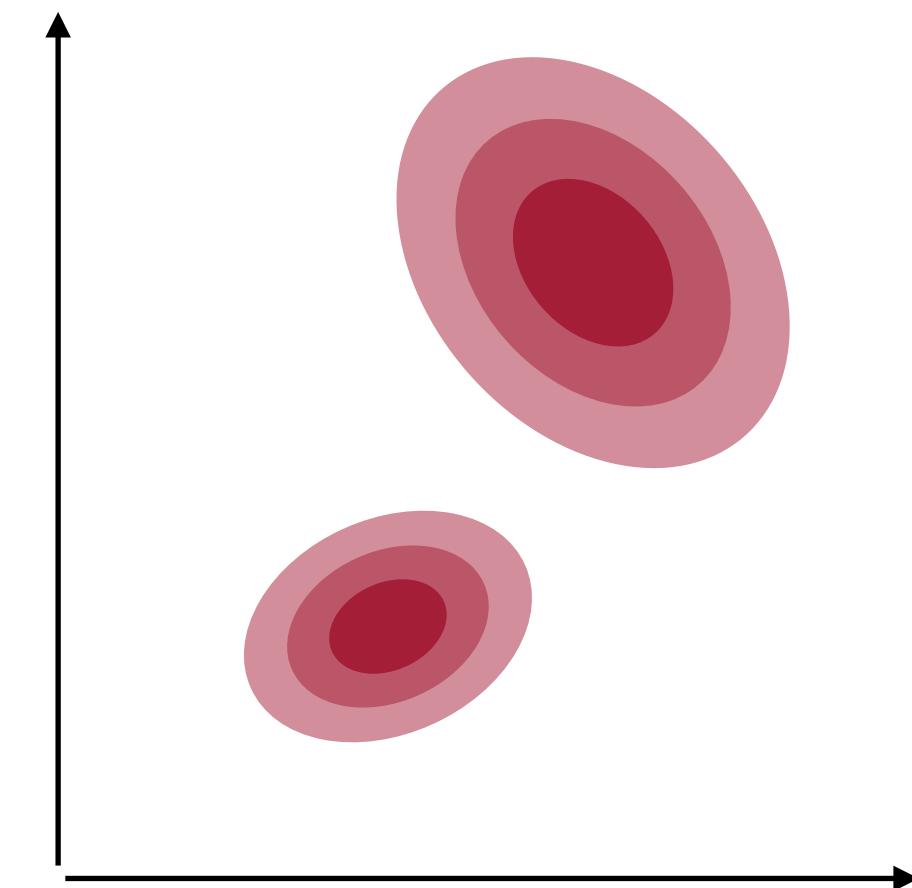
- find all parameters
- quantify uncertainty
- quantify relations between parameters



Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters

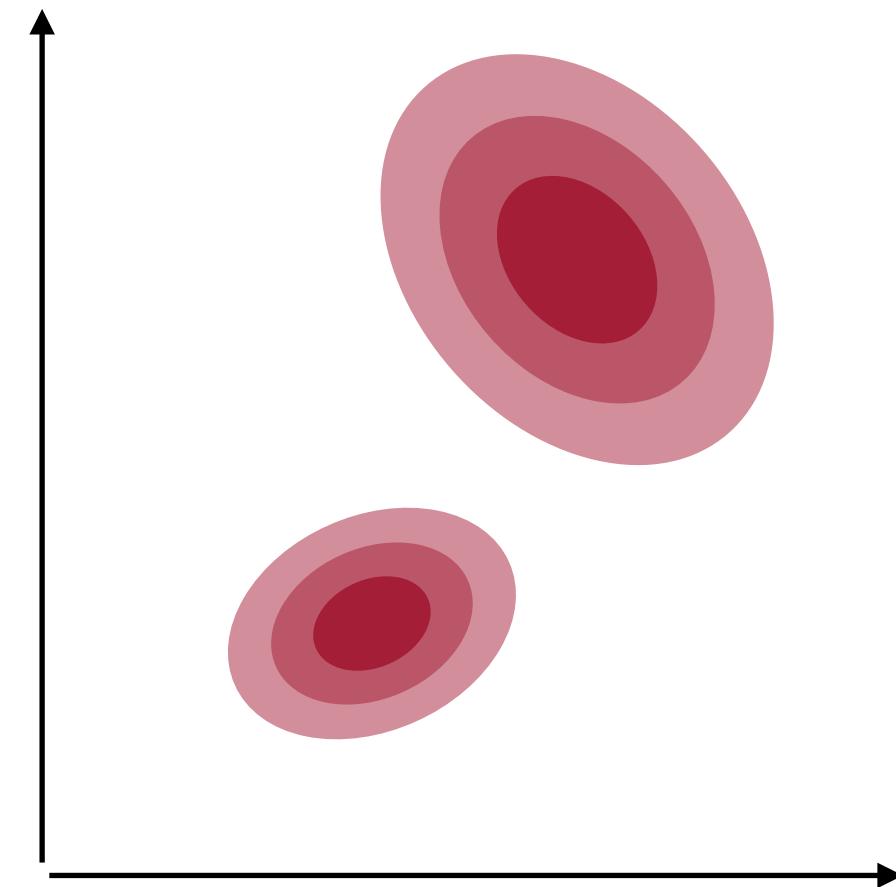


- data x and parameters θ are *random variables*

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters

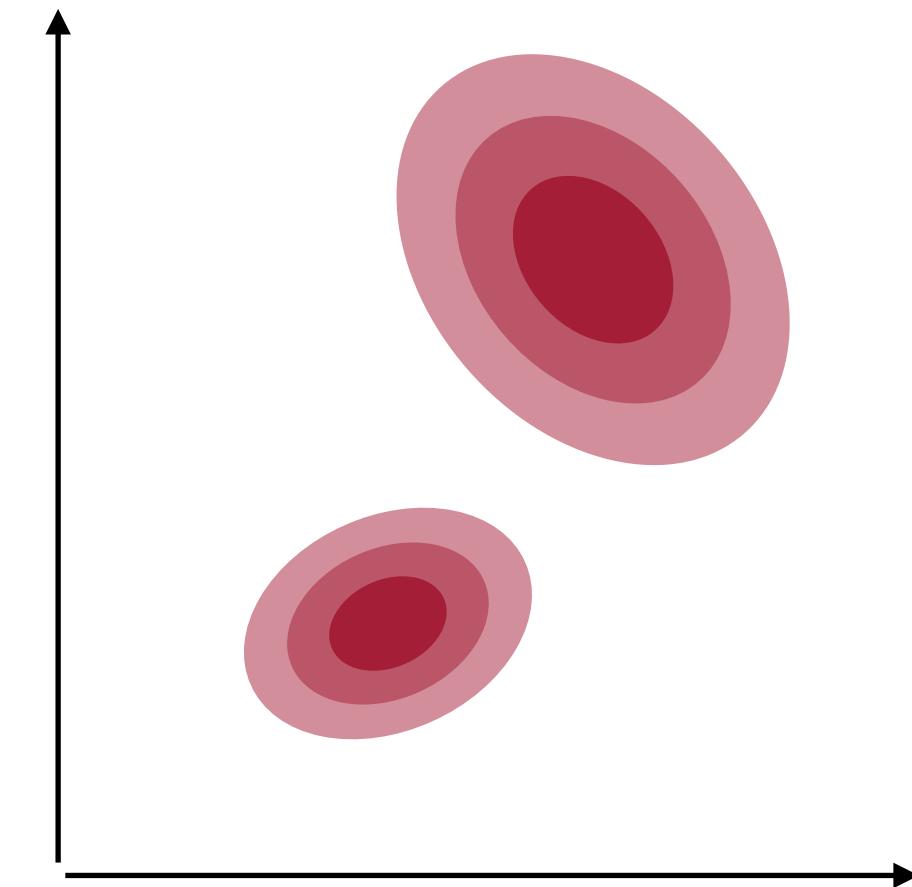


- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters

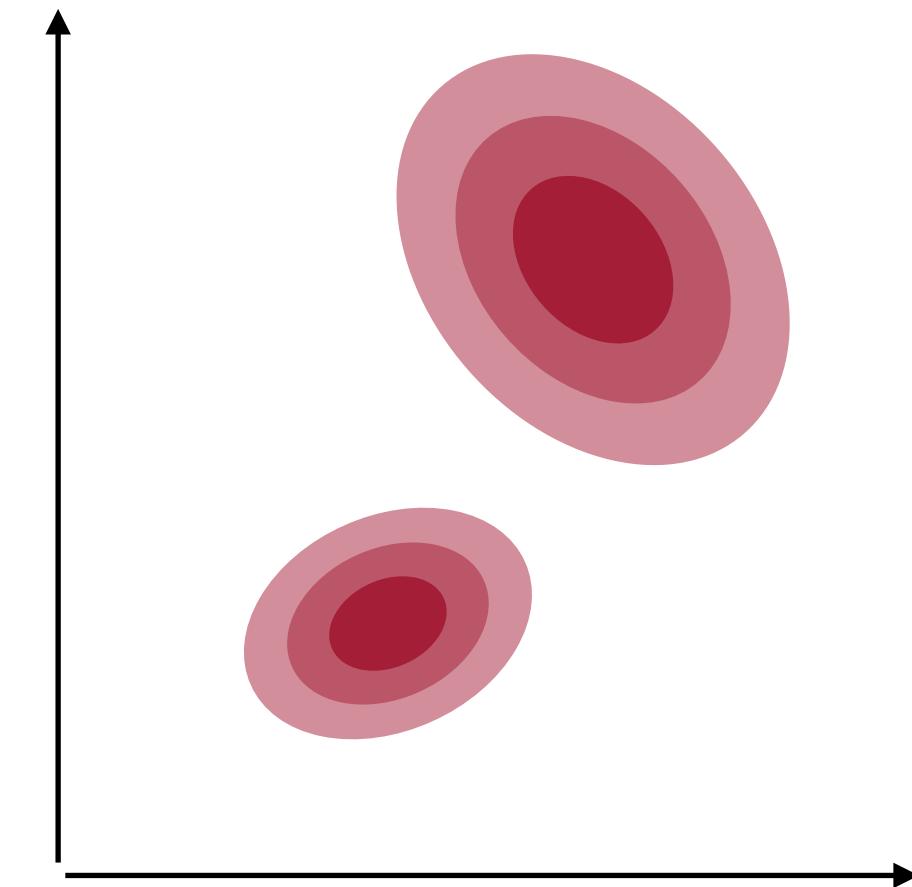


- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters

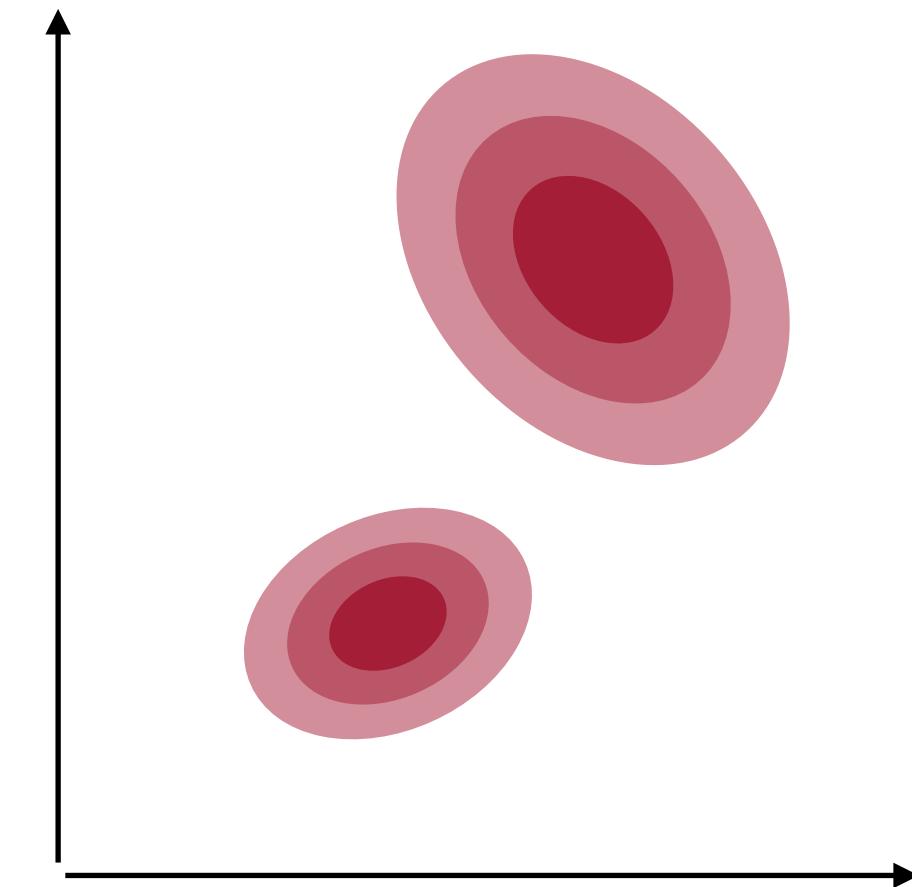


- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters
 - quantifies uncertainty

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters



- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters
 - quantifies uncertainty

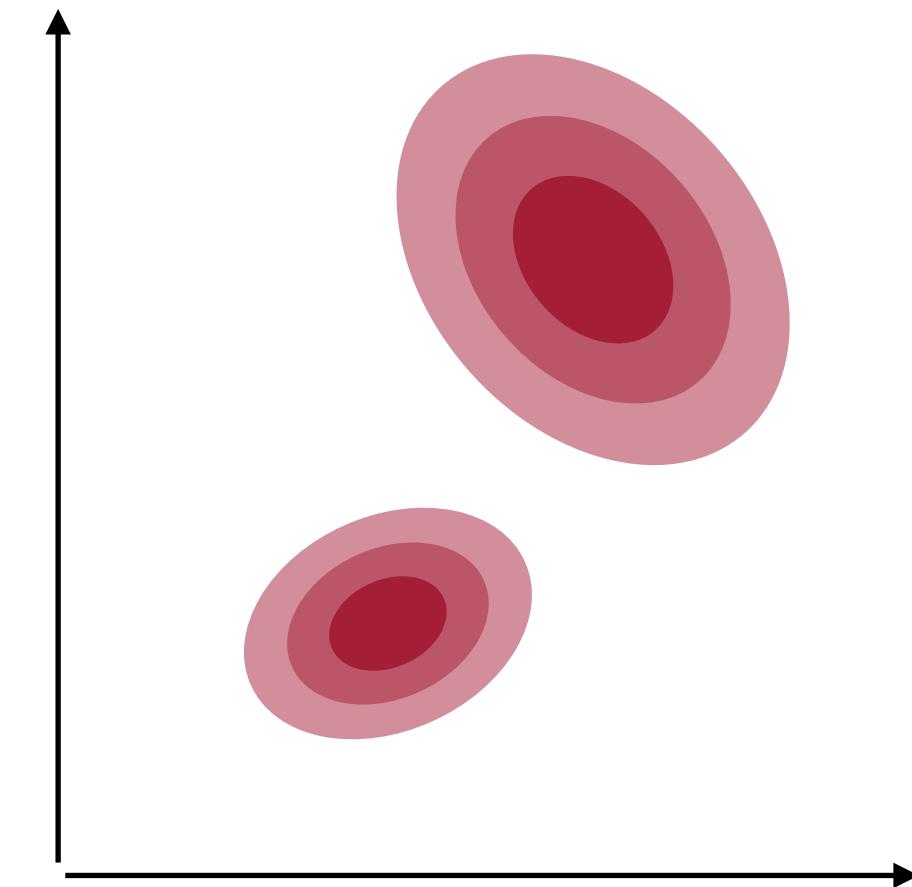
Bayes' Theorem

$$p(\theta | x) \propto p(x | \theta) p(\theta)$$

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters



- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters
 - quantifies uncertainty

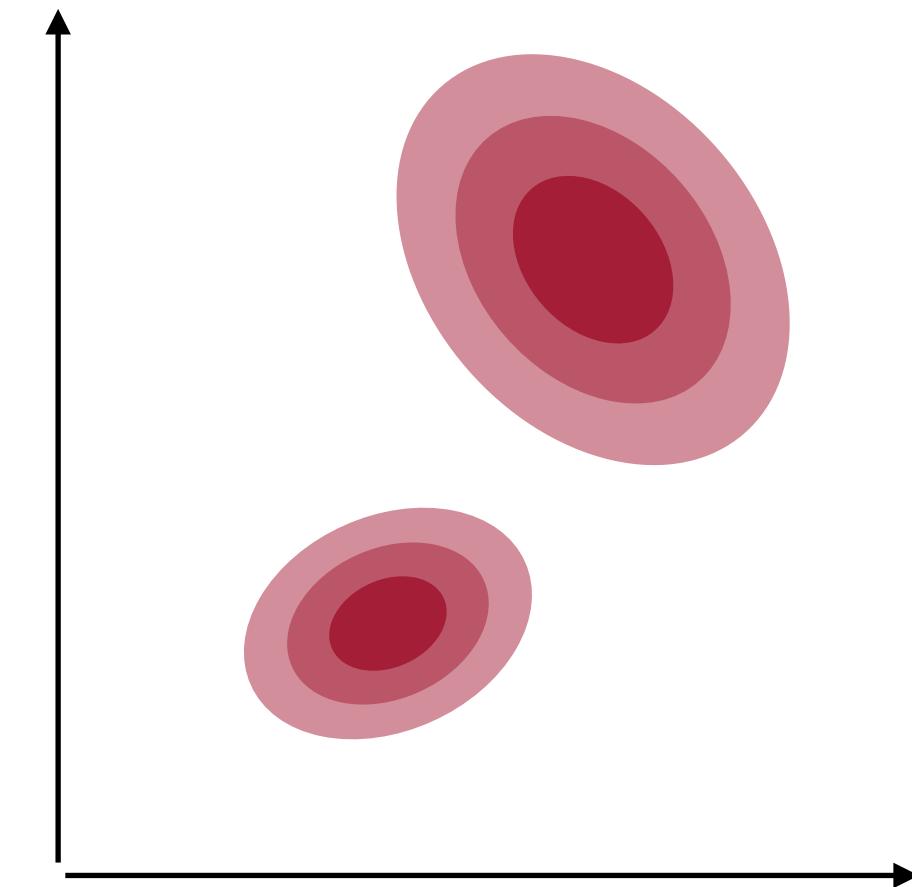
Bayes' Theorem

“posterior” → $p(\theta | x) \propto p(x | \theta) p(\theta)$

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters



- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters
 - quantifies uncertainty

Bayes' Theorem

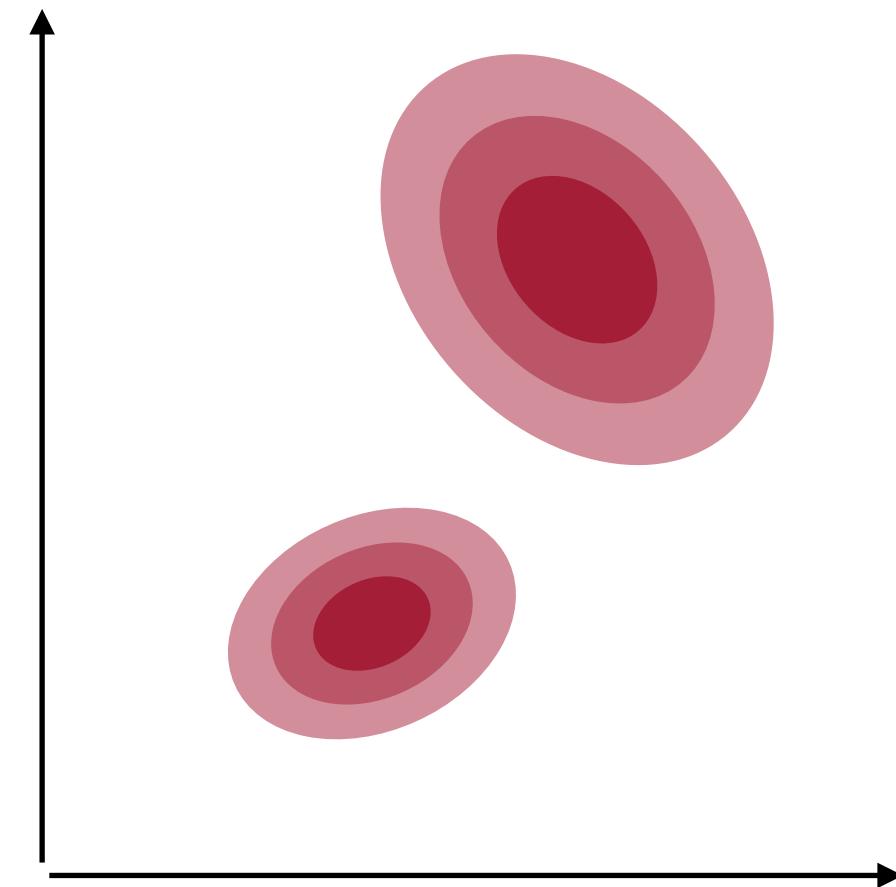
$$\text{“posterior”} \rightarrow p(\theta | x) \propto p(x | \theta) p(\theta)$$

“prior”
↓

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters



- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters
 - quantifies uncertainty

Bayes' Theorem

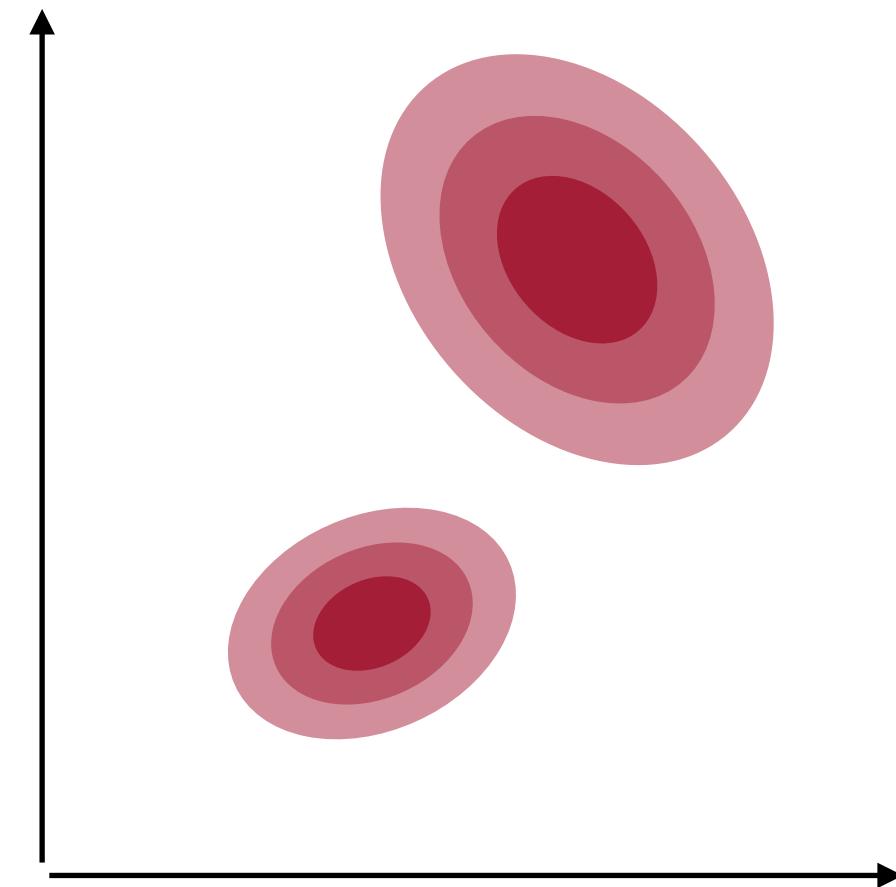
$$\text{“posterior”} \rightarrow p(\theta | x) \propto p(x | \theta) p(\theta)$$

“prior”
↑
“likelihood”

Bayesian parameter inference

Ideal approach

- find all parameters
- quantify uncertainty
- quantify relations between parameters



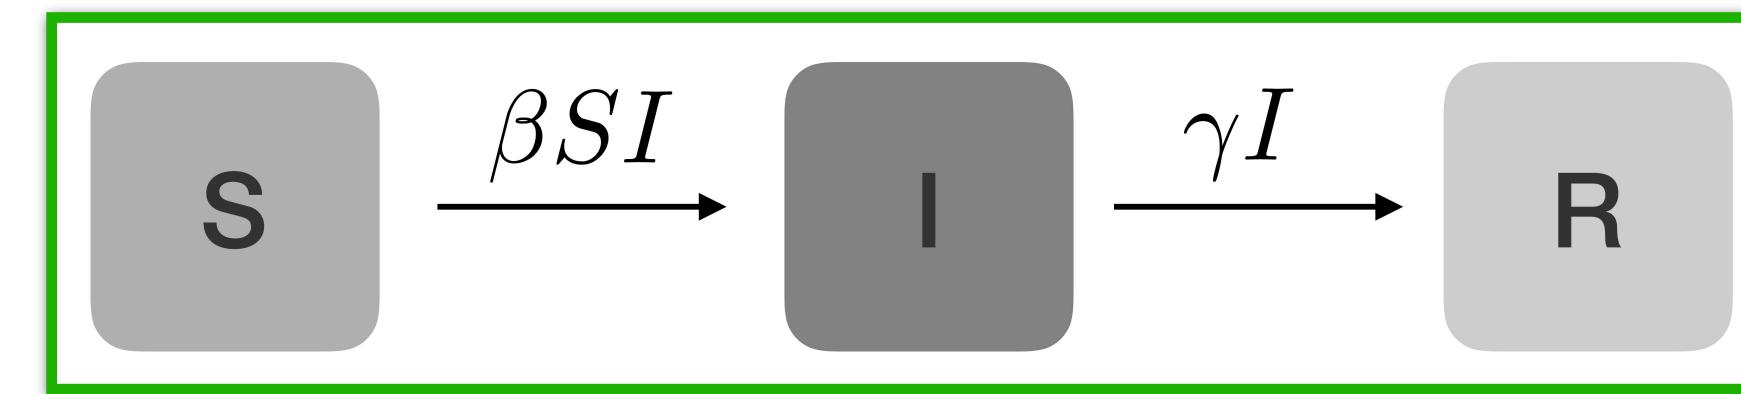
- data x and parameters θ are *random variables*
- **goal:** infer a distribution over parameters $p(\theta | x)$
 - characterizes **all** suitable parameters
 - quantifies uncertainty

Bayes' Theorem

$$\text{“posterior”} \rightarrow p(\theta | x) \propto p(x | \theta) p(\theta)$$

“prior”
“likelihood” = simulator

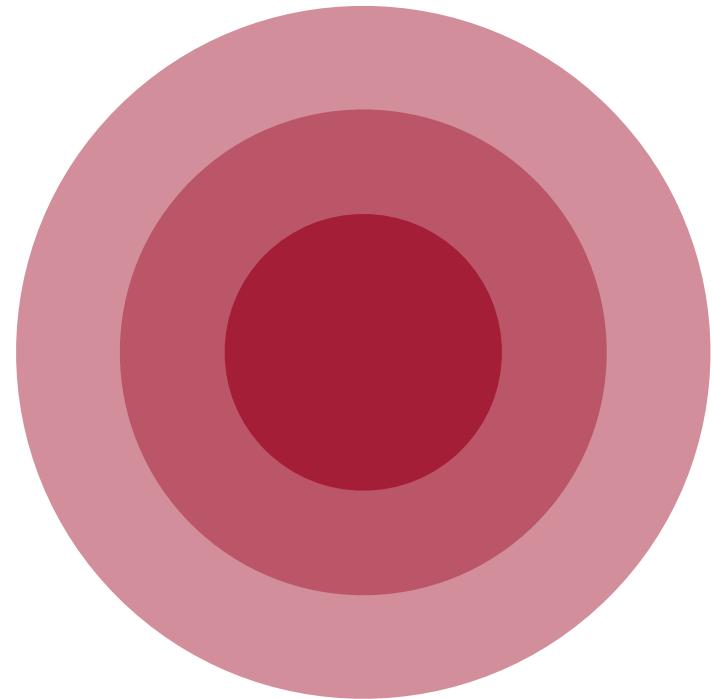
Bayesian parameter inference for simulators



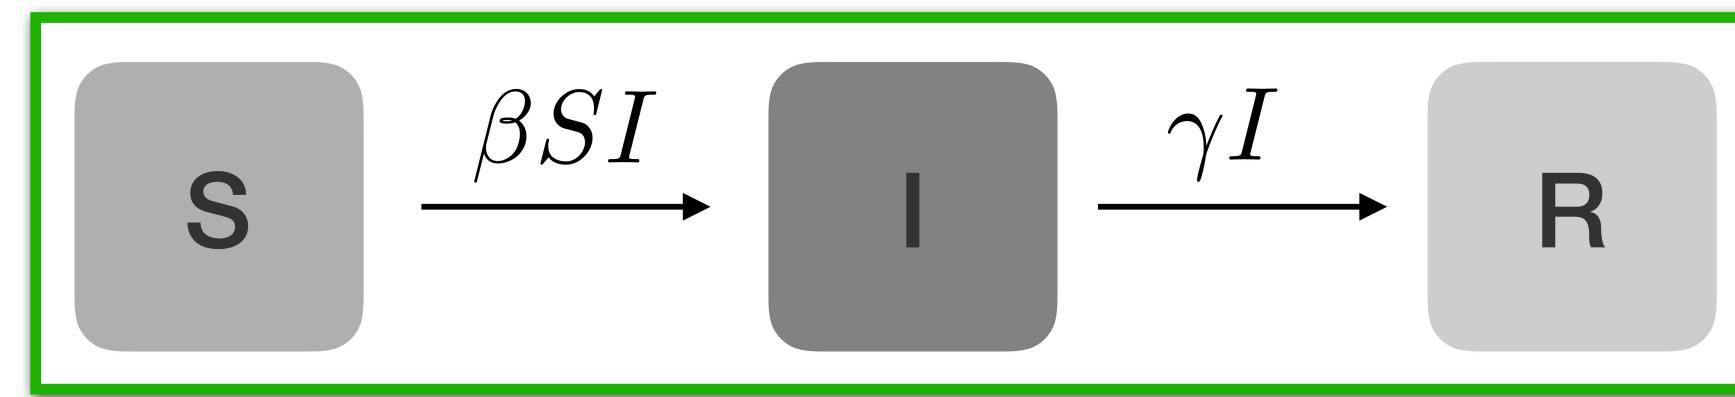
$$p(\theta | x) \propto p(x | \theta) p(\theta)$$

Bayesian parameter inference for simulators

parameters θ

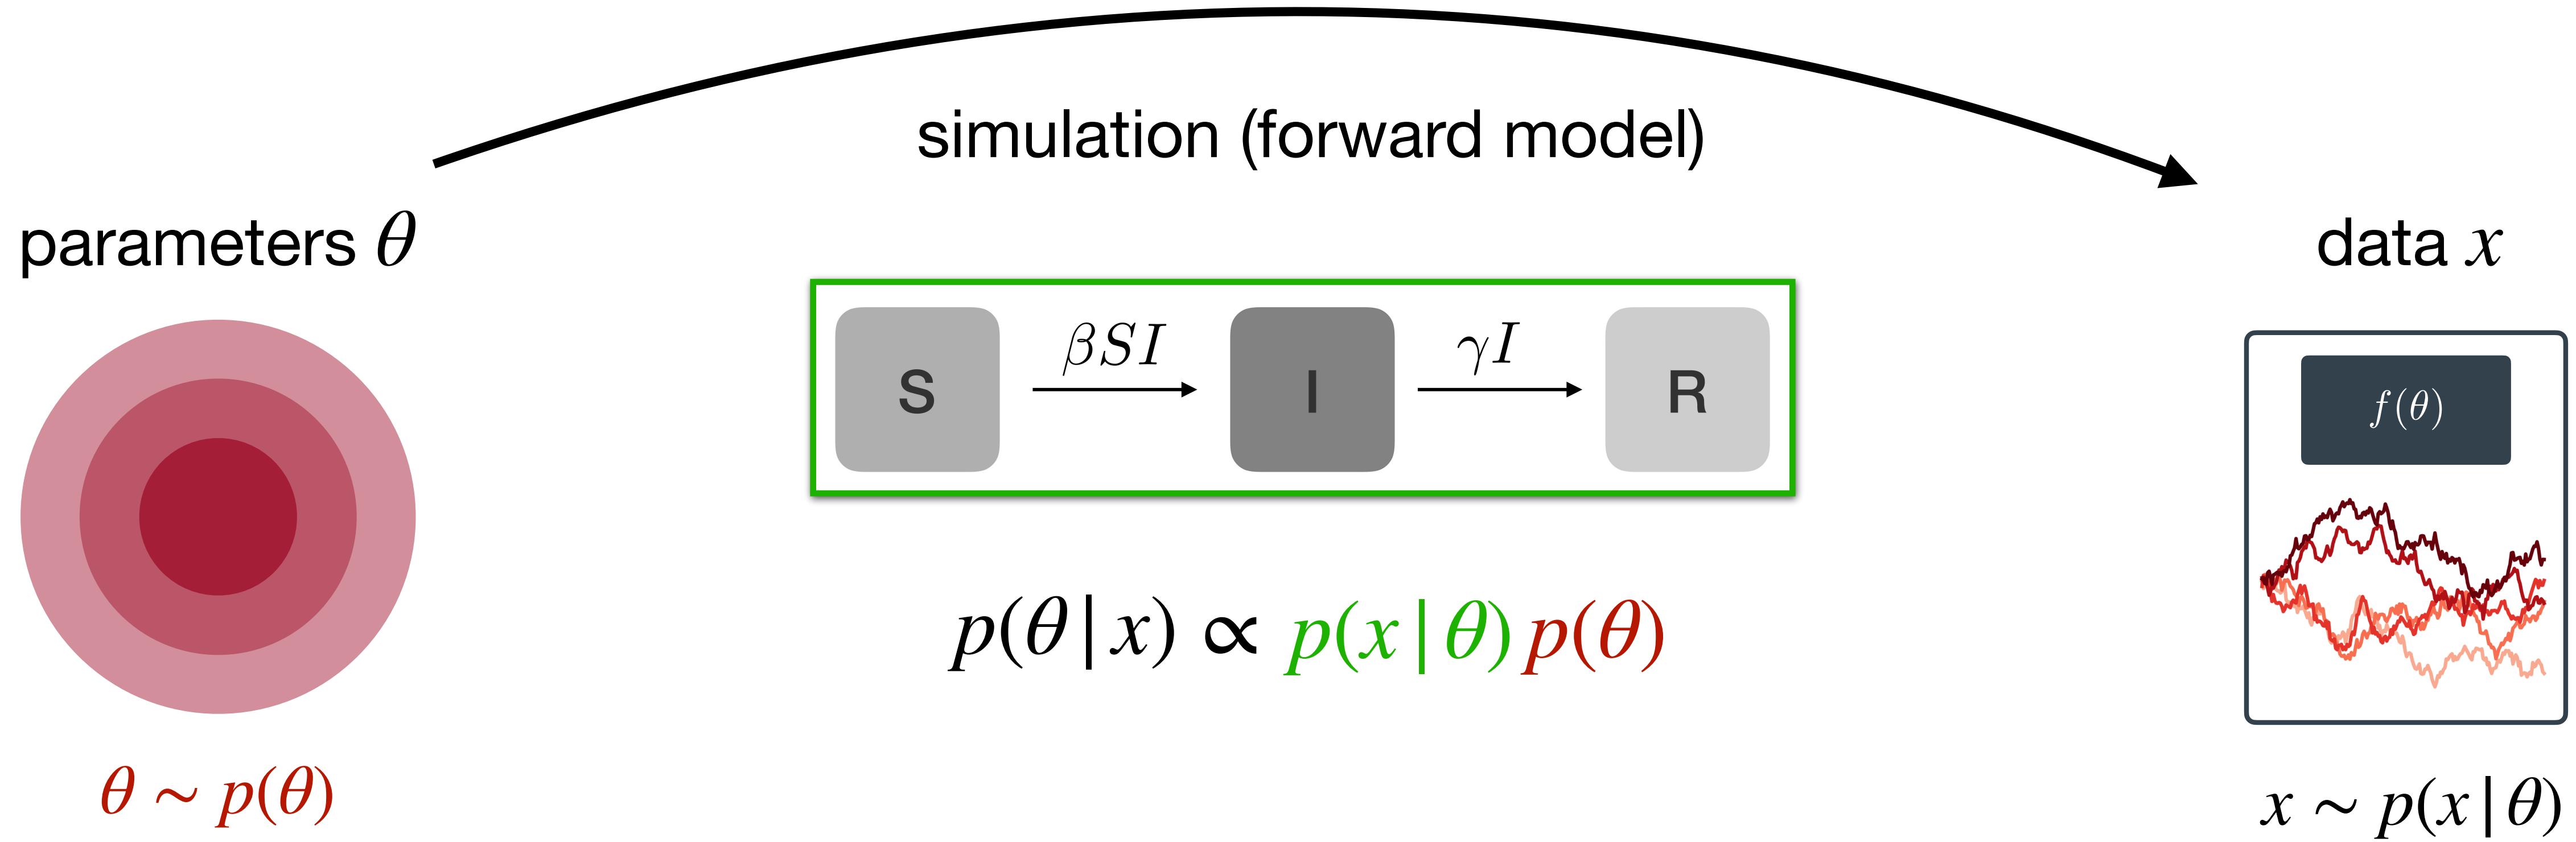


$$\theta \sim p(\theta)$$

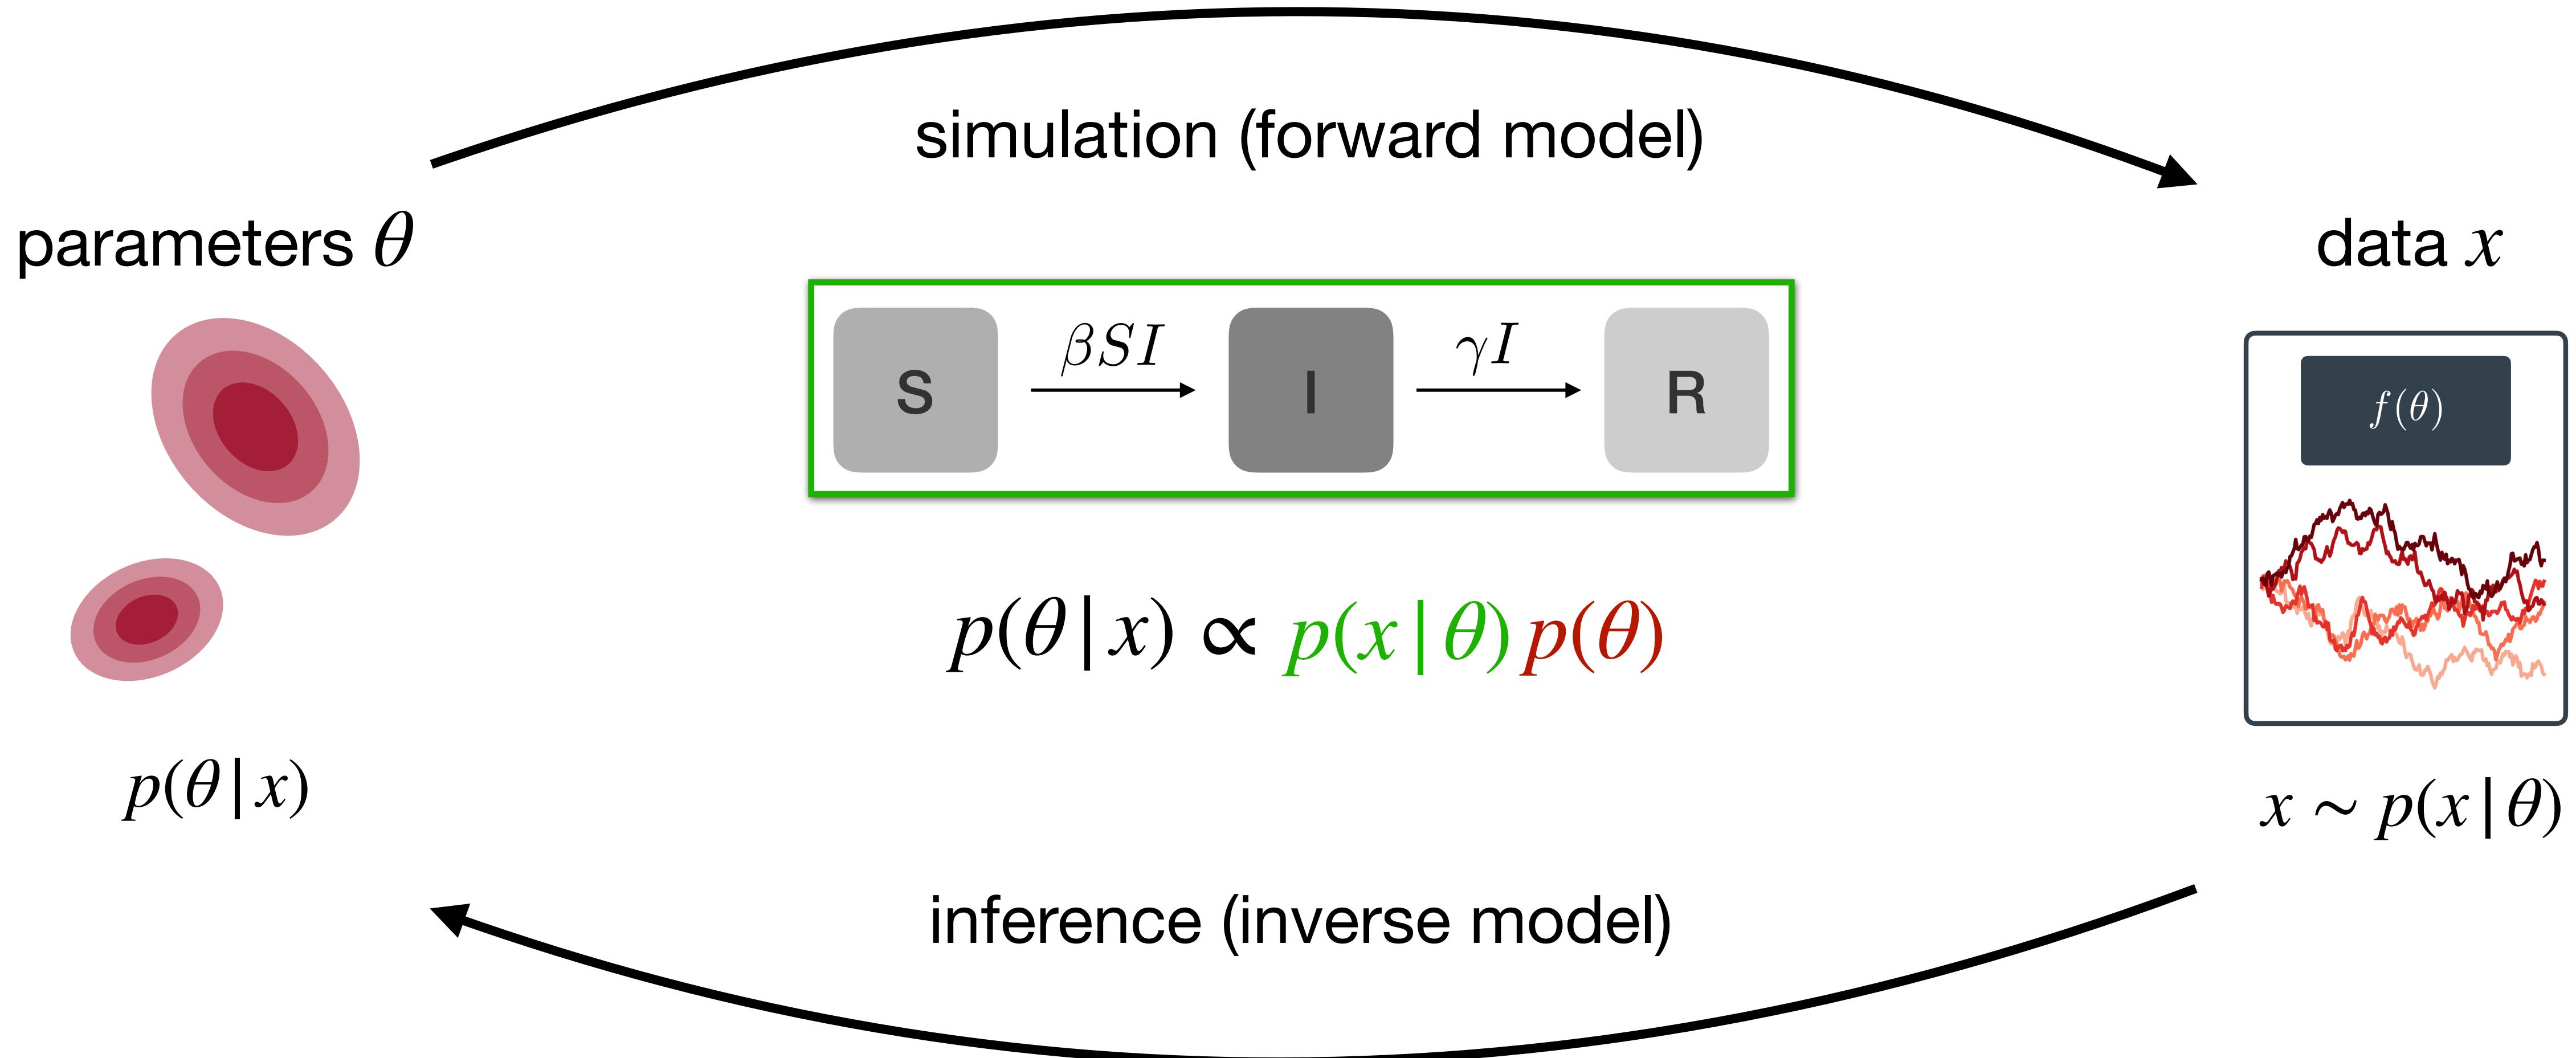


$$p(\theta | x) \propto p(x | \theta) p(\theta)$$

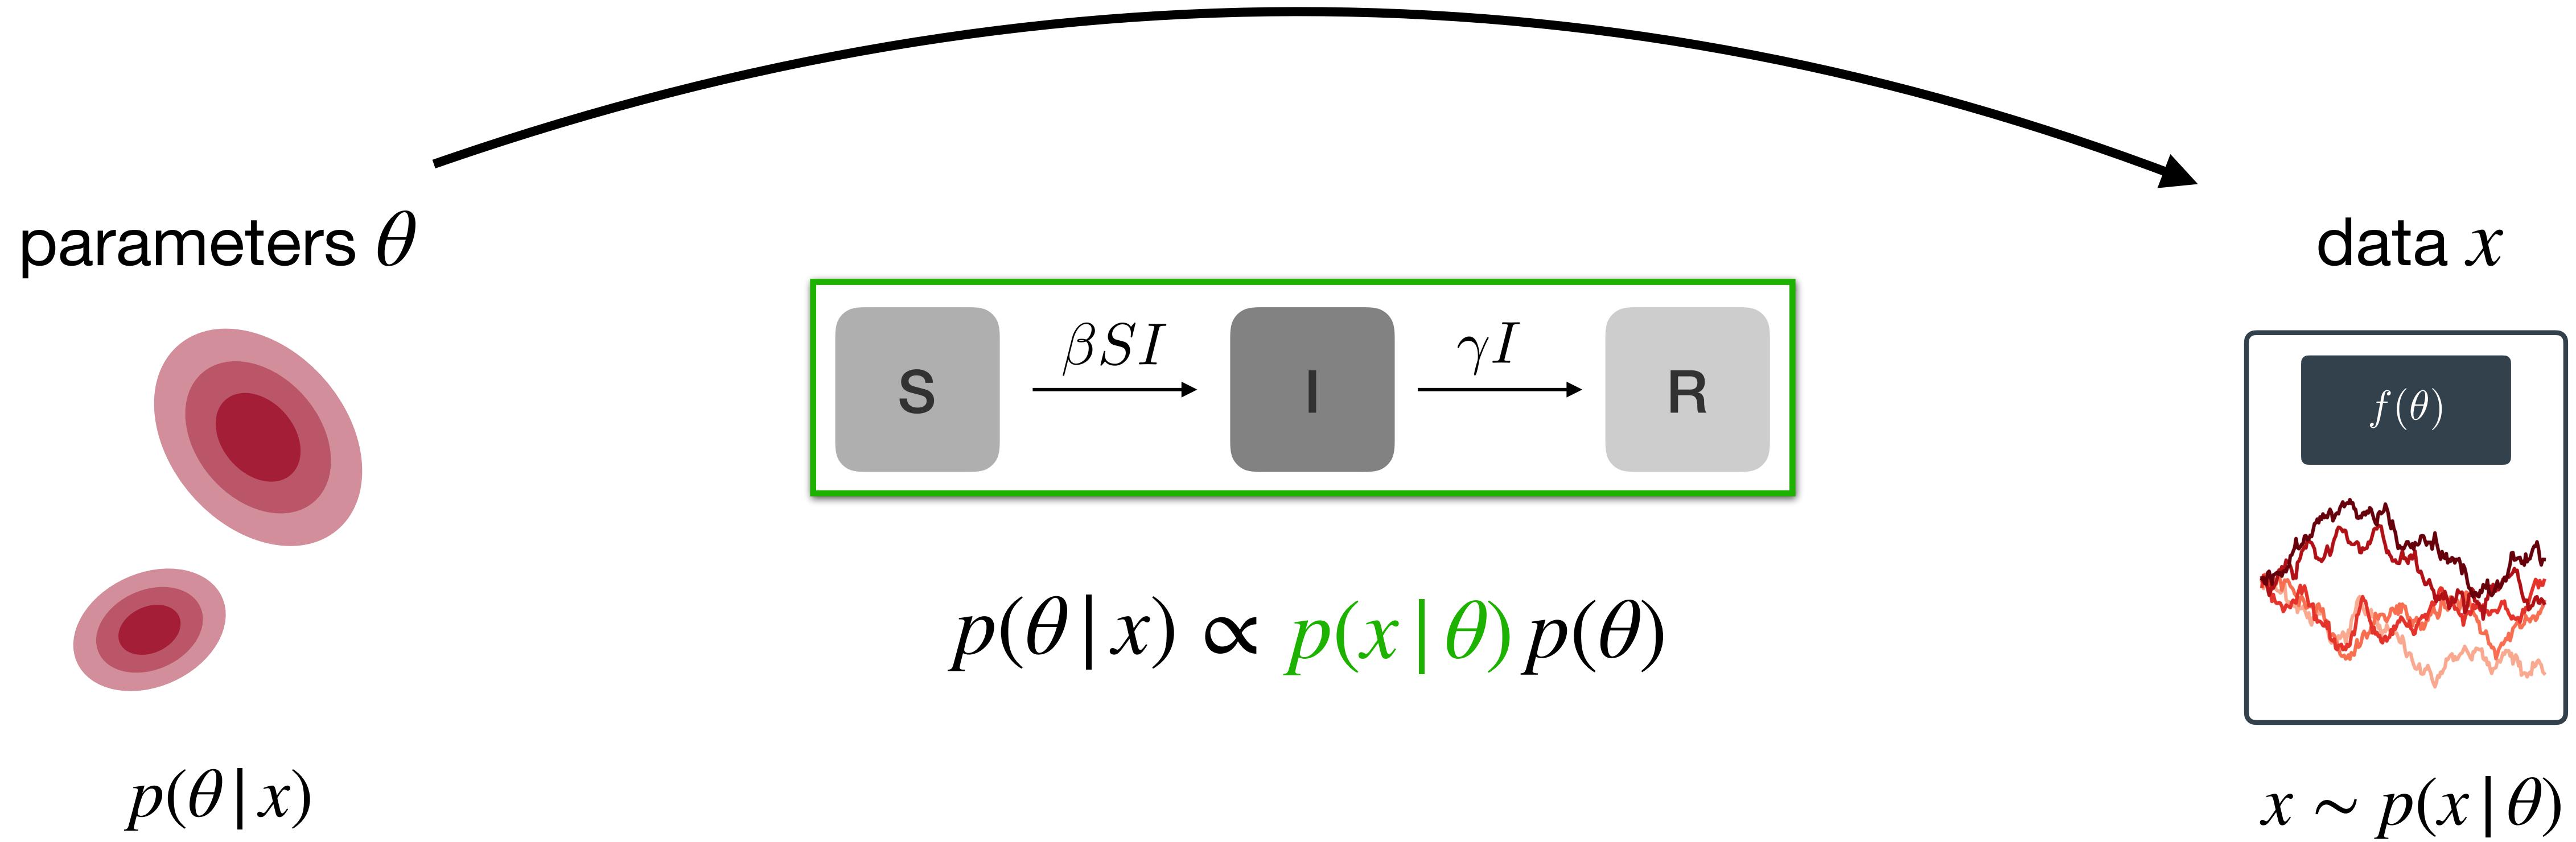
Bayesian parameter inference for simulators



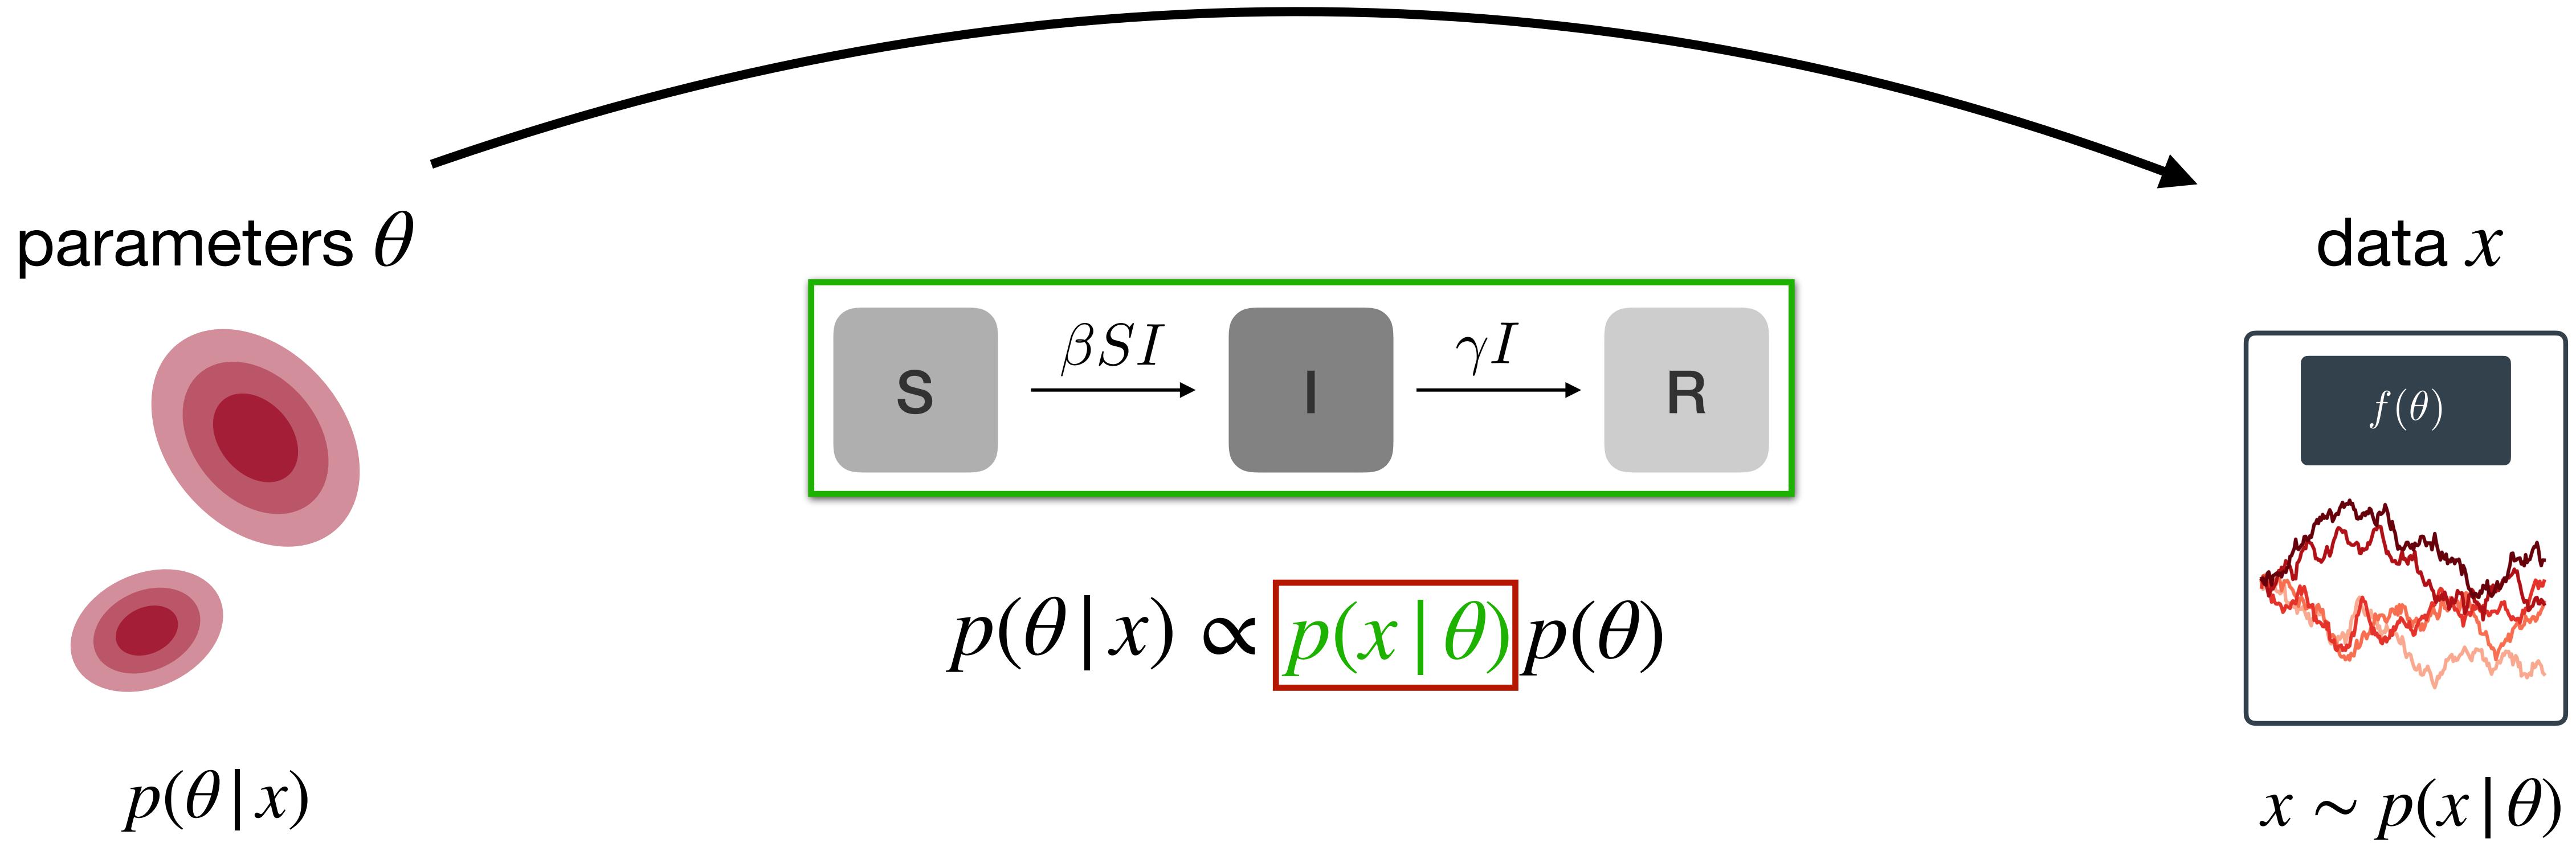
Bayesian parameter inference for simulators



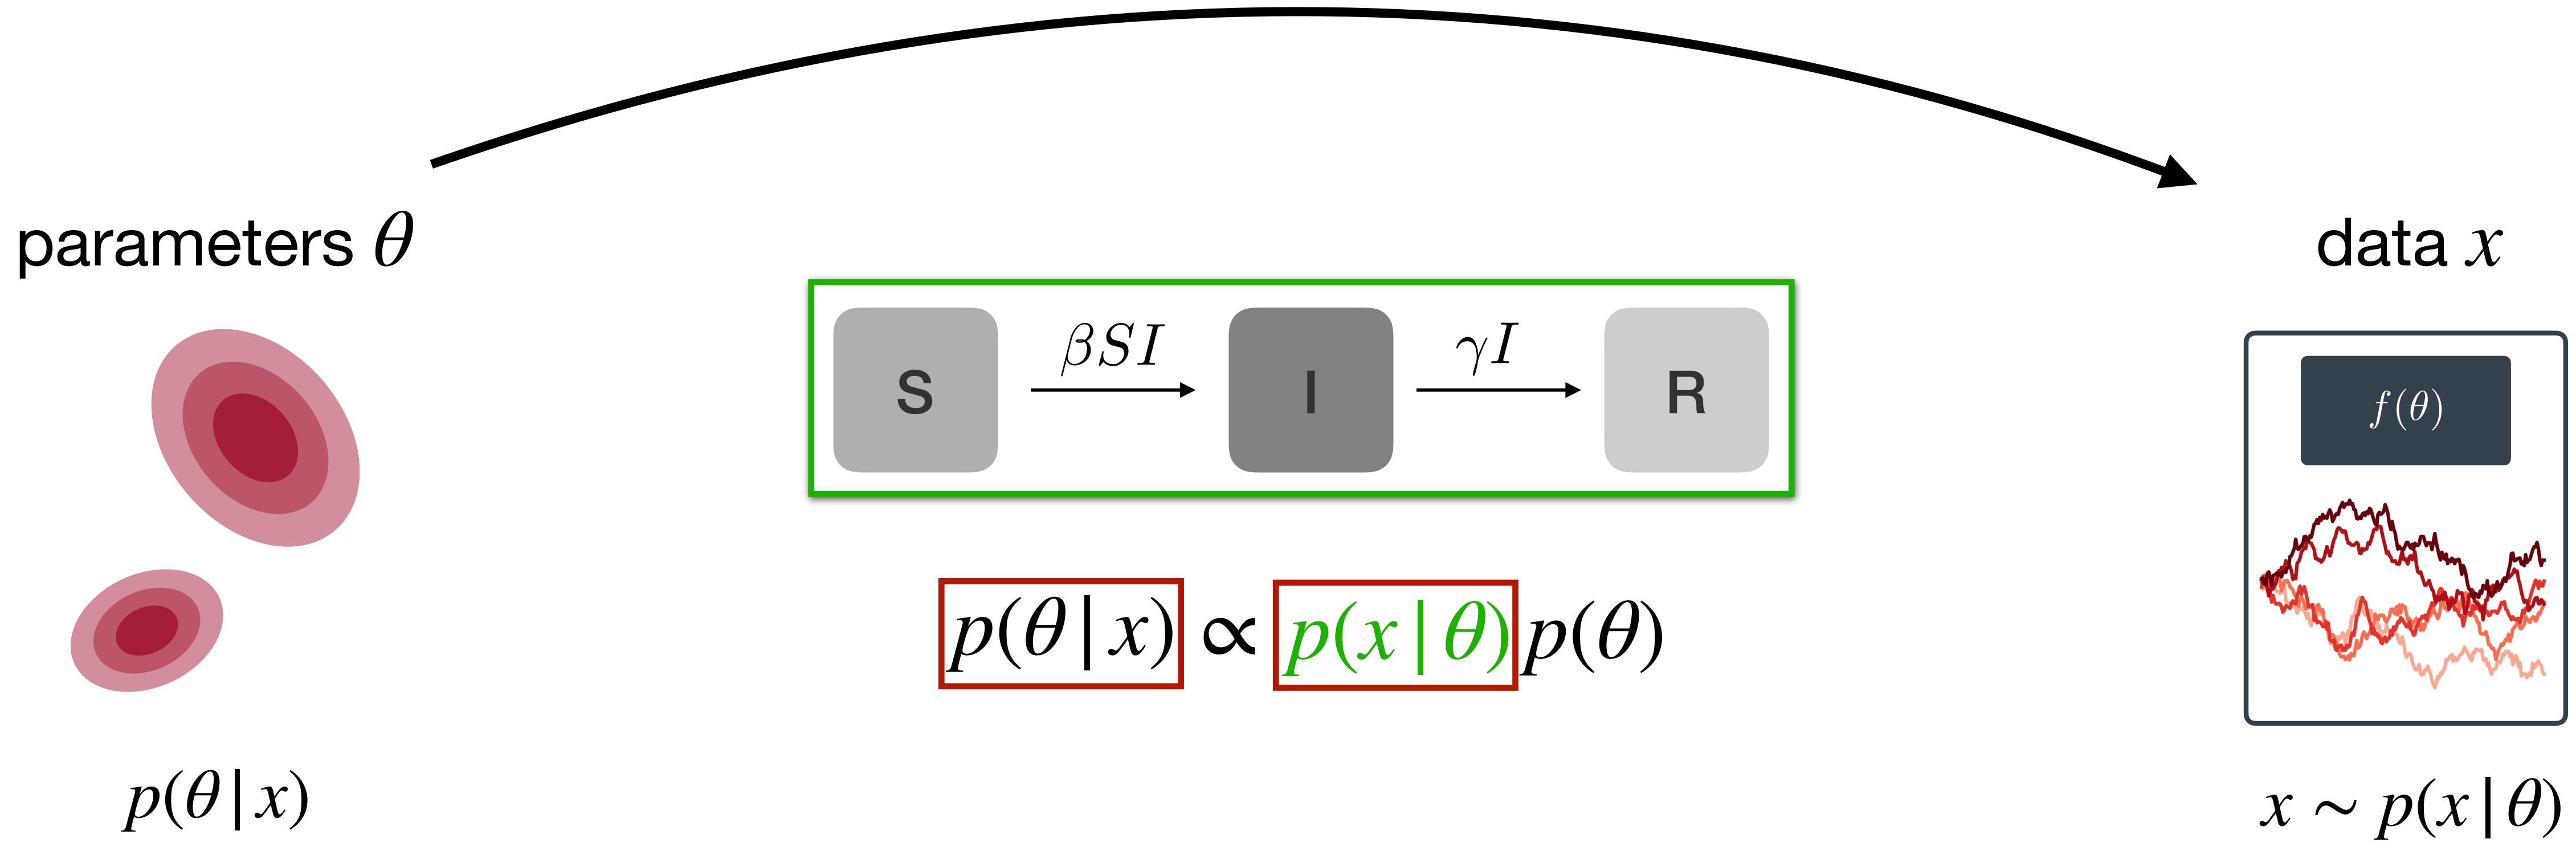
Simulation-based Bayesian inference (SBI)



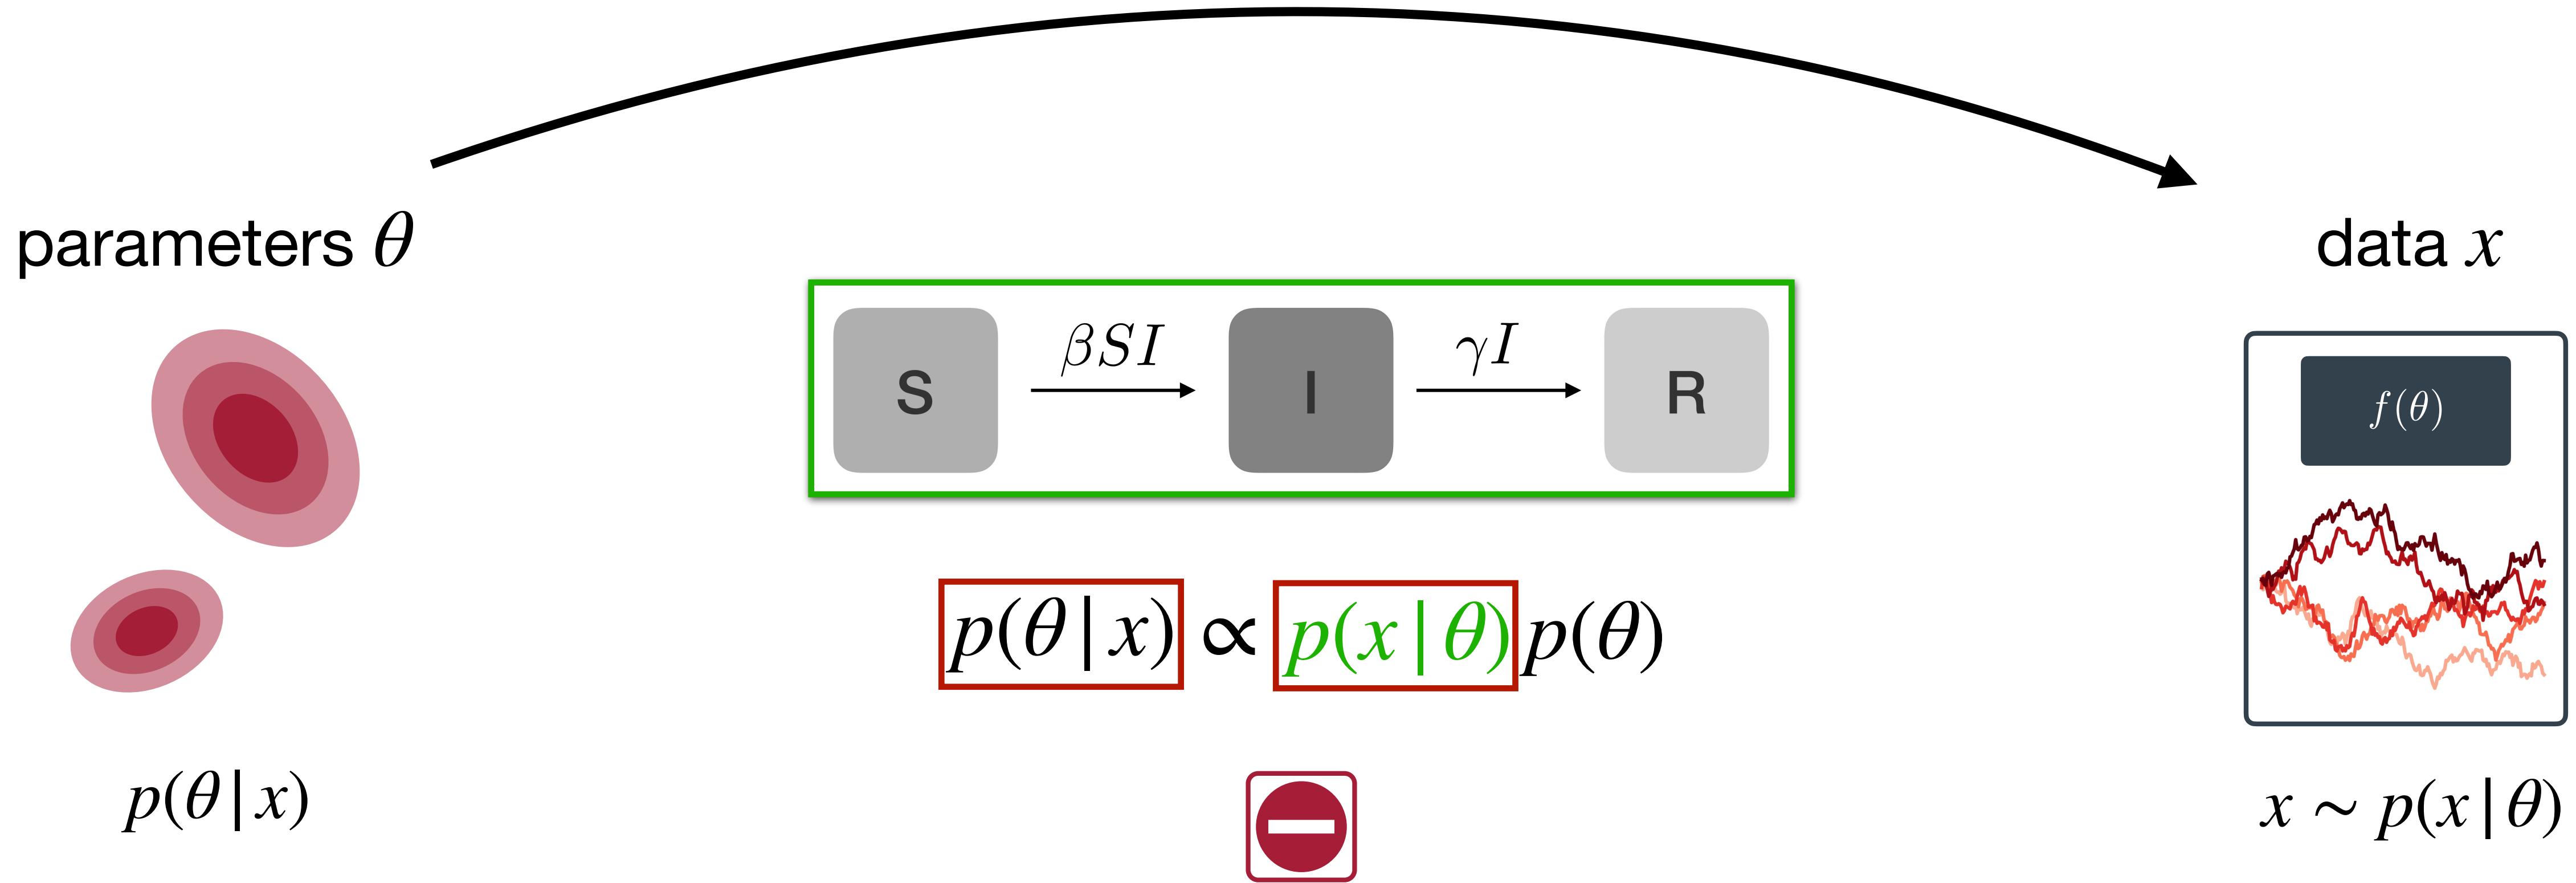
Simulation-based Bayesian inference (SBI)



Simulation-based Bayesian inference (SBI)

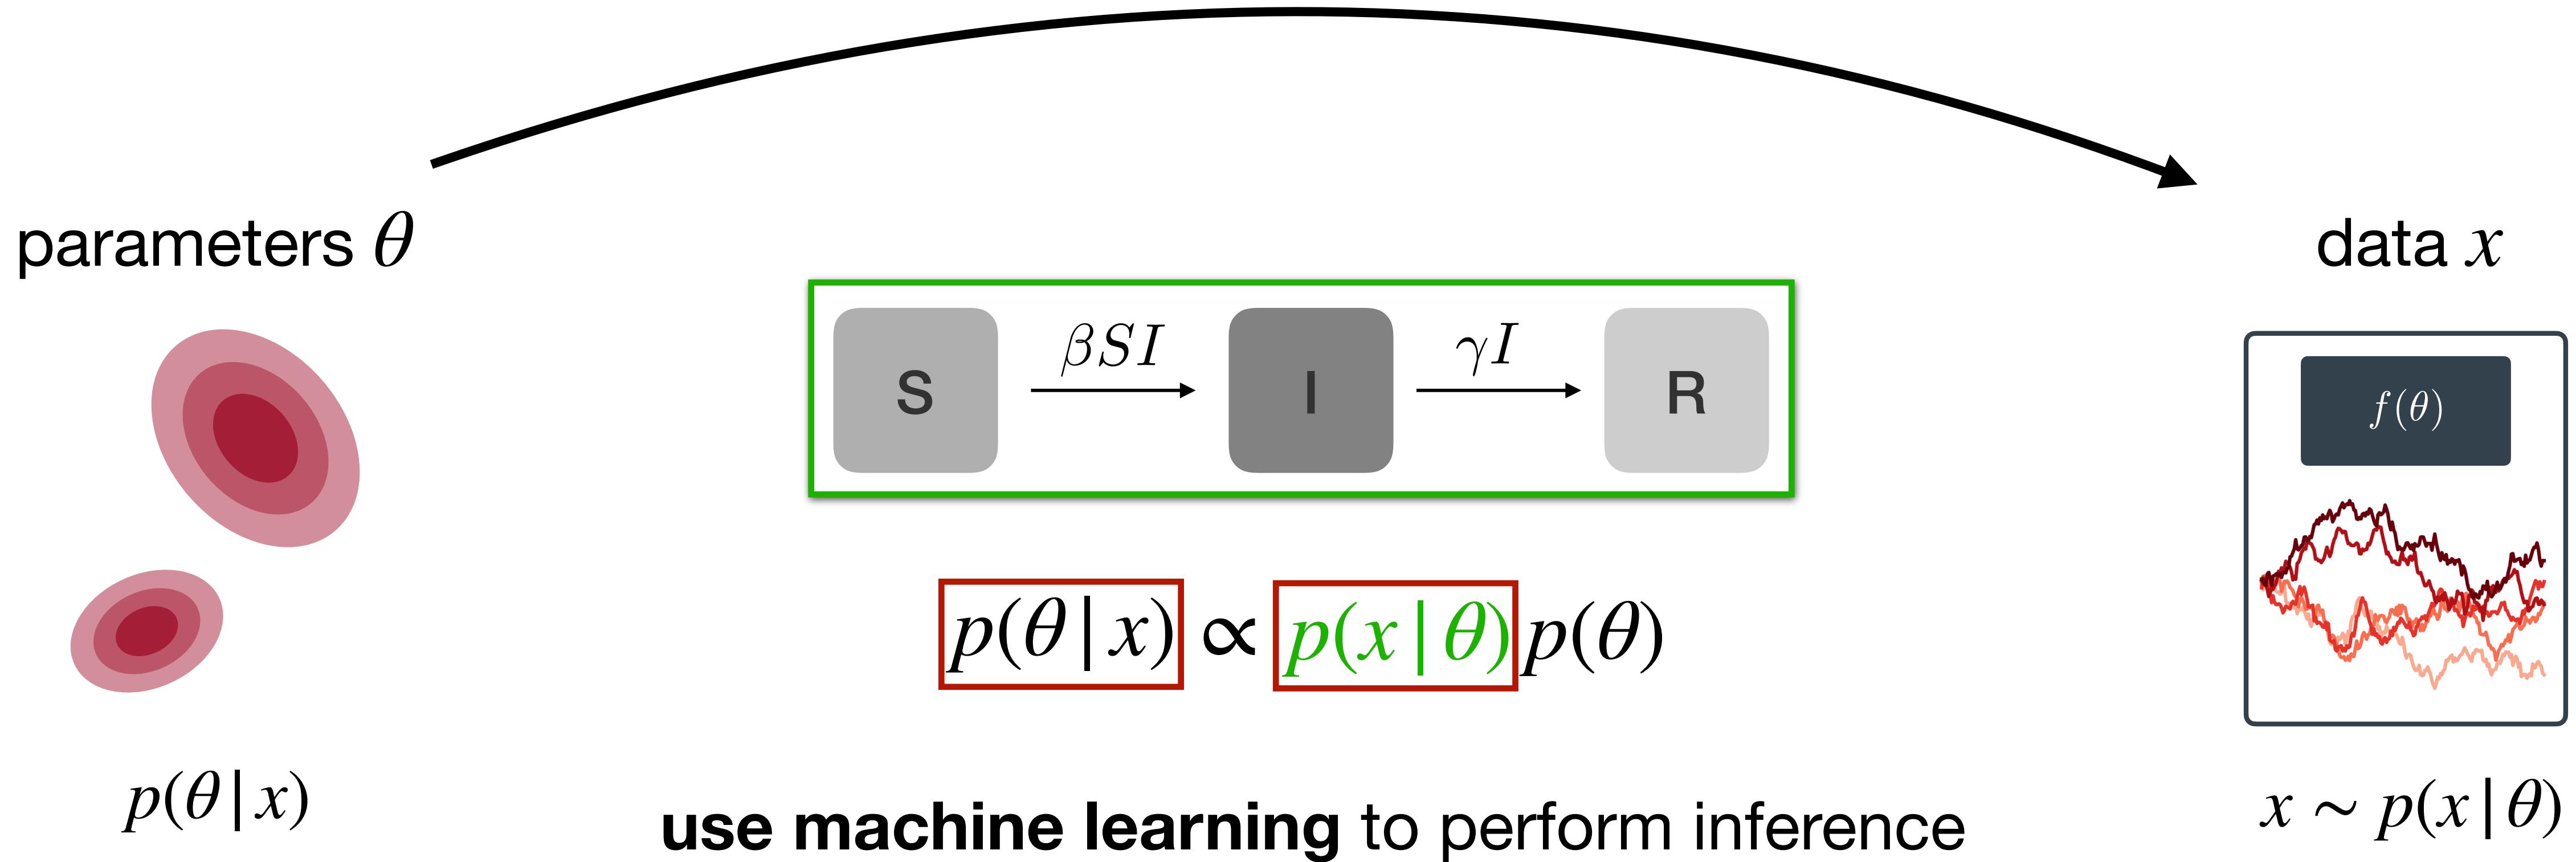


Simulation-based Bayesian inference (SBI)

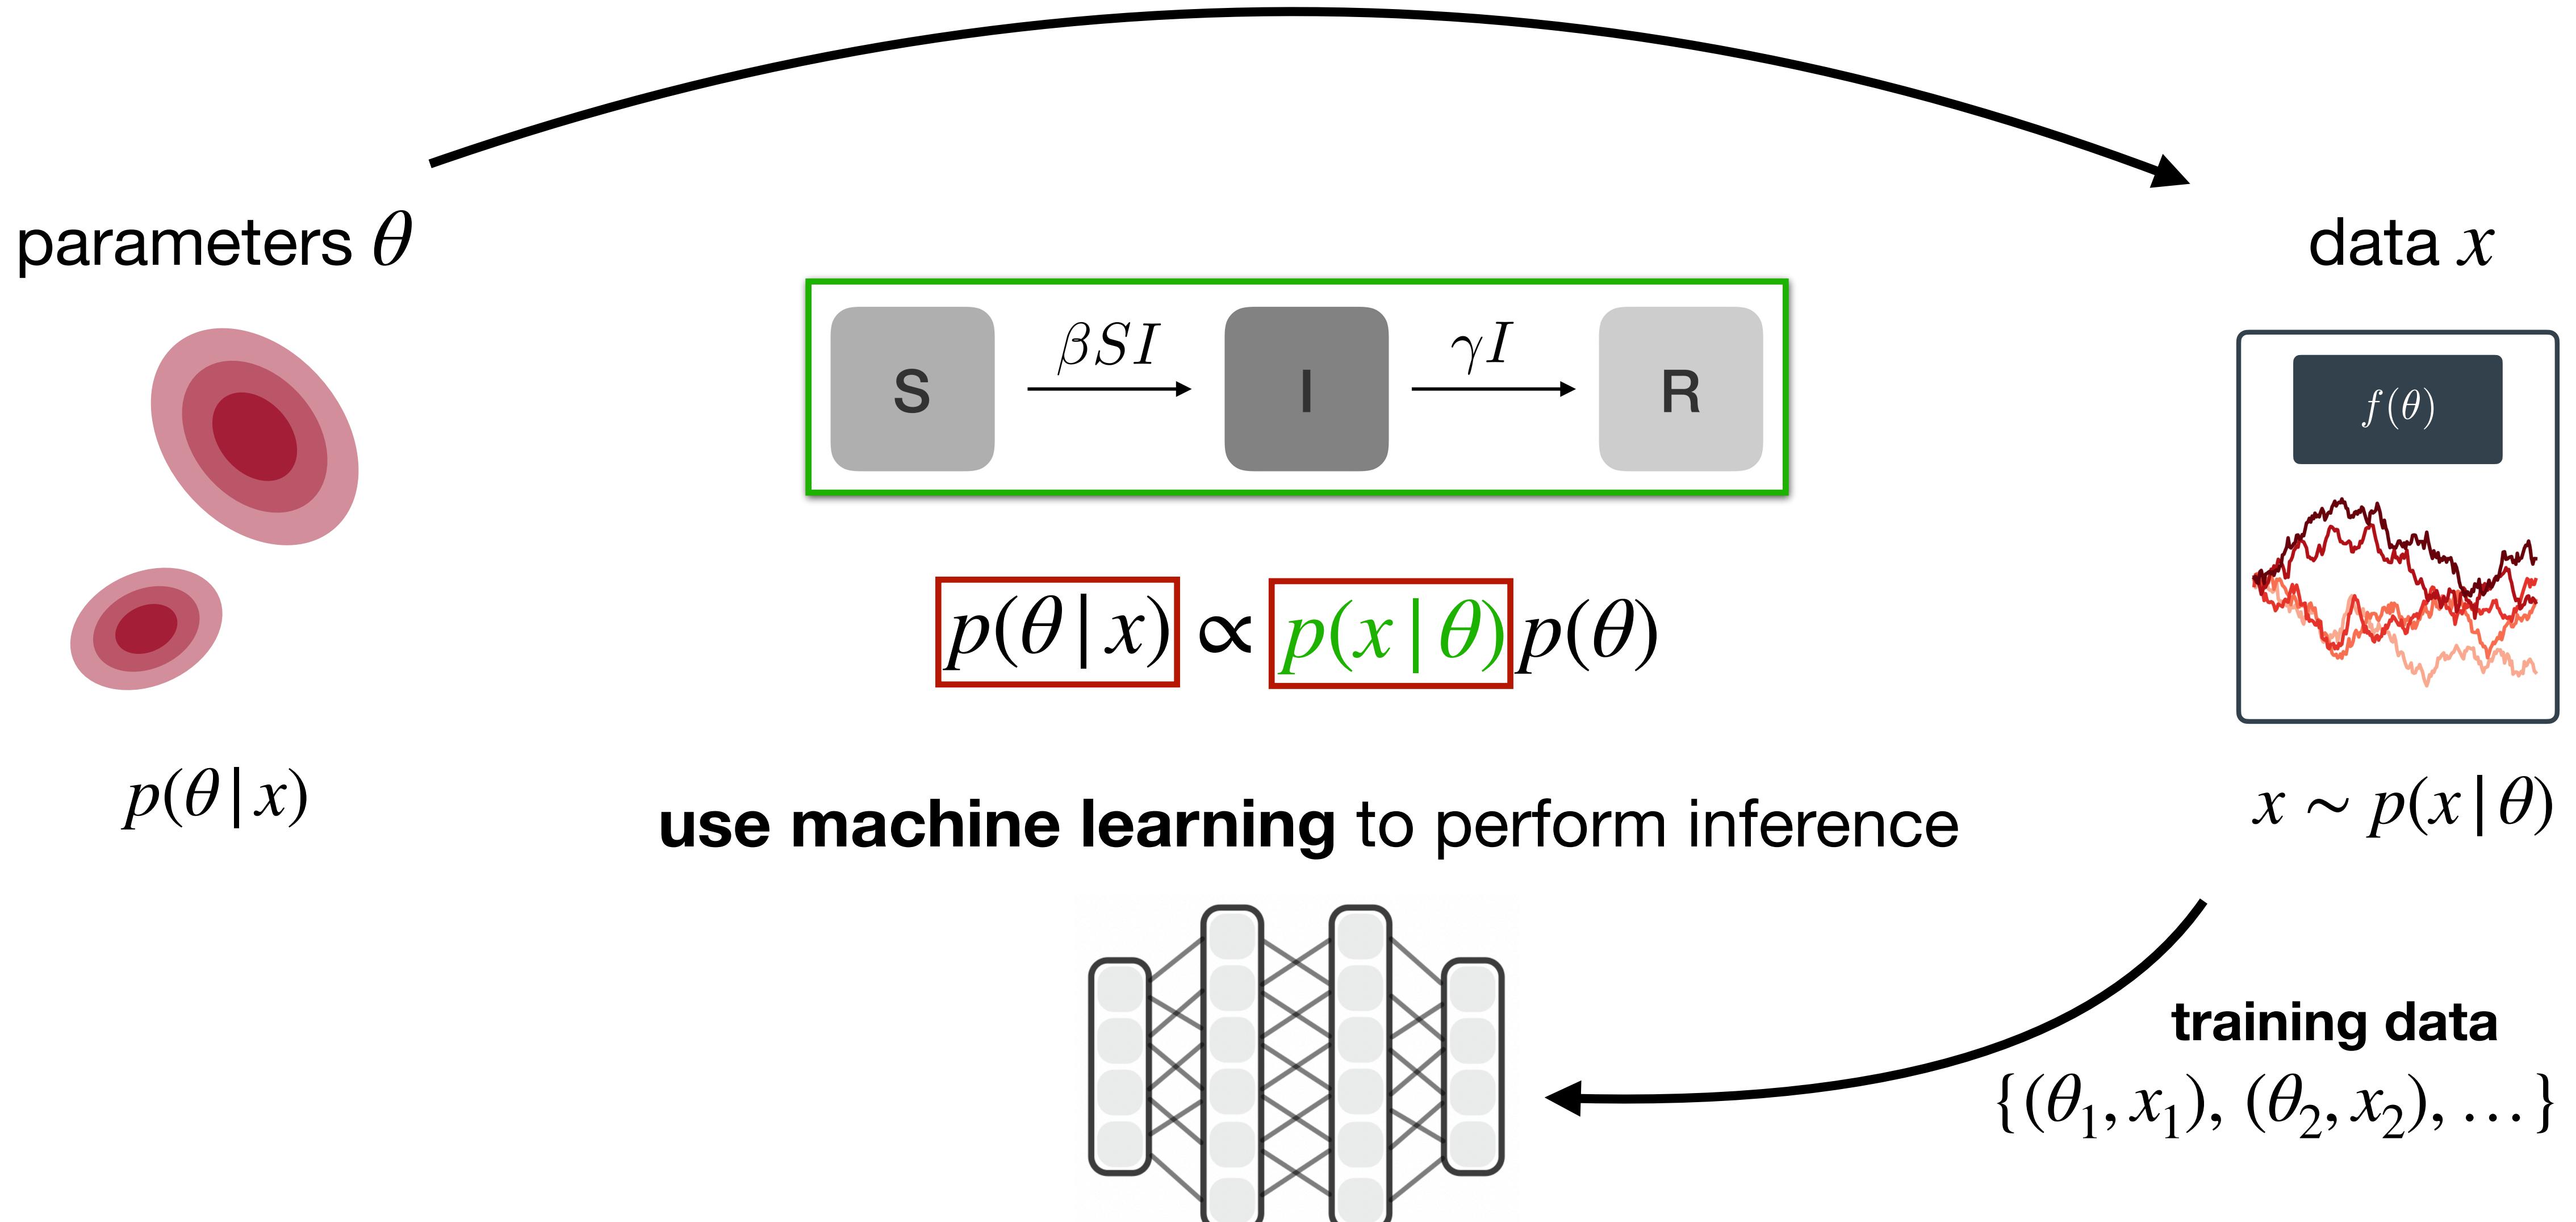


No Markov Chain Monte Carlo (MCMC)
No variational inference
No Bayesian optimization

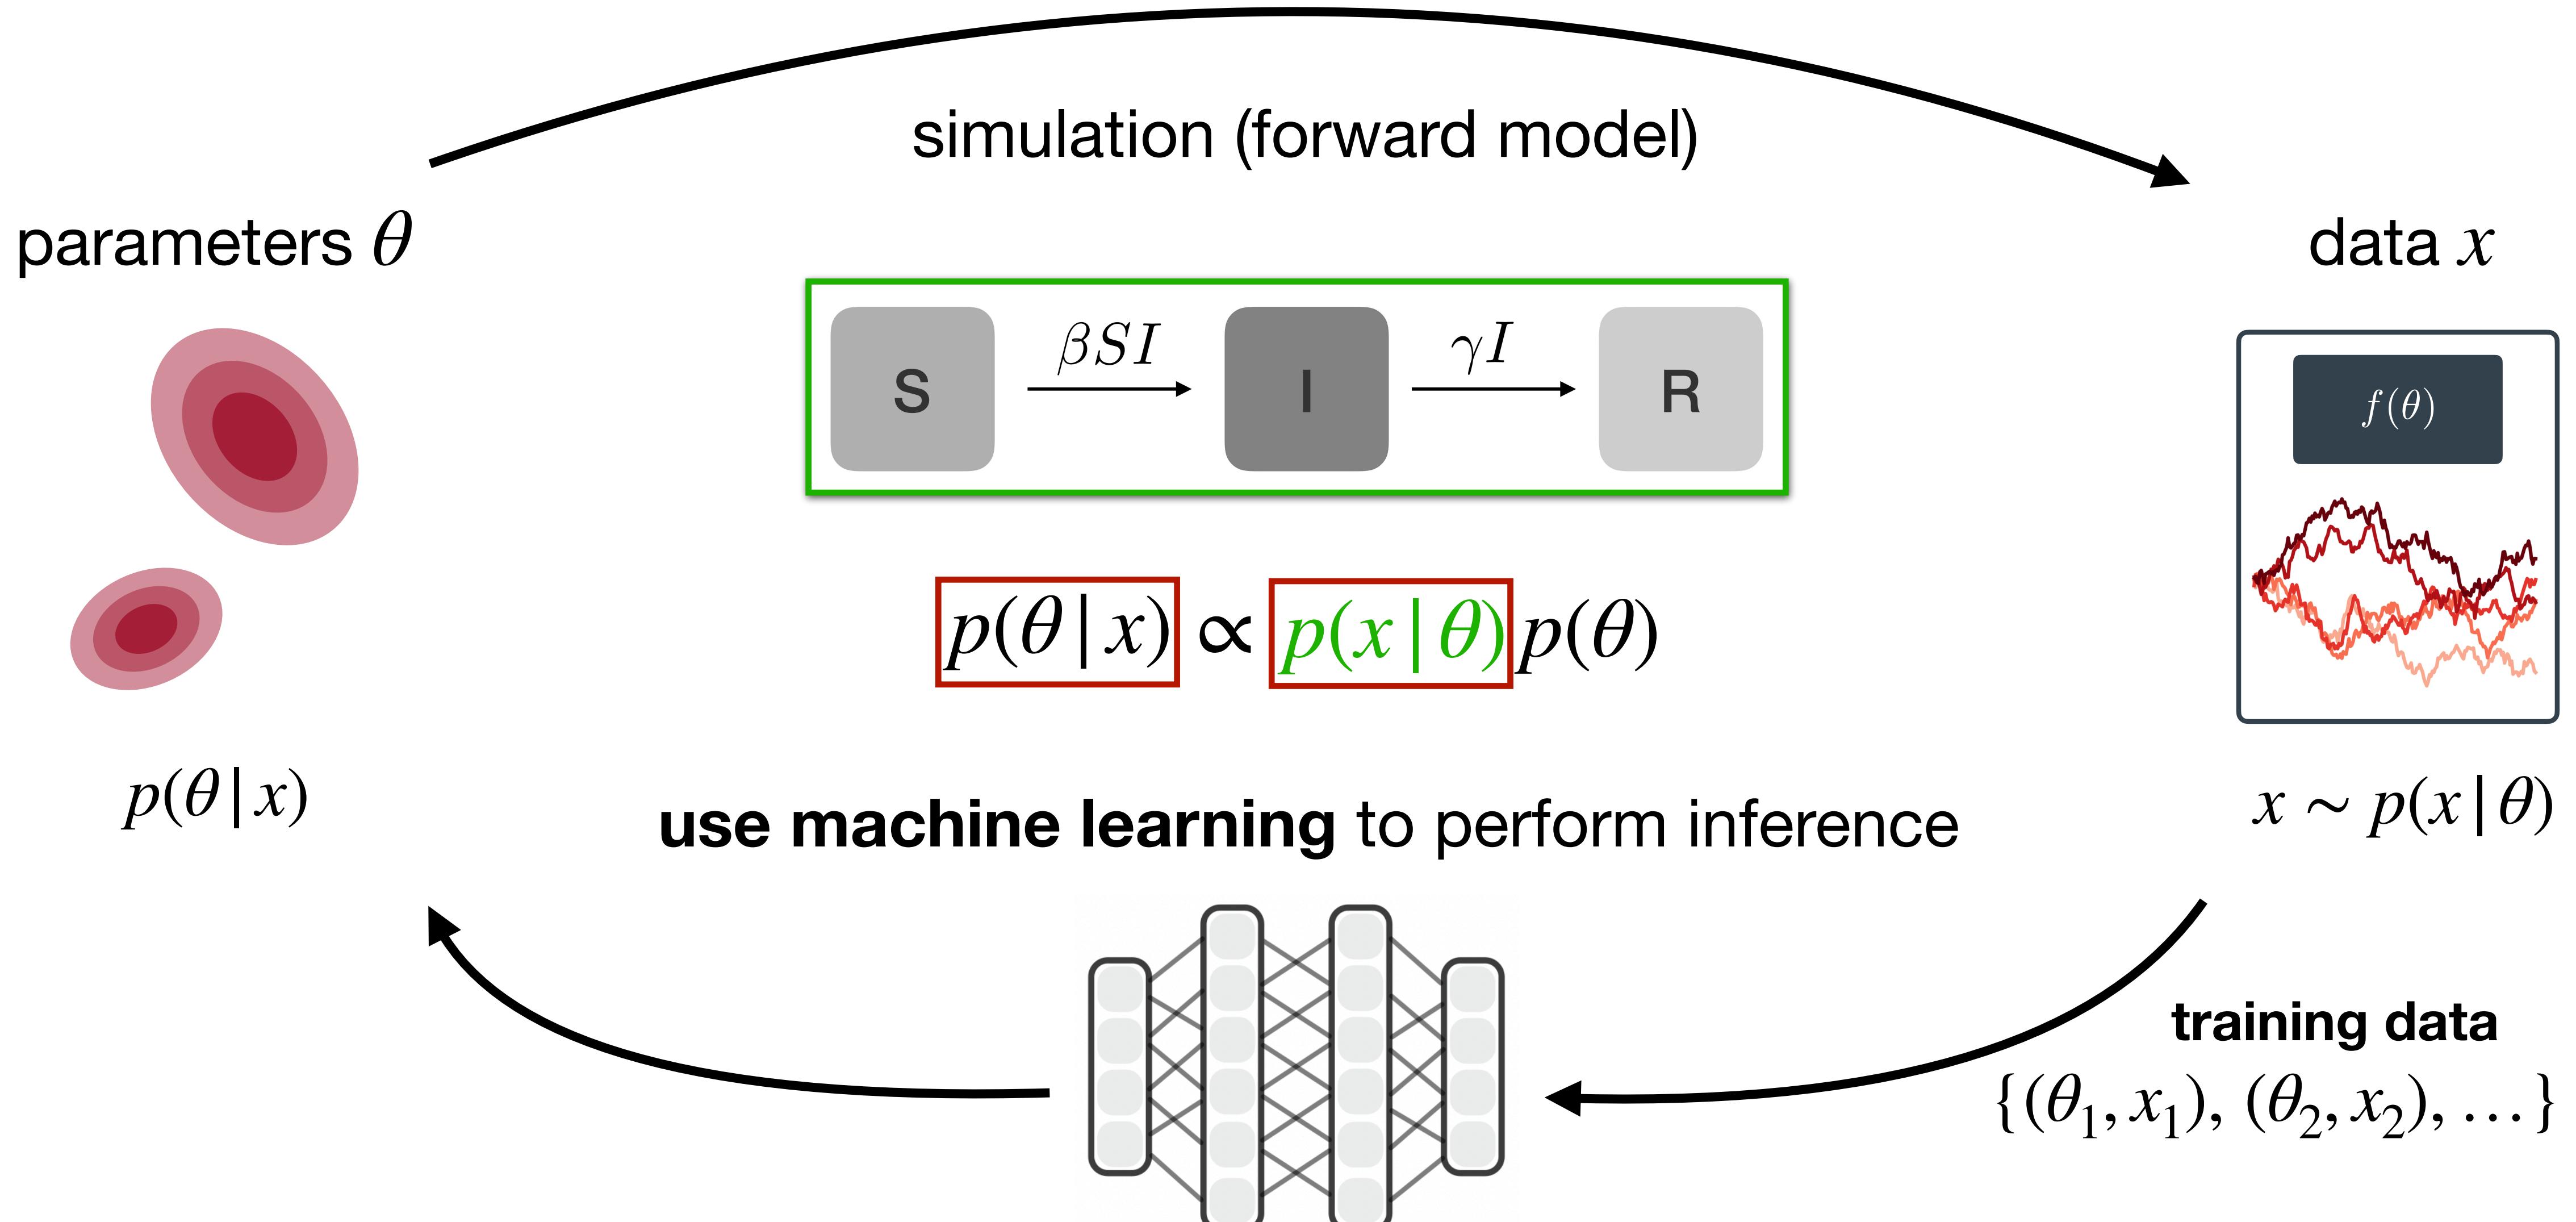
Simulation-based Bayesian inference (SBI)



Simulation-based Bayesian inference (SBI)



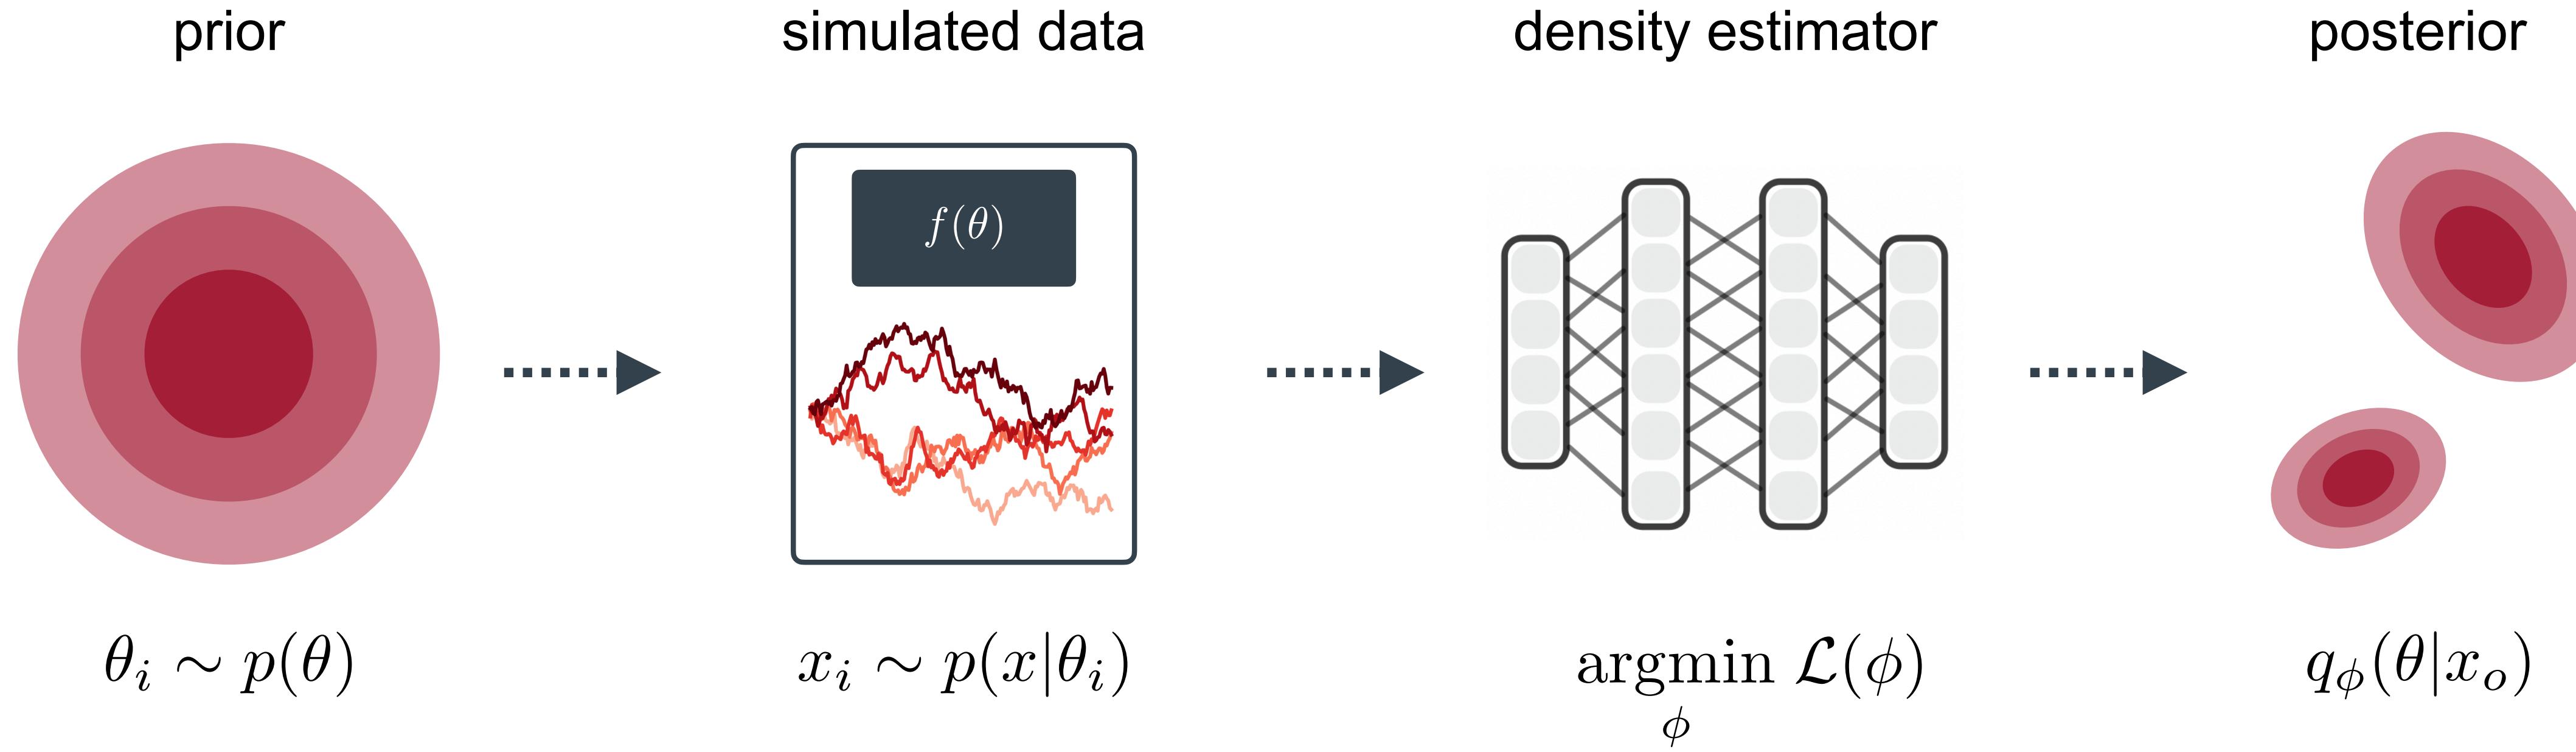
Simulation-based Bayesian inference (SBI)



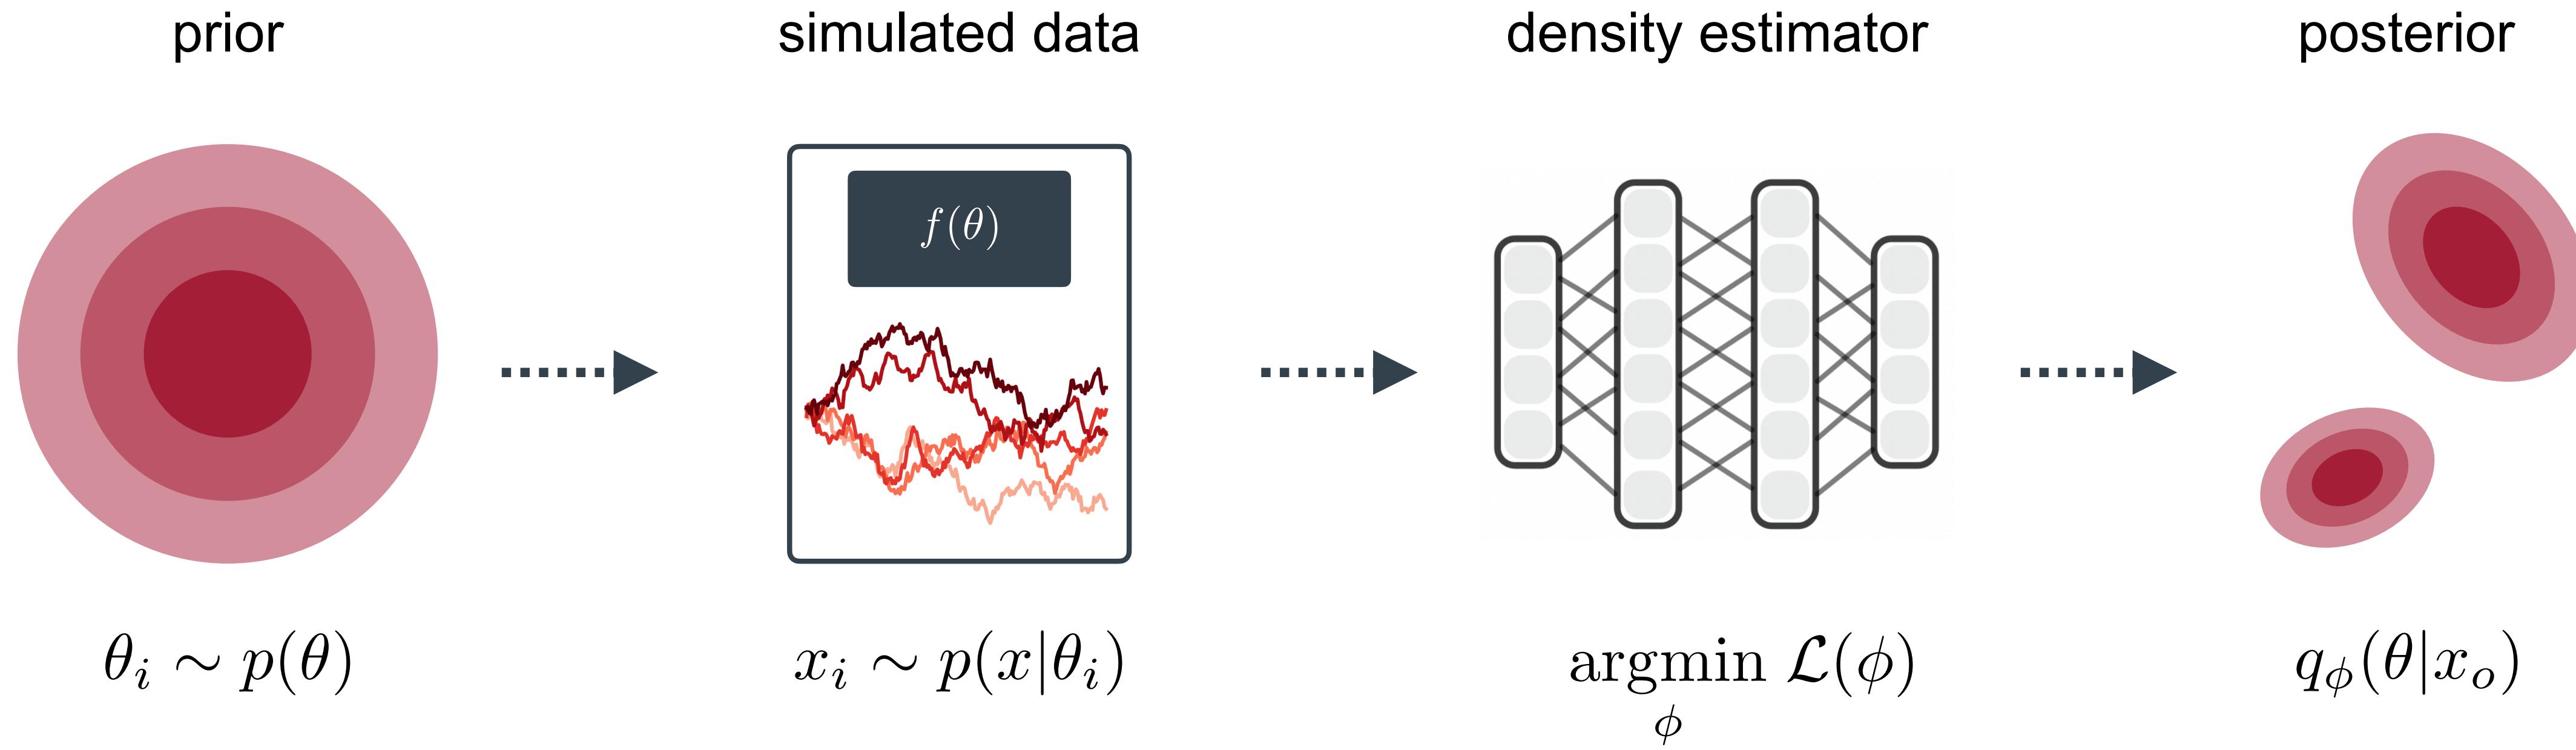
Outline

1. **Why** simulators?
2. **Why** Bayesian parameter inference?
3. **How** to do Bayesian inference for simulators?
4. **How** can you apply SBI to your simulator?

Neural Simulation-Based Inference

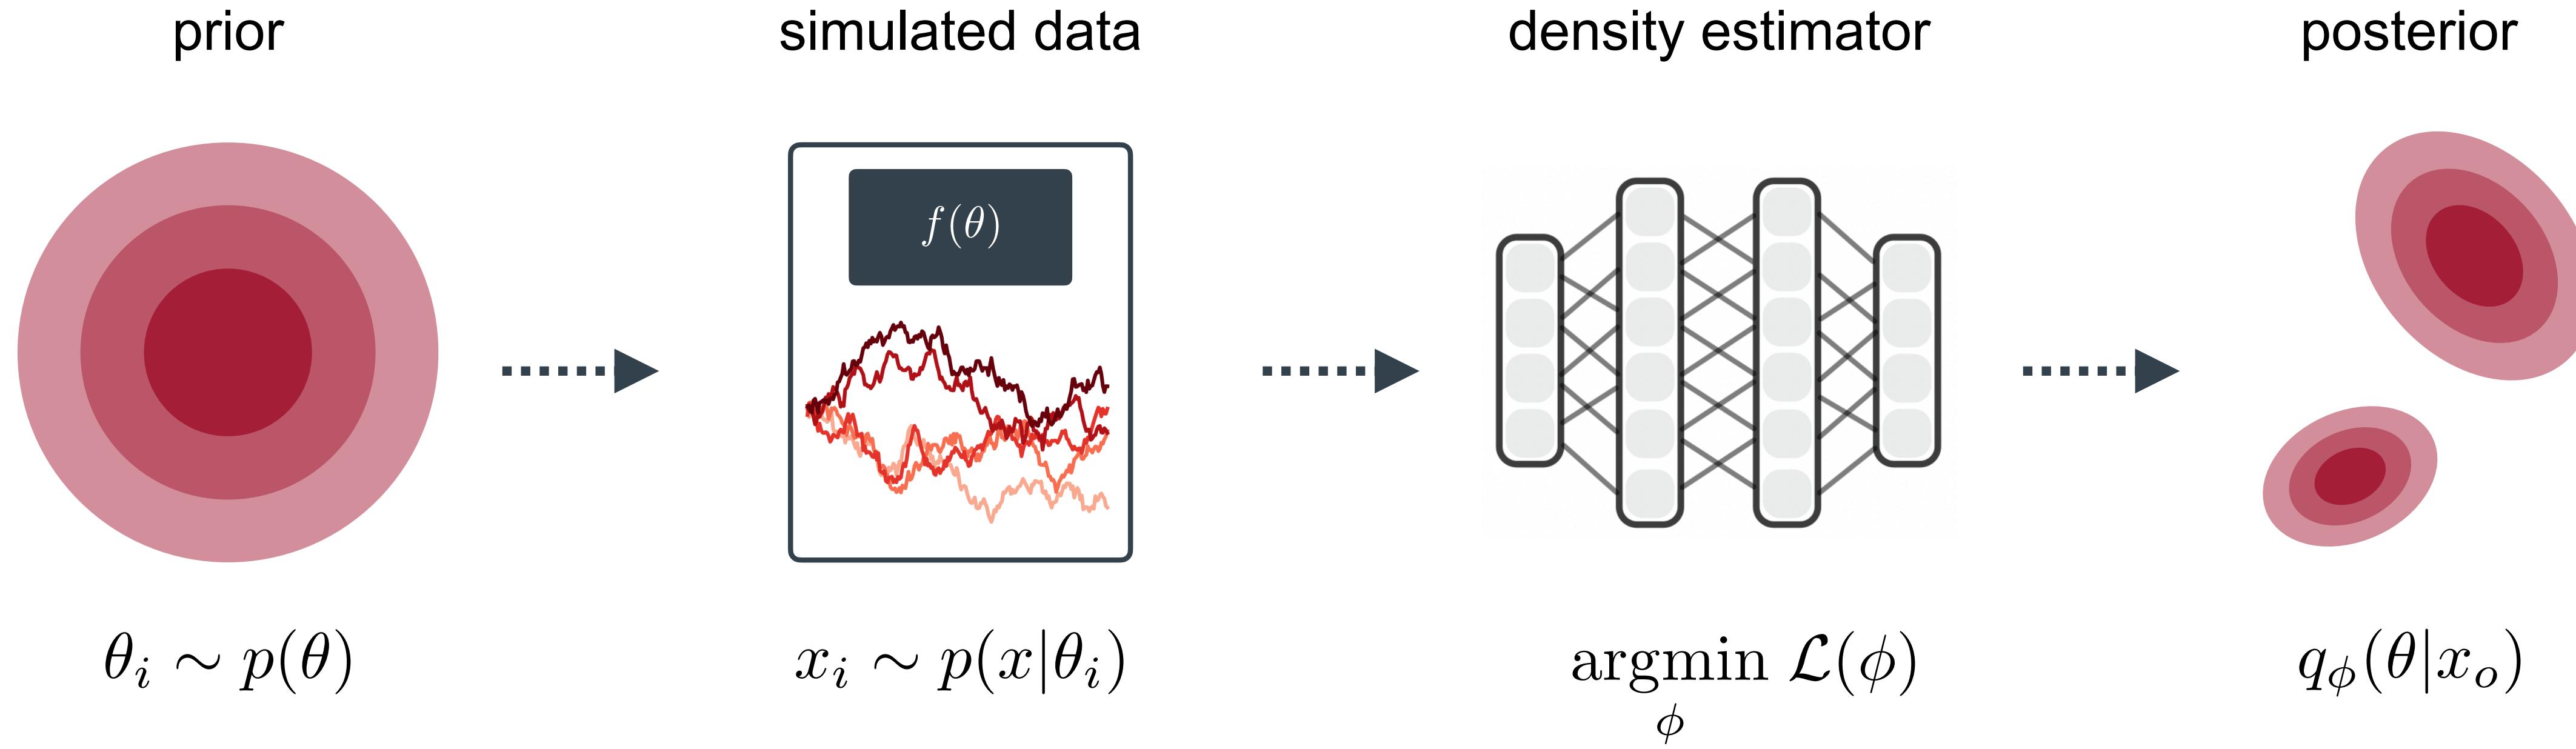


Neural Simulation-Based Inference



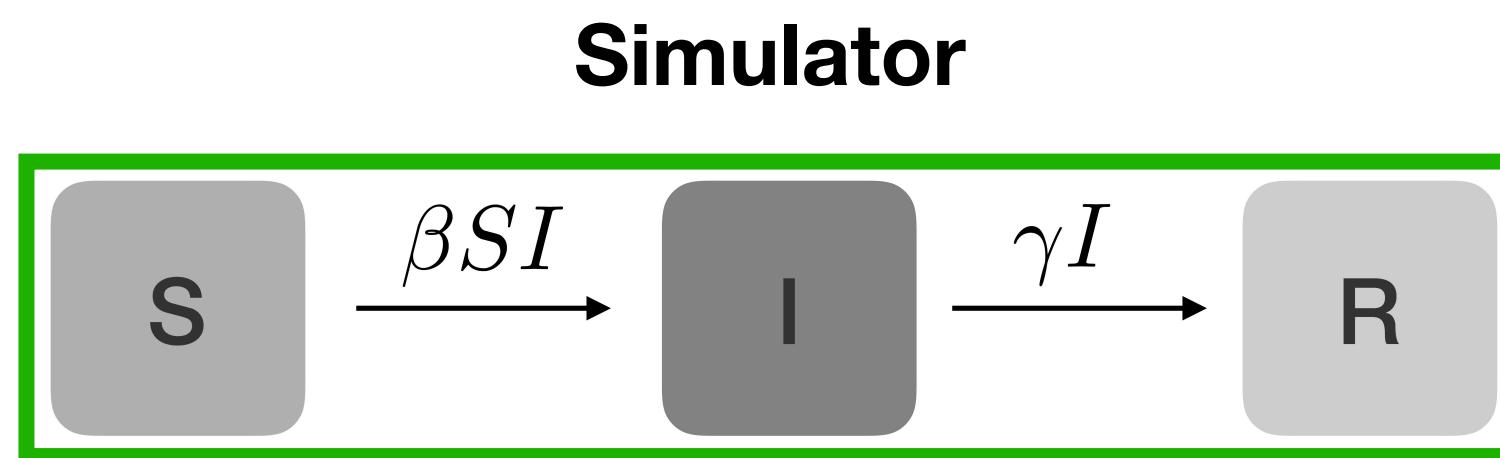
1. **Training:** train neural network (NN) with simulated data to learn $q_{\phi}(\theta | x)$

Neural Simulation-Based Inference

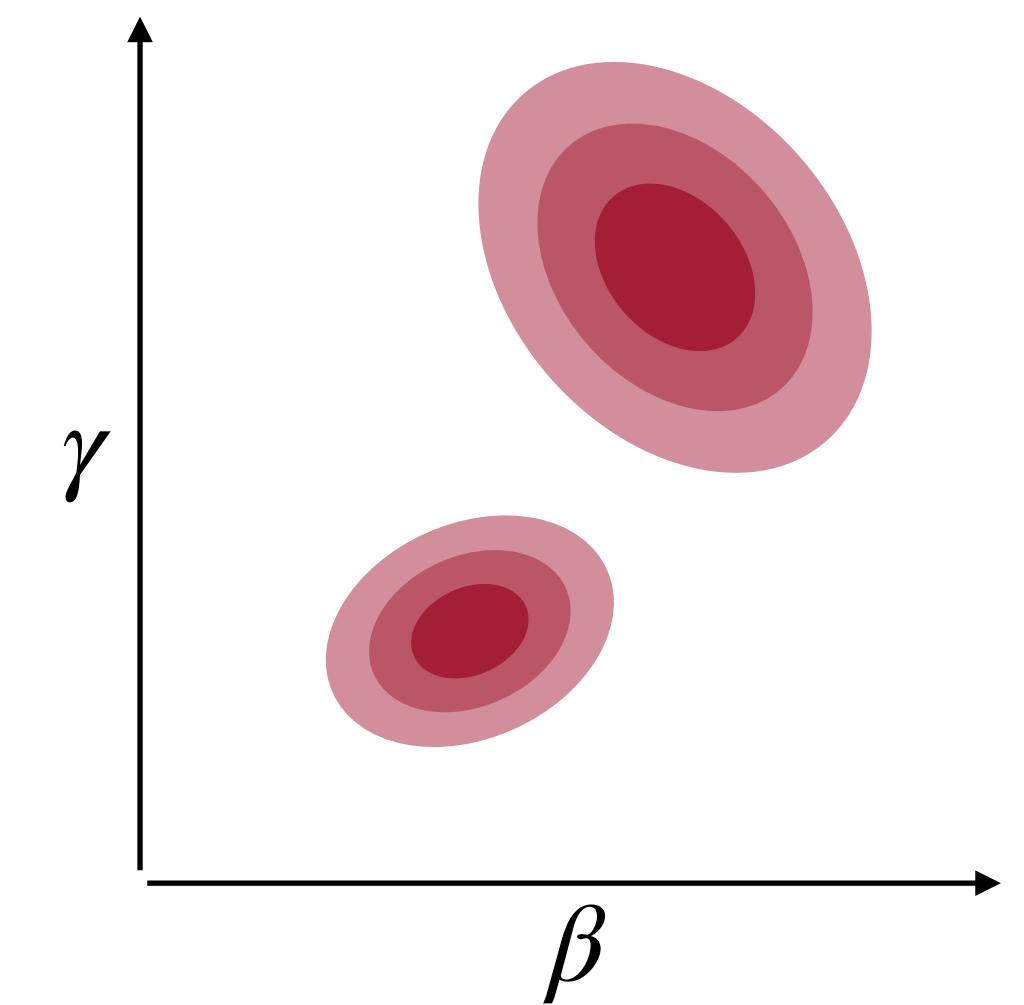


- Training:** train neural network (NN) with simulated data to learn $q_{\phi}(\theta | x)$
- Inference:** plug x_o into NN to obtain $q_{\phi}(\theta | x_o)$

Neural Simulation-Based Inference



Conditional Density Estimation



Prior

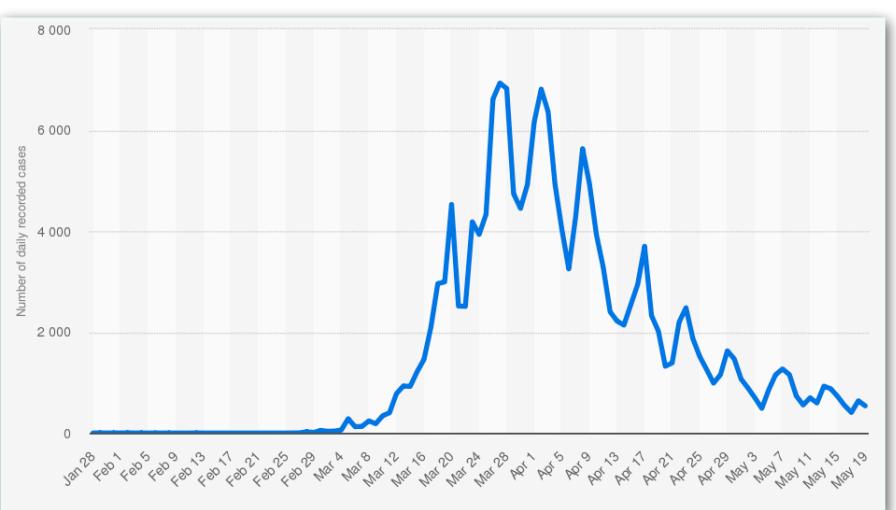
$$p(\theta) = p(\beta, \gamma)$$

Simulated data

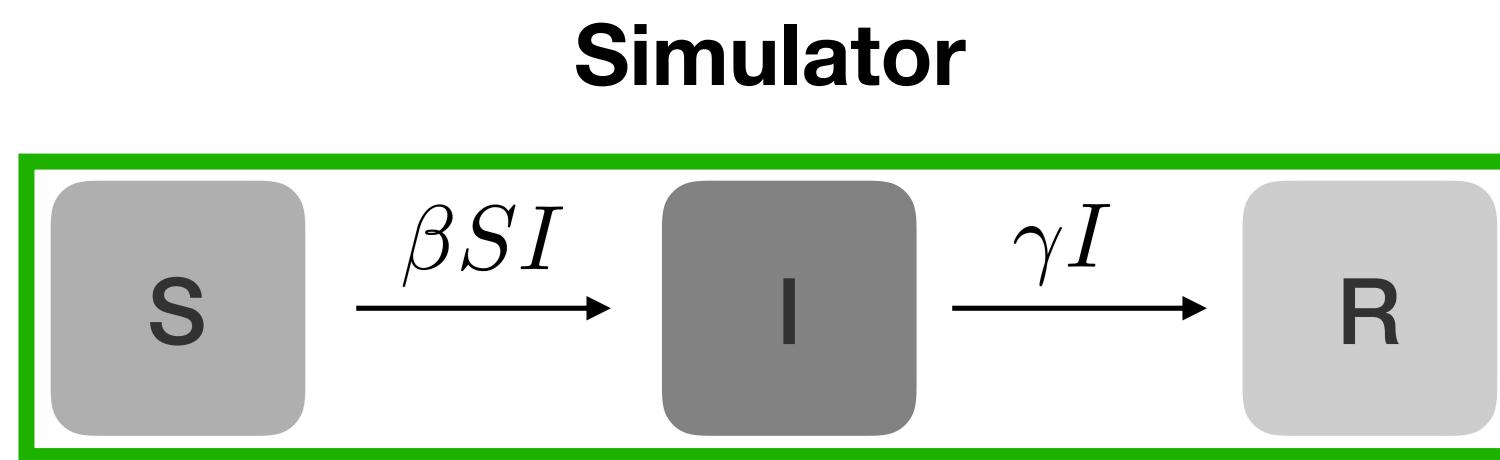
$$\theta_i \sim p(\theta)$$

$$x_i = \mathbf{SIR}(\theta_i)$$

Observed data x_o



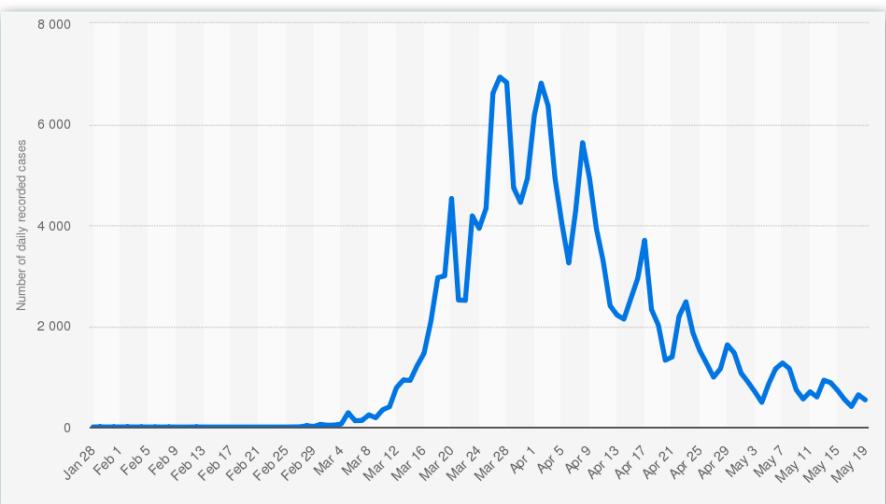
Neural Simulation-Based Inference



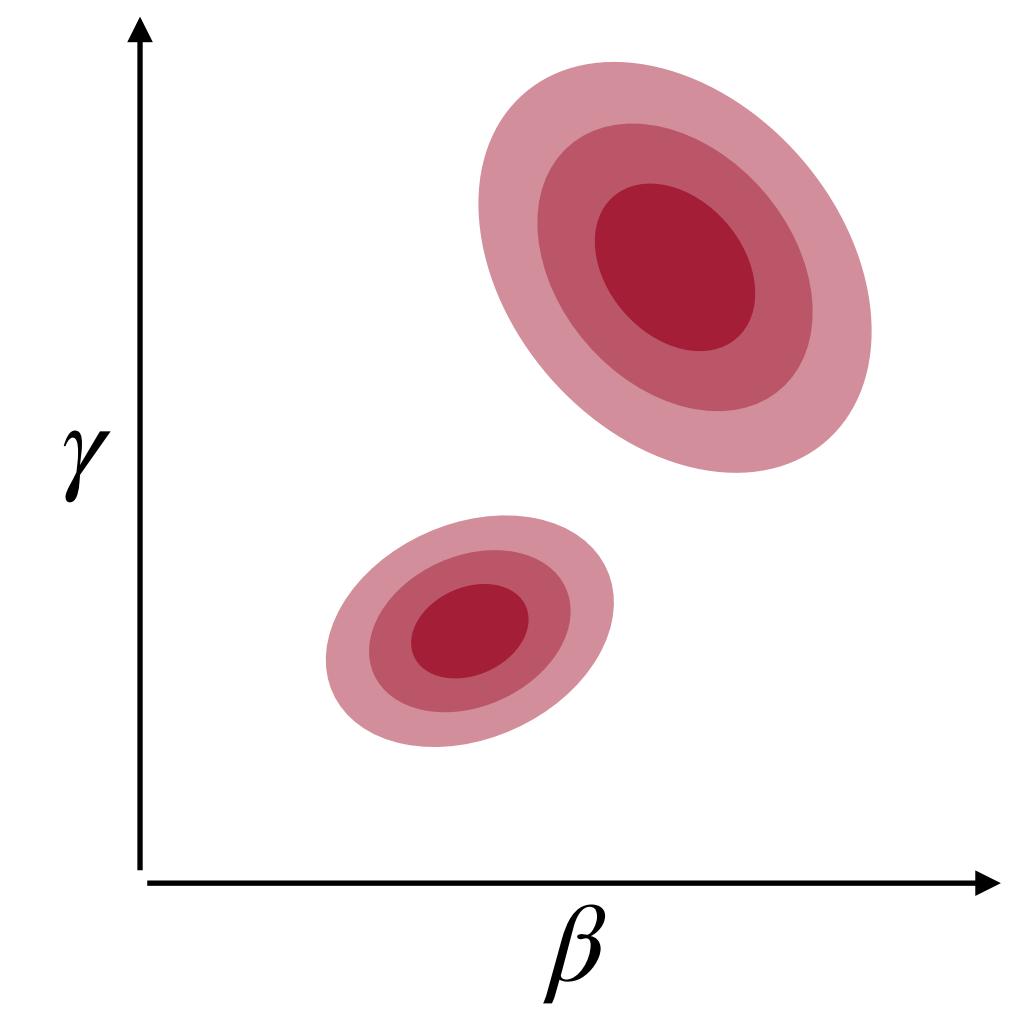
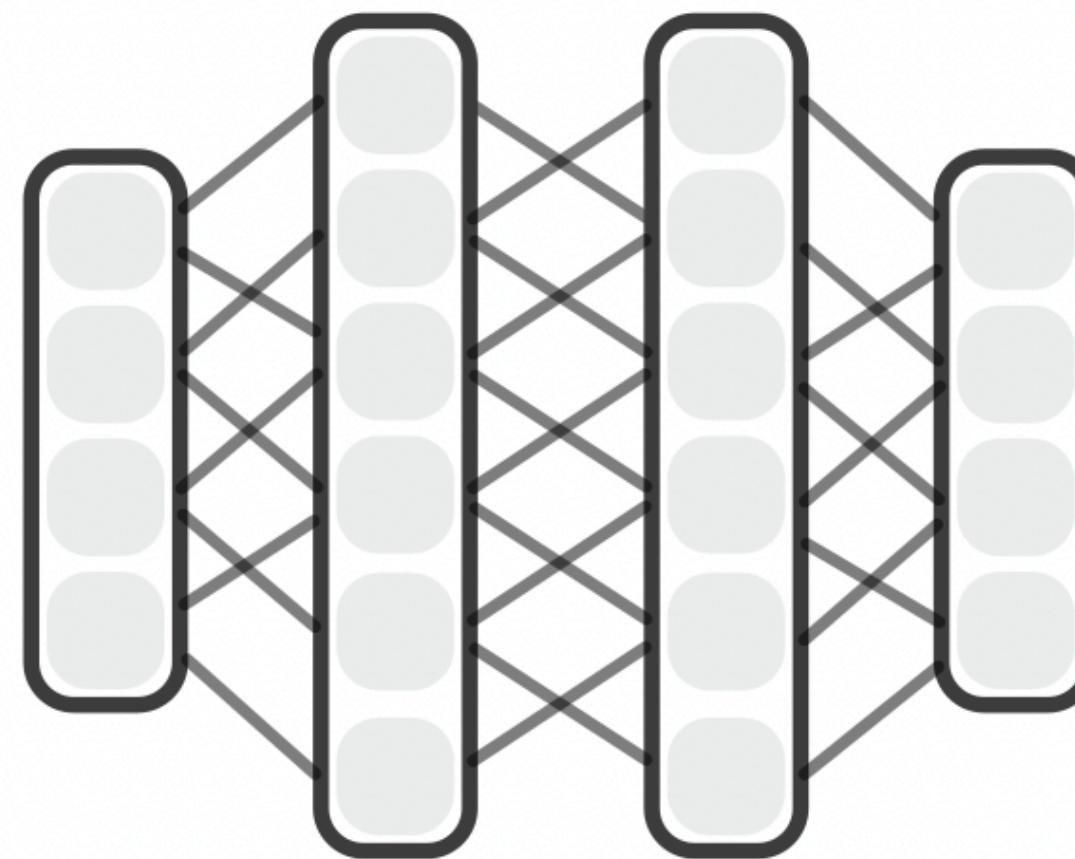
Prior
 $p(\theta) = p(\beta, \gamma)$

Simulated data
 $\theta_i \sim p(\theta)$
 $x_i = \mathbf{SIR}(\theta_i)$

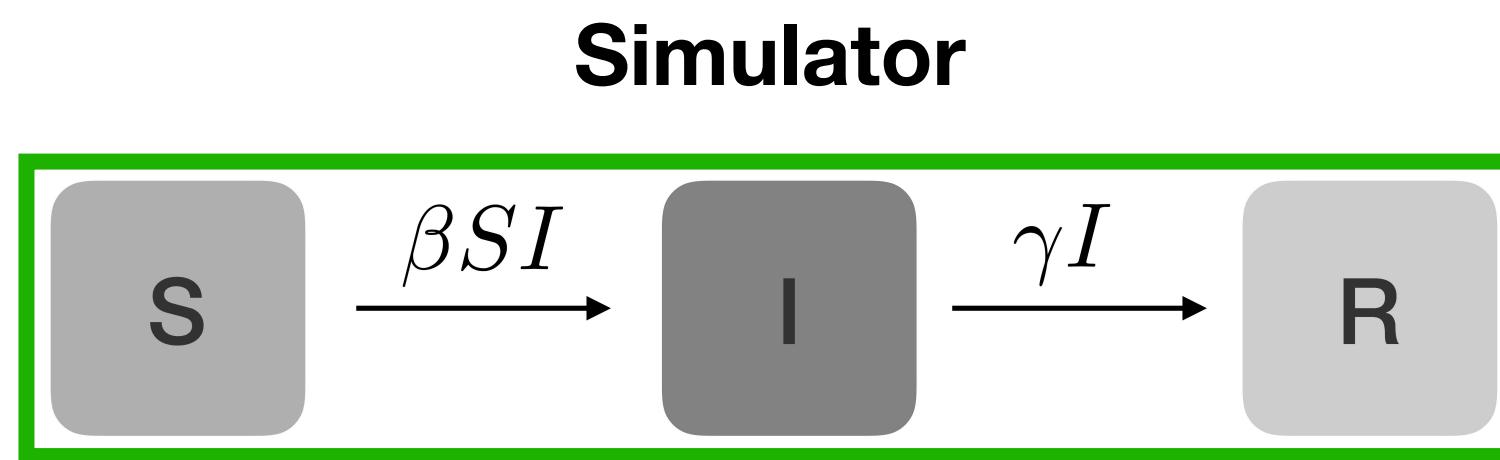
Observed data x_o



Conditional Density Estimation



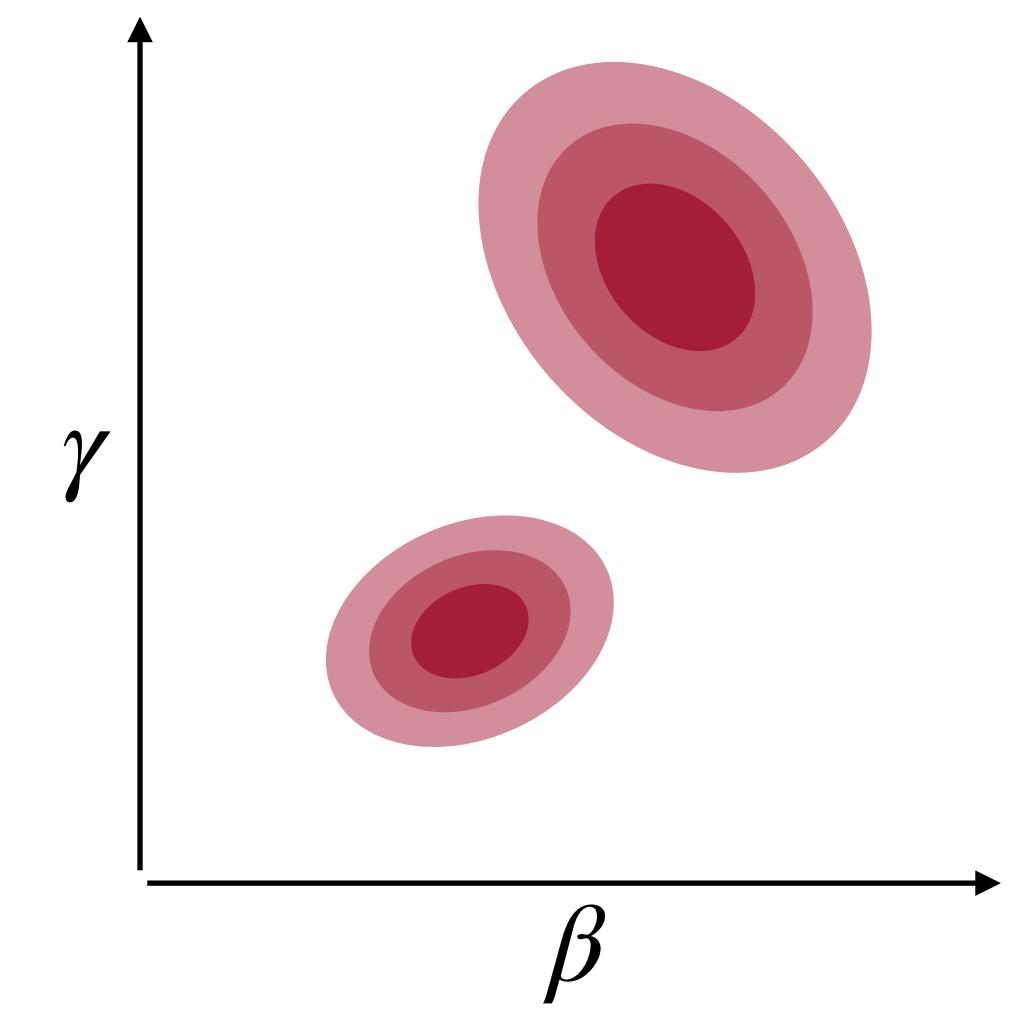
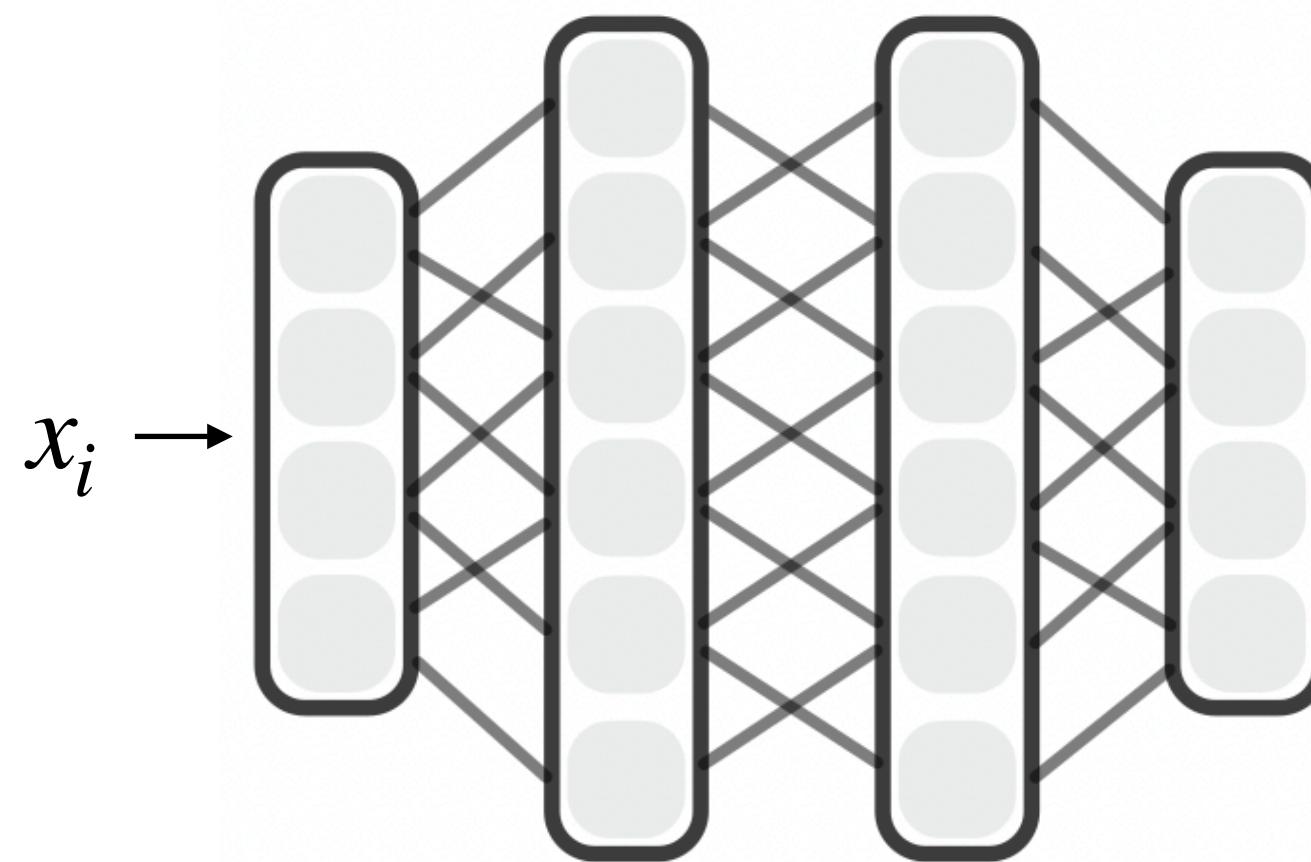
Neural Simulation-Based Inference



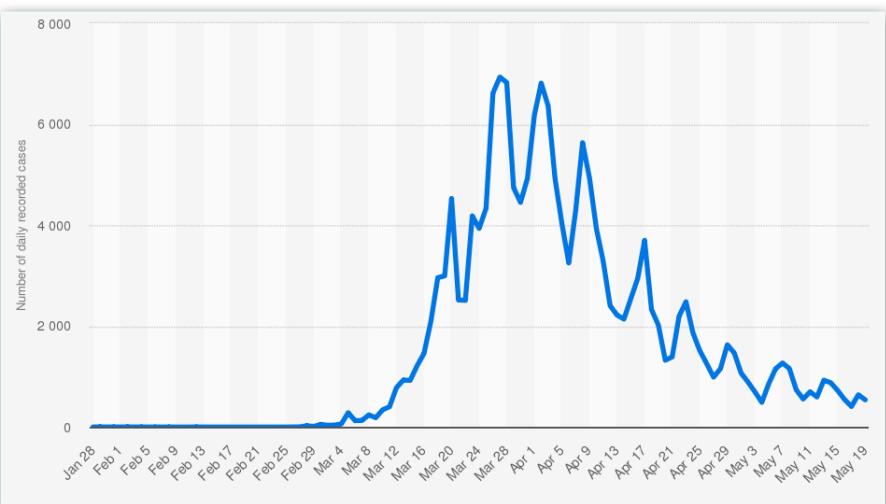
Prior
 $p(\theta) = p(\beta, \gamma)$

Simulated data
 $\theta_i \sim p(\theta)$
 $x_i = \mathbf{SIR}(\theta_i)$

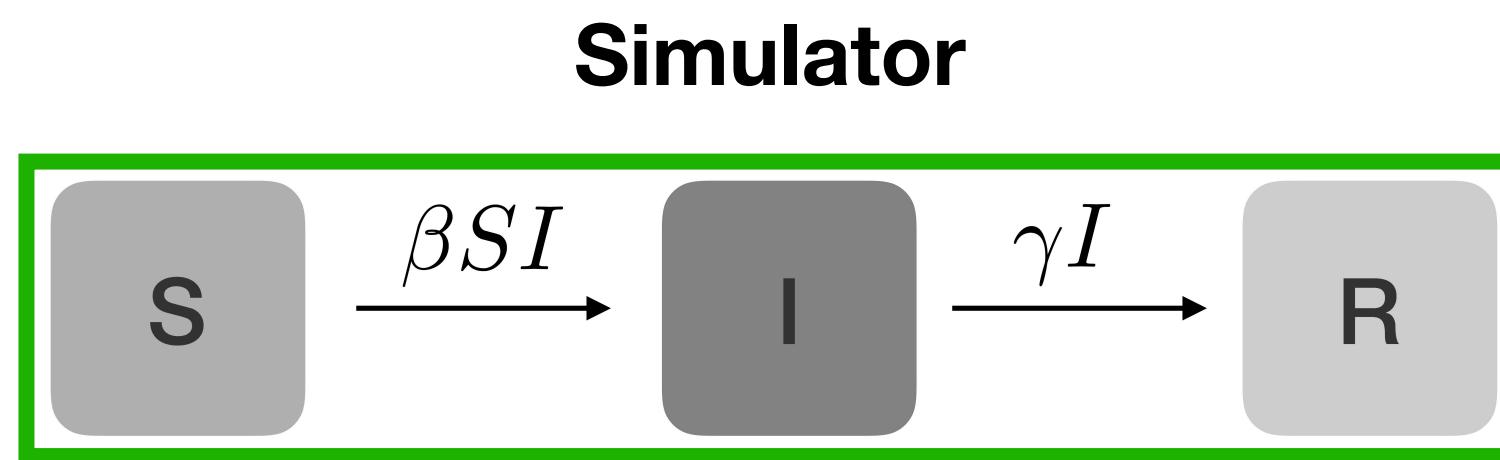
Conditional Density Estimation



Observed data x_o



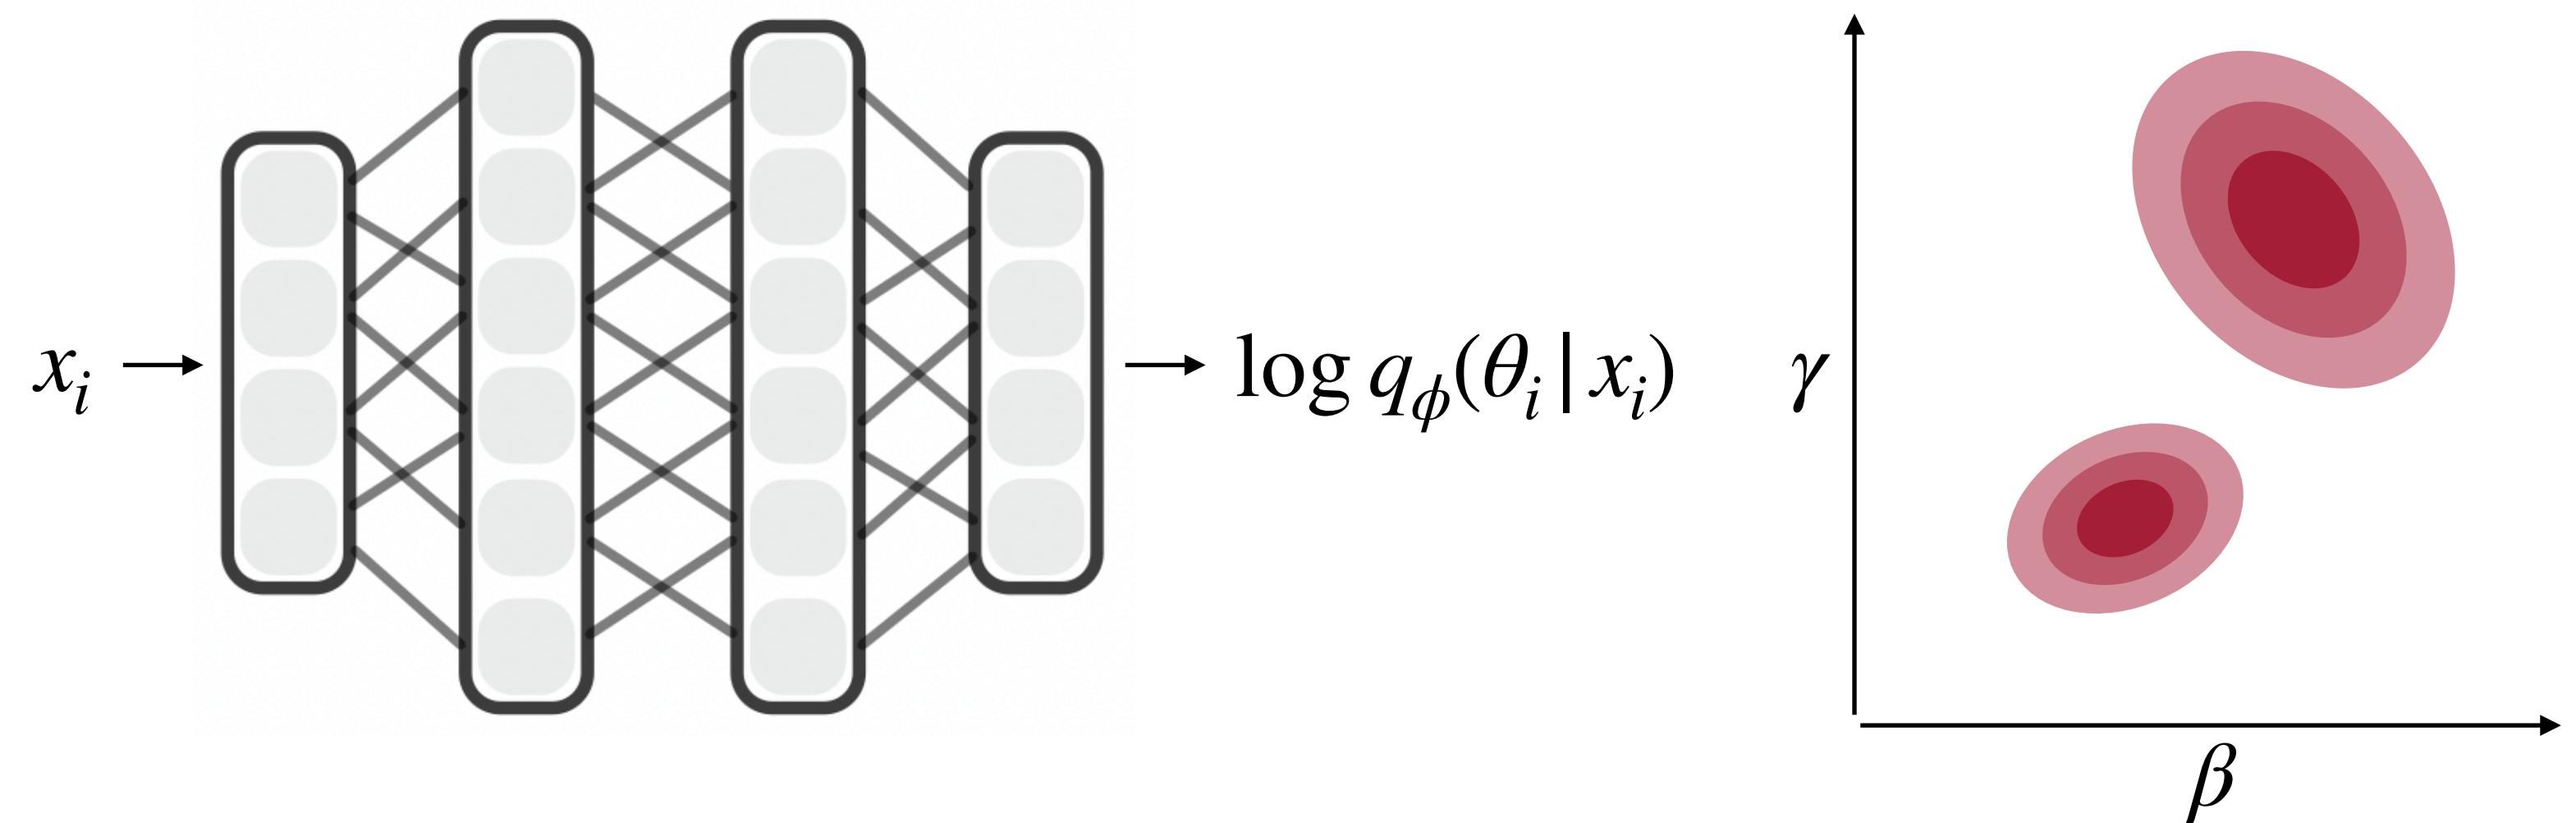
Neural Simulation-Based Inference



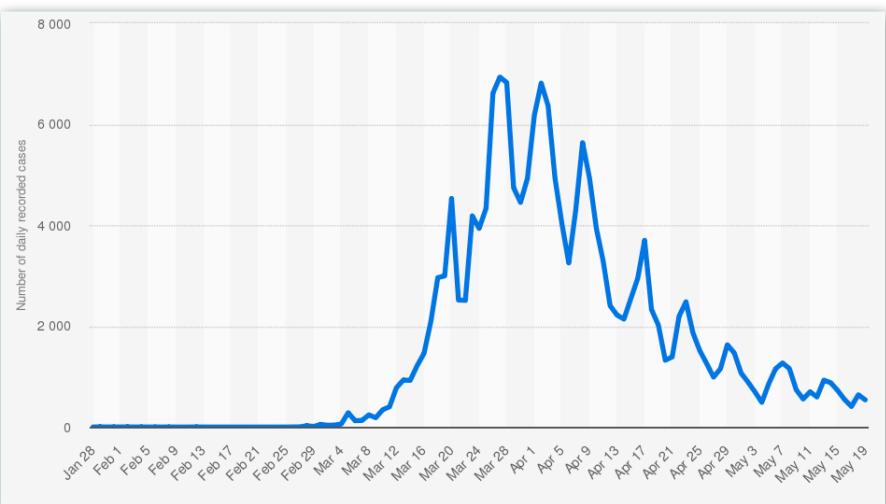
Prior
 $p(\theta) = p(\beta, \gamma)$

Simulated data
 $\theta_i \sim p(\theta)$
 $x_i = \text{SIR}(\theta_i)$

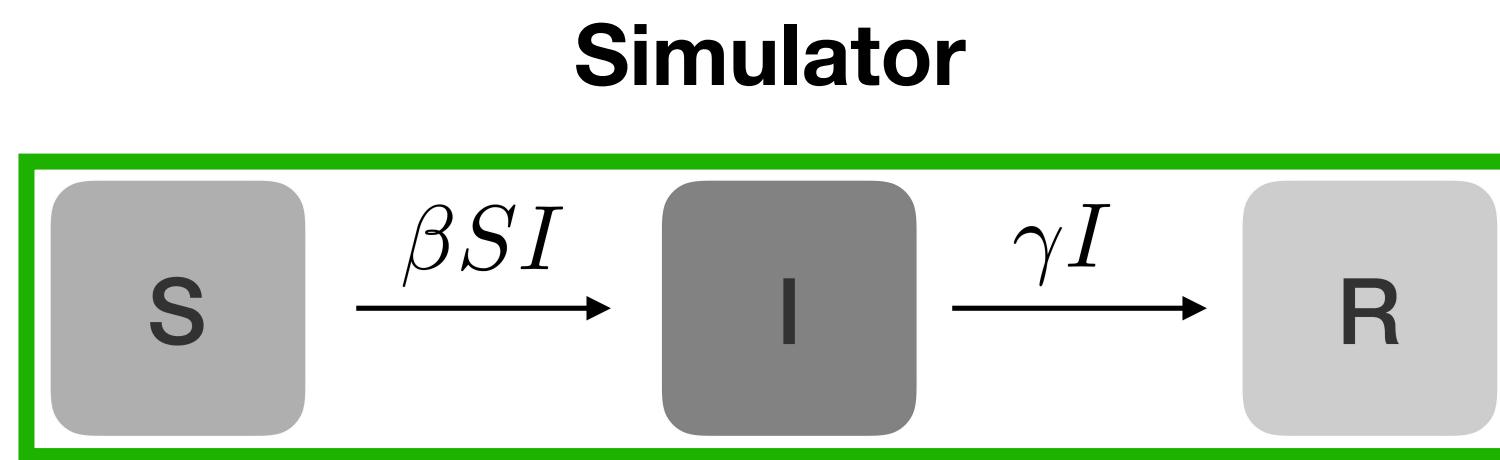
Conditional Density Estimation



Observed data x_o



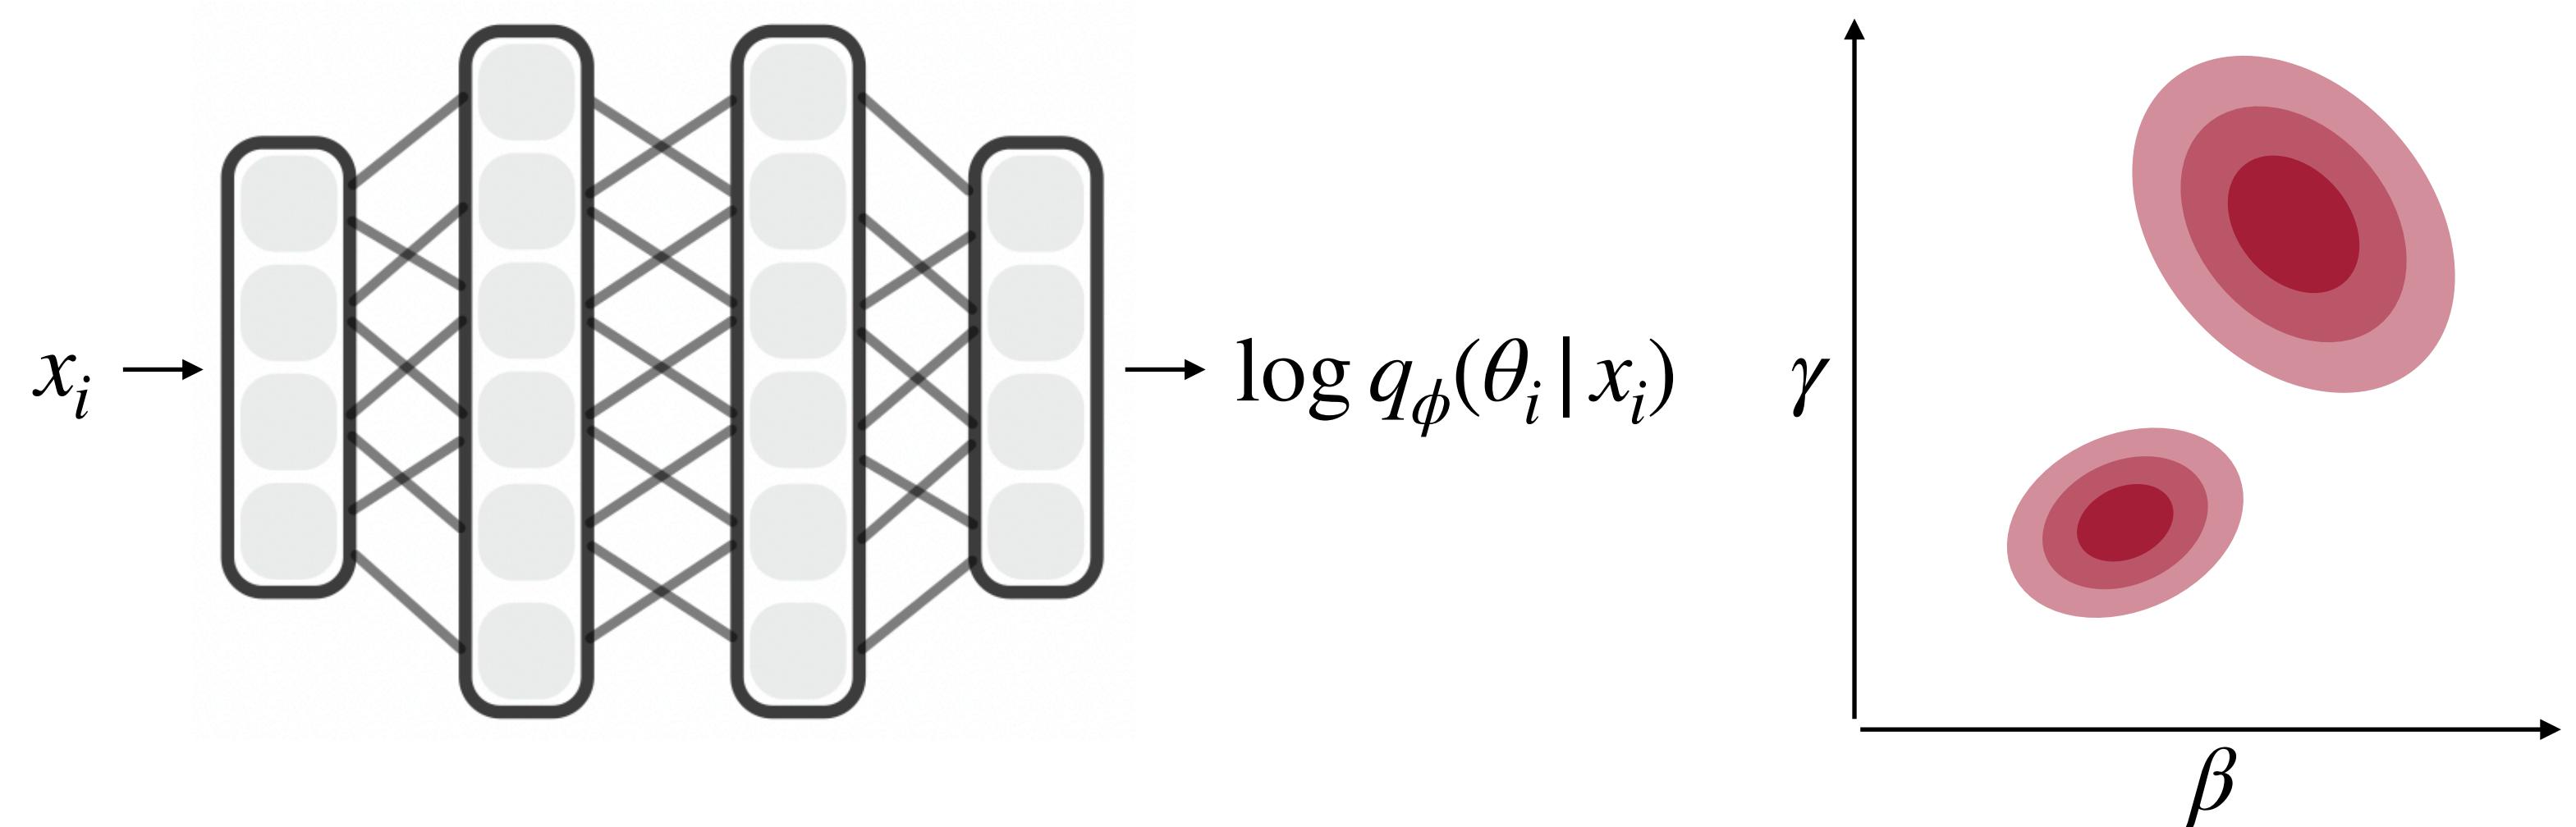
Neural Simulation-Based Inference



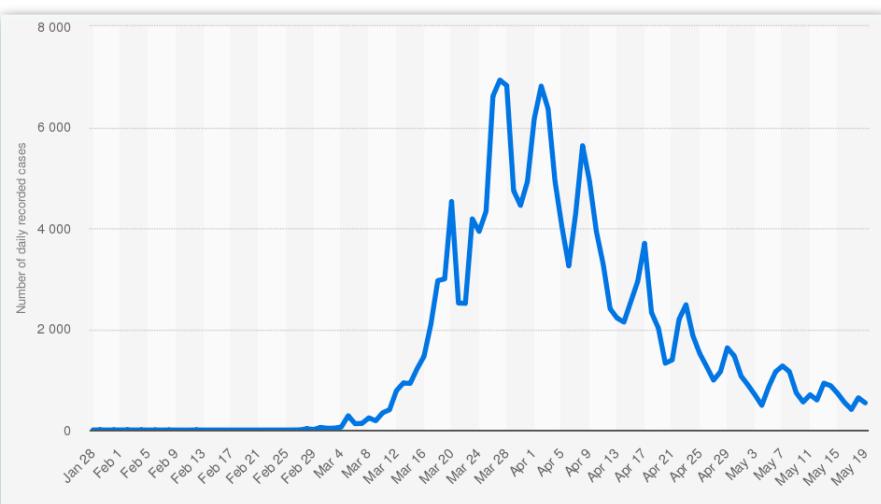
Prior
 $p(\theta) = p(\beta, \gamma)$

Simulated data
 $\theta_i \sim p(\theta)$
 $x_i = \text{SIR}(\theta_i)$

Conditional Density Estimation

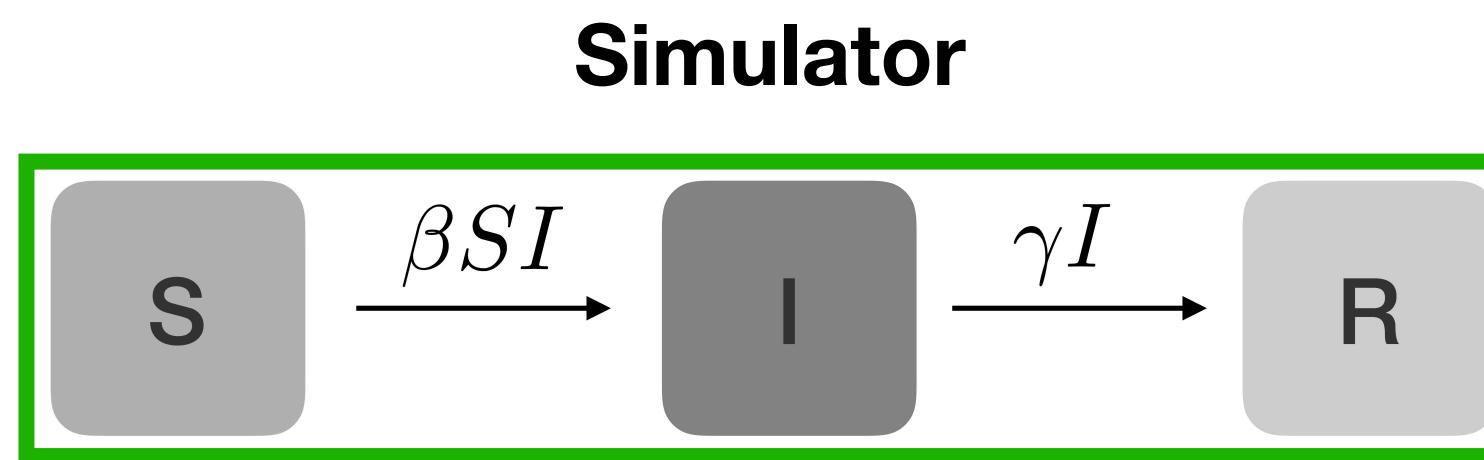


Observed data x_o



1. Train with simulated data

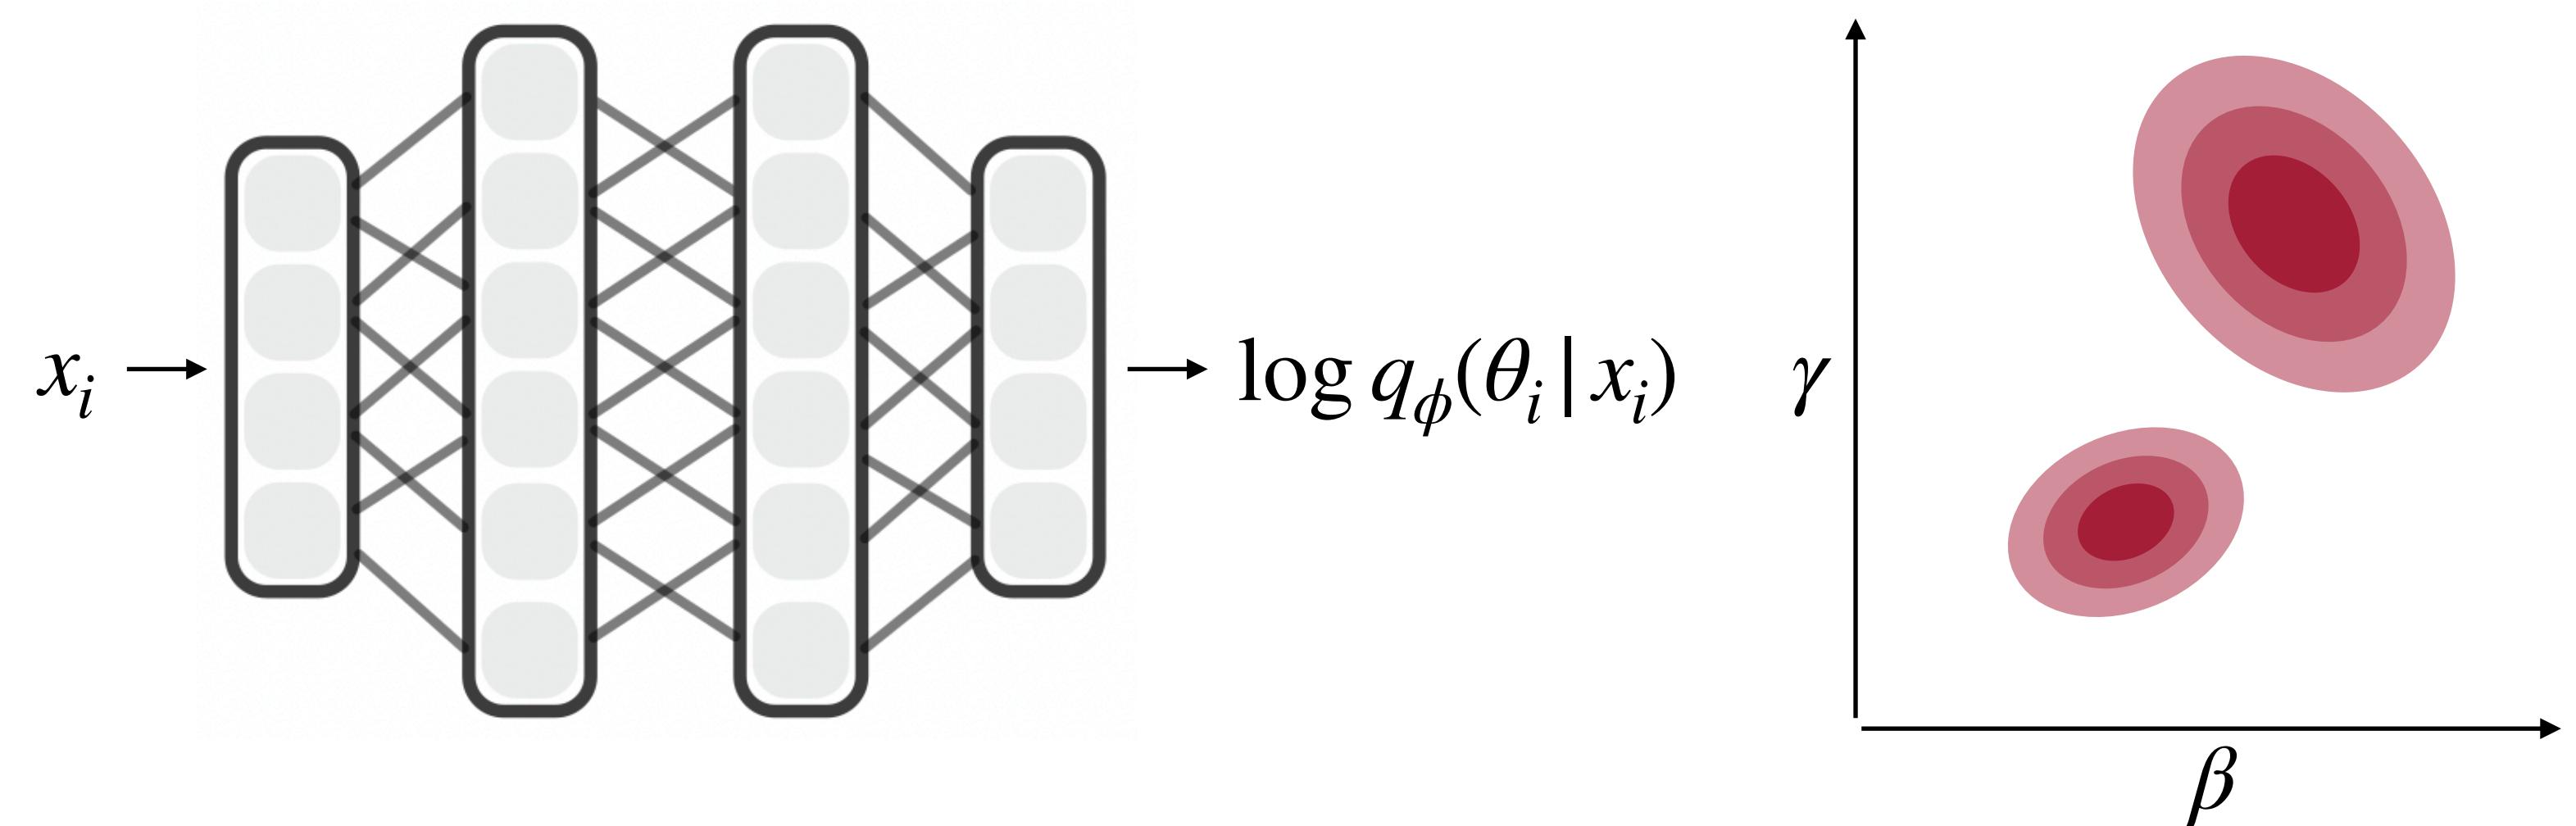
Neural Simulation-Based Inference



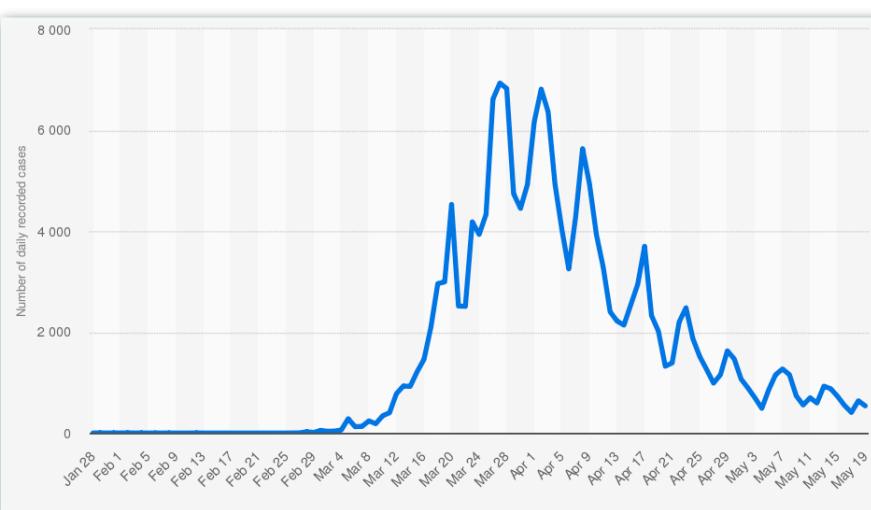
Prior
 $p(\theta) = p(\beta, \gamma)$

Simulated data
 $\theta_i \sim p(\theta)$
 $x_i = \text{SIR}(\theta_i)$

Conditional Density Estimation



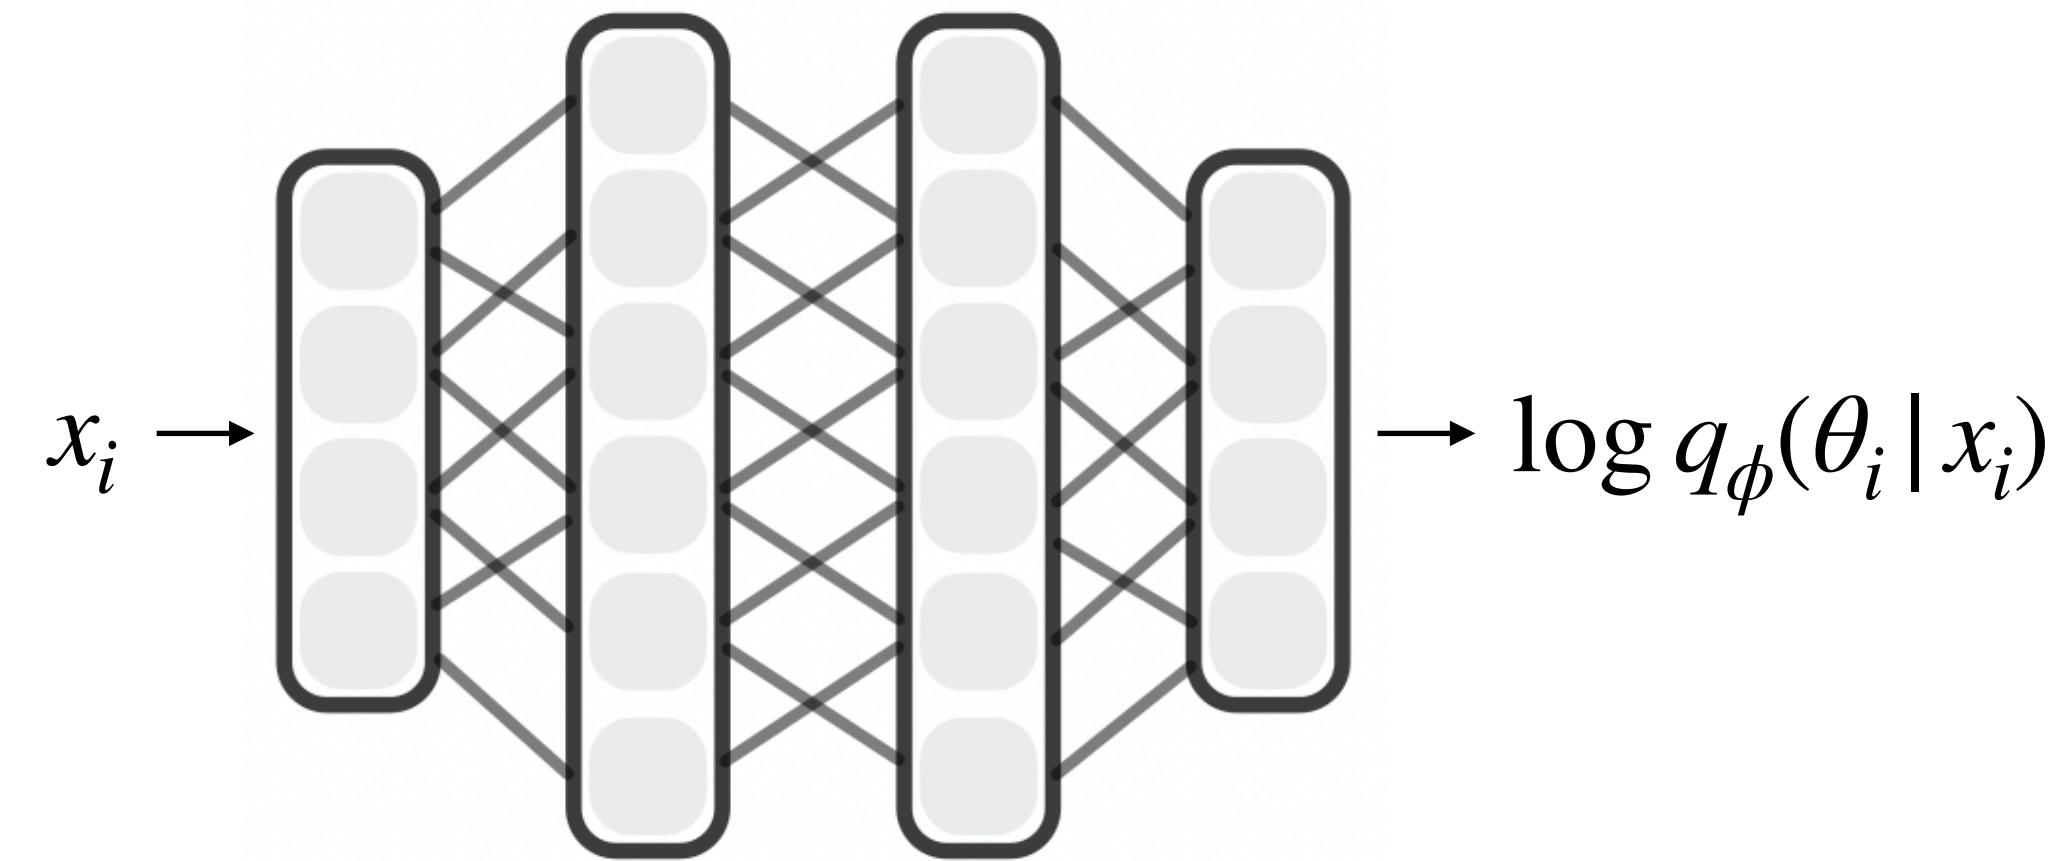
Observed data x_o



1. Train with simulated data
2. Infer given observed data

Conditional density estimation in a nutshell

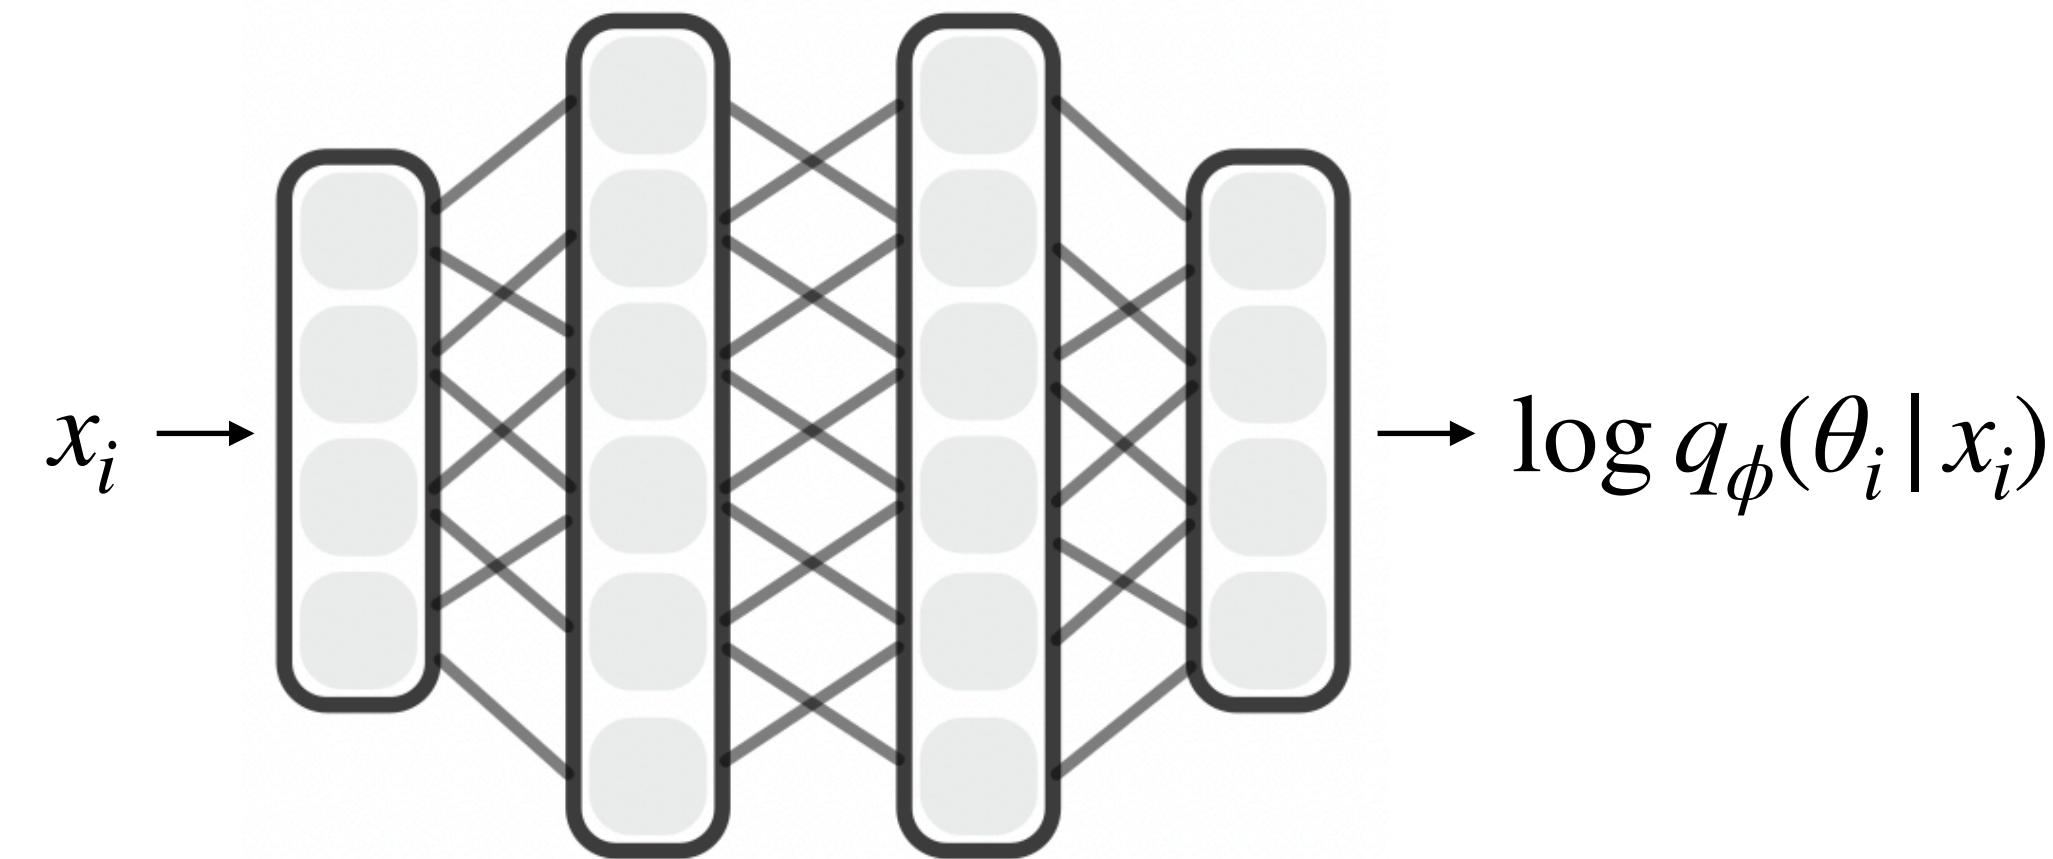
Conditional Density Estimation



Conditional density estimation in a nutshell

Conditional Density Estimation

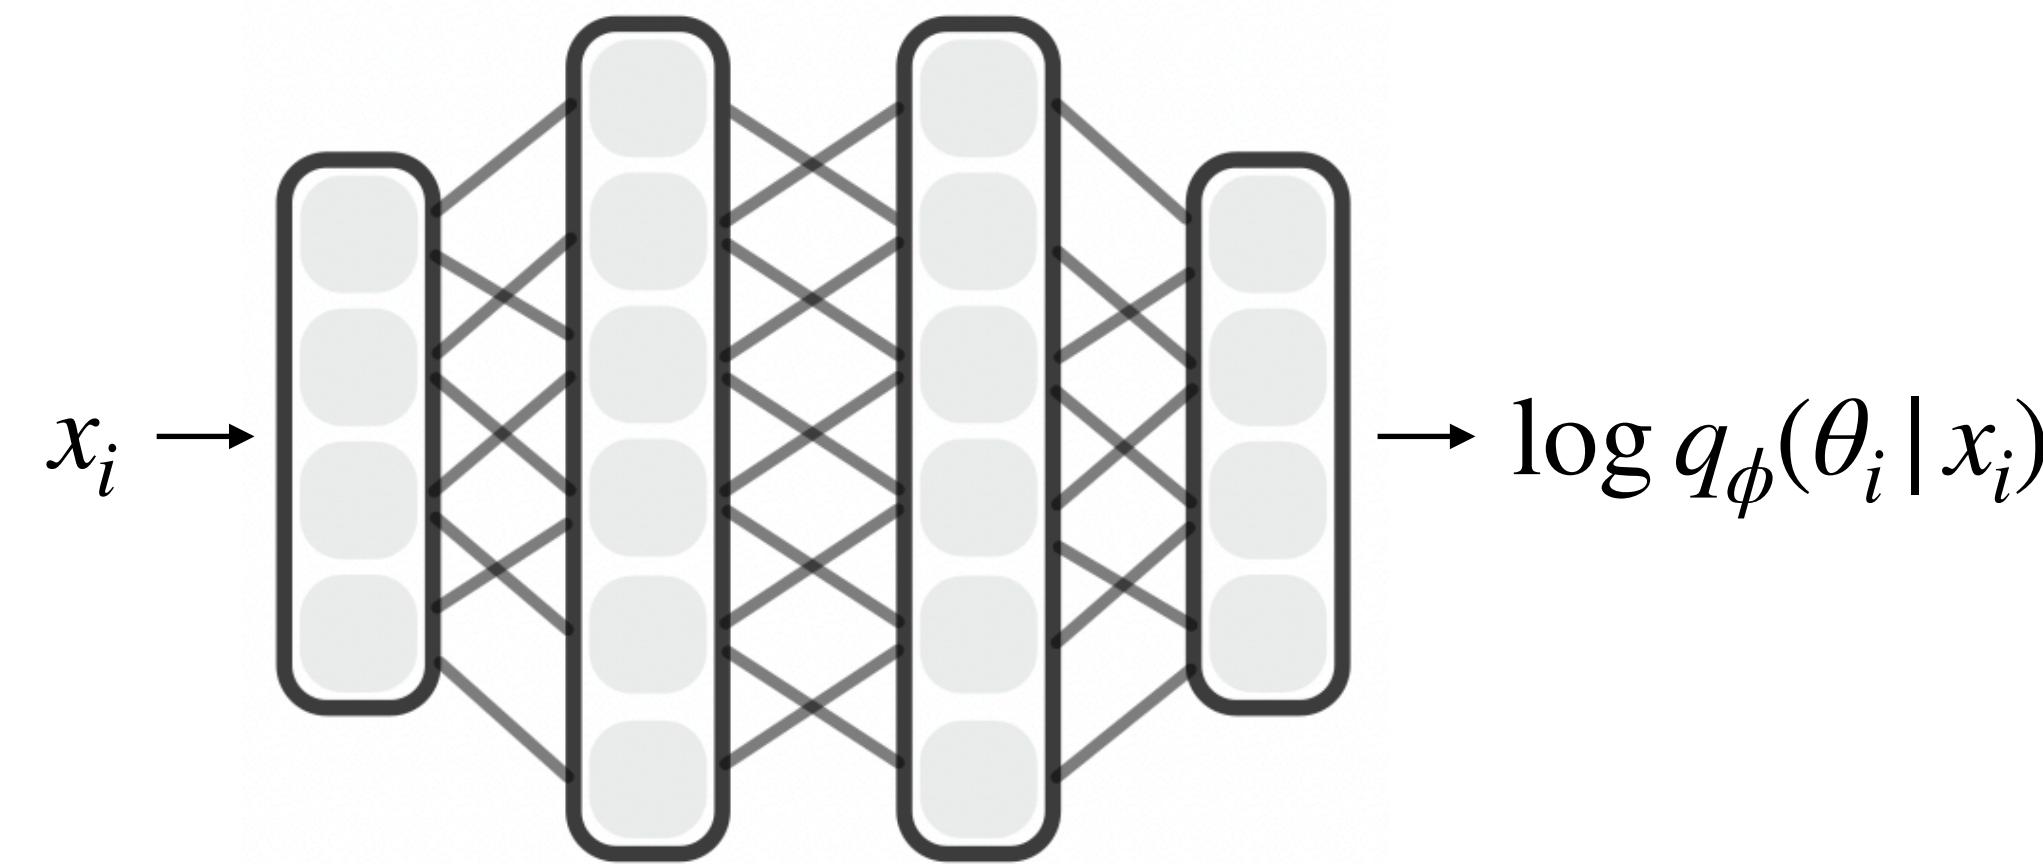
- Assume a parametric model for the posterior, e.g., a Gaussian



Conditional density estimation in a nutshell

Conditional Density Estimation

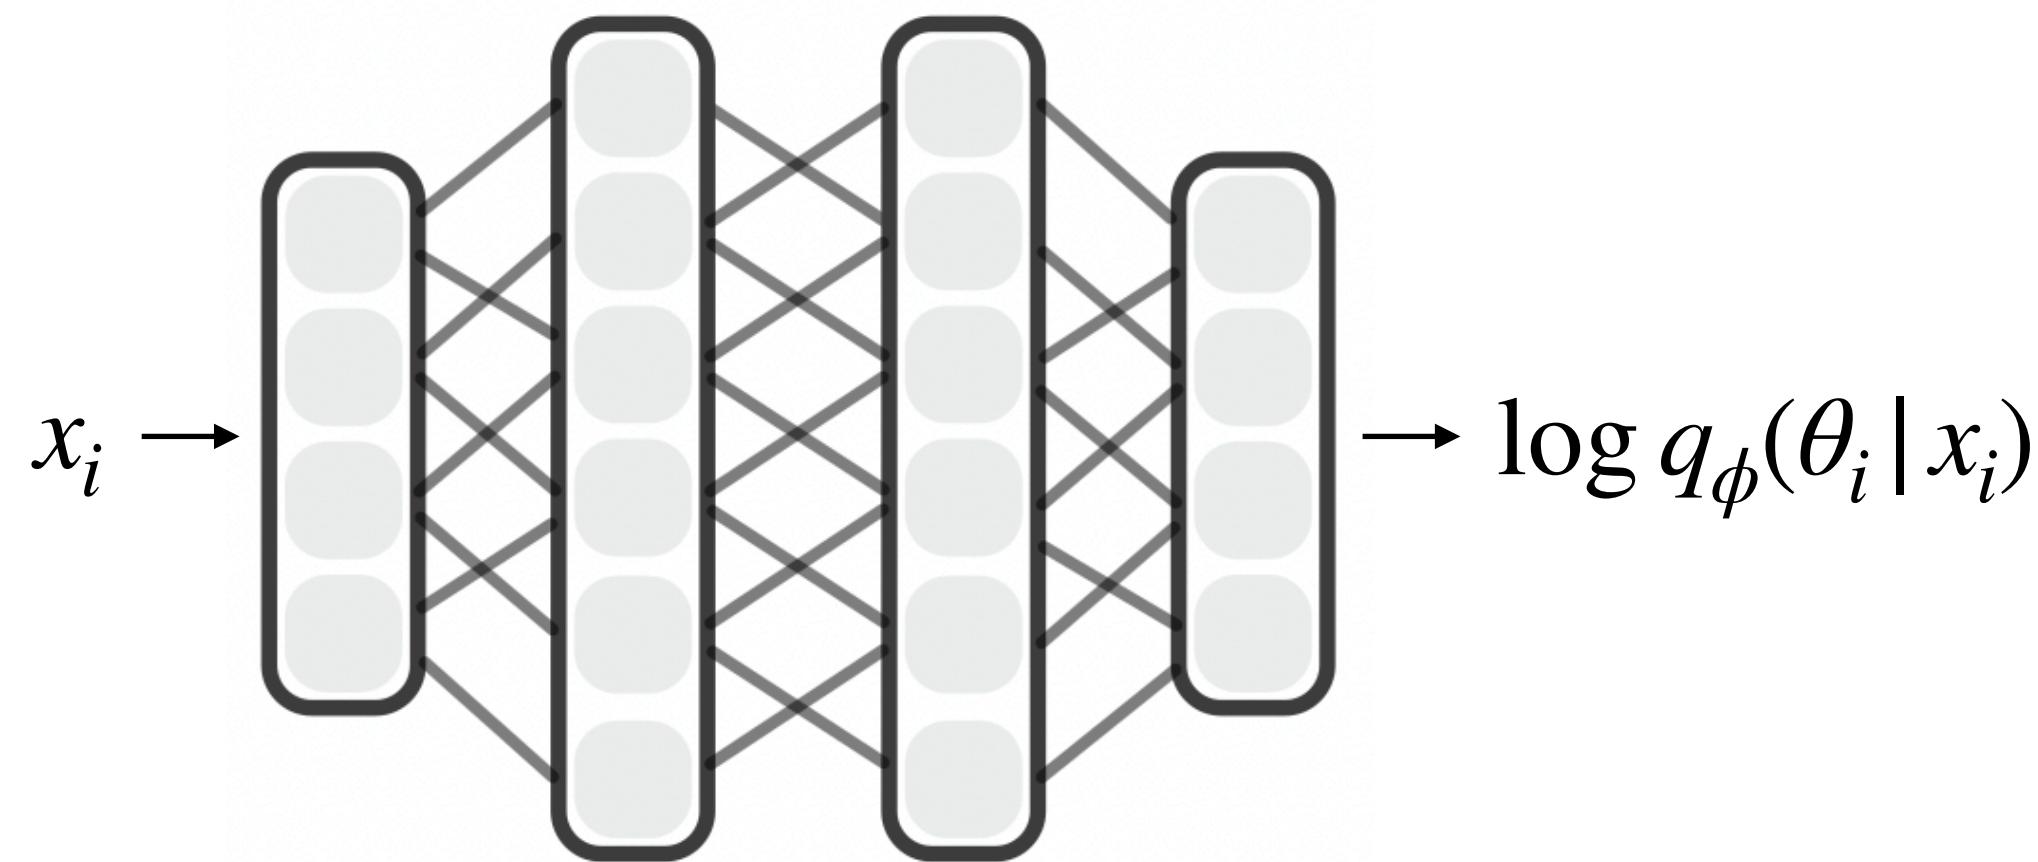
- Assume a parametric model for the posterior, e.g., a Gaussian
- NN predicts Gaussian mean and variance



Conditional density estimation in a nutshell

Conditional Density Estimation

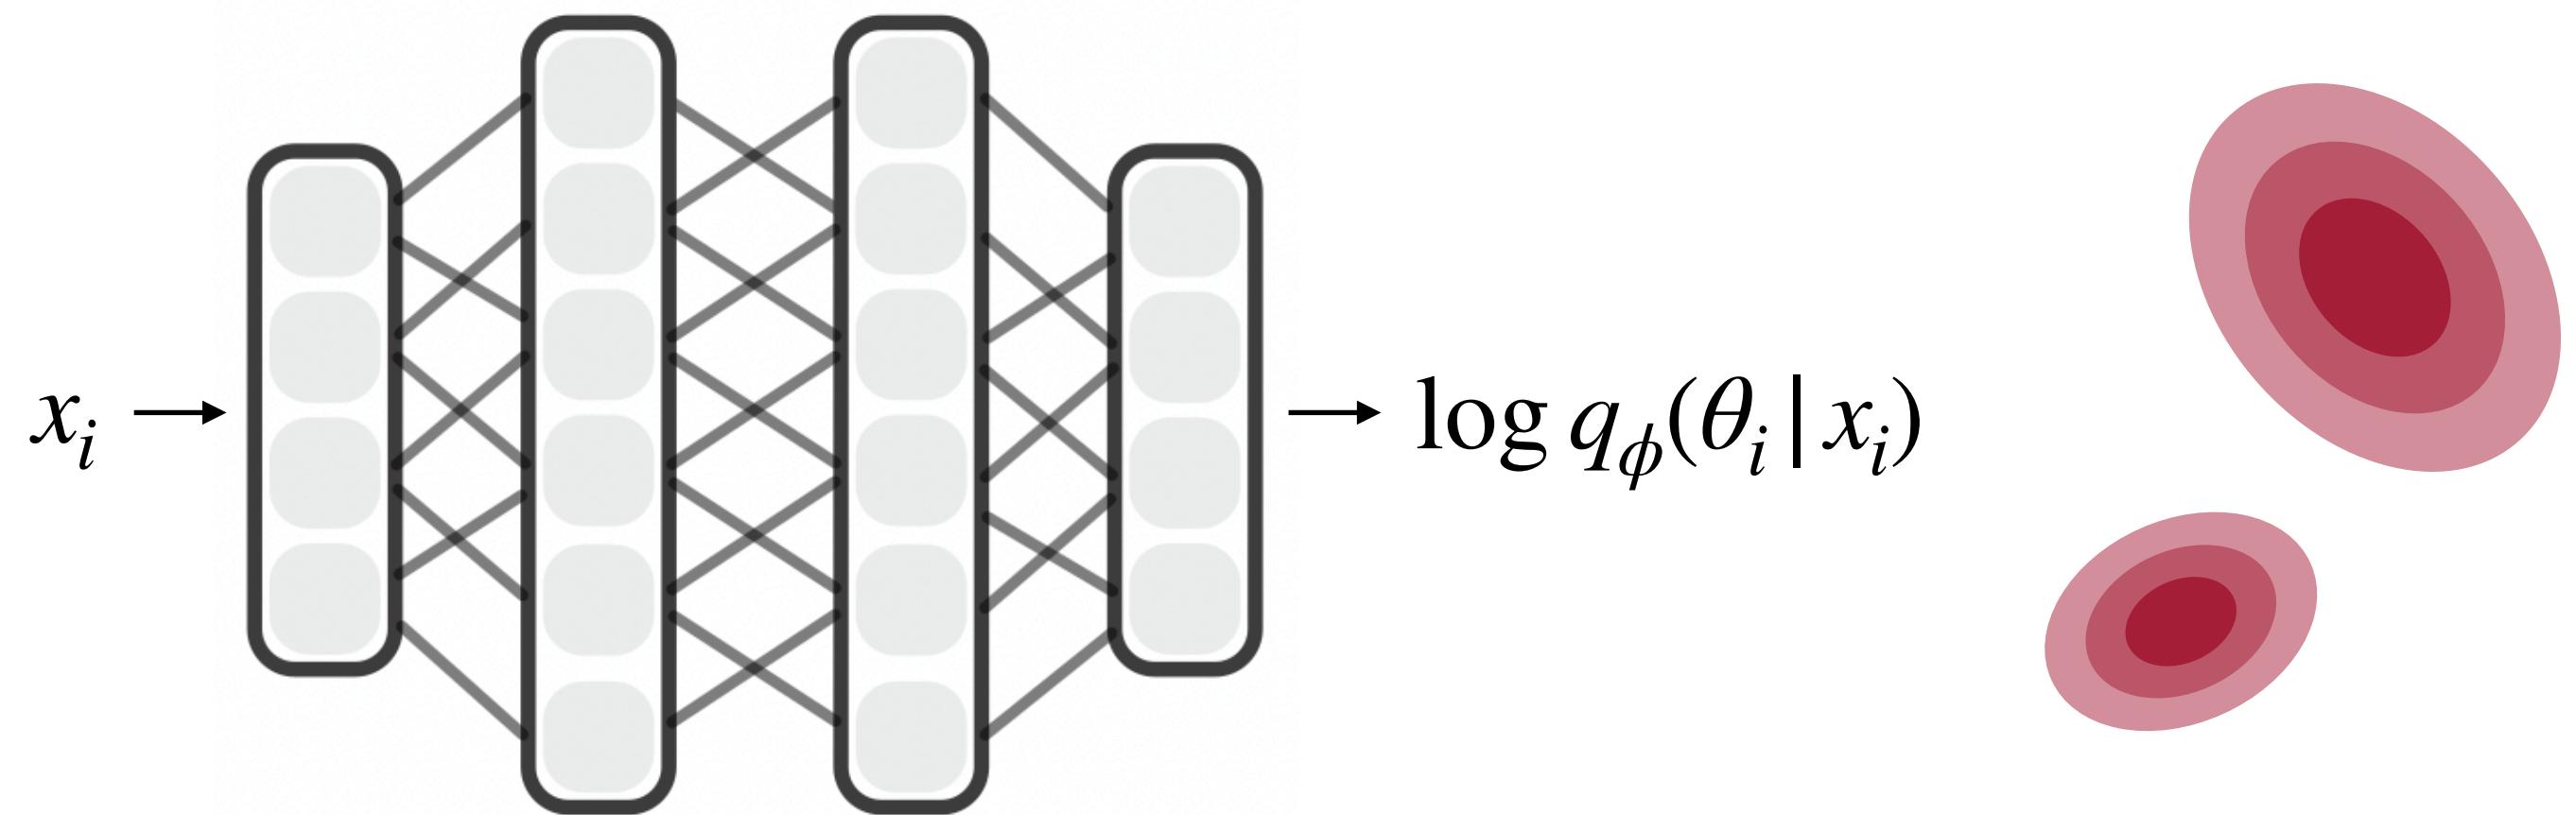
- Assume a parametric model for the posterior, e.g., a Gaussian
- NN predicts Gaussian mean and variance
- Loss signal: **model probability**



Conditional density estimation in a nutshell

Conditional Density Estimation

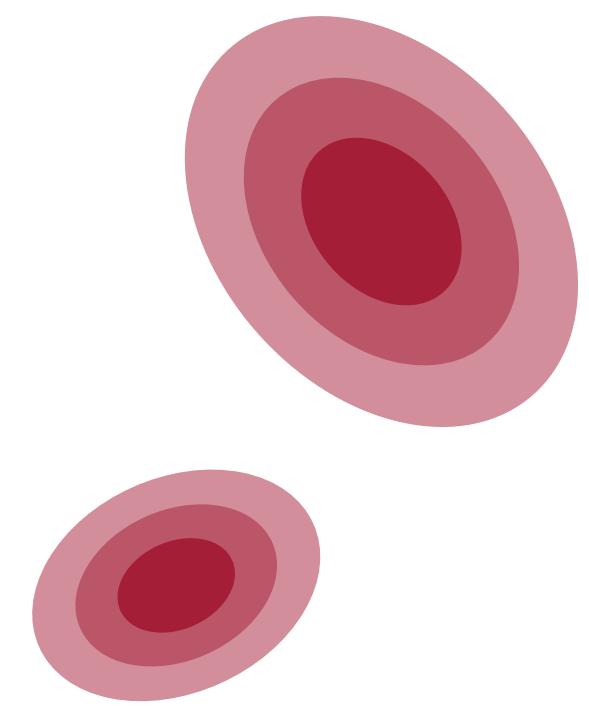
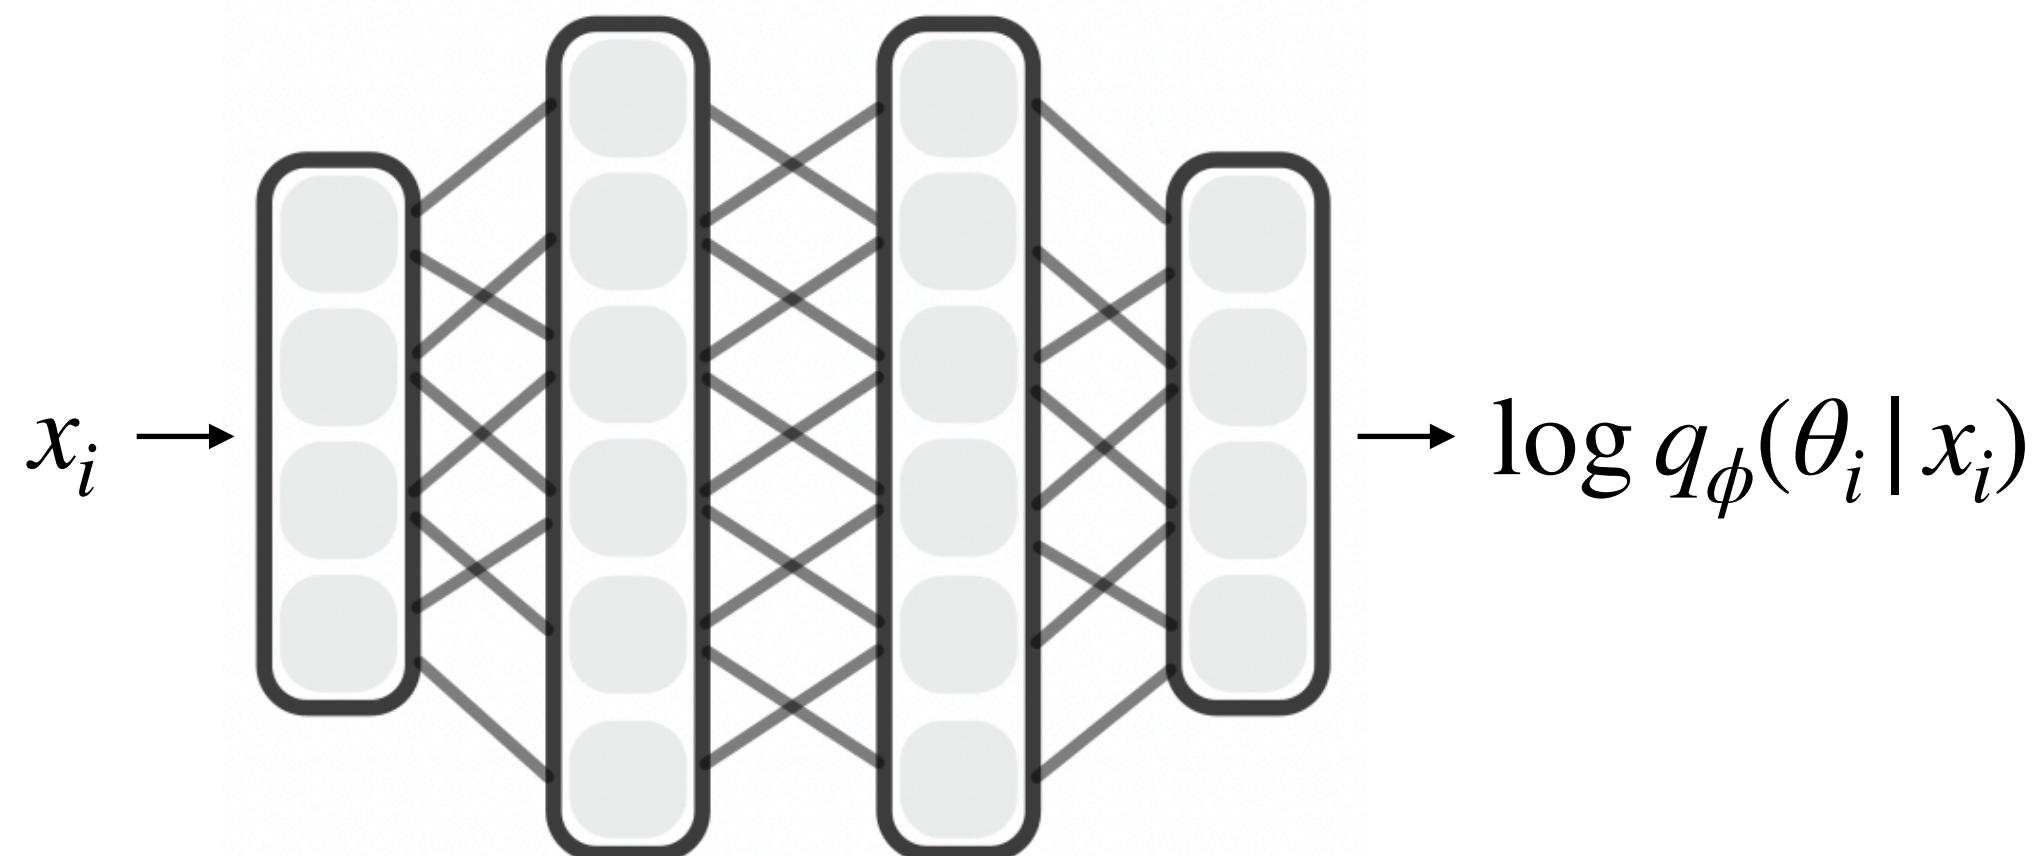
- Assume a parametric model for the posterior, e.g., a Gaussian
- NN predicts Gaussian mean and variance
- Loss signal: **model probability**



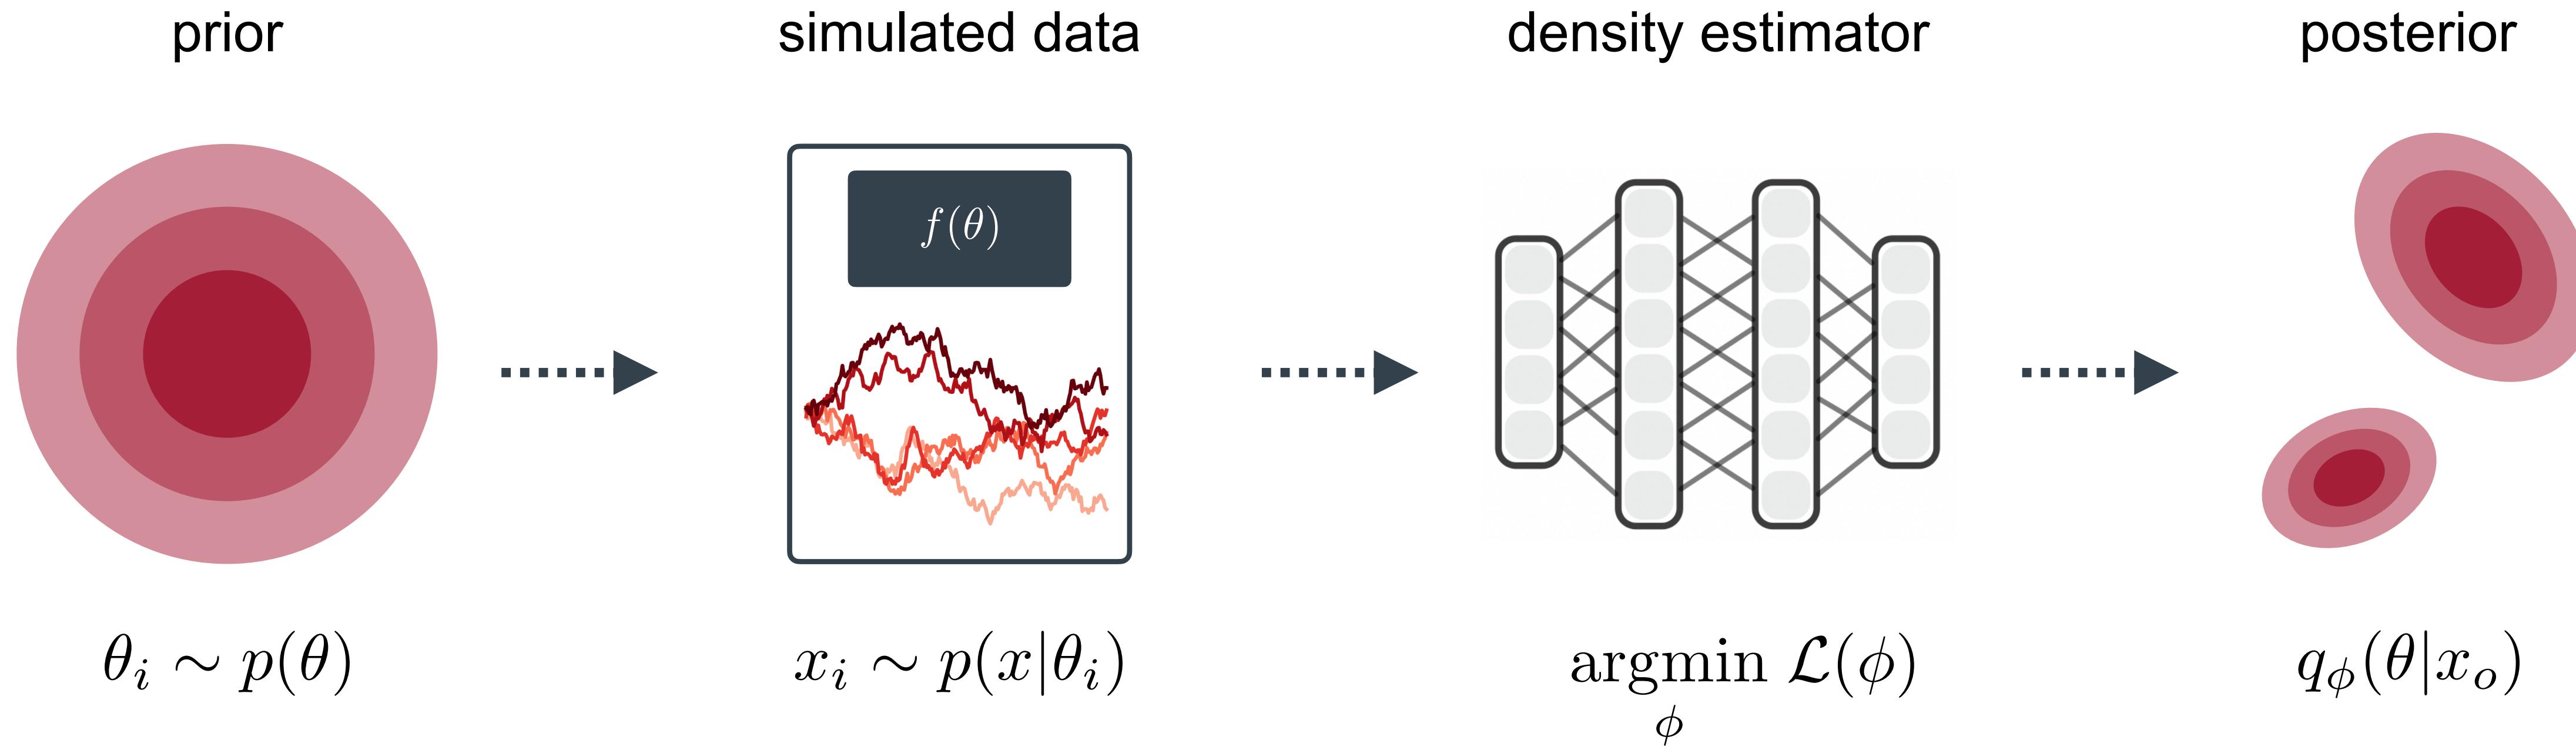
Conditional density estimation in a nutshell

Conditional Density Estimation

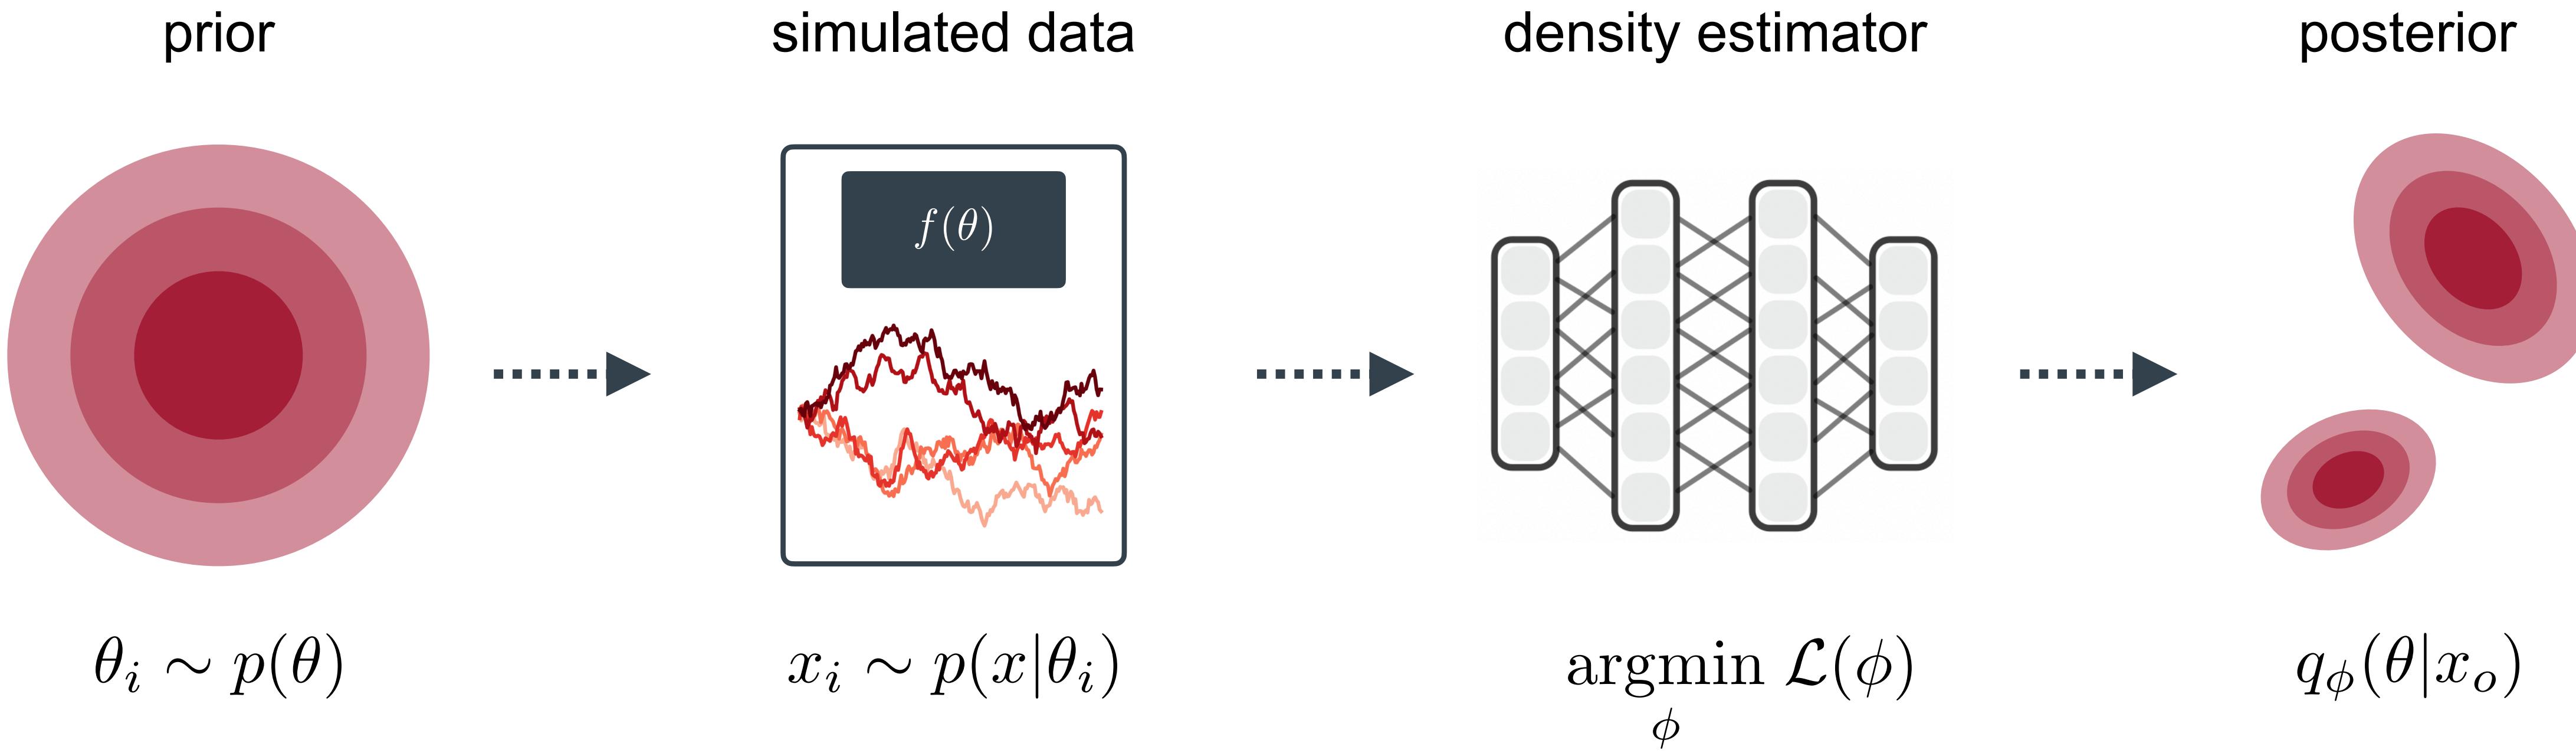
- Assume a parametric model for the posterior, e.g., a Gaussian
- NN predicts Gaussian mean and variance
- Loss signal: **model probability**
- Used in practice:
 - Mixture of Gaussians
 - Normalizing Flows
 - (Flow Matching, Score Matching)



Three main neural SBI approaches

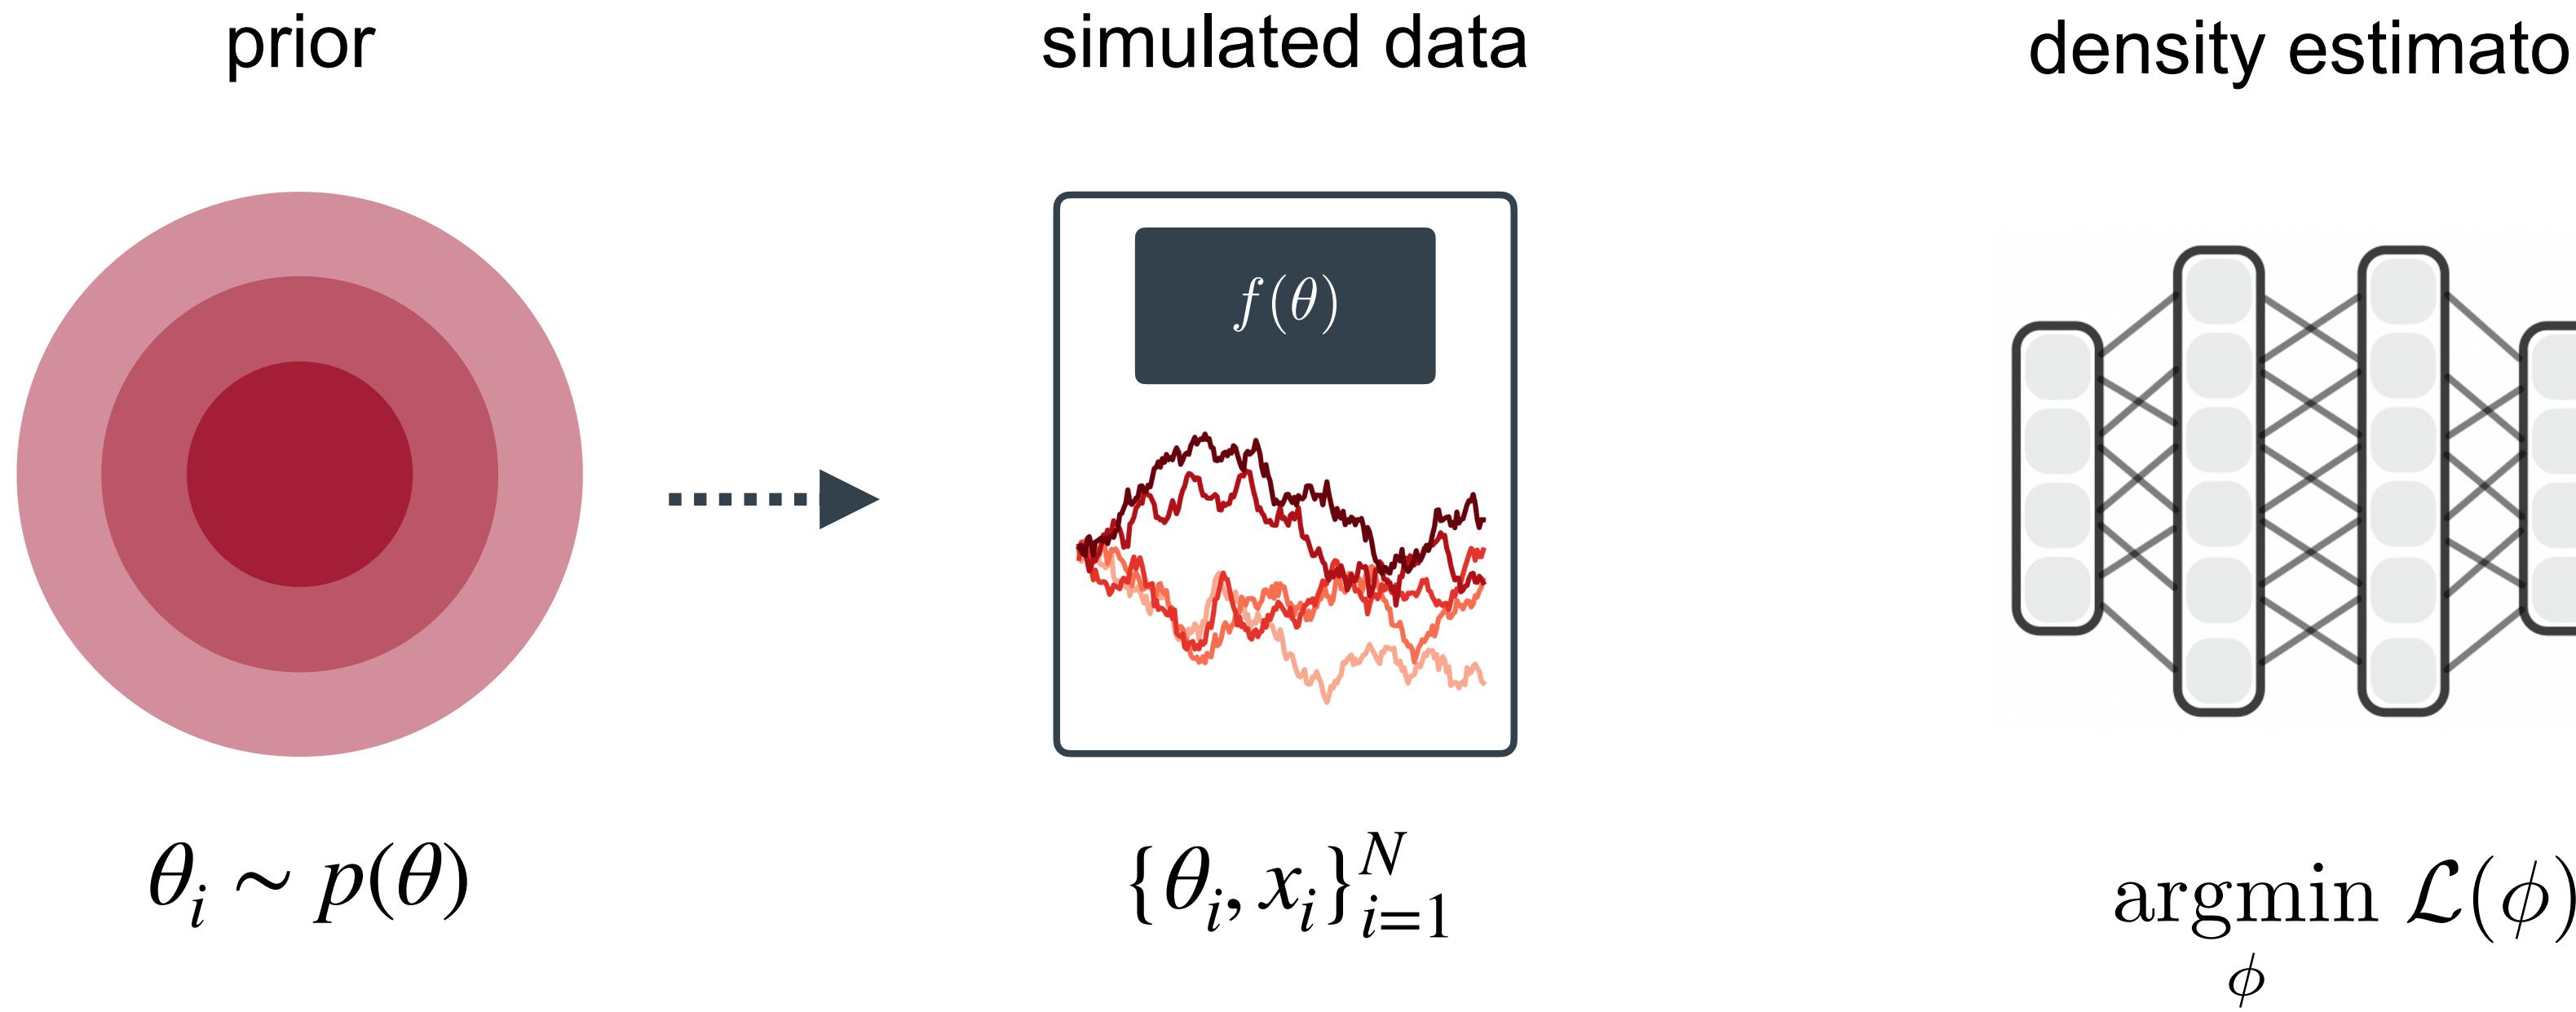


Three main neural SBI approaches



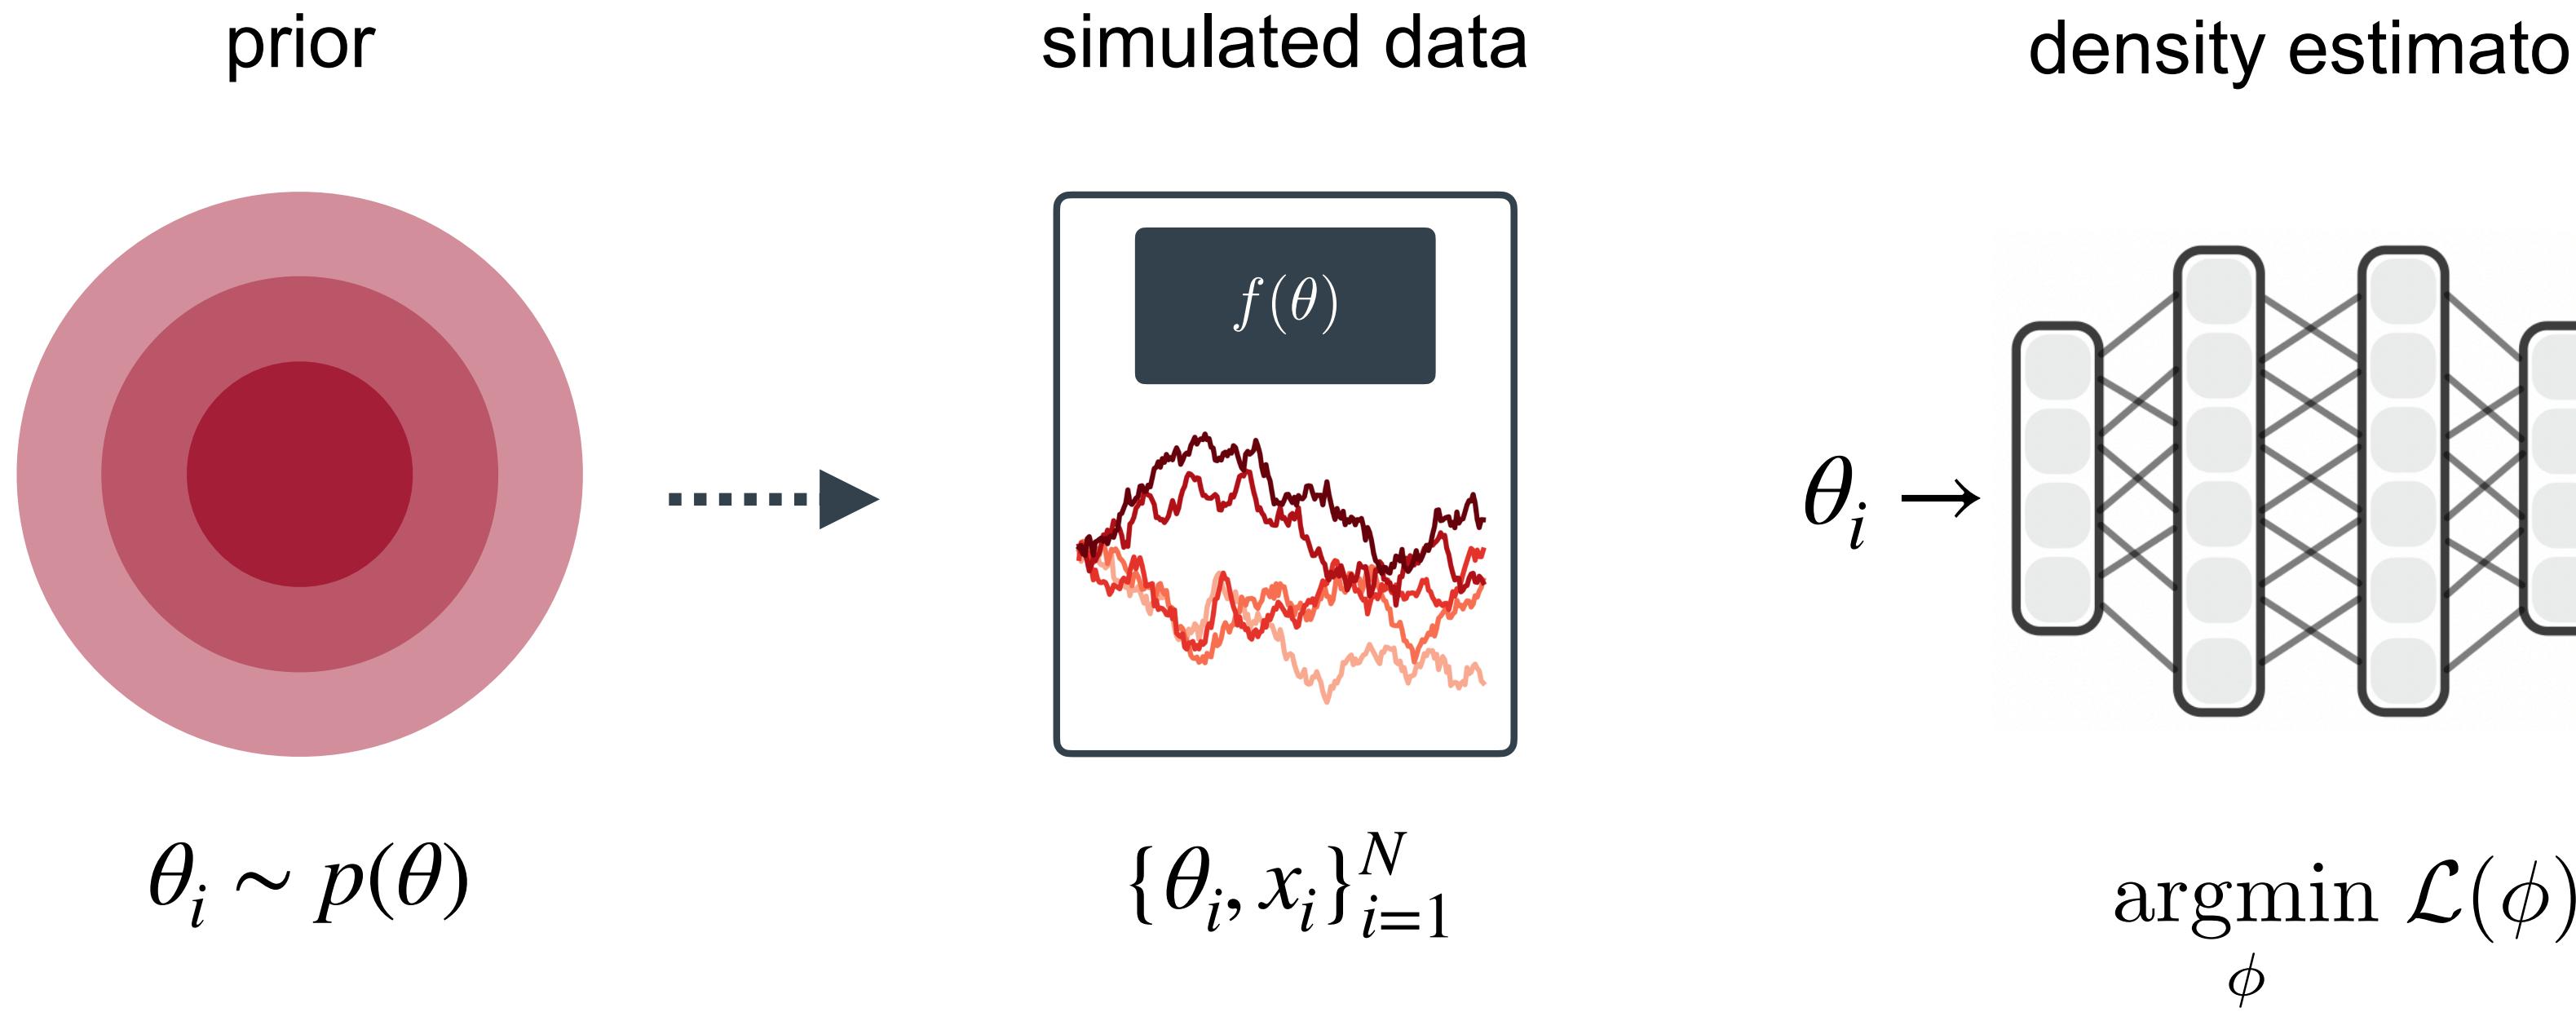
1. **Neural Posterior Estimation (NPE):** directly learn the posterior $q_{\phi}(\theta|x)$

Three main neural SBI approaches



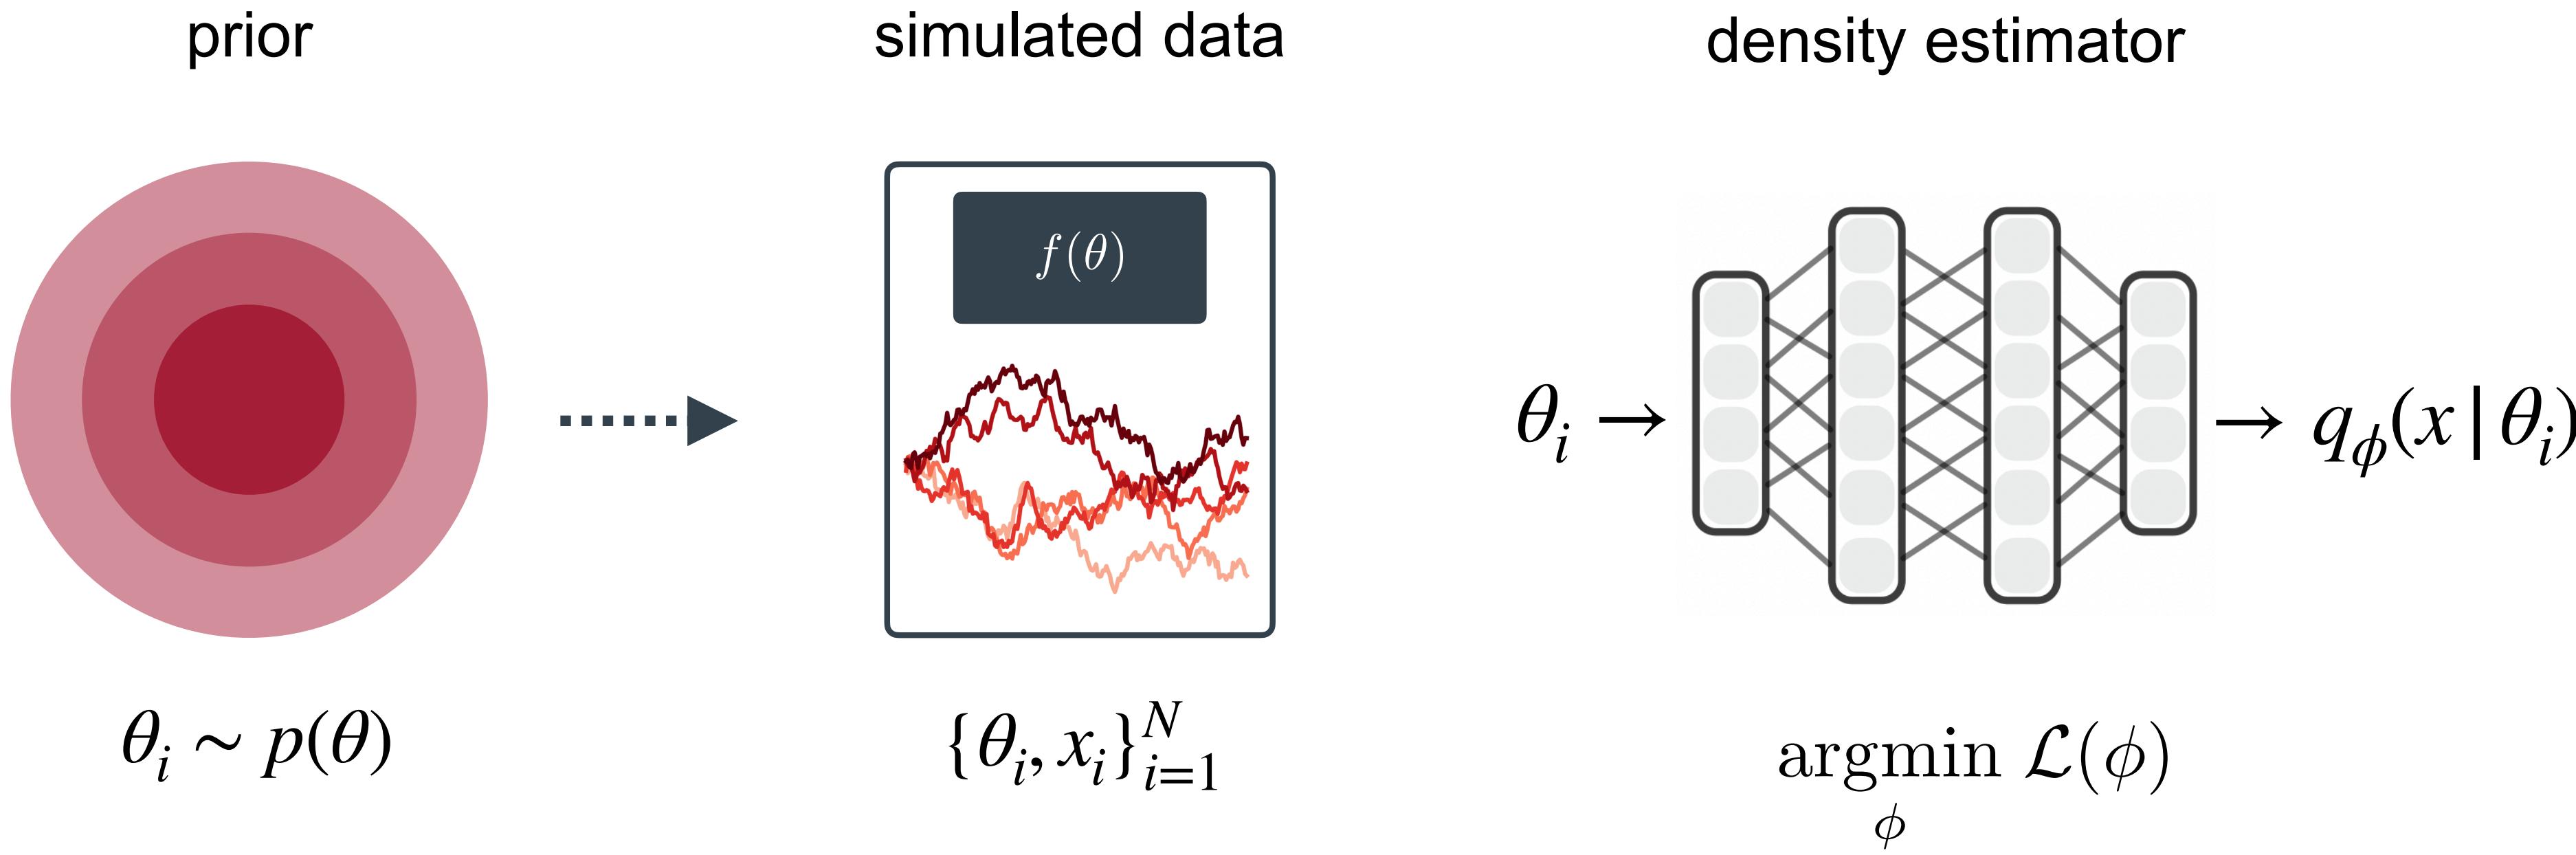
1. **Neural Posterior Estimation (NPE)**: directly learn the posterior $q_\phi(\theta | x)$
2. **Neural Likelihood Estimation (NLE)**: learn the likelihood $q_\phi(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$

Three main neural SBI approaches



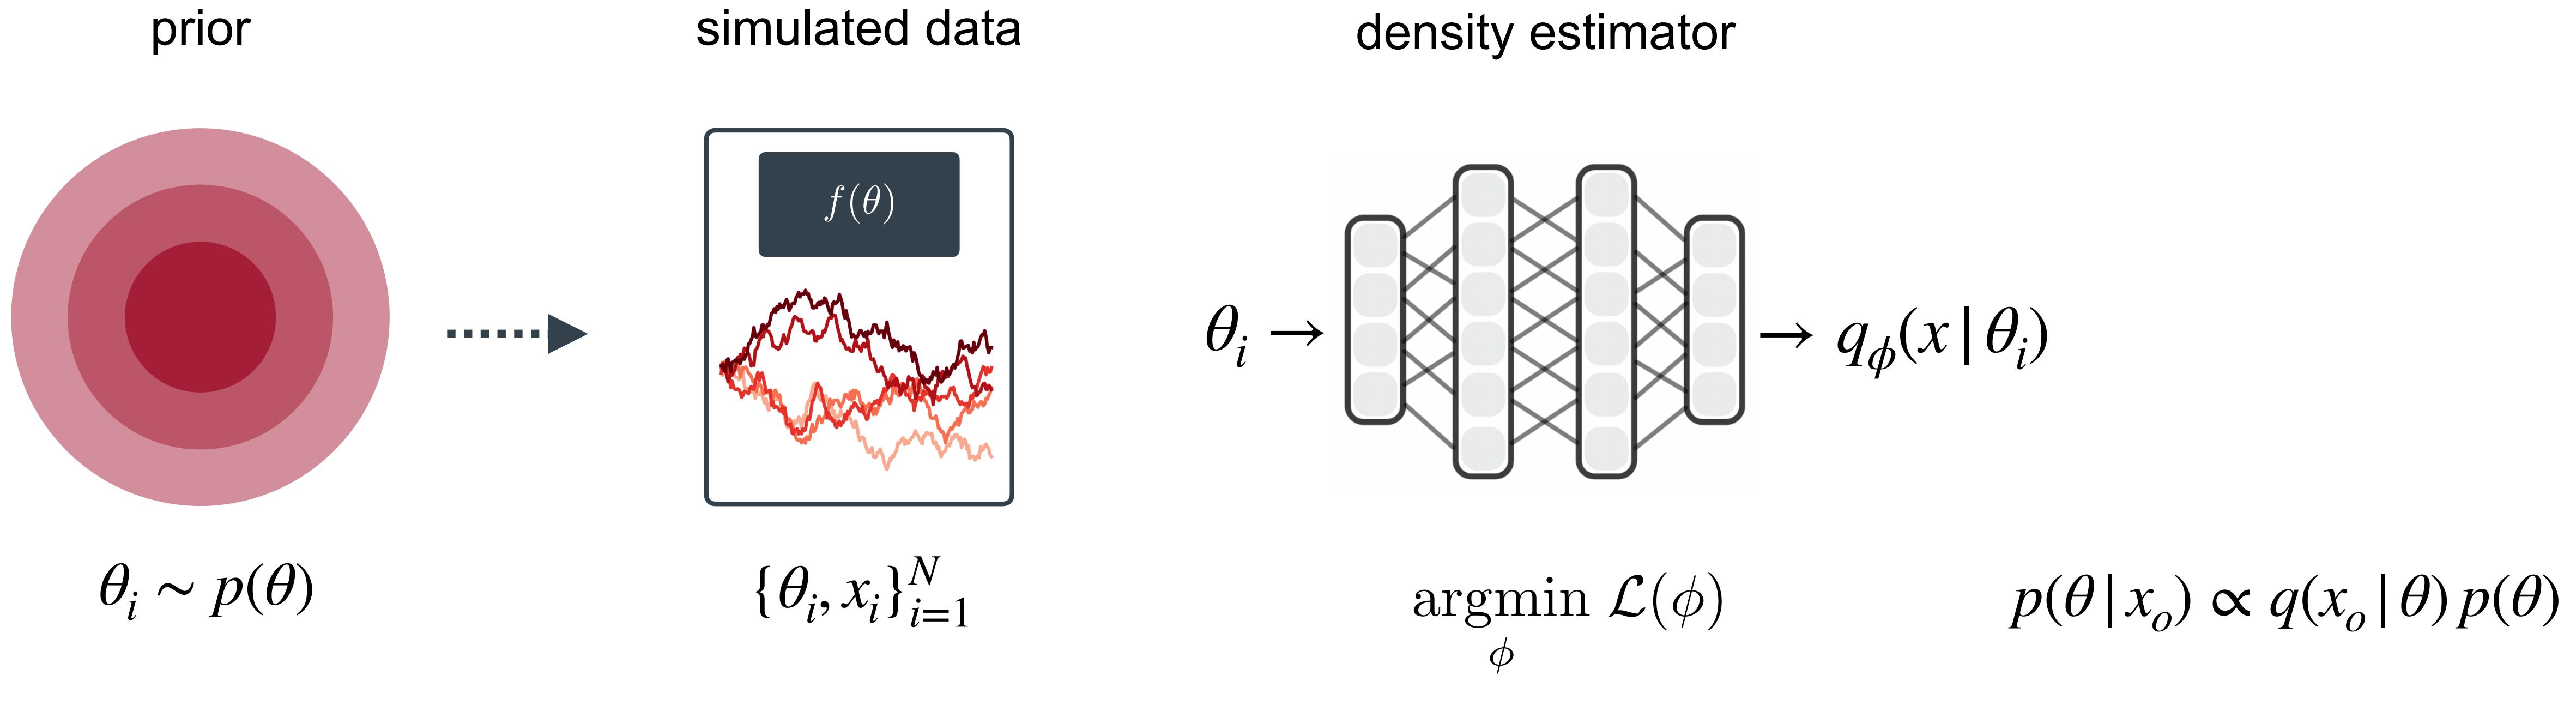
1. **Neural Posterior Estimation (NPE)**: directly learn the posterior $q_\phi(\theta | x)$
2. **Neural Likelihood Estimation (NLE)**: learn the likelihood $q_\phi(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$

Three main neural SBI approaches



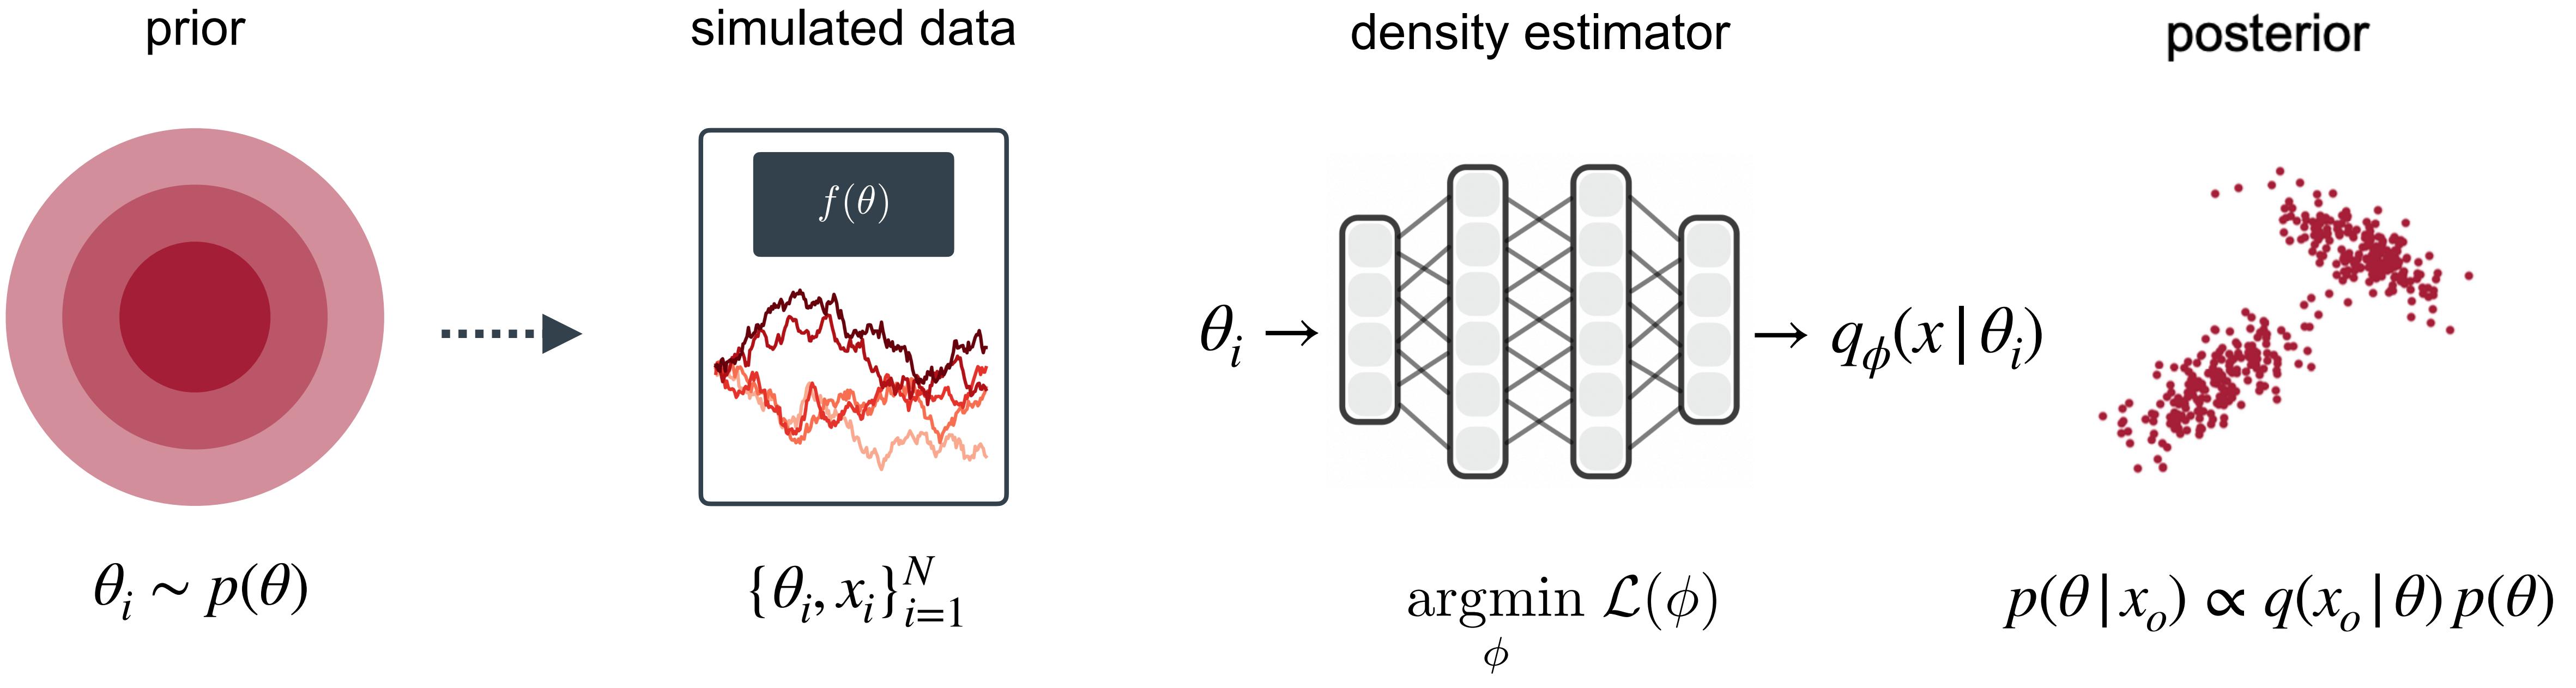
1. **Neural Posterior Estimation (NPE)**: directly learn the posterior $q_\phi(\theta | x)$
2. **Neural Likelihood Estimation (NLE)**: learn the likelihood $q_\phi(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$

Three main neural SBI approaches



1. **Neural Posterior Estimation (NPE)**: directly learn the posterior $q_\phi(\theta | x)$
2. **Neural Likelihood Estimation (NLE)**: learn the likelihood $q_\phi(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$

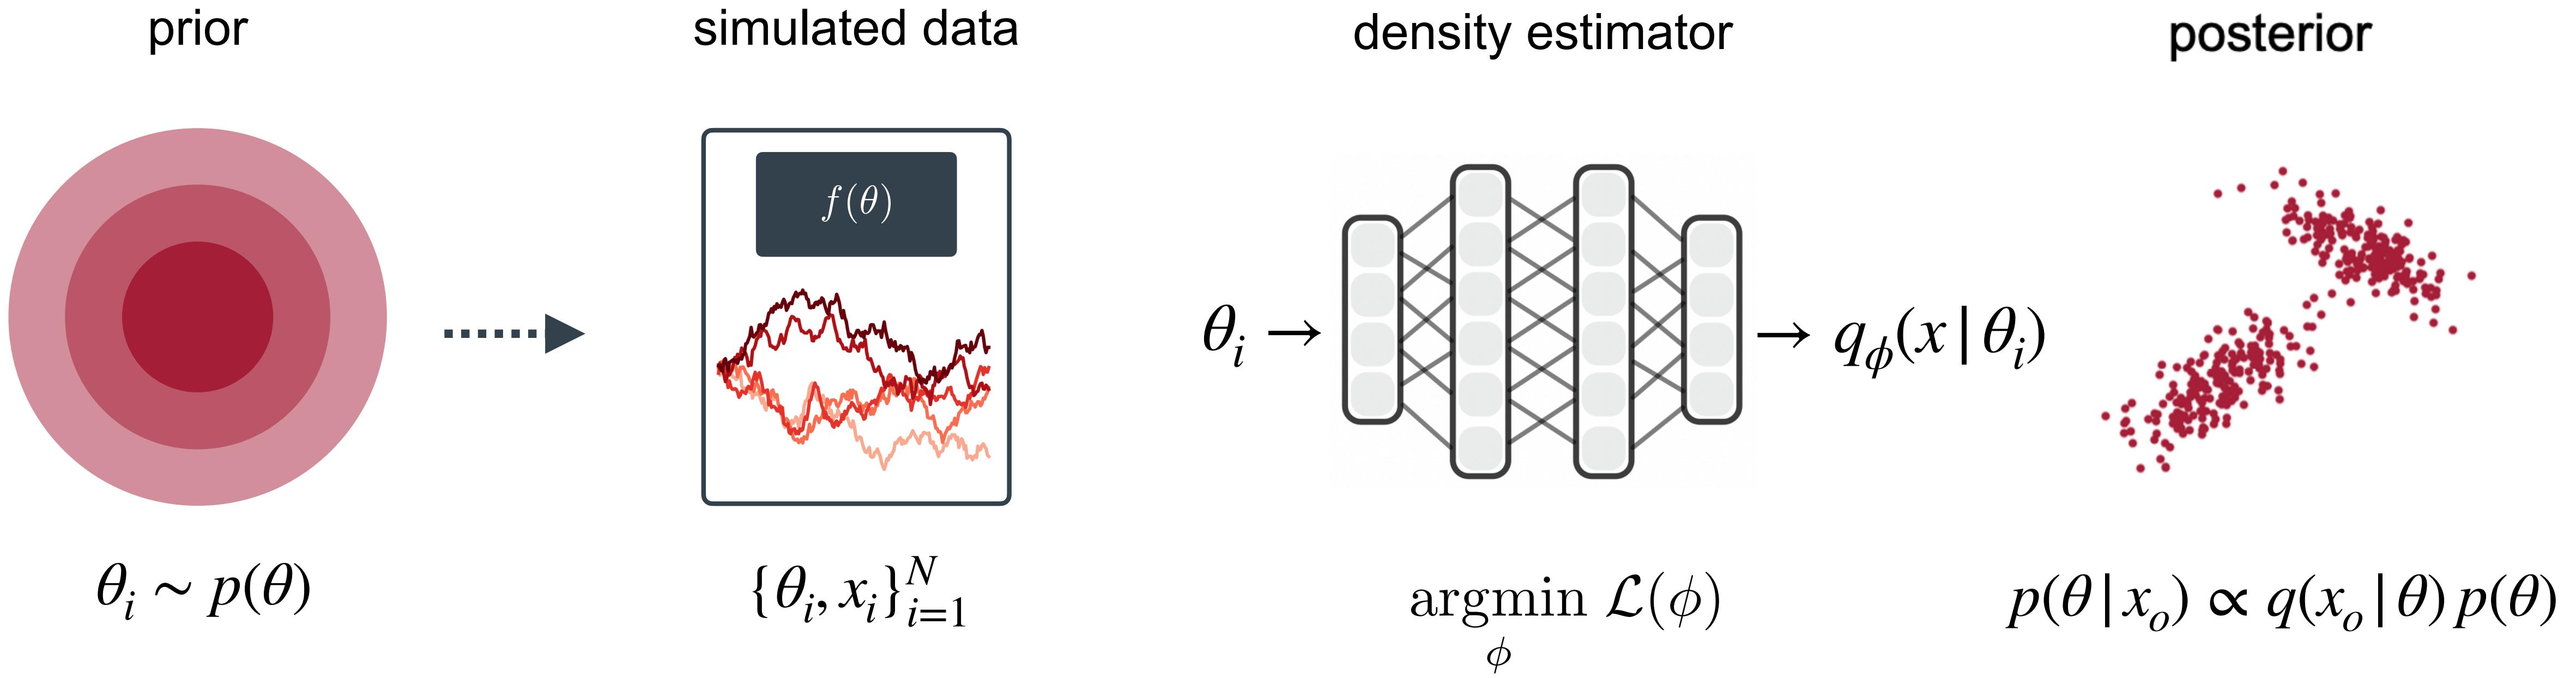
Three main neural SBI approaches



1. **Neural Posterior Estimation (NPE)**: directly learn the posterior $q_\phi(\theta | x)$

2. **Neural Likelihood Estimation (NLE)**: learn the likelihood $q_\phi(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$

Three main neural SBI approaches

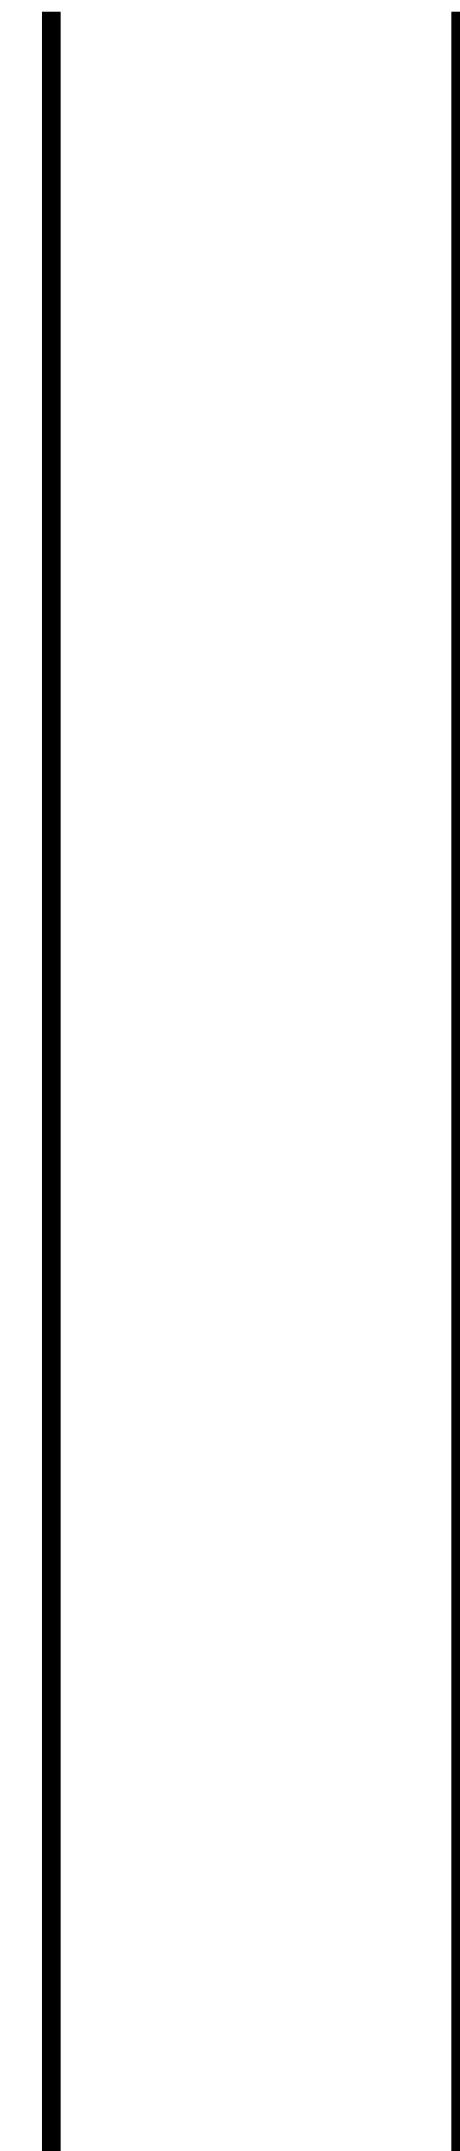


1. **Neural Posterior Estimation (NPE)**: directly learn the posterior $q_\phi(\theta | x)$
2. **Neural Likelihood Estimation (NLE)**: learn the likelihood $q_\phi(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$
3. **Neural Ratio Estimation (NRE)**: learn the likelihood ratio $r(x | \theta)$, run MCMC to sample from $p(\theta | x_o)$

Outline

1. **Why** simulators?
2. **Why** Bayesian parameter inference?
3. **How** to do Bayesian inference for simulators?
4. **How** can you apply SBI to your simulator?

Mind the Gap: Methods and Applicability of SBI



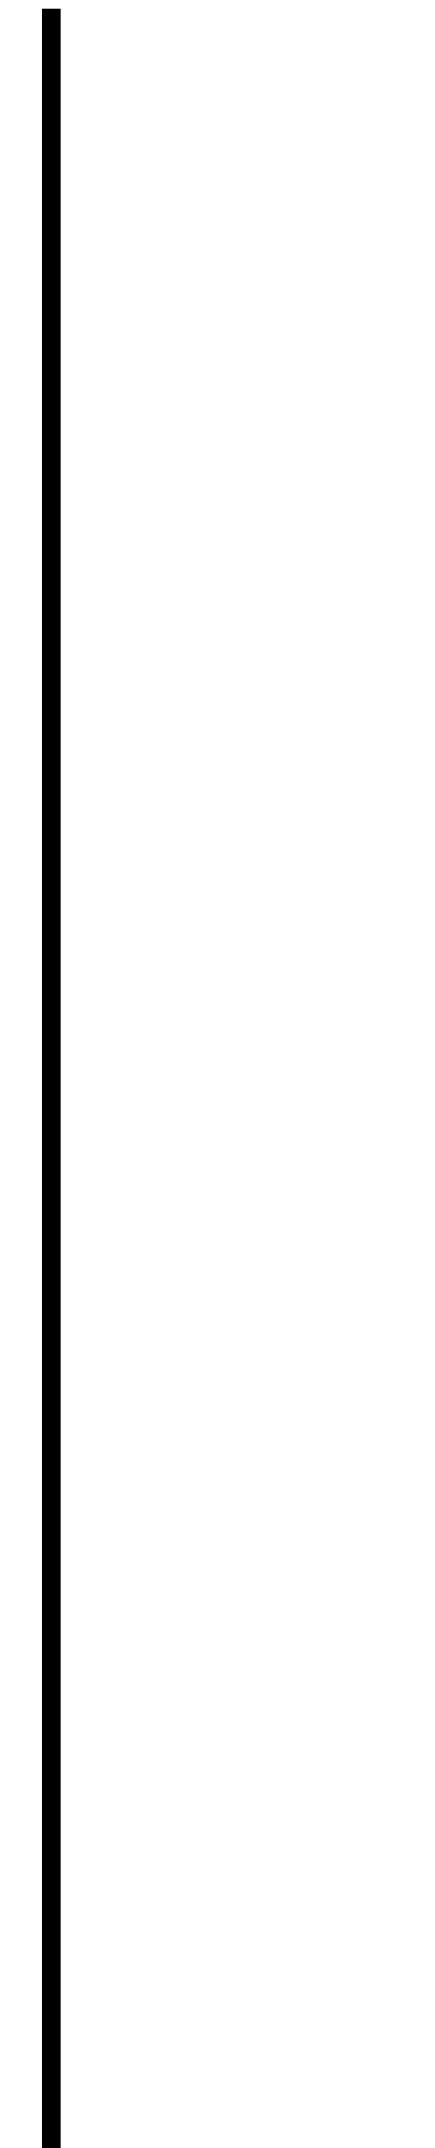
Mind the Gap: Methods and Applicability of SBI

NPE, NLE, NRE

SNLE, SNPE, SNRE

diffusion models

flow matching



Mind the Gap: Methods and Applicability of SBI

NPE, NLE, NRE

SNLE, SNPE, SNRE

diffusion models

flow matching

Fast ϵ -free Inference of Simulation Models with Bayesian Conditional Density Estimation

Sequential Neural Likelihood:
Fast Likelihood-free Inference with Autoregressive Flows

Likelihood-free MCMC with Amortized Approximate Ratio Estimators

Flow Matching for Scalable Simulation-Based Inference

Mind the Gap: Methods and Applicability of SBI

NPE, NLE, NRE

SNLE, SNPE, SNRE

diffusion models

flow matching

Fast ϵ -free Inference of Simulation Models with Bayesian Conditional Density Estimation

Sequential Neural Likelihood:
Fast Likelihood-free Inference with Autoregressive Flows

Likelihood-free MCMC with Amortized Approximate Ratio Estimators

Flow Matching for Scalable Simulation-Based Inference

- Which method to use?
- How to use it?
- Is there usable code?
- How does it compare to other methods
- ...

Mind the Gap: Methods and Applicability of SBI

NPE, NLE, NRE
SNLE, SNPE, SNRE
diffusion models
flow matching



sbi

Fast ϵ -free Inference of Simulation Models with Bayesian Conditional Density Estimation

Sequential Neural Likelihood:
Fast Likelihood-free Inference with Autoregressive Flows

Likelihood-free MCMC with Amortized Approximate Ratio Estimators

Flow Matching for Scalable Simulation-Based Inference

Mind the Gap: Methods and Applicability of SBI

NPE, NLE, NRE
SNLE, SNPE, SNRE
diffusion models
flow matching

Fast ϵ -free Inference of Simulation Models with Bayesian Conditional Density Estimation

Sequential Neural Likelihood:
Fast Likelihood-free Inference with Autoregressive Flows

Likelihood-free MCMC with Amortized Approximate Ratio Estimators

Flow Matching for Scalable Simulation-Based Inference

The sbi logo consists of the lowercase letters "sbi" in a bold, black font, enclosed within a blue circular arrow that forms a loop.

```
import torch
from sbi.inference import NPE

# define shifted Gaussian simulator.
def simulator(theta):
    return theta + torch.randn_like(theta)

# draw parameters from Gaussian prior.
theta = torch.randn(1000, 2)
# simulate data
x = simulator(theta)

# choose sbi method and train
inference = NPE()
inference.append_simulations(theta, x).train()

# do inference given observed data
x_o = torch.ones(2)
posterior = inference.build_posterior()
samples = posterior.sample((1000,), x=x_o)
```



An accessible toolkit for applying SBI

Home Search

sbi
[Home](#)
[Tutorials and Examples](#)
[API Reference](#)
[FAQ](#)
[Contributing](#) >
[Citation](#)
[Credits](#)

Overview

To get started, install the `sbi` package with:

```
pip install sbi
```

for more advanced install options, see our [Install Guide](#).

Then, check out our material:

[Motivation and approach](#)
General motivation for the SBI framework and methods included in sbi.

[Tutorials and Examples](#)
Various examples illustrating how to get started or use the sbi package.

[Reference API](#)
The detailed description of the package classes and functions.

[Citation](#)
How to cite the sbi package.



An accessible toolkit for applying SBI

- Off-the-shelf SBI methods
- Embedding networks
- Validation techniques
- Plotting tools

The screenshot shows the homepage of the sbi documentation. At the top, there is a dark blue header bar with a white navigation bar below it. The navigation bar includes a "Home" button with a house icon, a search bar with a magnifying glass icon, and a "sbi" logo. The main content area has a light gray background. On the left, there is a sidebar with links to various sections: Home, Tutorials and Examples, API Reference, FAQ, Contributing, Citation, and Credits. The main content area starts with the title "Overview". Below it, there is a section titled "To get started, install the `sbi` package with:" followed by a command line interface (CLI) command: "pip install sbi". There is also a small icon of a rocket ship. To the right of this, there is a section titled "for more advanced install options, see our [Install Guide](#)". Below these sections, there are four cards: "Motivation and approach" (with a gear icon), "Tutorials and Examples" (with a rocket ship icon), "Reference API" (with a book icon), and "Citation" (with a book icon). Each card has a brief description of its purpose.

Home

Search

sbi

Home

Tutorials and Examples

API Reference

FAQ

Contributing

Citation

Credits

Overview

To get started, install the `sbi` package with:

```
pip install sbi
```

for more advanced install options, see our [Install Guide](#).

Then, check out our material:

[Motivation and approach](#)
General motivation for the SBI framework and methods included in `sbi`.

[Tutorials and Examples](#)
Various examples illustrating how to [get started](#) or use the `sbi` package.

[Reference API](#)
The detailed description of the package classes and functions.

[Citation](#)
How to cite the `sbi` package.



An accessible toolkit for applying SBI

- Off-the-shelf SBI methods
- Embedding networks
- Validation techniques
- Plotting tools
- Thorough documentation
- Tutorials and examples
- Issues, discussions

The screenshot shows the homepage of the sbi package documentation. At the top, there is a dark blue header bar with a white navigation bar below it. The navigation bar includes a 'Home' button with a house icon, a search bar with a magnifying glass icon, and a 'sbi' logo. The main content area has a light gray background. On the left, there is a sidebar with the 'sbi' logo and links to 'Home', 'Tutorials and Examples', 'API Reference', 'FAQ', 'Contributing', 'Citation', and 'Credits'. To the right of the sidebar, the word 'Overview' is displayed. Below it, a section titled 'To get started, install the `sbi` package with:' contains the command `pip install sbi`. A small icon of a rocket ship is positioned to the right of the command. Below this, a note says 'for more advanced install options, see our [Install Guide](#)'. Further down, another note says 'Then, check out our material:' followed by four cards: 'Motivation and approach' (with a gear icon), 'Reference API' (with a book icon), 'Tutorials and Examples' (with a rocket ship icon), and 'Citation' (with a book icon). Each card has a brief description.

Home

Search

sbi

Home

Tutorials and Examples

API Reference

FAQ

Contributing

Citation

Credits

Overview

To get started, install the `sbi` package with:

```
pip install sbi
```

for more advanced install options, see our [Install Guide](#).

Then, check out our material:

[Motivation and approach](#)
General motivation for the SBI framework and methods included in `sbi`.

[Reference API](#)
The detailed description of the package classes and functions.

[Tutorials and Examples](#)
Various examples illustrating how to [get started](#) or use the `sbi` package.

[Citation](#)
How to cite the `sbi` package.

Demo: SBI on the SIR model

SBI toolkit has been adopted by the community

SBI toolkit has been adopted by the community

Article

A biophysical account of multiplication by a single neuron

PNAS

RESEARCH ARTICLE

NEUROSCIENCE

Energy-efficient network activity from disparate circuit parameters

Michael Deistler^a, Jakob H. Macke^{a,b,1,2}, and Pedro J. Gonçalves^{a,c,1,2}

Using simulation-based inference to determine the parameters of an integrated hydrologic model: a case study from the upper Colorado River basin

Robert Hull , Elena Leonardiuzzi, Luis De la Fuente, Hoang Viet Tran, Andrew Bennett, Peter Melchior, Reed D

and I
THE ASTROPHYSICAL JOURNAL

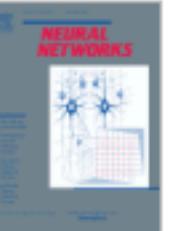
OPEN ACCESS

Calibrating Cosmological Simulations with Implicit Likelihood Inference Using Galaxy Growth Observables



Neural Networks

Volume 163, June 2023, Pages 178-194



Amortized Bayesian inference on generative dynamical network models of epilepsy using deep neural density estimators

Journal of Neural Engineering

 , Viktor Sip ^a 
 irsa ^a  

PAPER

Fast inference of spinal neuromodulation for motor control using amortized neural networks

Calibrating Agent-based Models to Microdata with Graph Neural Networks

Bayesian Object Models for Robotic Interaction with Differentiable Probabilistic Programming

Krishna Murthy Jatavallabhula, Miles
Proceedings of The 6th Conference on

Farmer ^{1,2,4} Sebastian M. Schmon ^{3,5}

Neural simulation-based inference approach for characterizing the Galactic Center γ -ray excess

Siddharth Mishra-Sharma and Kyle Cranmer
Phys. Rev. D **105**, 063017 – Published 23 March 2022

SBI toolkit has been adopted by the community

Article

A biophysical account of multiplication by a single neuron

PNAS

RESEARCH ARTICLE

NEUROSCIENCE

Energy-efficient network activity from circuit parameters

Michael Deistler^a, Jakob H. Macke^{a,b,1,2}, and Pedro J. Gonçalves^{a,c,1,2}

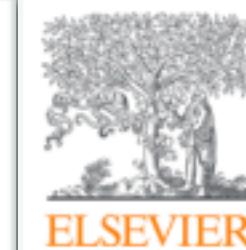
Using simulation-based inference to determine the of an integrated hydrologic model: a case study from the upper Colorado River basin

Robert Hull¹, Elena Leonardiuzzi¹, Luis De la Fuente¹, Hoang Viet Tran¹, Andrew Bennett¹, Peter Melchior¹, Reed Tsoo¹, and I

THE ASTROPHYSICAL JOURNAL

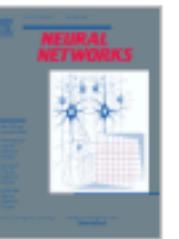
OPEN ACCESS

Calibrating Cosmological Simulations with Implicit Likelihood Inference Using Galaxy Growth Observables



Neural Networks

Volume 163, June 2023, Pages 178-194



Amortized Bayesian inference on generative dynamical network models of epilepsy using deep neural density estimators

Journal of Neural Engineering

Viktor Sip^a,
Irsa^a

>50 contributors

>500 GitHub stars

of spinal neuromodulation for motor control
ed neural networks

Bayesian Object Models for Robotic Interaction with Differentiable Probabilistic Programming

Krishna Murthy Jatavallabhula, Miles
Proceedings of The 6th Conference on

Farmer^{1,2,4}, Sebastian M. Schmon^{3,5}

Neural simulation-based inference approach for characterizing the Galactic Center γ -ray excess

Siddharth Mishra-Sharma and Kyle Cranmer
Phys. Rev. D **105**, 063017 – Published 23 March 2022



New contributors welcome!

NUMFOCUS
OPEN CODE = BETTER SCIENCE



New contributors welcome!



- **A lot of potential** in scientific and industrial applications



New contributors welcome!



- **A lot of potential** in scientific and industrial applications
- **A lot to improve:** new methods, API design, documentation



New contributors welcome!



- **A lot of potential** in scientific and industrial applications
- **A lot to improve:** new methods, API design, documentation
- Actively maintained



New contributors welcome!



- **A lot of potential** in scientific and industrial applications
- **A lot to improve:** new methods, API design, documentation
- Actively maintained
- Upcoming events:



New contributors welcome!



- **A lot of potential** in scientific and industrial applications
- **A lot to improve:** new methods, API design, documentation
- Actively maintained
- Upcoming events:
 - **Tomorrow** 9am-11am “New Contributor Sprint”



New contributors welcome!



- **A lot of potential** in scientific and industrial applications
- **A lot to improve:** new methods, API design, documentation
- Actively maintained
- Upcoming events:
 - **Tomorrow** 9am-11am “New Contributor Sprint”
 - Spring 2025: SBI Hackathon & Workshop

Acknowledgments

Contributors 57

+ 43 contributors

Used by 122

+ 114

MLS

EBERHARD KARLS
UNIVERSITÄT
TÜBINGEN

ai TransferLab



Tübingen AI Center

ai appliedAI
institute
for europe

Thank you!

Questions?

Try it out: **pip install sbi**

GitHub: github.com/sbi-dev/sbi

X, Mastodon: @sbi_devs

TransferLab website: transferlab.ai

sbi sbi developers
@sbi_devs ...

We just released a new version of `sbi`, and this one has _a ton_ of new features! Many of these features are thanks to more than 30 (mostly new) contributors. We are very excited about the growing community and the new release! 1/8

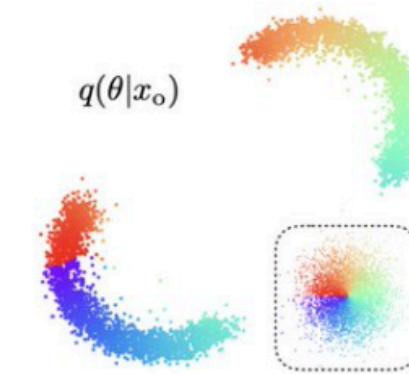
sbi v0.23

pip install sbi --upgrade

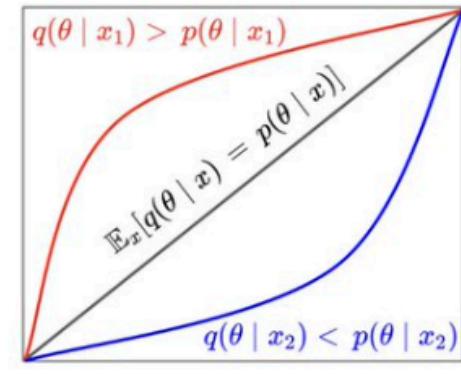
Full control over the training loop

```
net = build_nsf(theta, x)
opt = Adam(list(net.parameters()))
for _ in range(100):
    opt.zero_grad()
    losses = net.loss(theta, x)
    loss = torch.mean(losses)
    loss.backward()
    opt.step()
```

Flow-matching & score-matching



New diagnostics: L-C2ST & TARP



Dax et al., 2023 Linhart et al., 2023 & Lemos et al., 2023

9:34 AM · Aug 29, 2024 · 1,279 Views

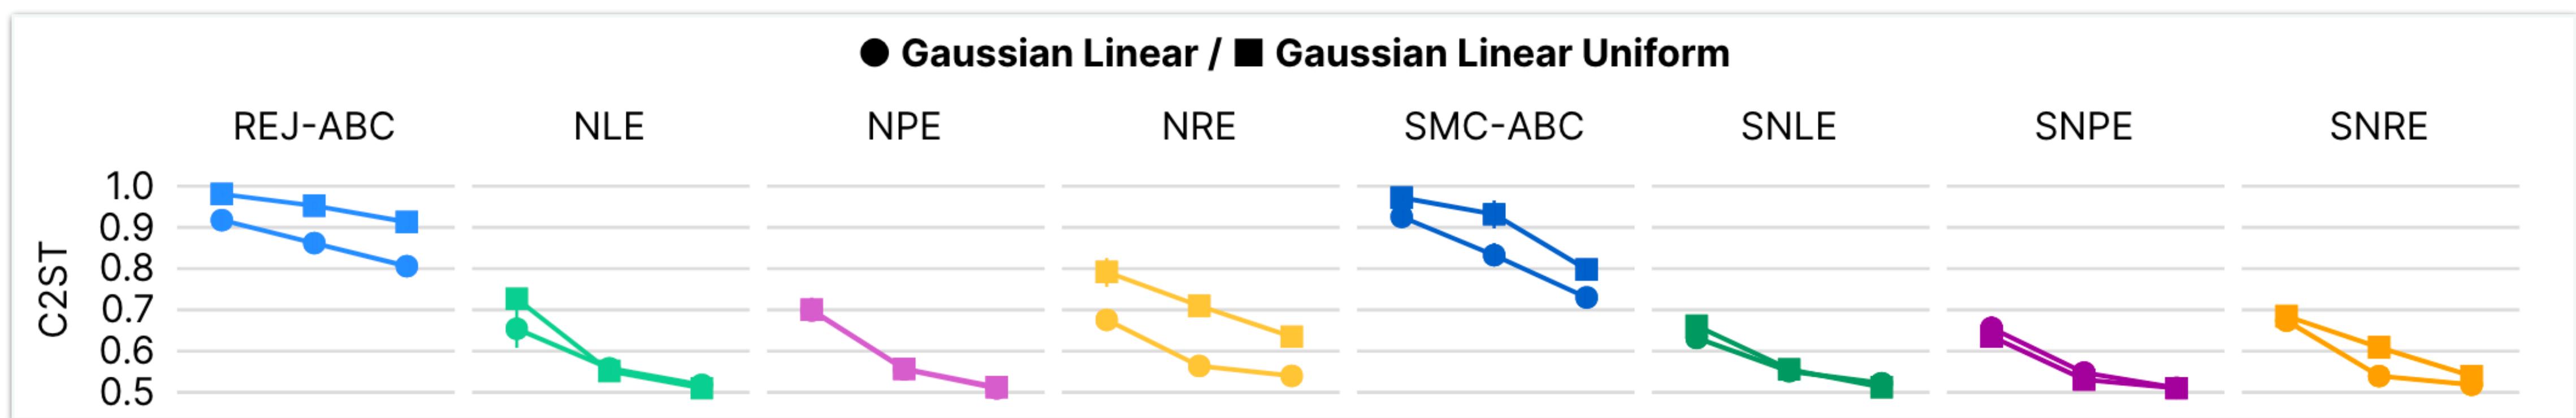
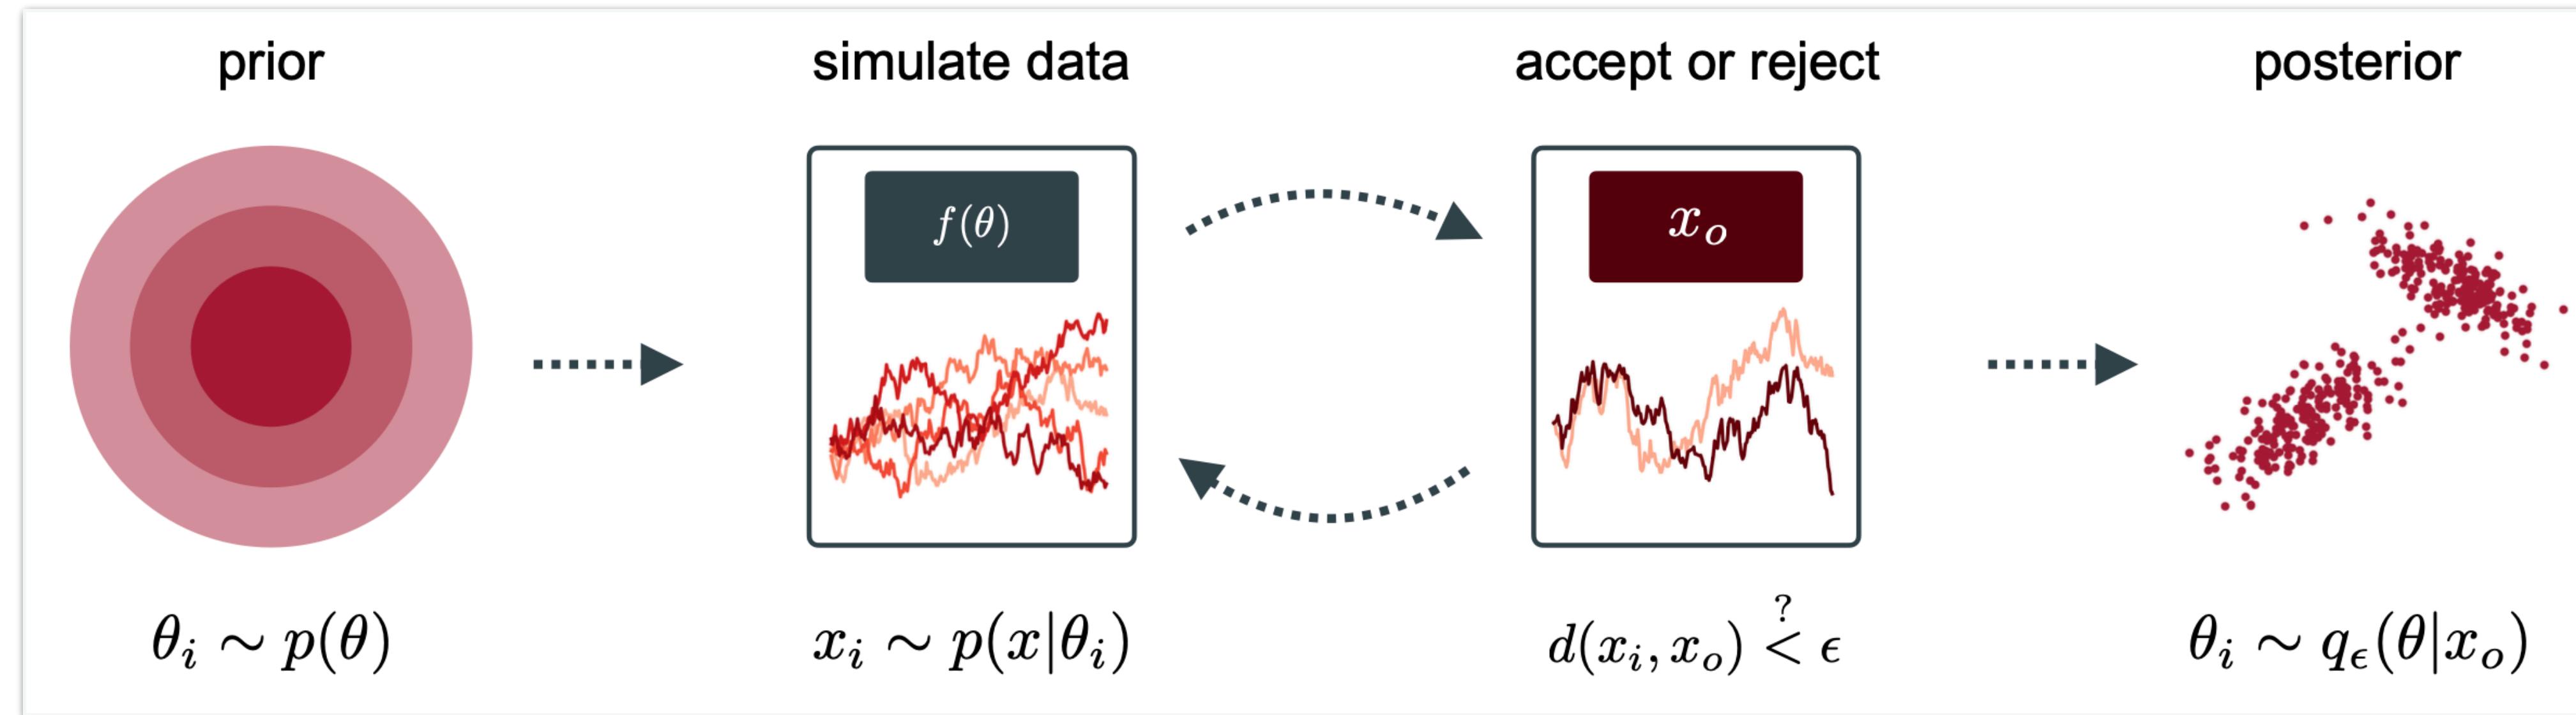
Supplementary slides

Is SBI for you?

Some heuristics

- Has been applied to simulators with ~10 parameters, ~100 in some cases
- Depends on how much data you can generate: runtime ~sec
- High-dimensional data can be embedded: CNNs, RNNs, transformers etc.
- Different methods for different settings: i.i.d. data, hierarchical models, amortized inference, sequential inference

Approximate Bayesian Computation (ABC)

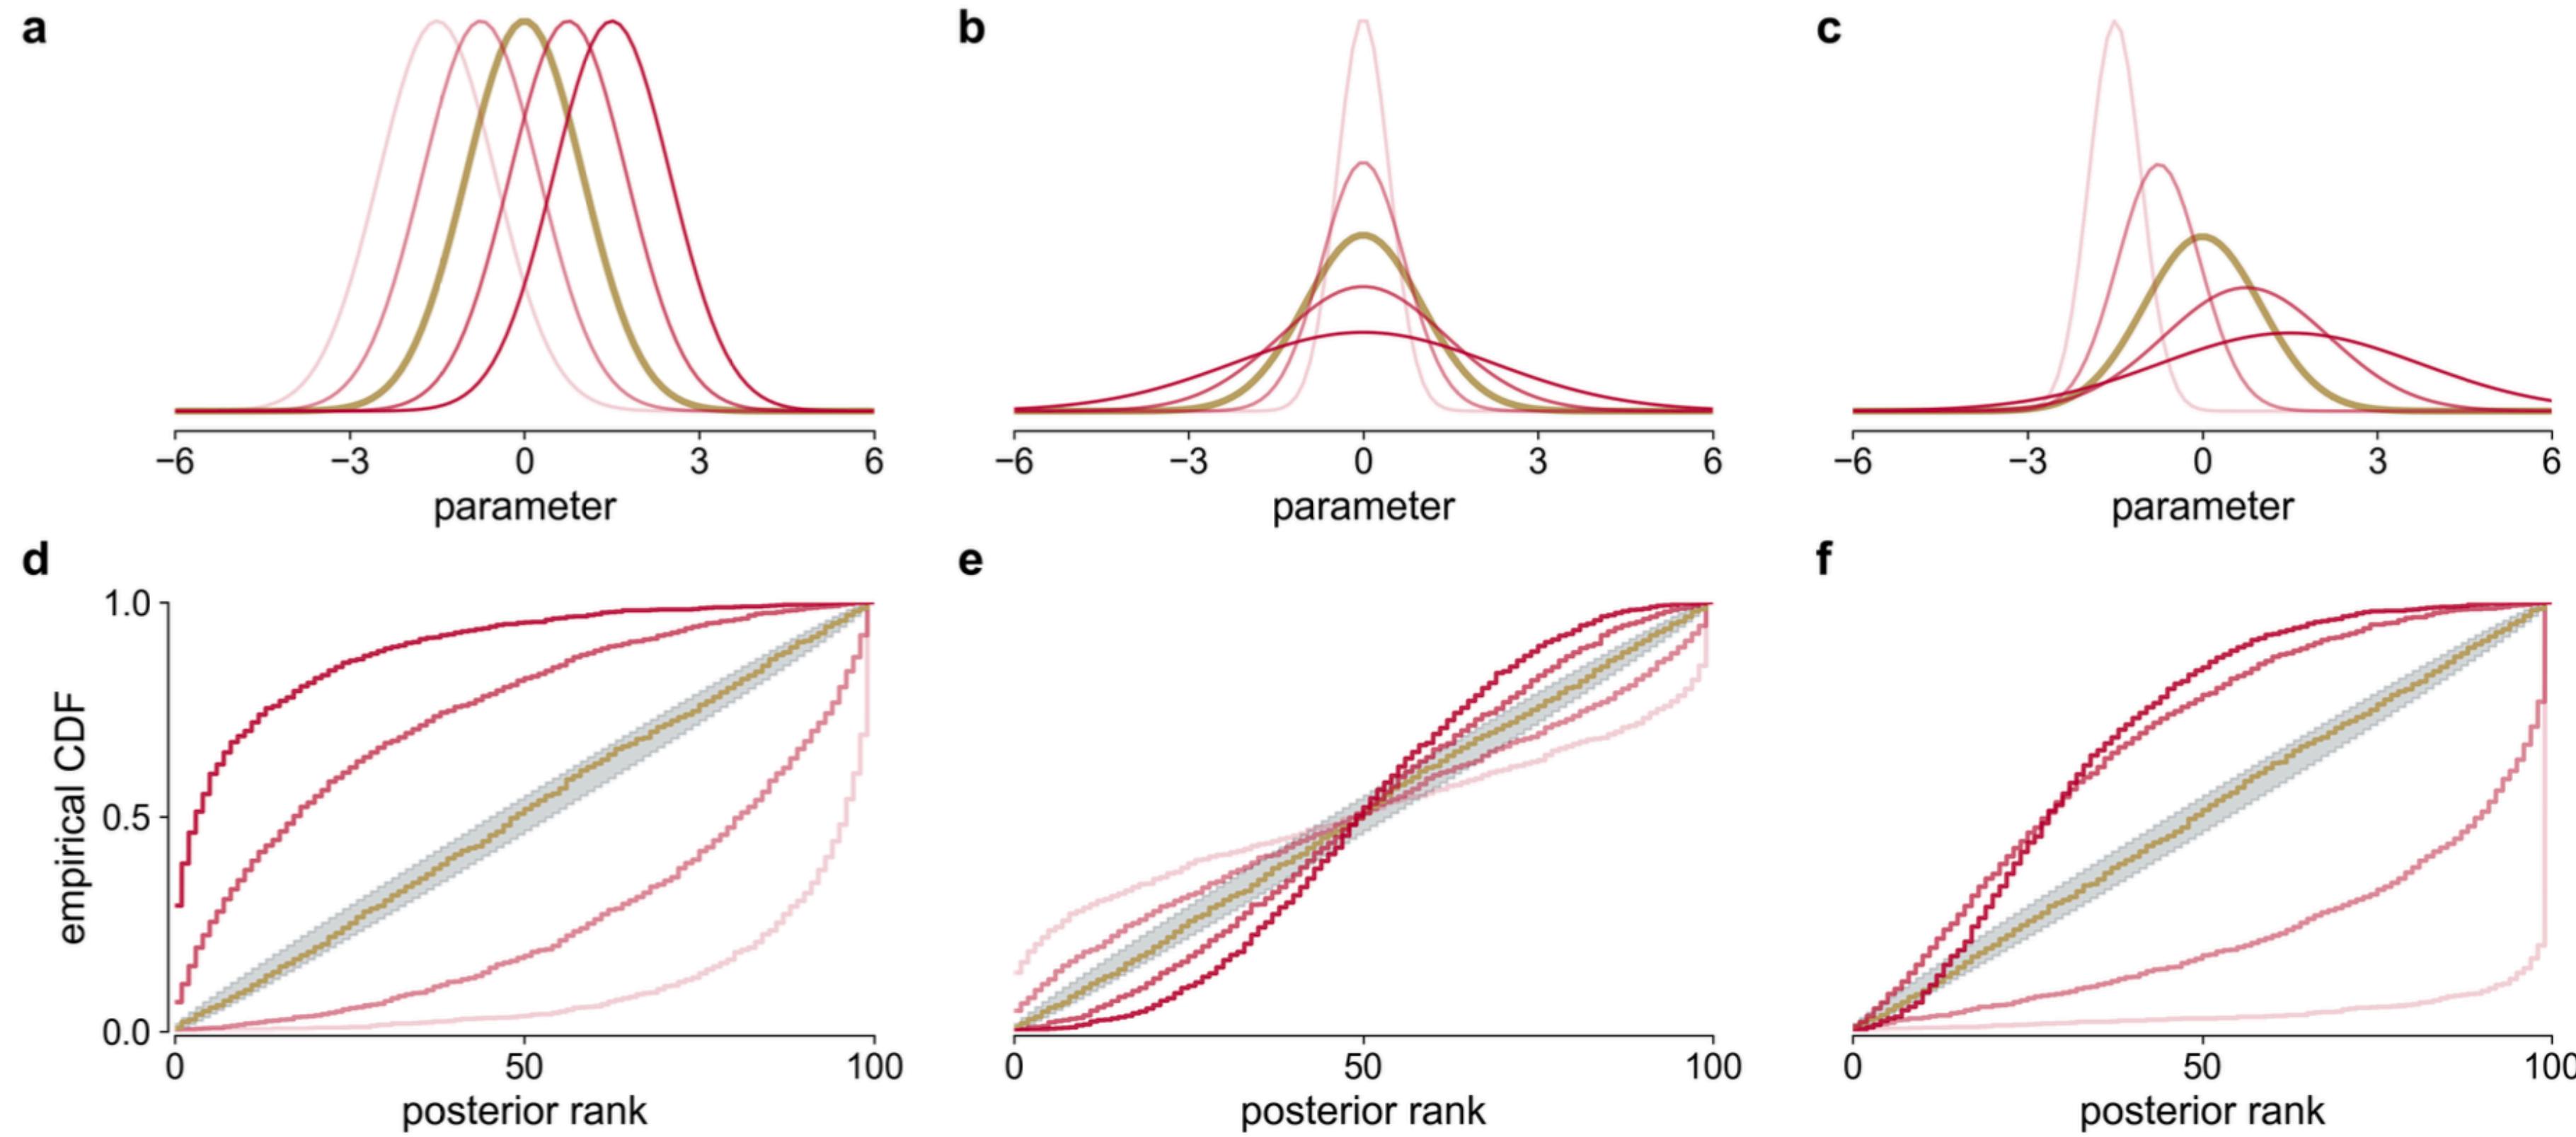


NLE / NPE loss functions

Algorithm 5: Single round Neural Posterior Estimation as in Papamakarios and Murray (2016)

```
for  $j = 1 : N$  do
    | Sample  $\boldsymbol{\theta}_j \sim p(\boldsymbol{\theta})$ 
    | Simulate  $\mathbf{x}_j \sim p(\mathbf{x}|\boldsymbol{\theta}_j)$ 
end
 $\boldsymbol{\phi} \leftarrow \arg \min \sum_j^N -\log q_{F(\mathbf{x}_j, \boldsymbol{\phi})}(\boldsymbol{\theta}_j)$ 
Set  $\hat{p}(\boldsymbol{\theta}|\mathbf{x}_o) = q_{F(\mathbf{x}_o, \boldsymbol{\phi})}(\boldsymbol{\theta})$ 
return Samples from  $\hat{p}(\boldsymbol{\theta}|\mathbf{x}_o); q_{F(\mathbf{x}, \boldsymbol{\phi})}(\boldsymbol{\theta})$ 
```

Simulation-based calibration



$$\theta^* \sim p(\theta)$$

$$x^* \sim p(x|\theta^*)$$

$$\{\theta_1, \dots, \theta_L\} \sim p_i(\theta|x^*)$$

$$r_i = \sum_{j=1}^L \mathbb{I}[f(\theta_j) < f(\theta_i^*)] \in [0, L],$$

Model misspecification in SBI

- **Notation:** $x \sim p(x | \theta)$ simulated data; $y \sim p^*(y)$ observed data
- **Problem:** simulator is misspecified
- **Approach:** assume an error model $\tilde{x} \sim p(x | y)$, e.g., “spike and slab”

Algorithm 1: Robust neural posterior estimation (RNPE)

```
for  $i$  in  $1 : N$  do
  1   Sample  $\theta_i \sim p(\theta)$ 
  2   Simulate  $x_i \sim p(x | \theta)$ 
end
3 Train NPE  $q(\theta | x)$  on  $\{(\theta_i, x_i)\}_{i=1}^N$ 
4 Train  $q(x)$  on  $\{x_i\}_{i=1}^N$ 
5 Sample  $\tilde{x}_m \sim \hat{p}(x | y_o) \propto p(y_o | x)q(x)$ ,  $m = 1, \dots, M$  using MCMC
6 Sample  $\tilde{\theta}_m \sim q(\theta | \tilde{x}_m)$ ,  $m = 1, \dots, M$ 
return  $\{(\tilde{\theta}_m, \tilde{x}_m)\}_{m=1}^M$ , samples drawn approximately from  $p(\theta, x | y_o)$ 
```
